

SPARSE RESERVOIR COMPUTING

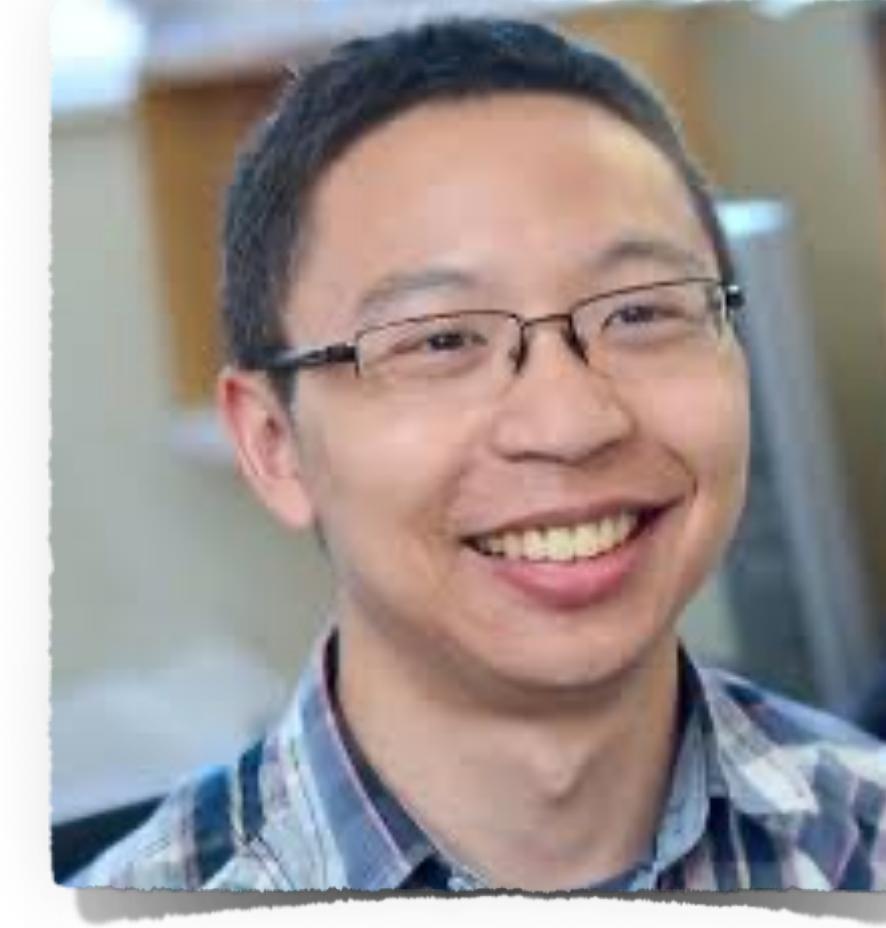
Eleni Vasilaki, University of Sheffield

WORK WITH



Luca Manneschi

Manneschi et al (2021) IEEE Trans. Neural Networks & Learning Systems <https://arxiv.org/abs/1907.08600>,



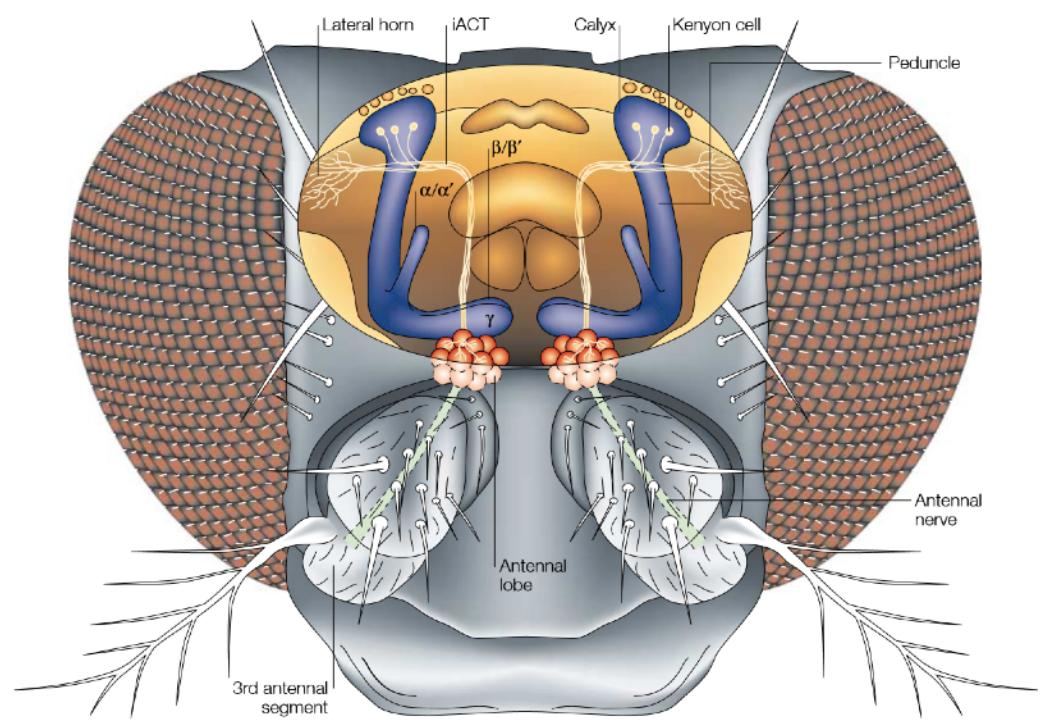
Andrew Lin



EPSRC

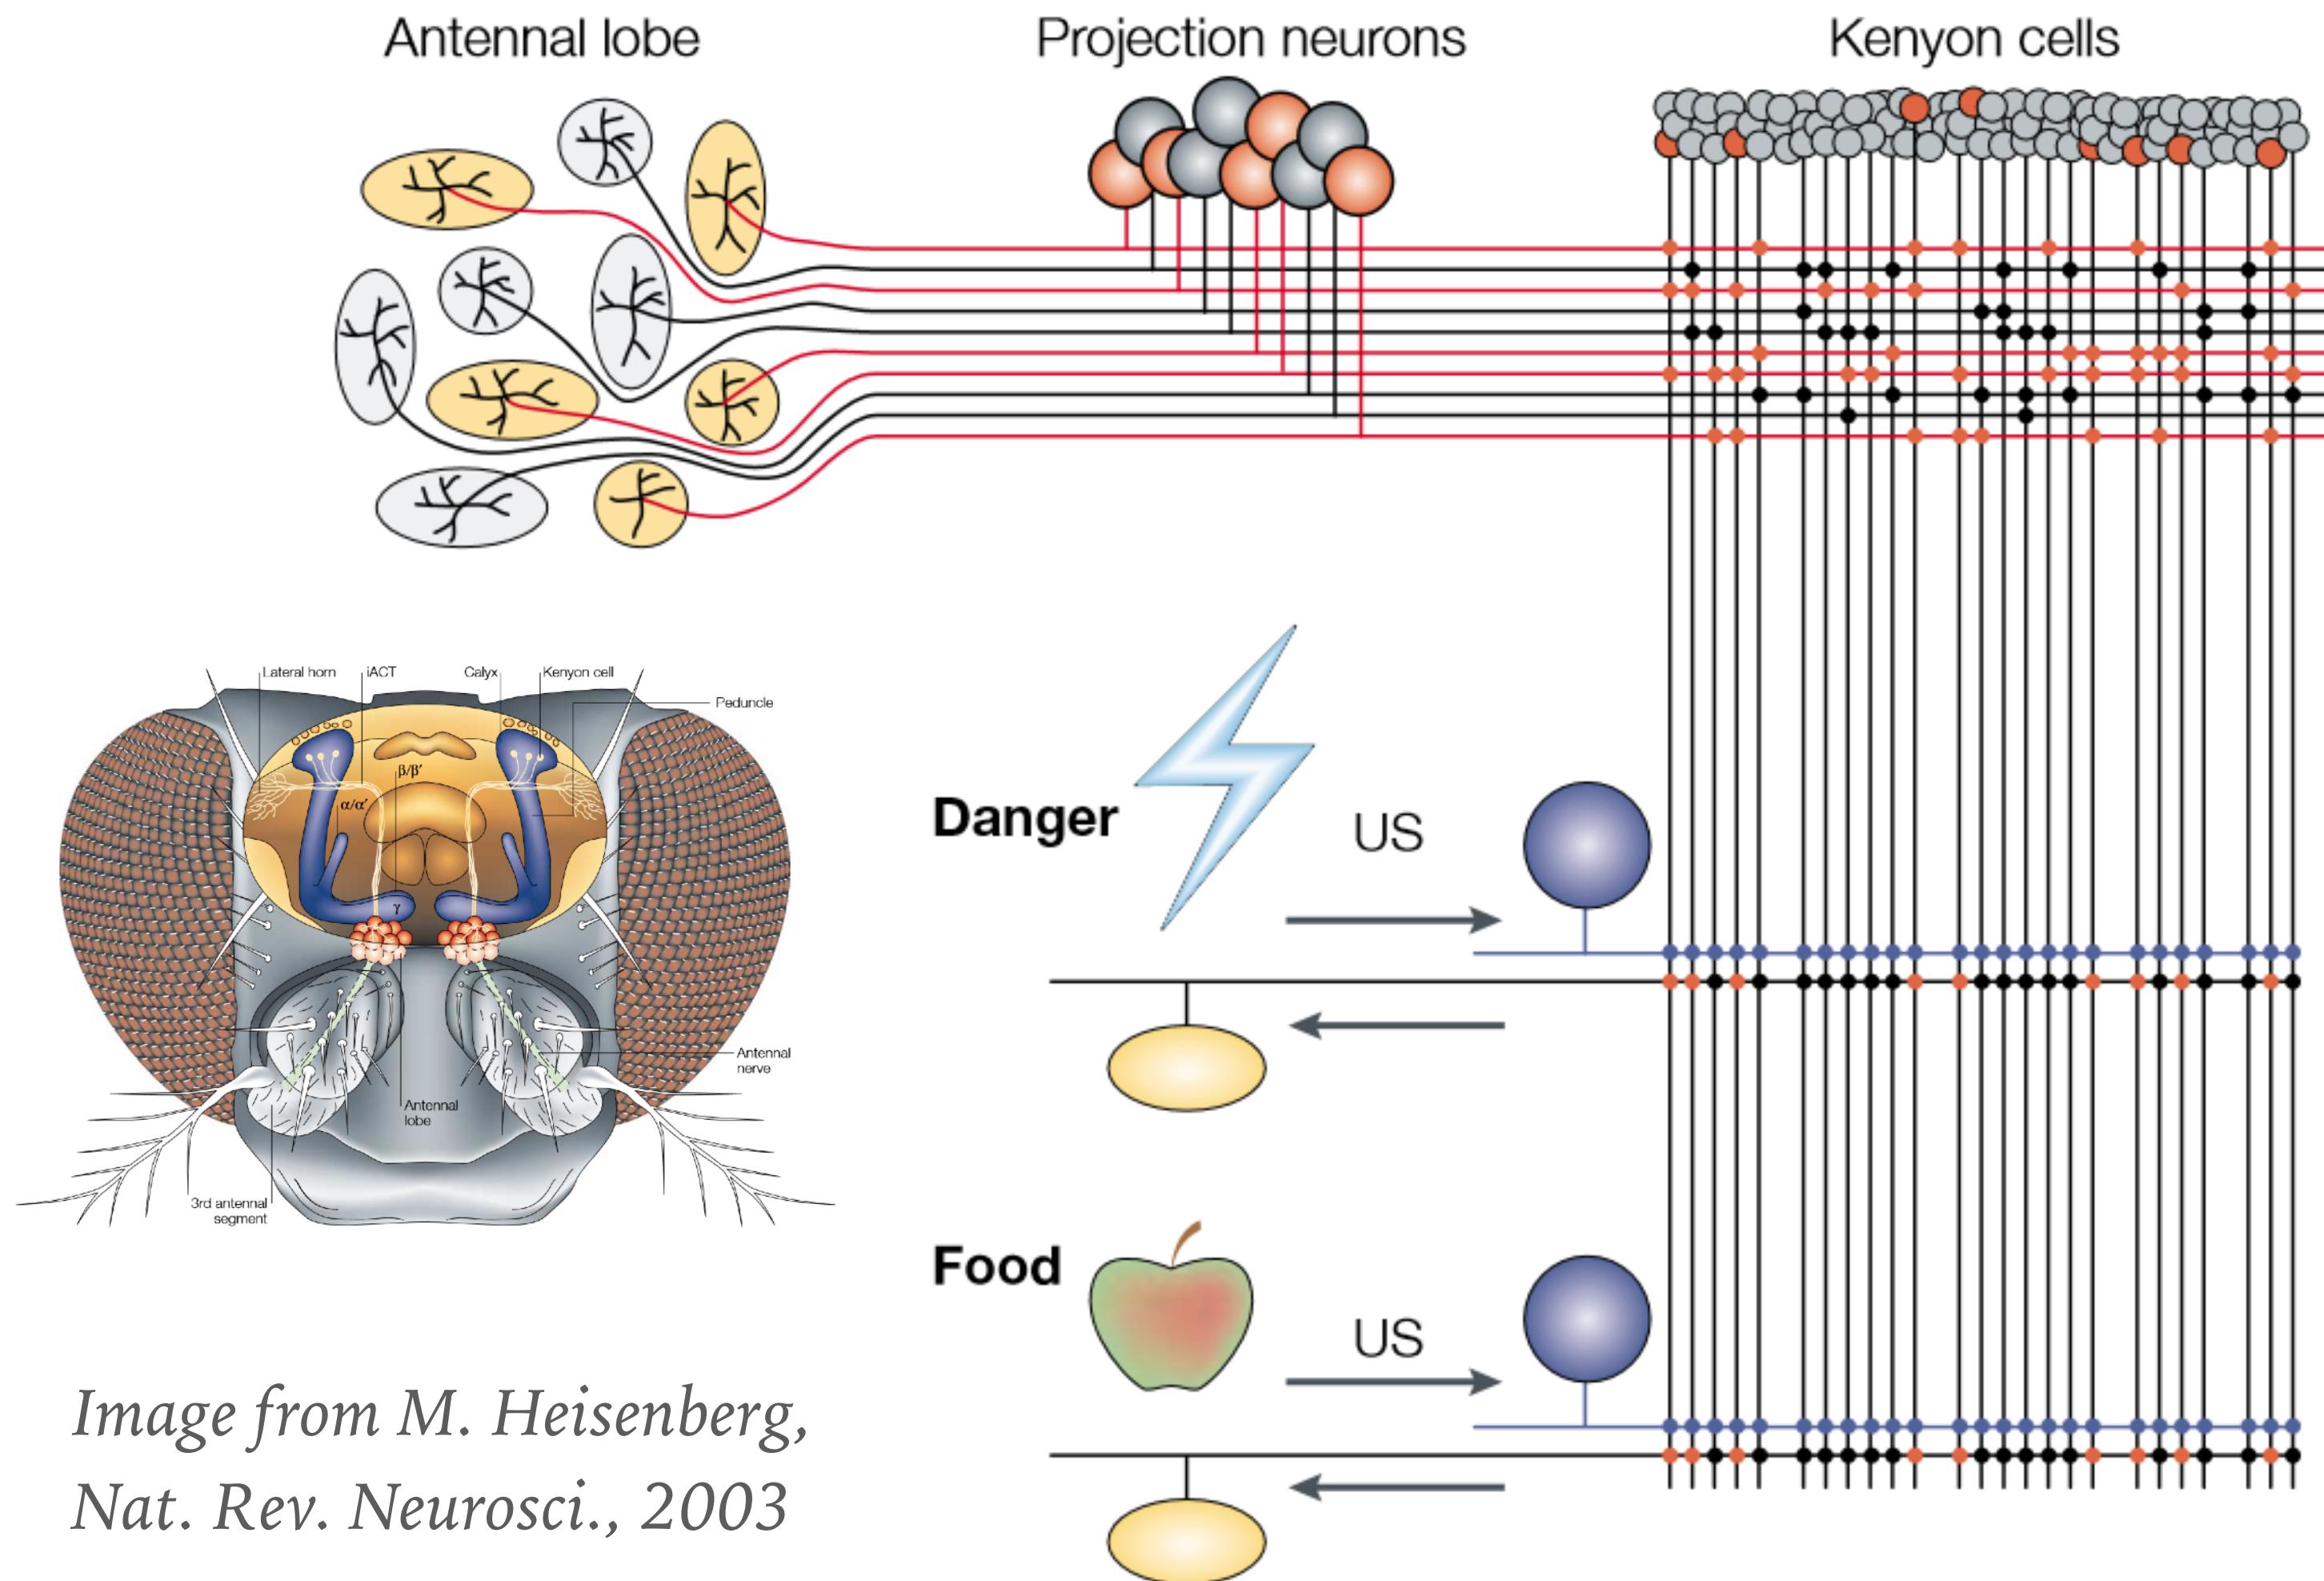
Engineering and Physical Sciences
Research Council

LEARNING ASSOCIATIONS: THE MUSHROOM BODY

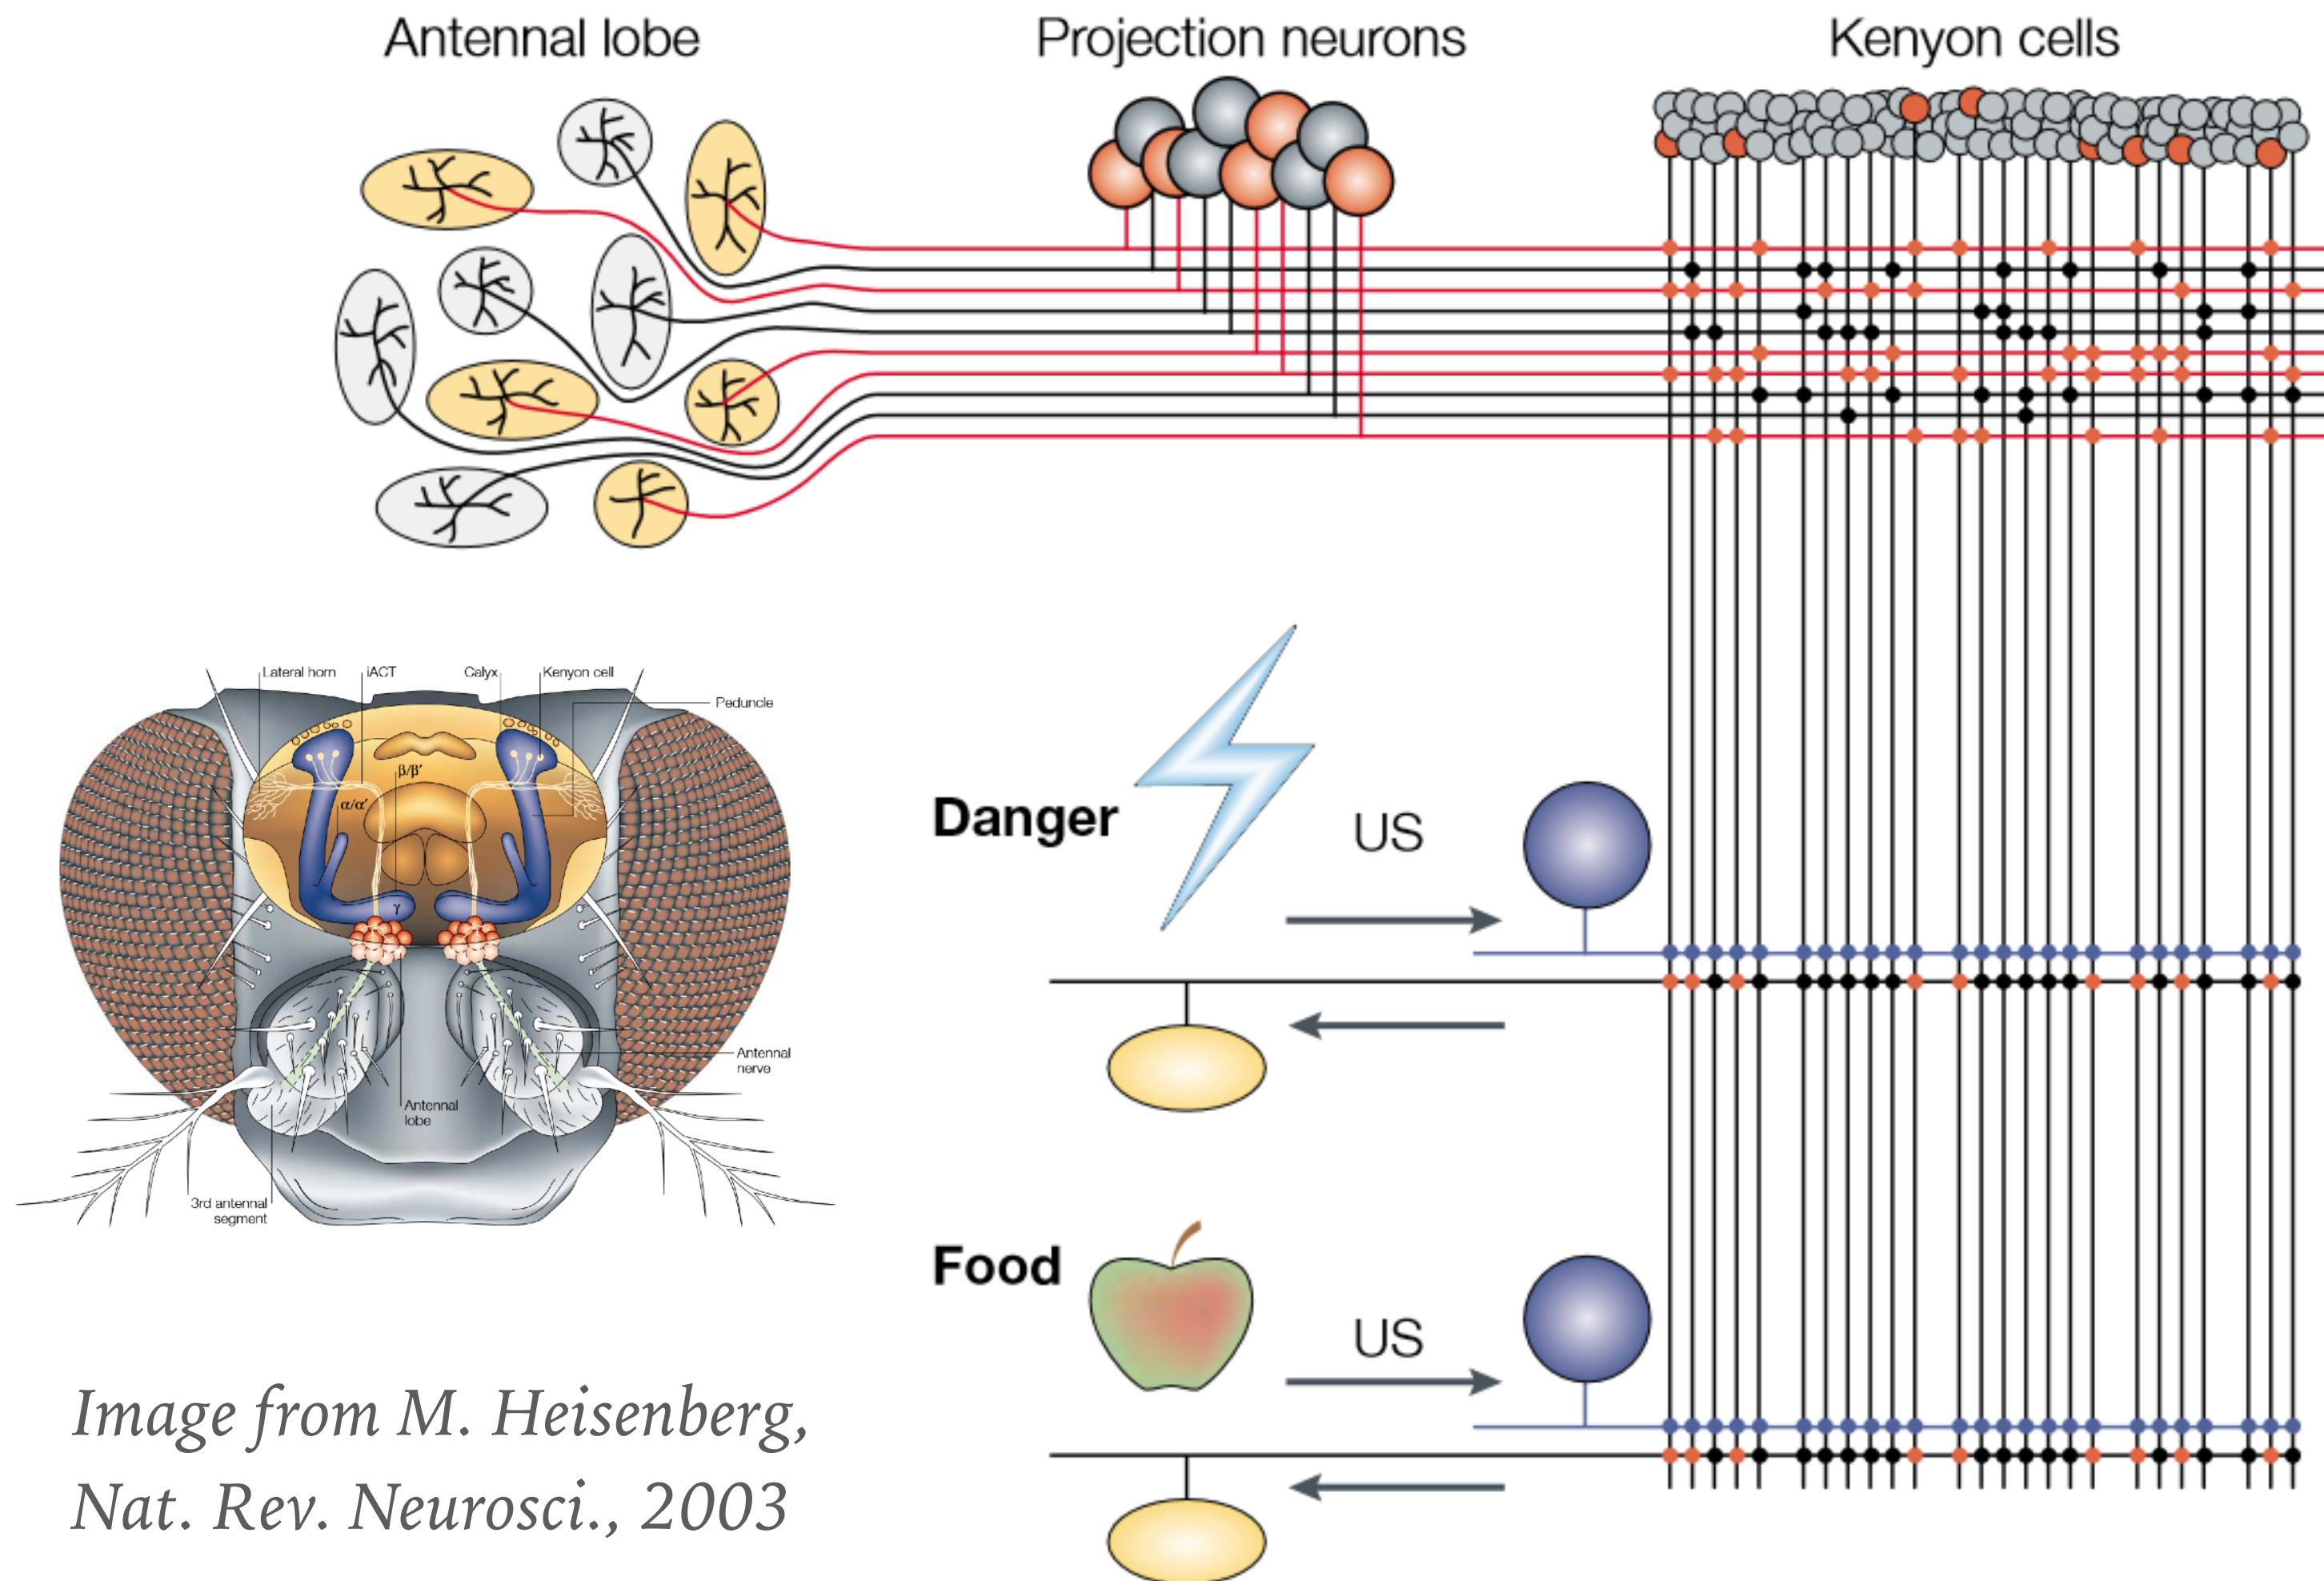


*Image from M. Heisenberg,
Nat. Rev. Neurosci., 2003*

LEARNING ASSOCIATIONS: THE MUSHROOM BODY

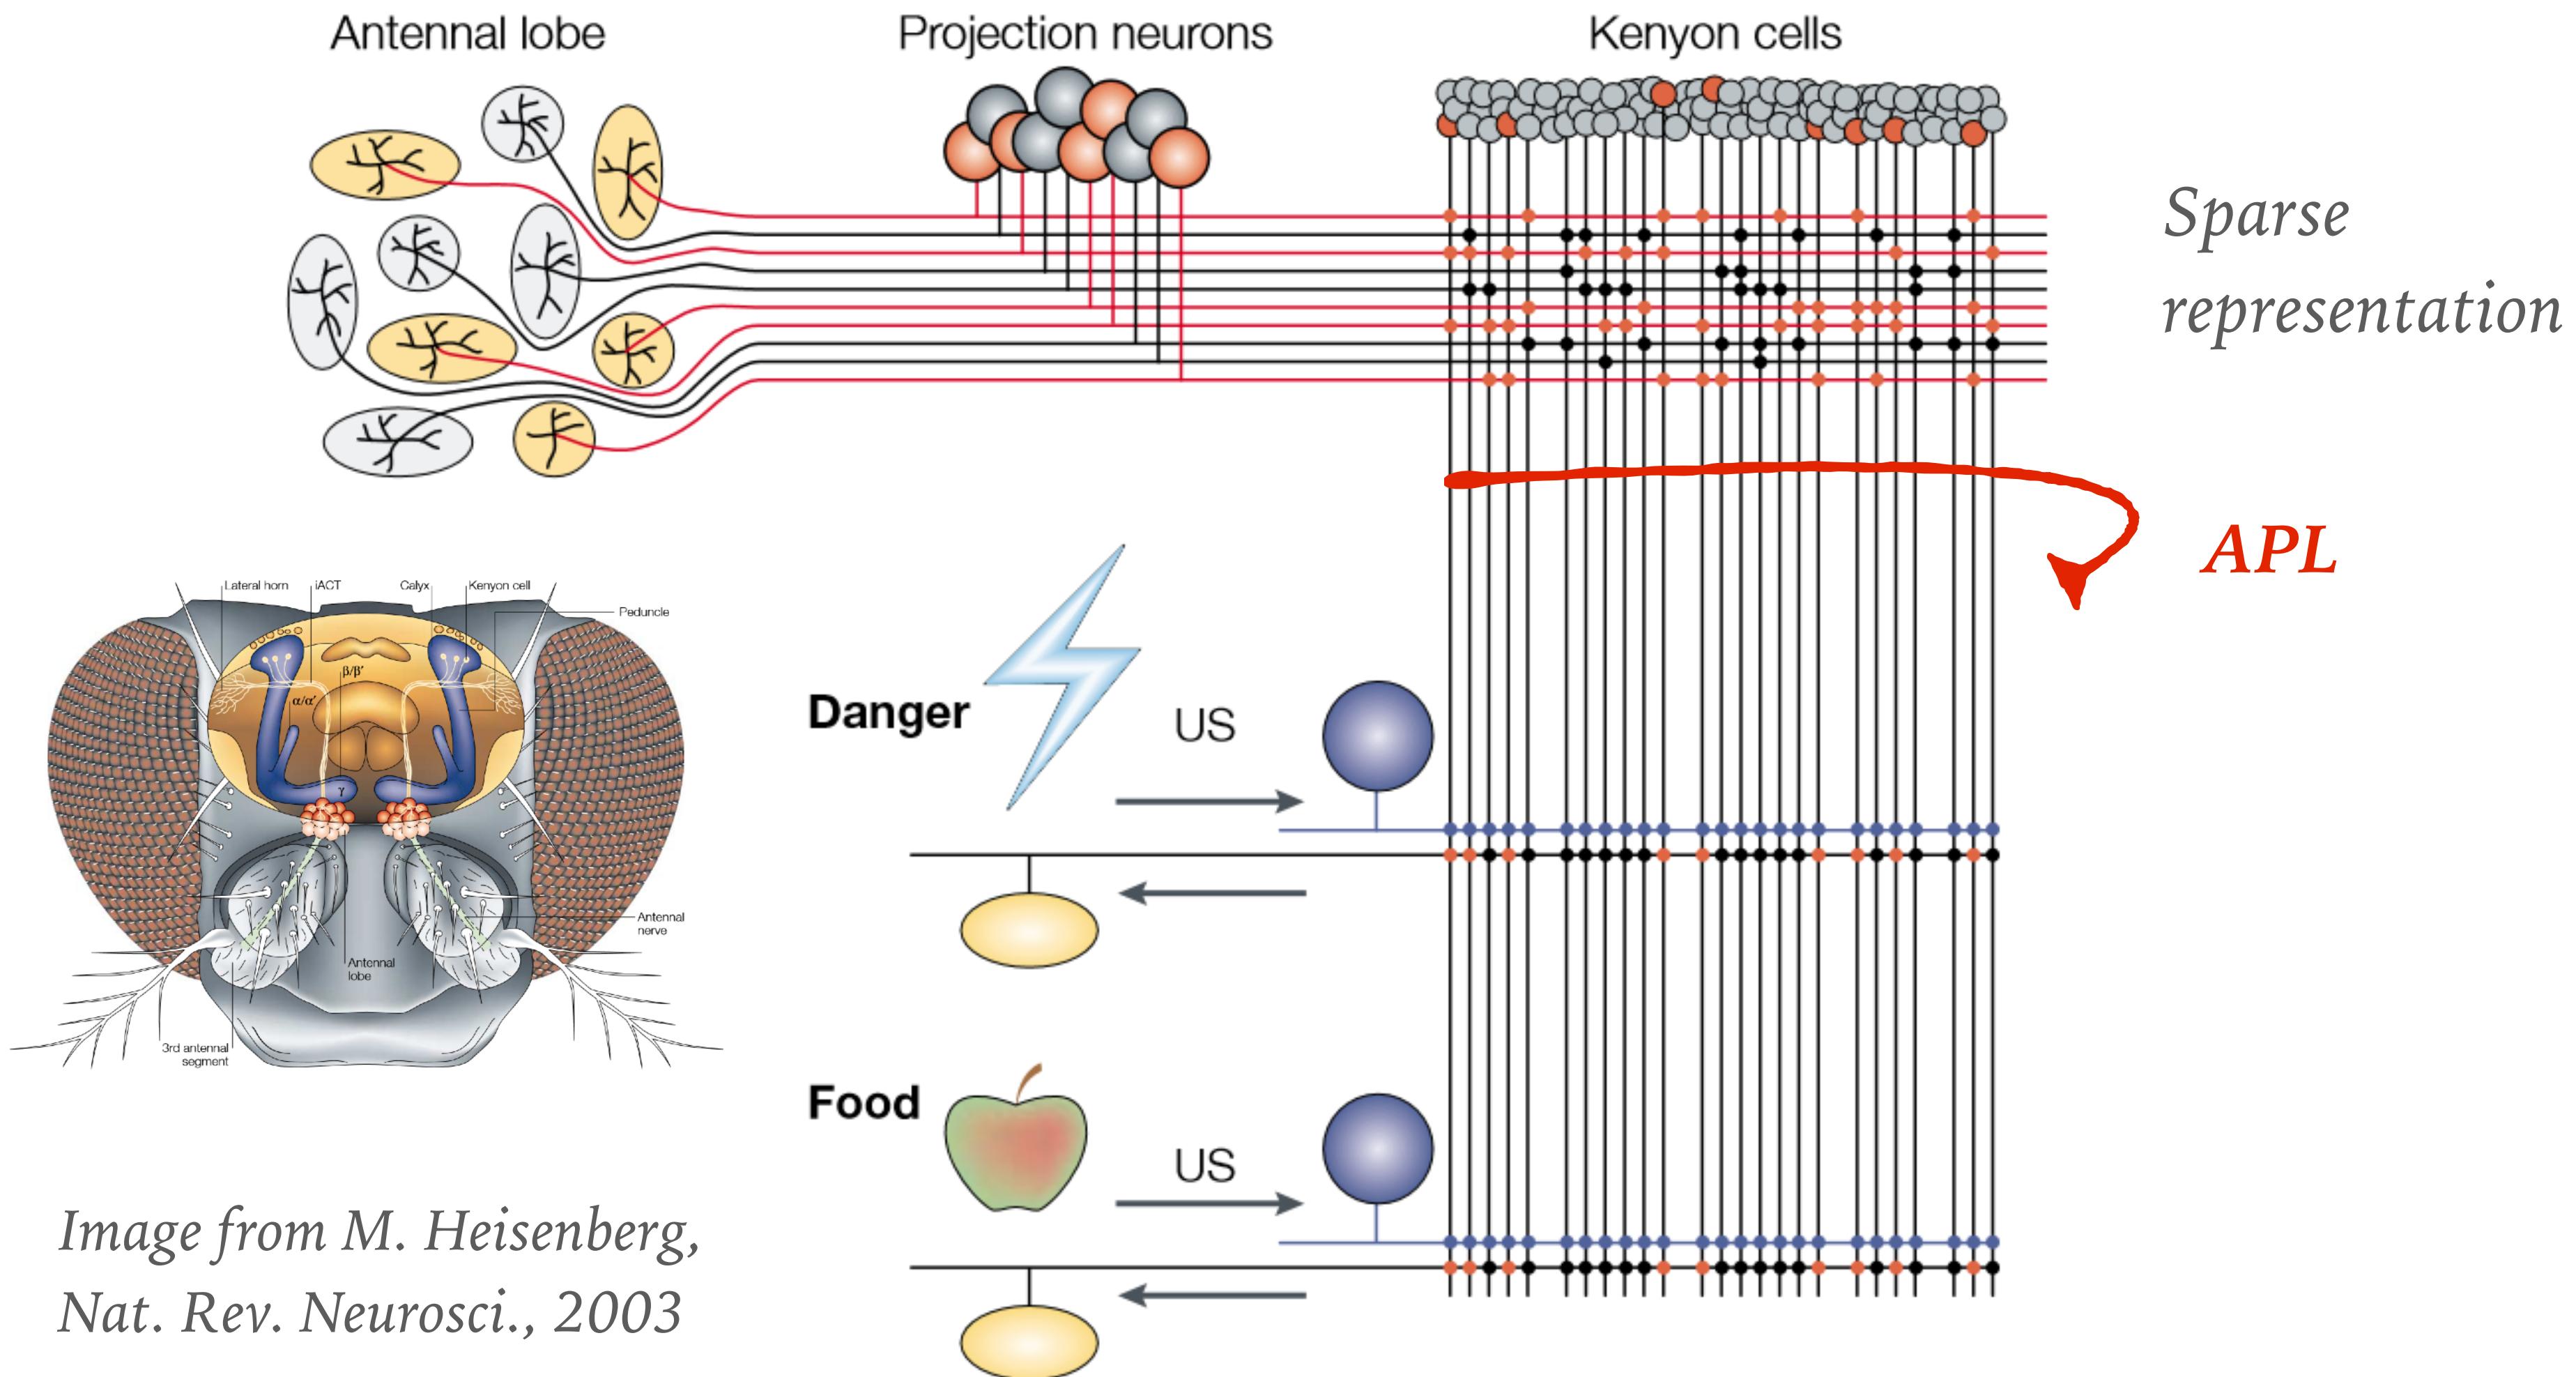


LEARNING ASSOCIATIONS: THE MUSHROOM BODY

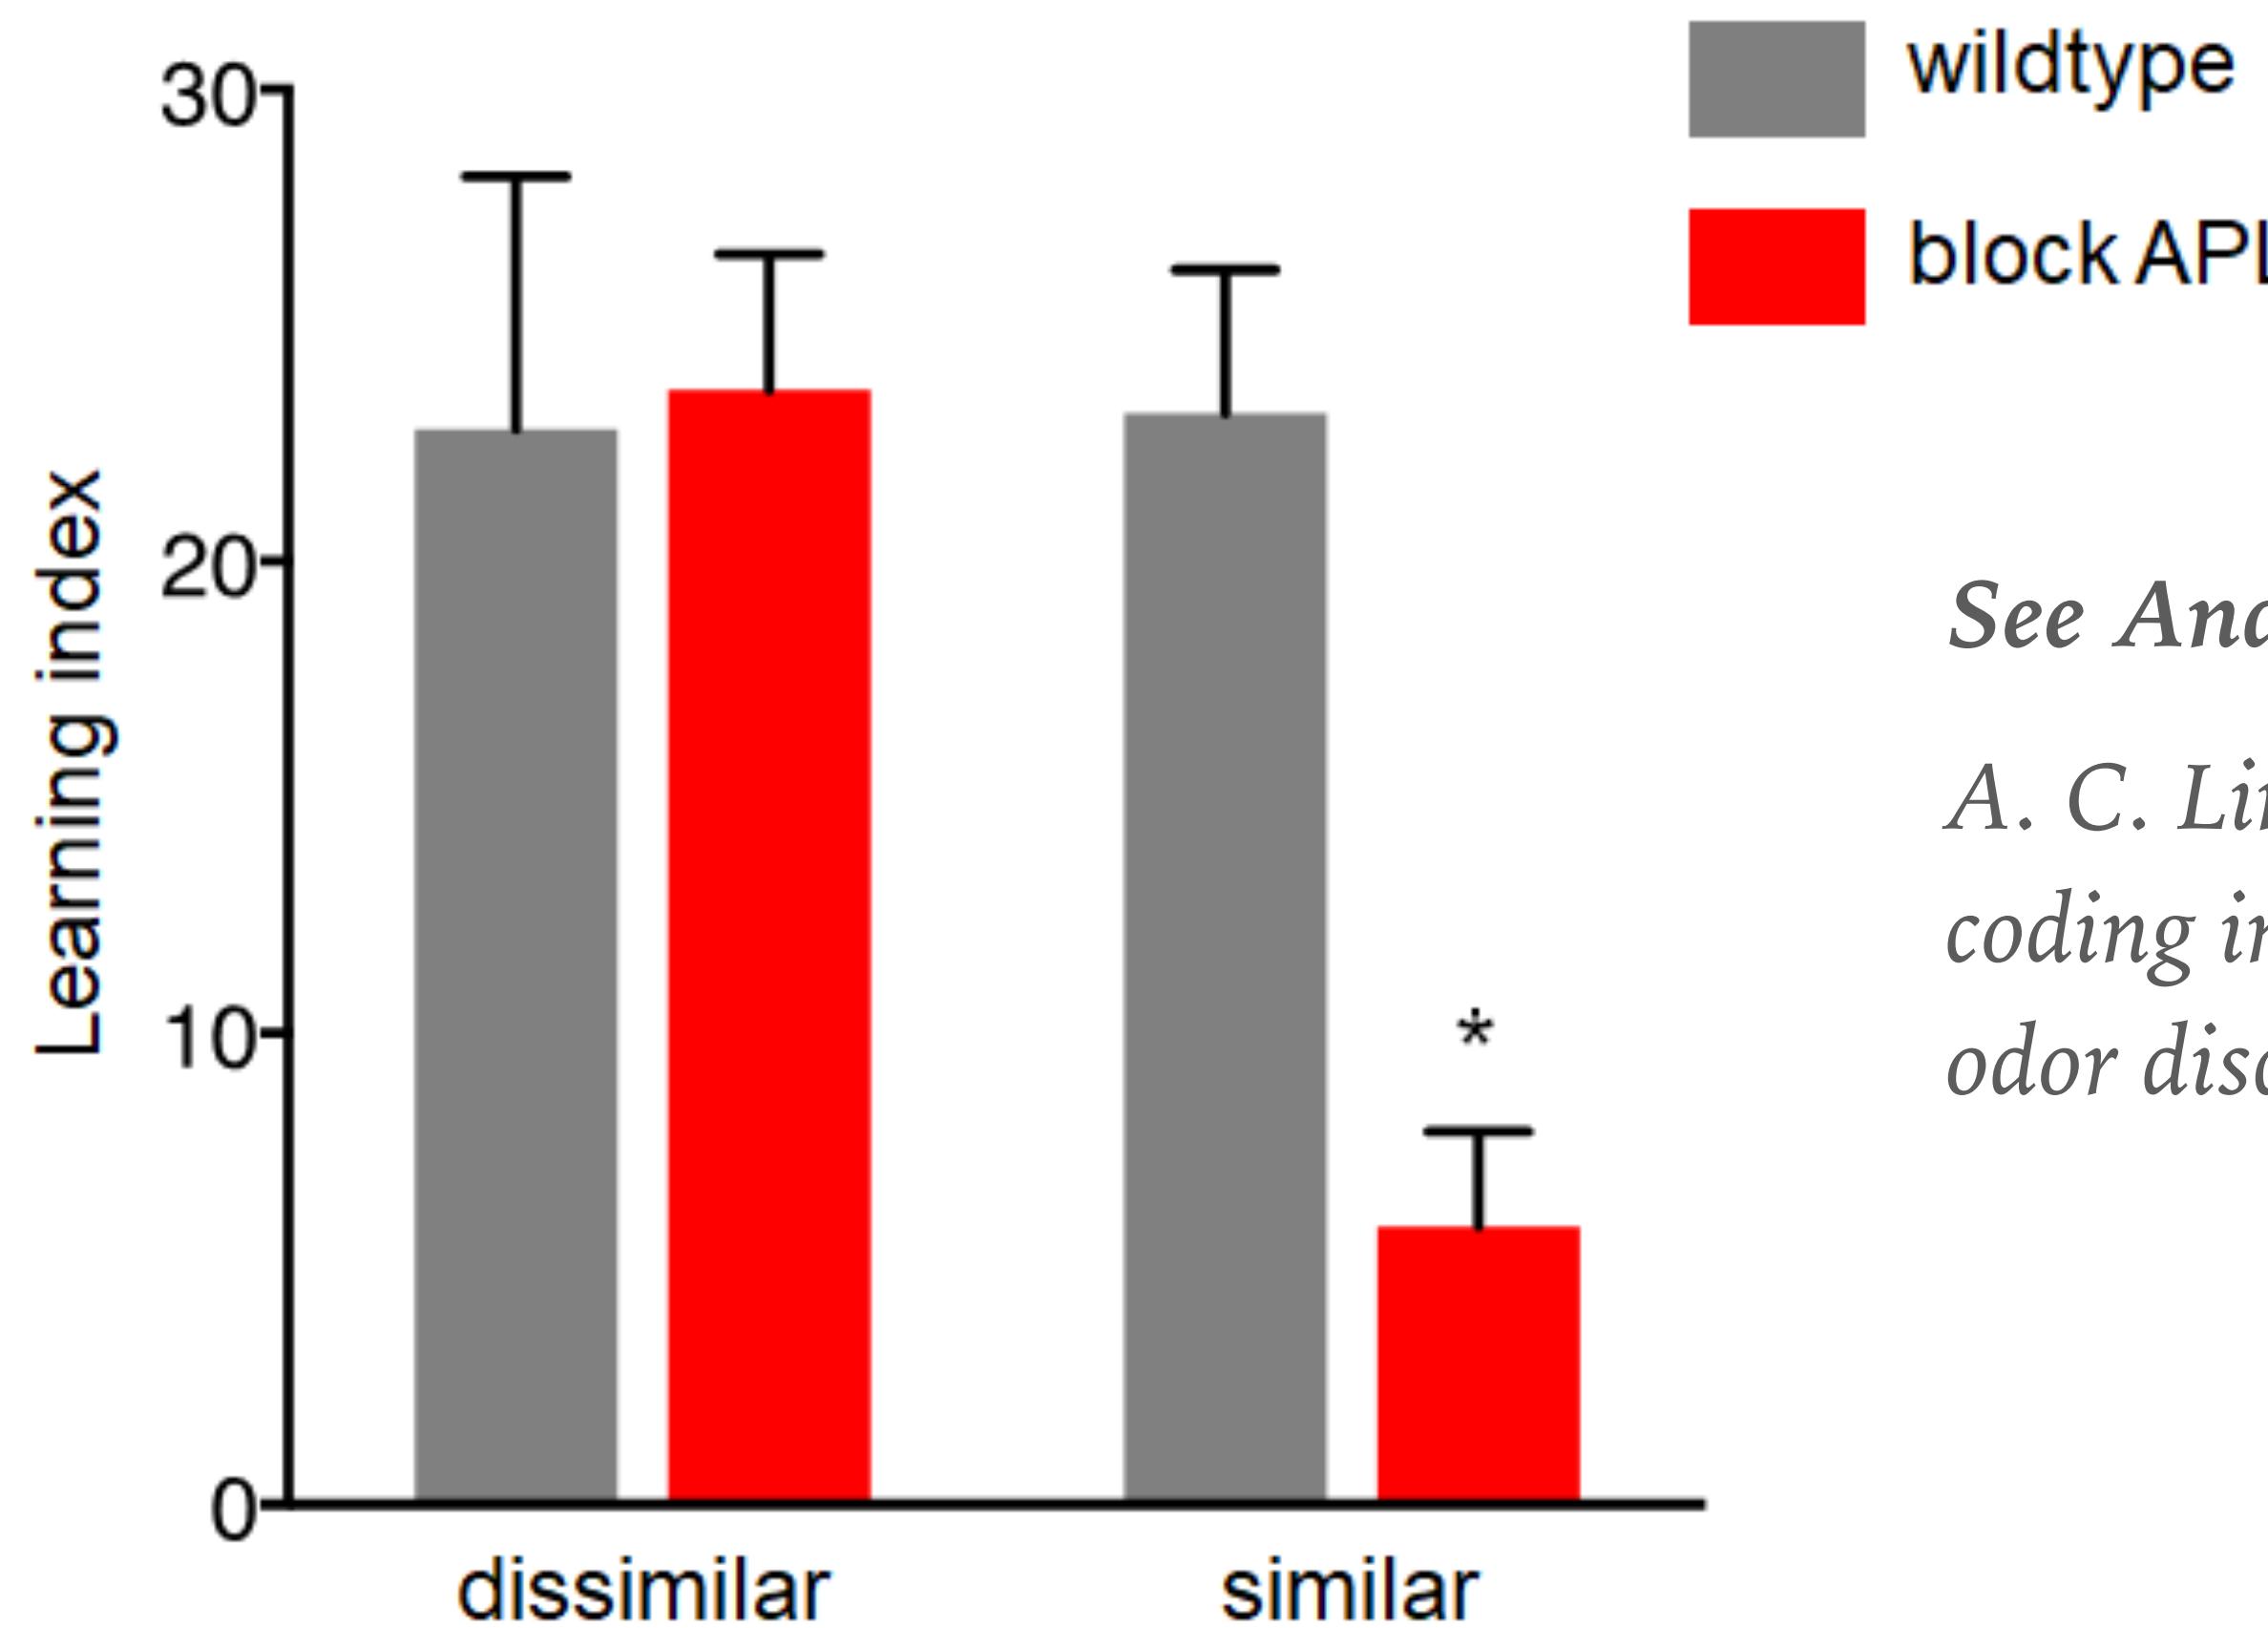


*Sparse
representation*

LEARNING ASSOCIATIONS: THE MUSHROOM BODY



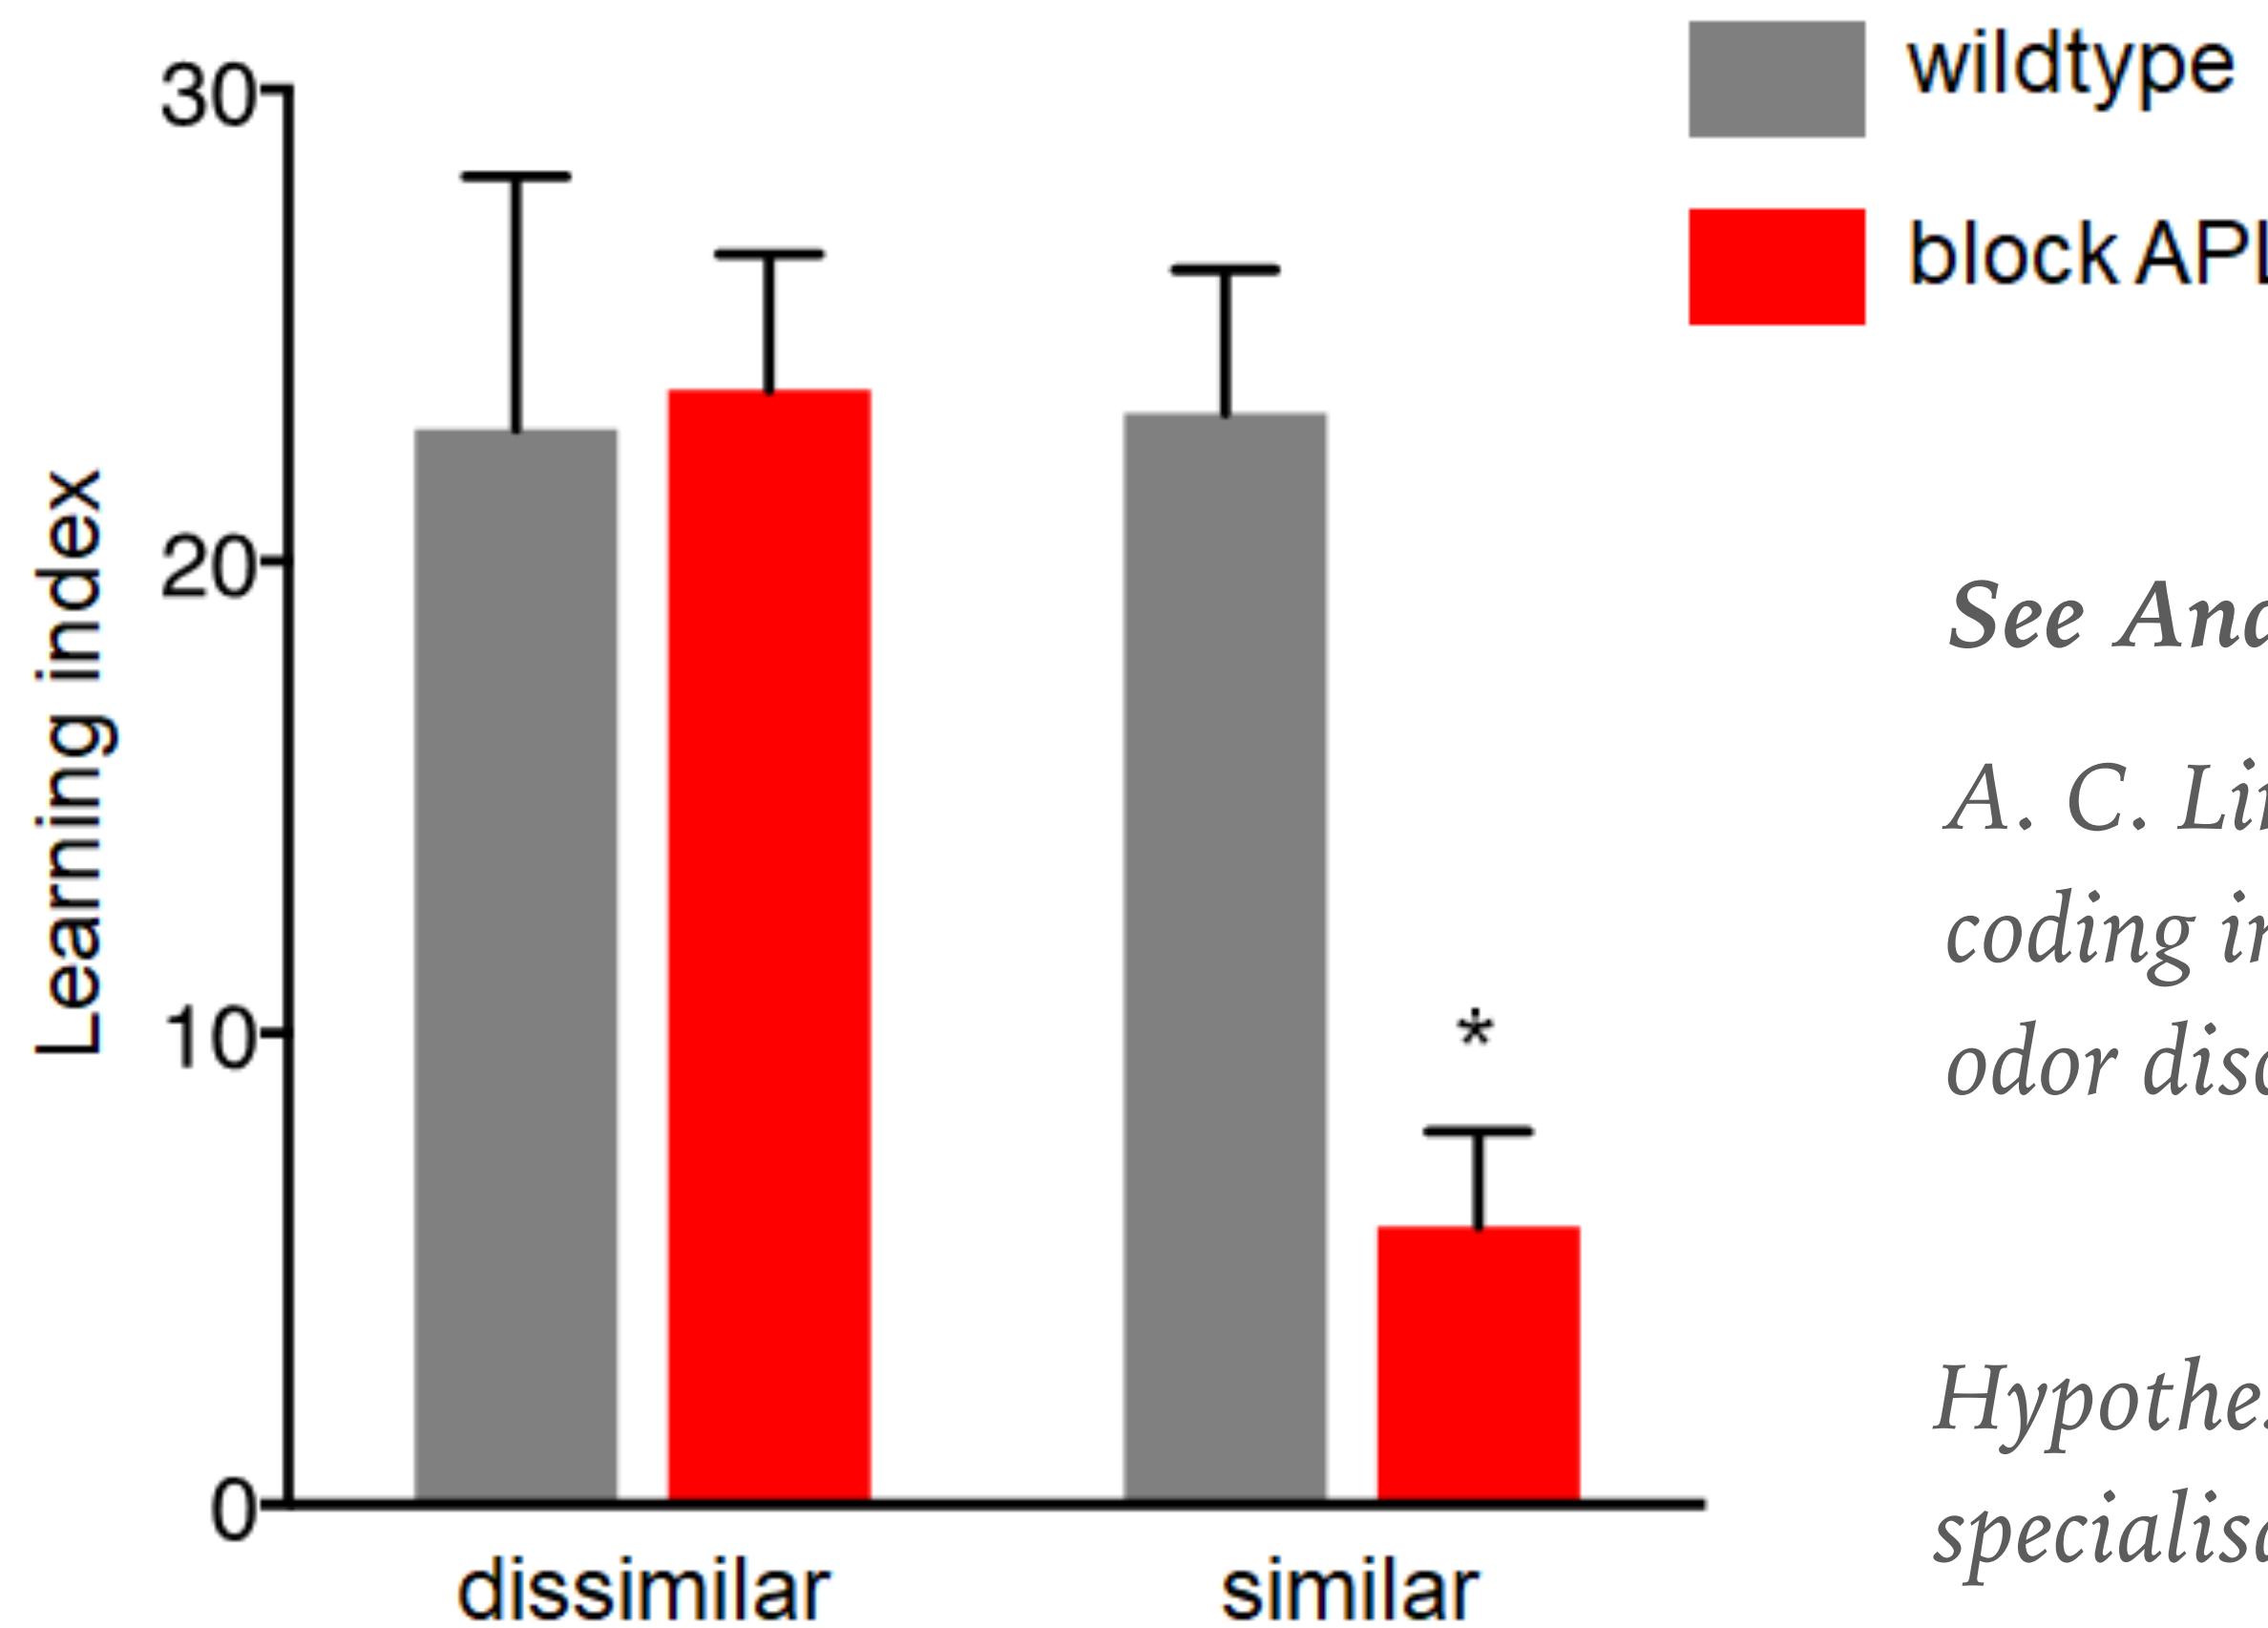
BLOCKING APL PREVENTS LEARNED DISCRIMINATION OF SIMILAR ODOURS BUT NOT DISSIMILAR ODOURS



See Andrew's work:

A. C. Lin et al “Sparse, decorrelated odor coding in the mushroom body enhances learned odor discrimination,” *Nature Neurosci.*, 2014.

BLOCKING APL PREVENTS LEARNED DISCRIMINATION OF SIMILAR ODOURS BUT NOT DISSIMILAR ODOURS

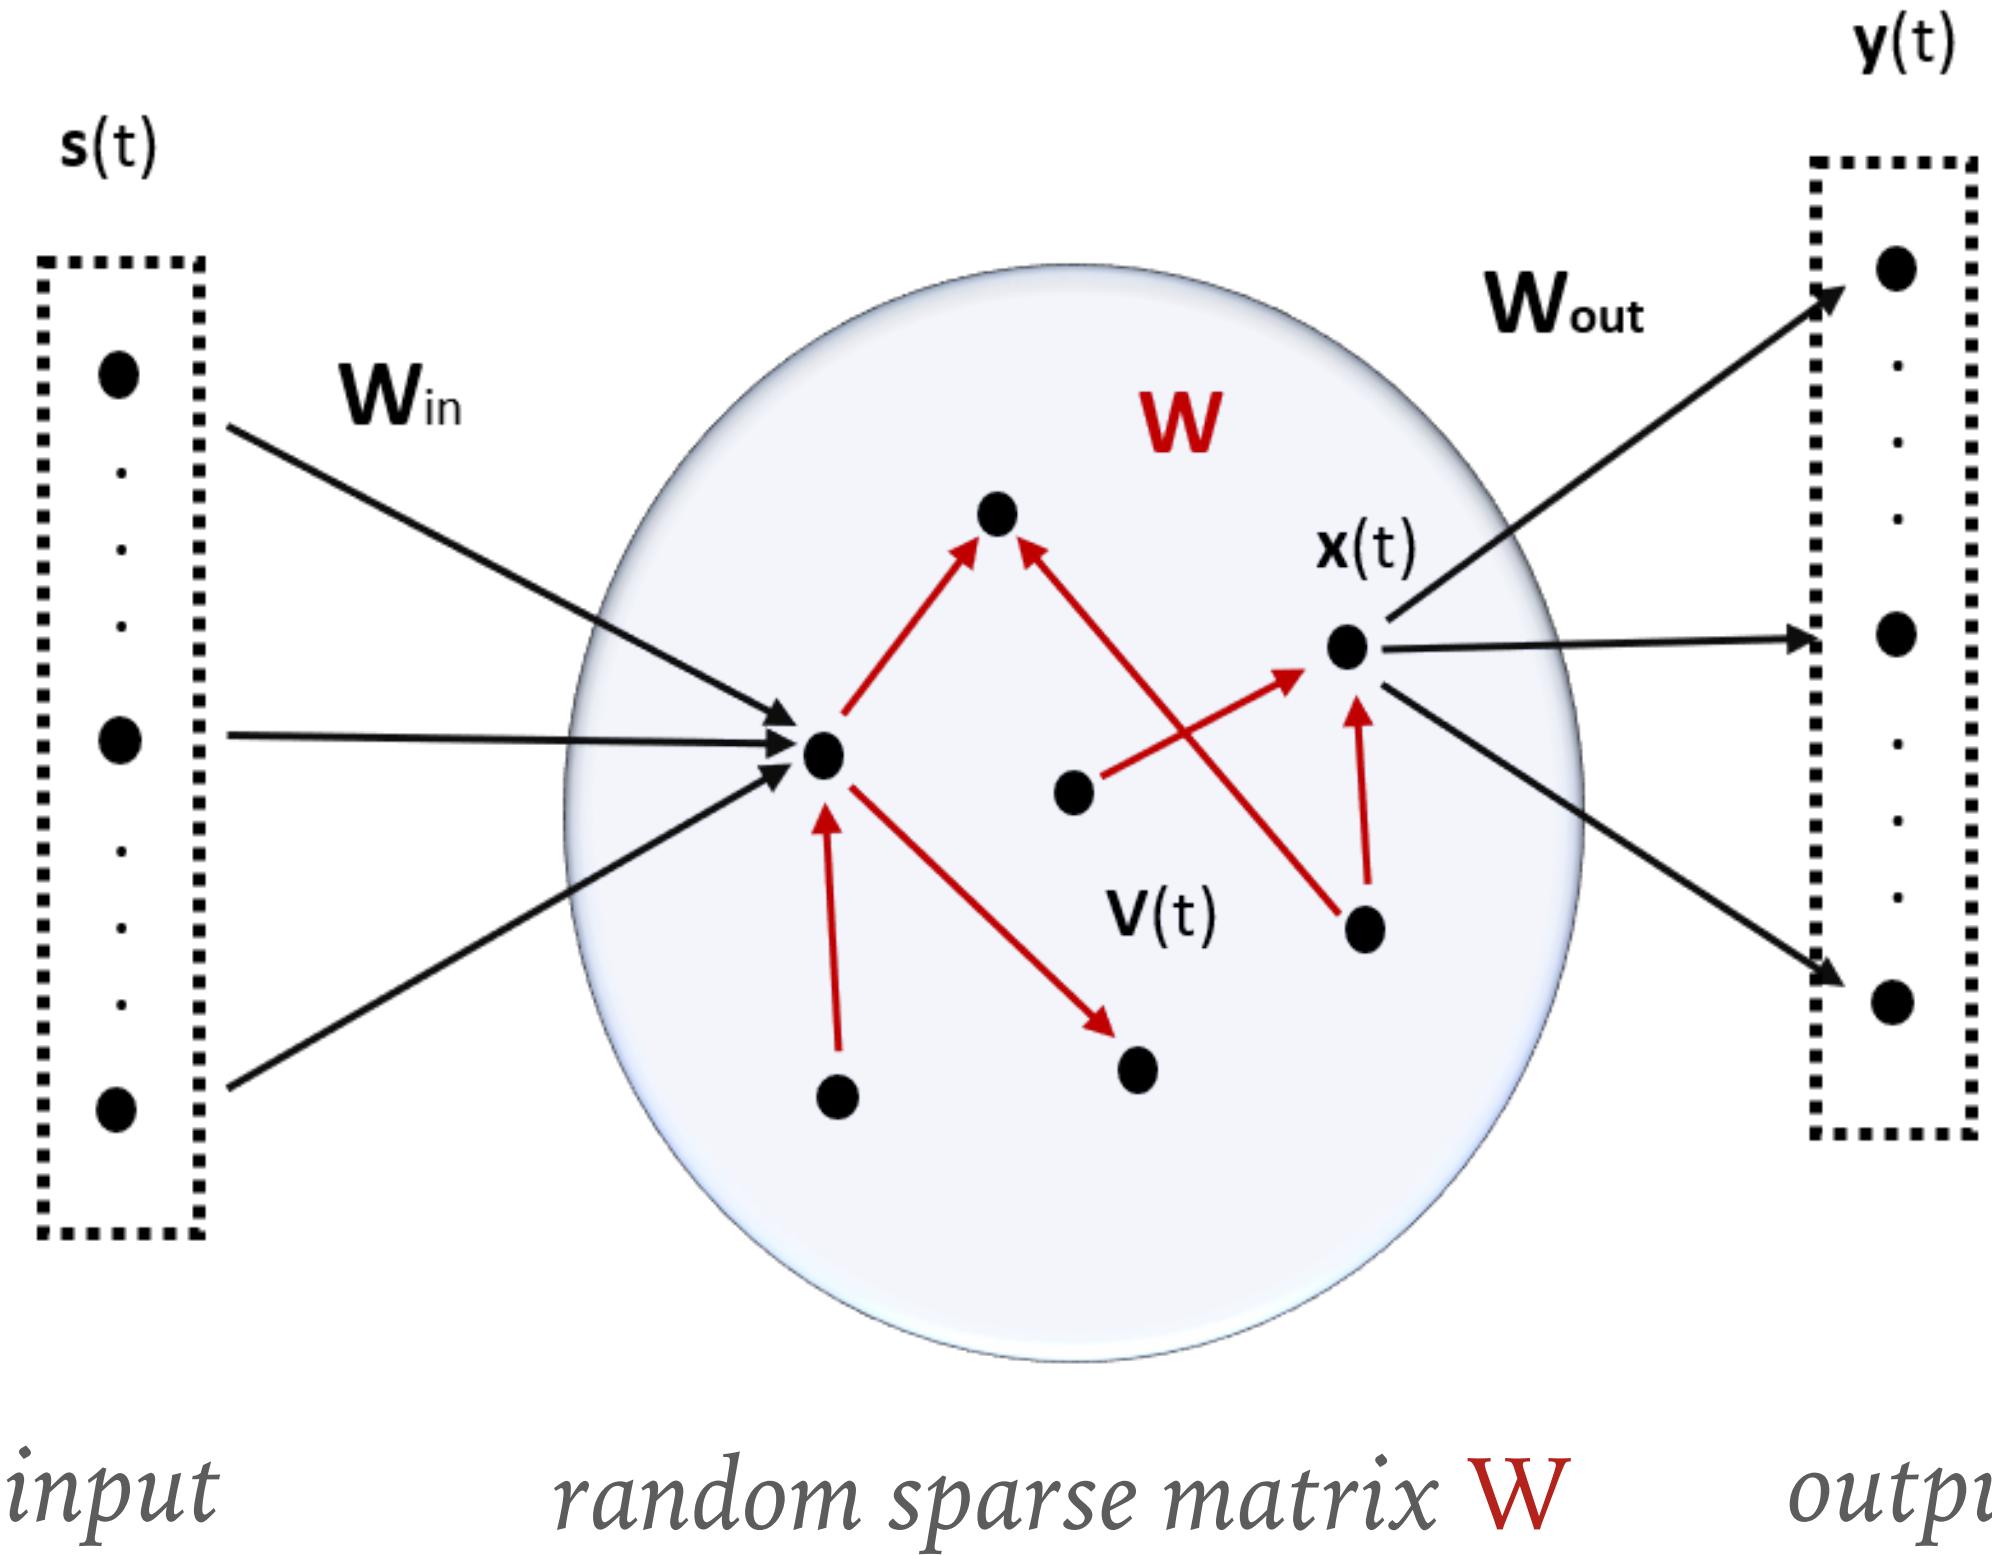


See Andrew's work:

A. C. Lin et al “Sparse, decorrelated odor coding in the mushroom body enhances learned odor discrimination,” *Nature Neurosci.*, 2014.

Hypothesis: Sparsity introduces neuronal specialisation

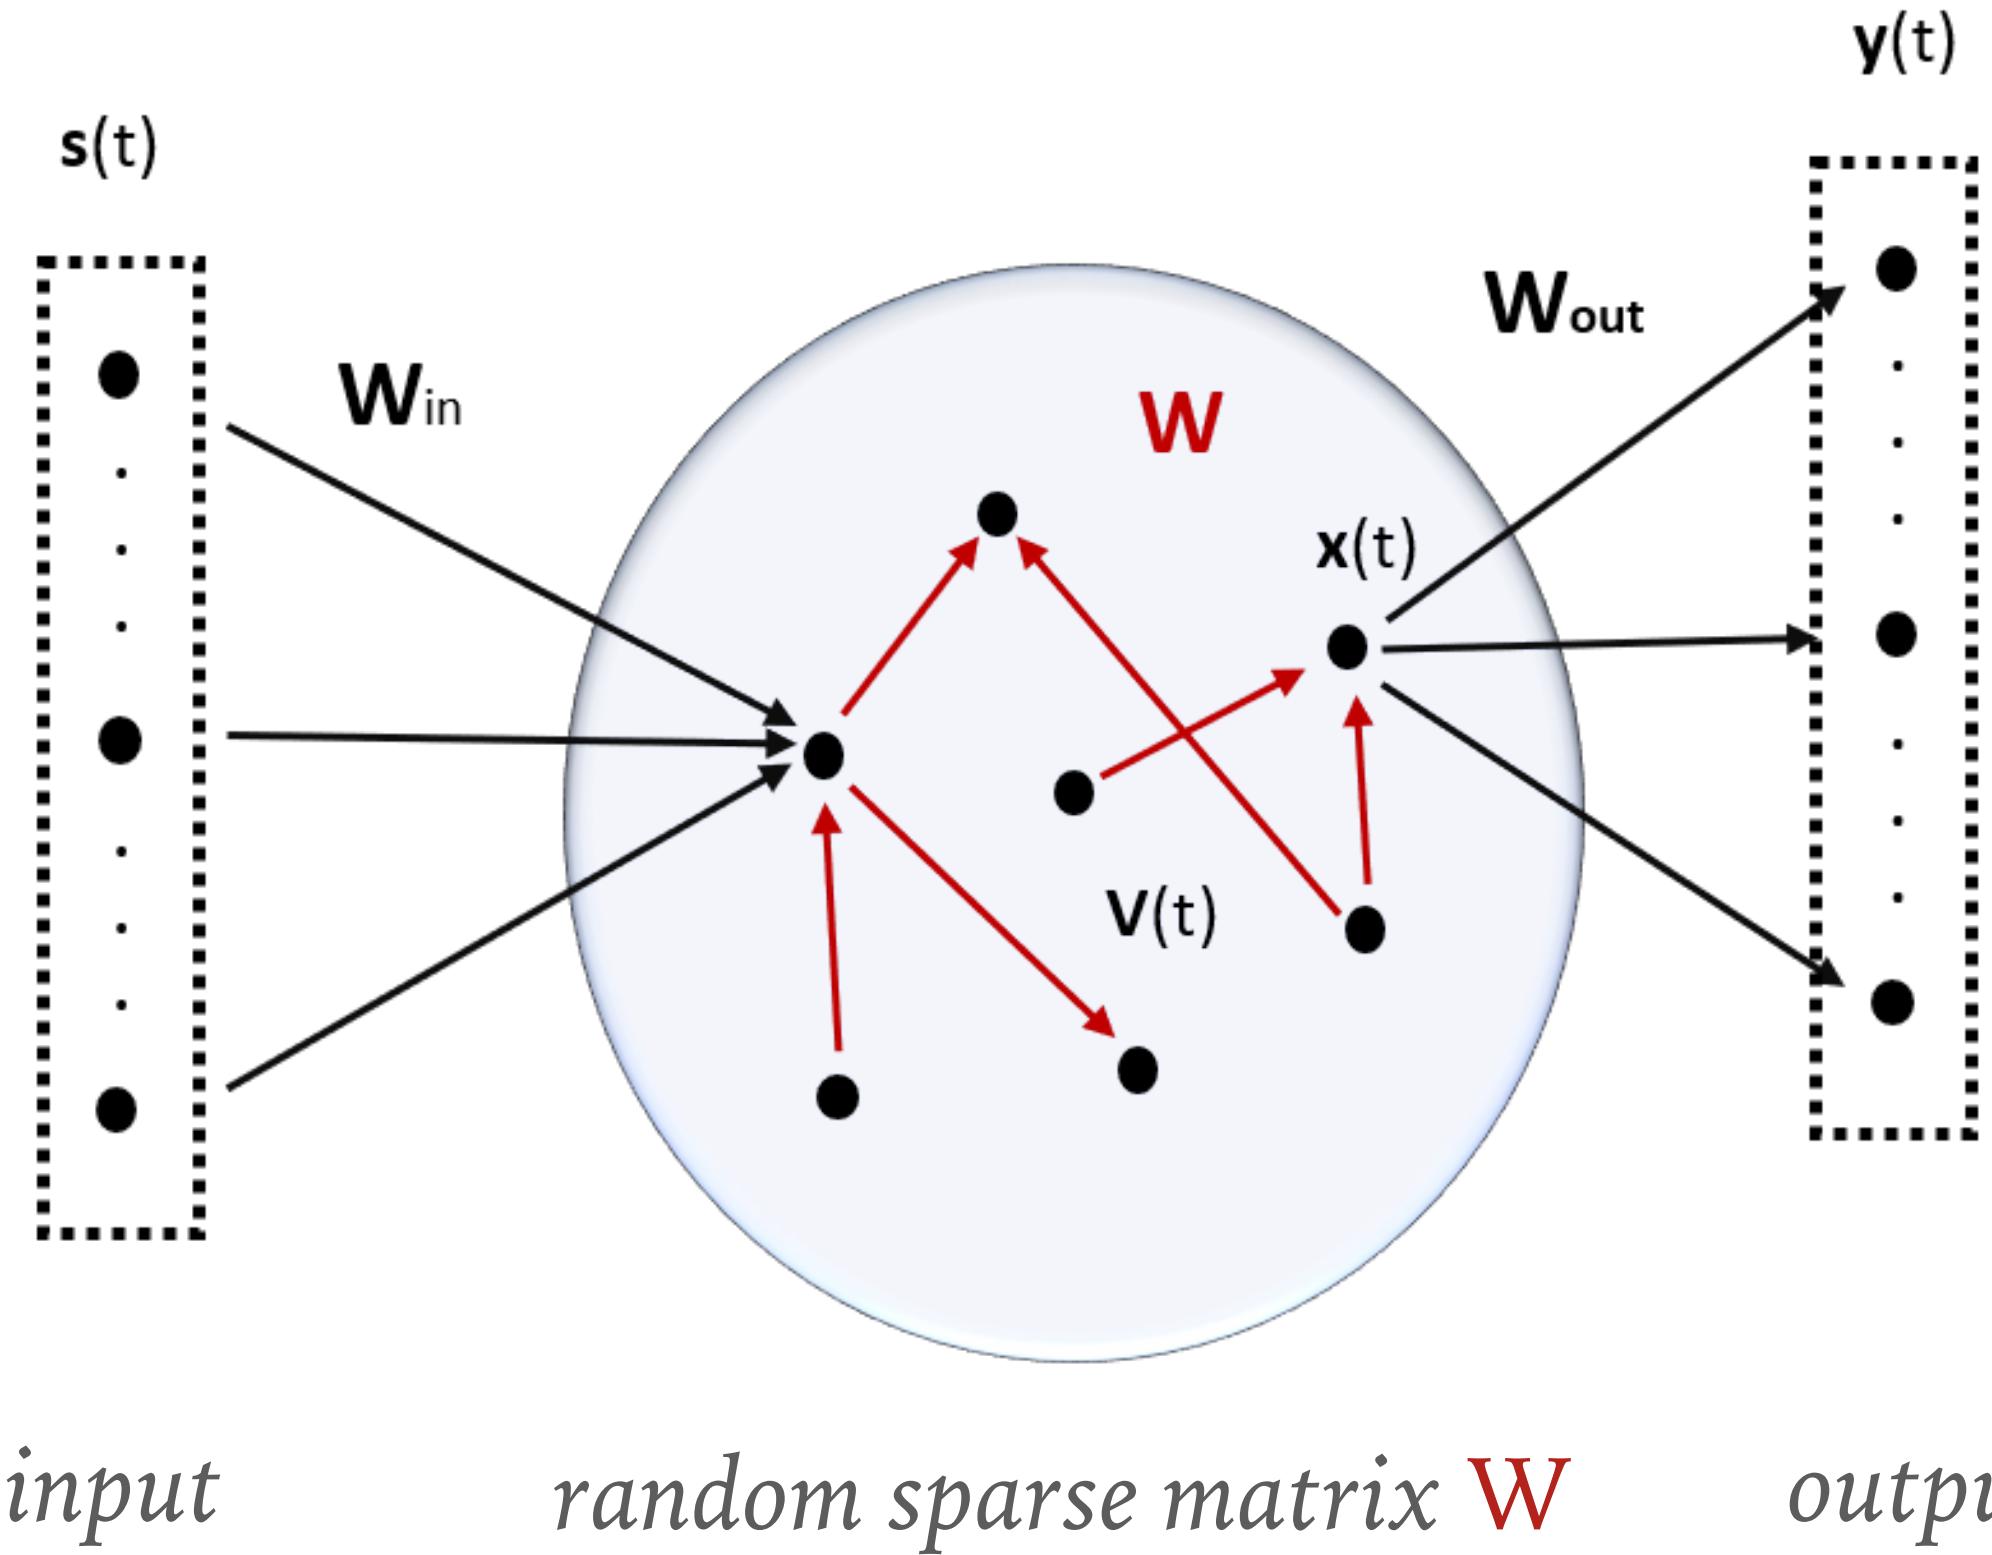
THE MUSHROOM BODY AS RESERVOIR



H. Jaeger, GMD-Report (2001), Neural Networks (2007)

W. Maass et al, Neural Comput. (2002) W. Maass & H. Markram, JCSS (2004)

THE MUSHROOM BODY AS RESERVOIR

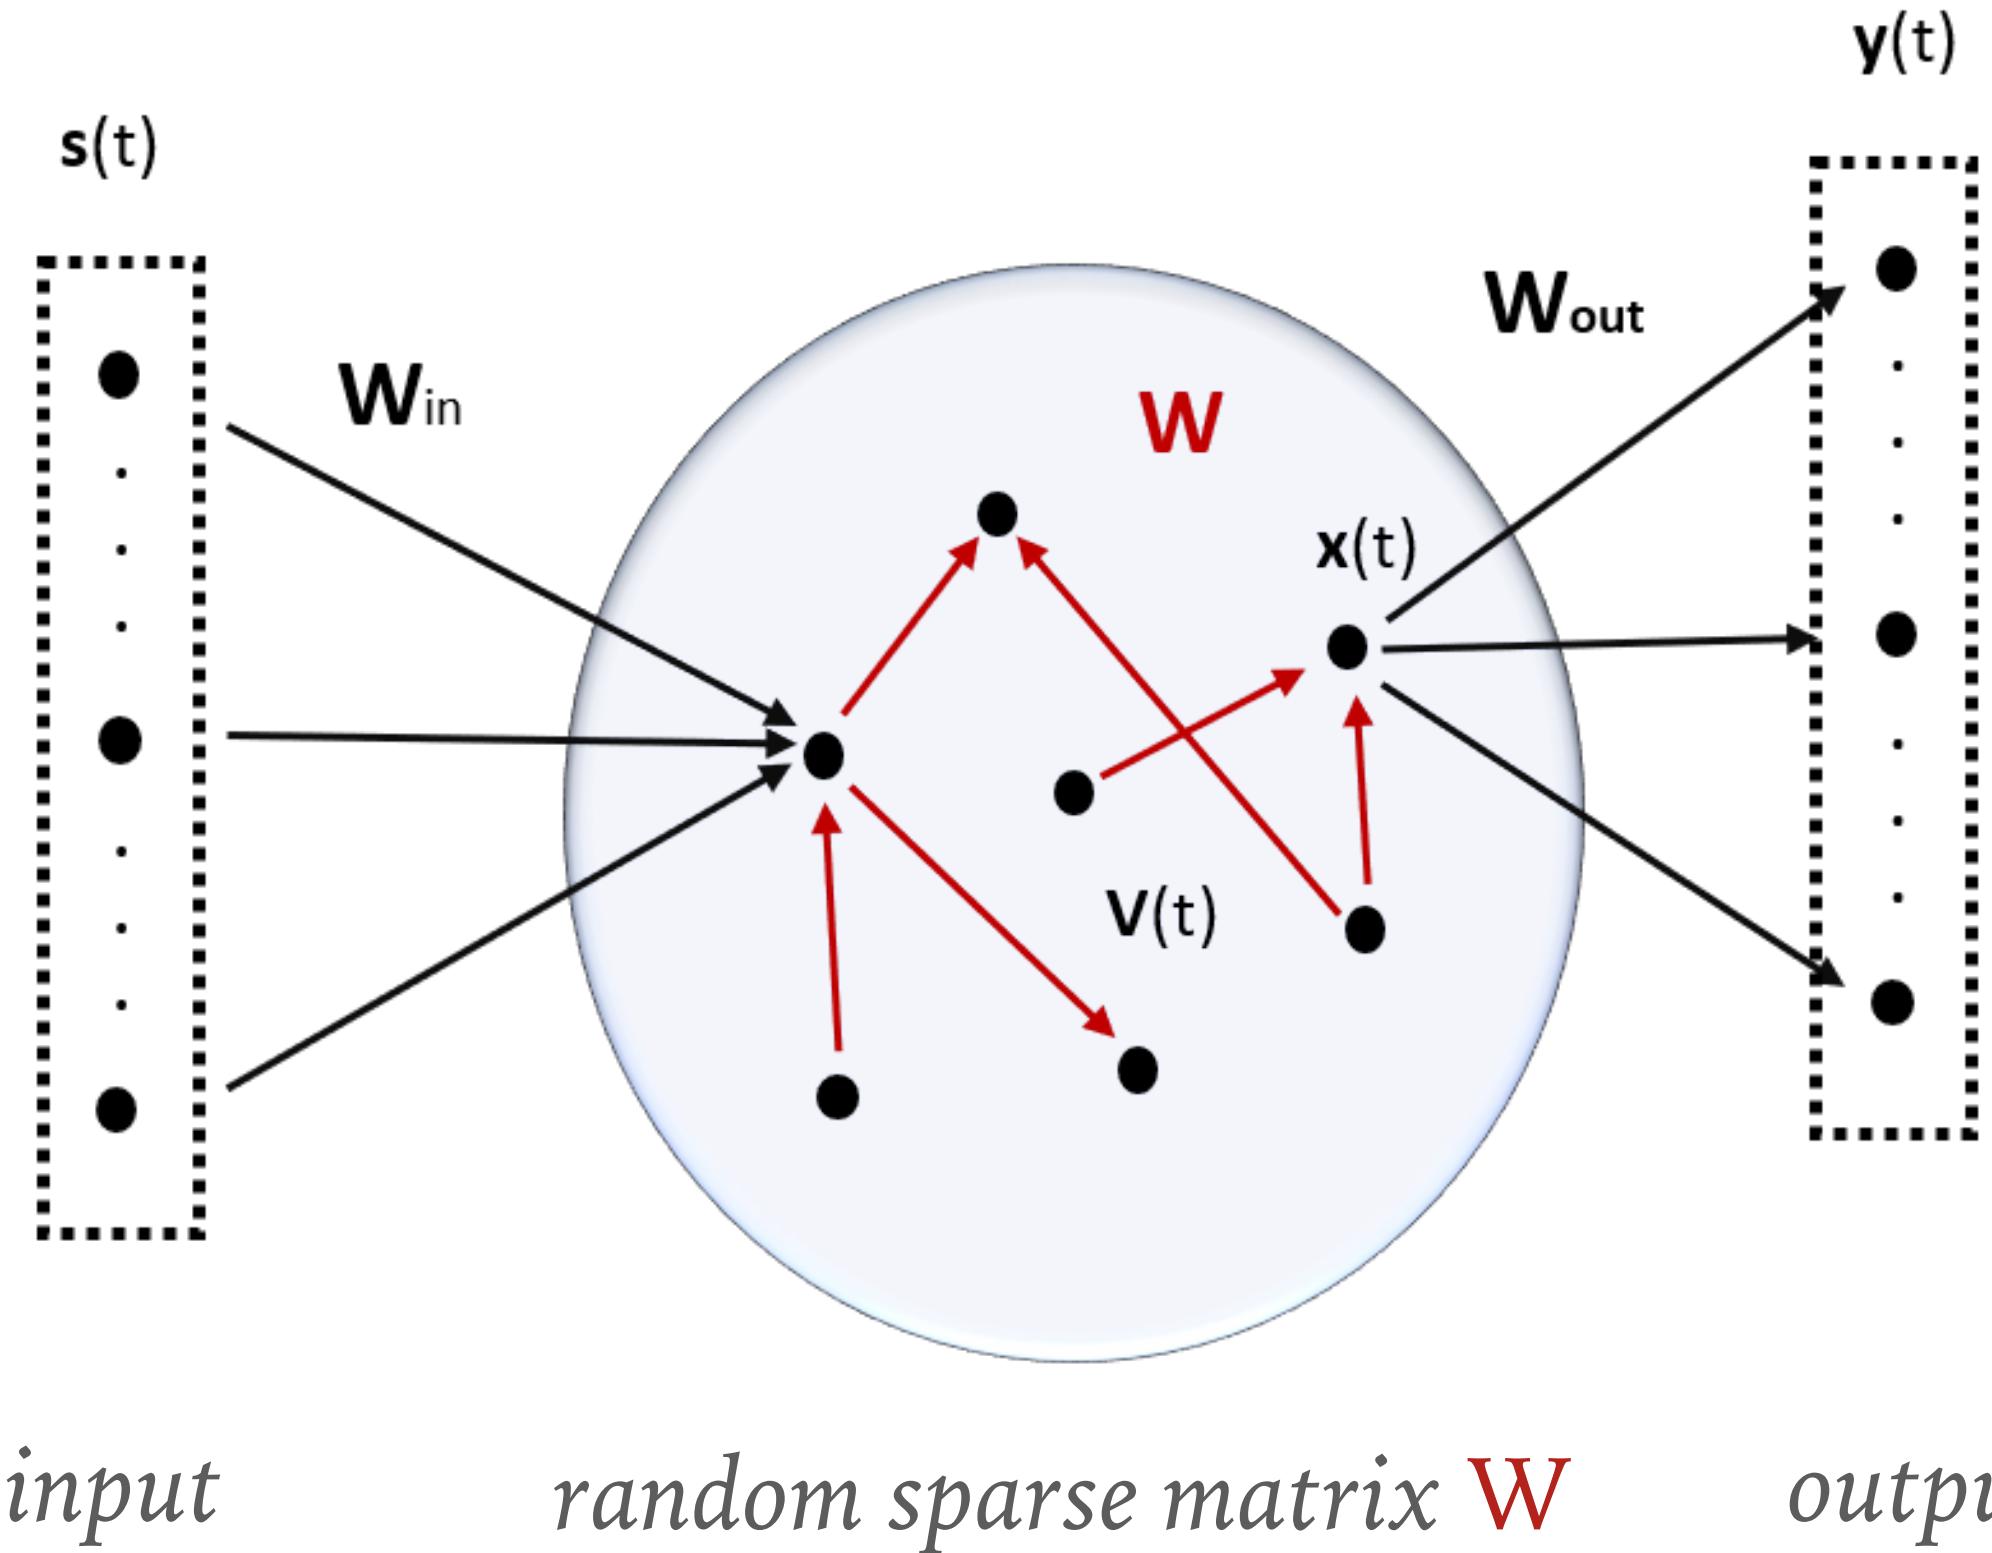


*Exploit randomness to achieve
a rich representation*

H. Jaeger, GMD-Report (2001), Neural Networks (2007)

W. Maass et al, Neural Comput. (2002) W. Maass & H. Markram, JCSS (2004)

THE MUSHROOM BODY AS RESERVOIR



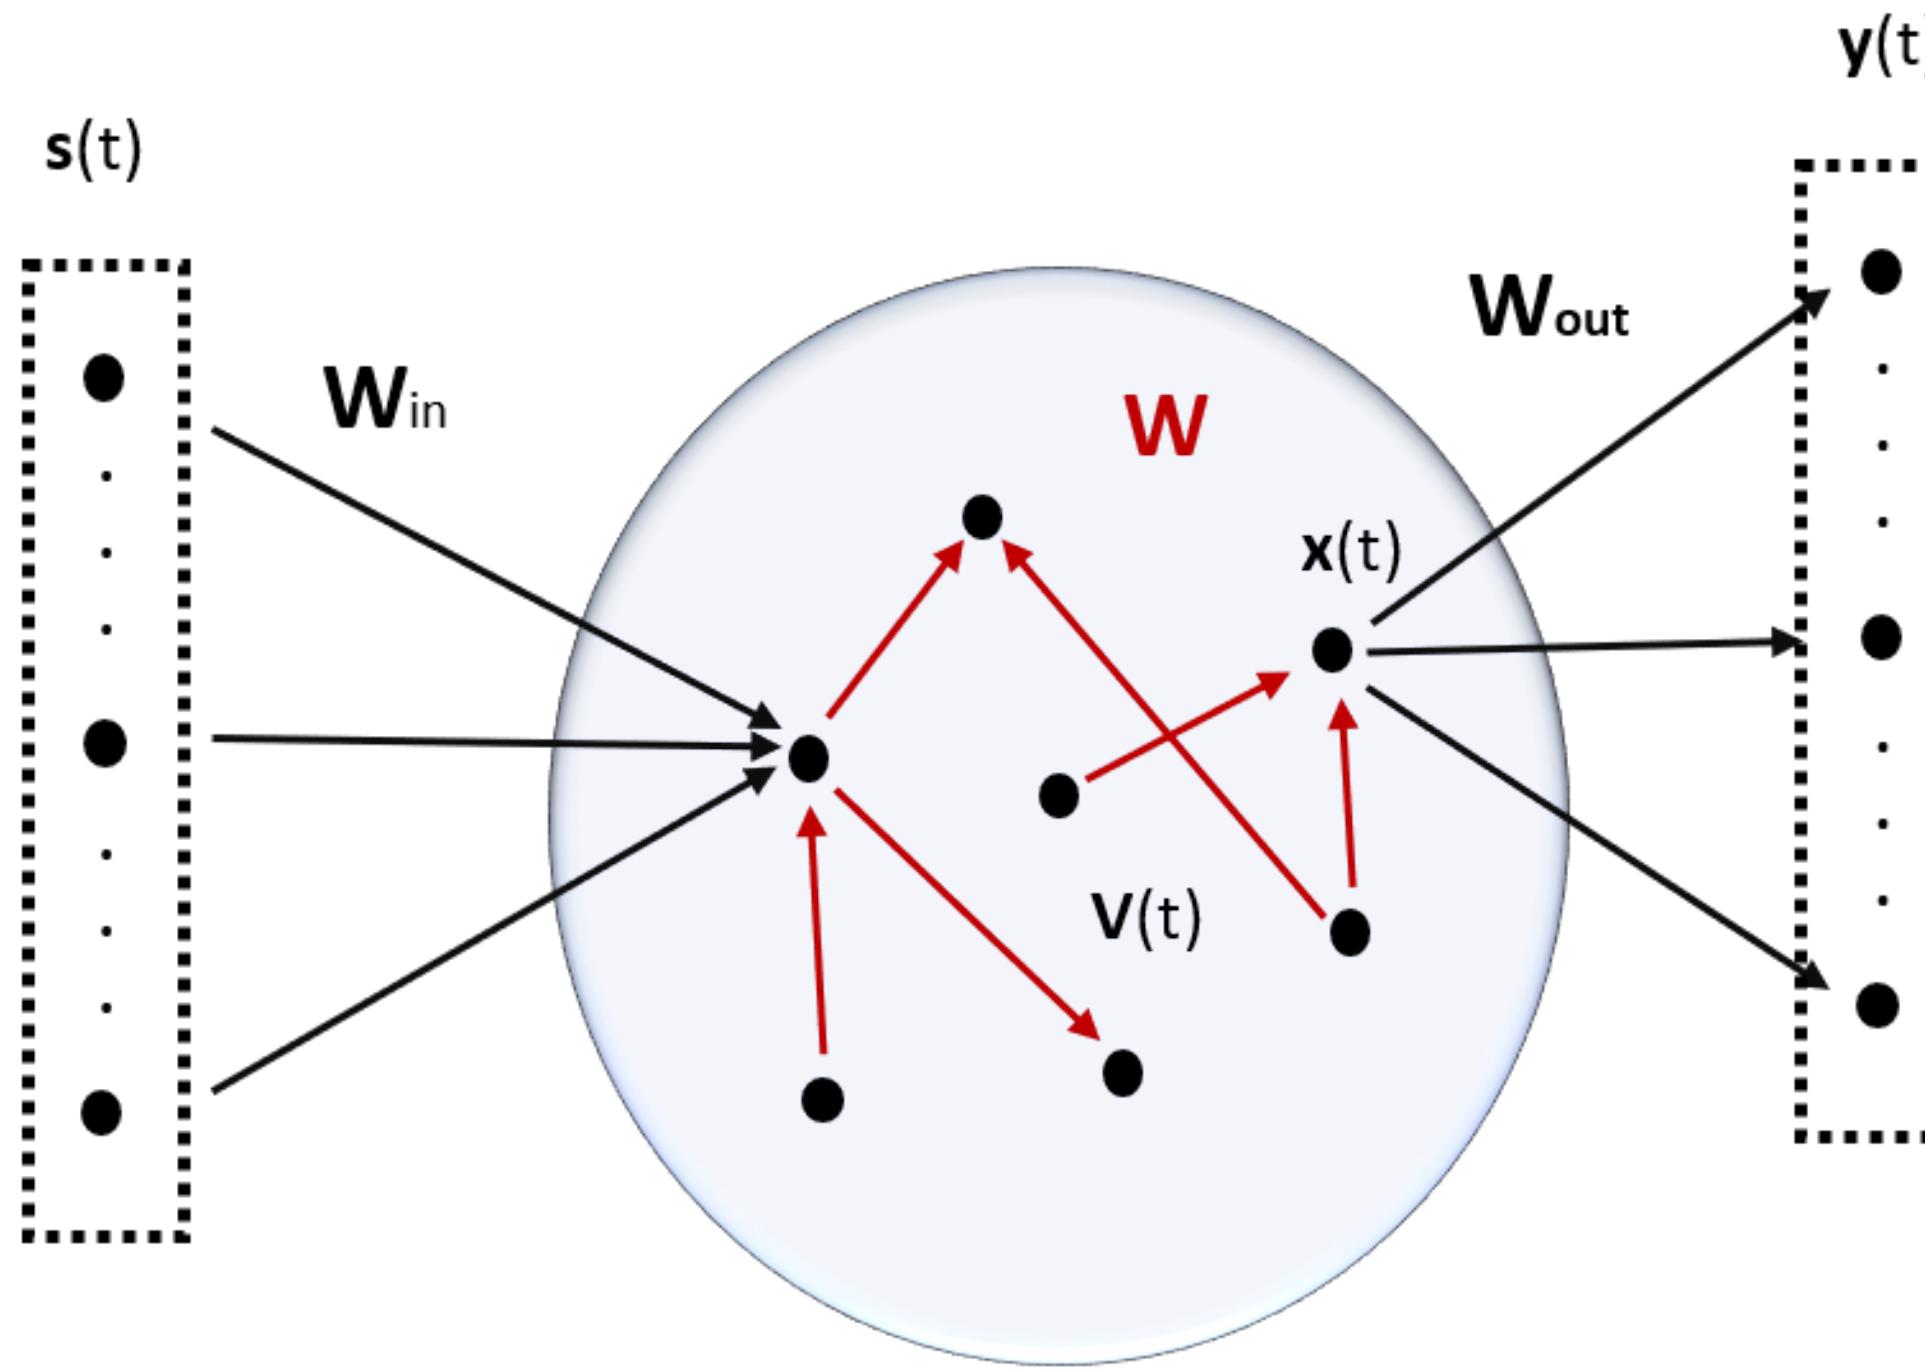
*Exploit randomness to achieve
a rich representation*

Learn only output matrix W_{out}

H. Jaeger, GMD-Report (2001), Neural Networks (2007)

W. Maass et al, Neural Comput. (2002) W. Maass & H. Markram, JCSS (2004)

RESERVOIR COMPUTING



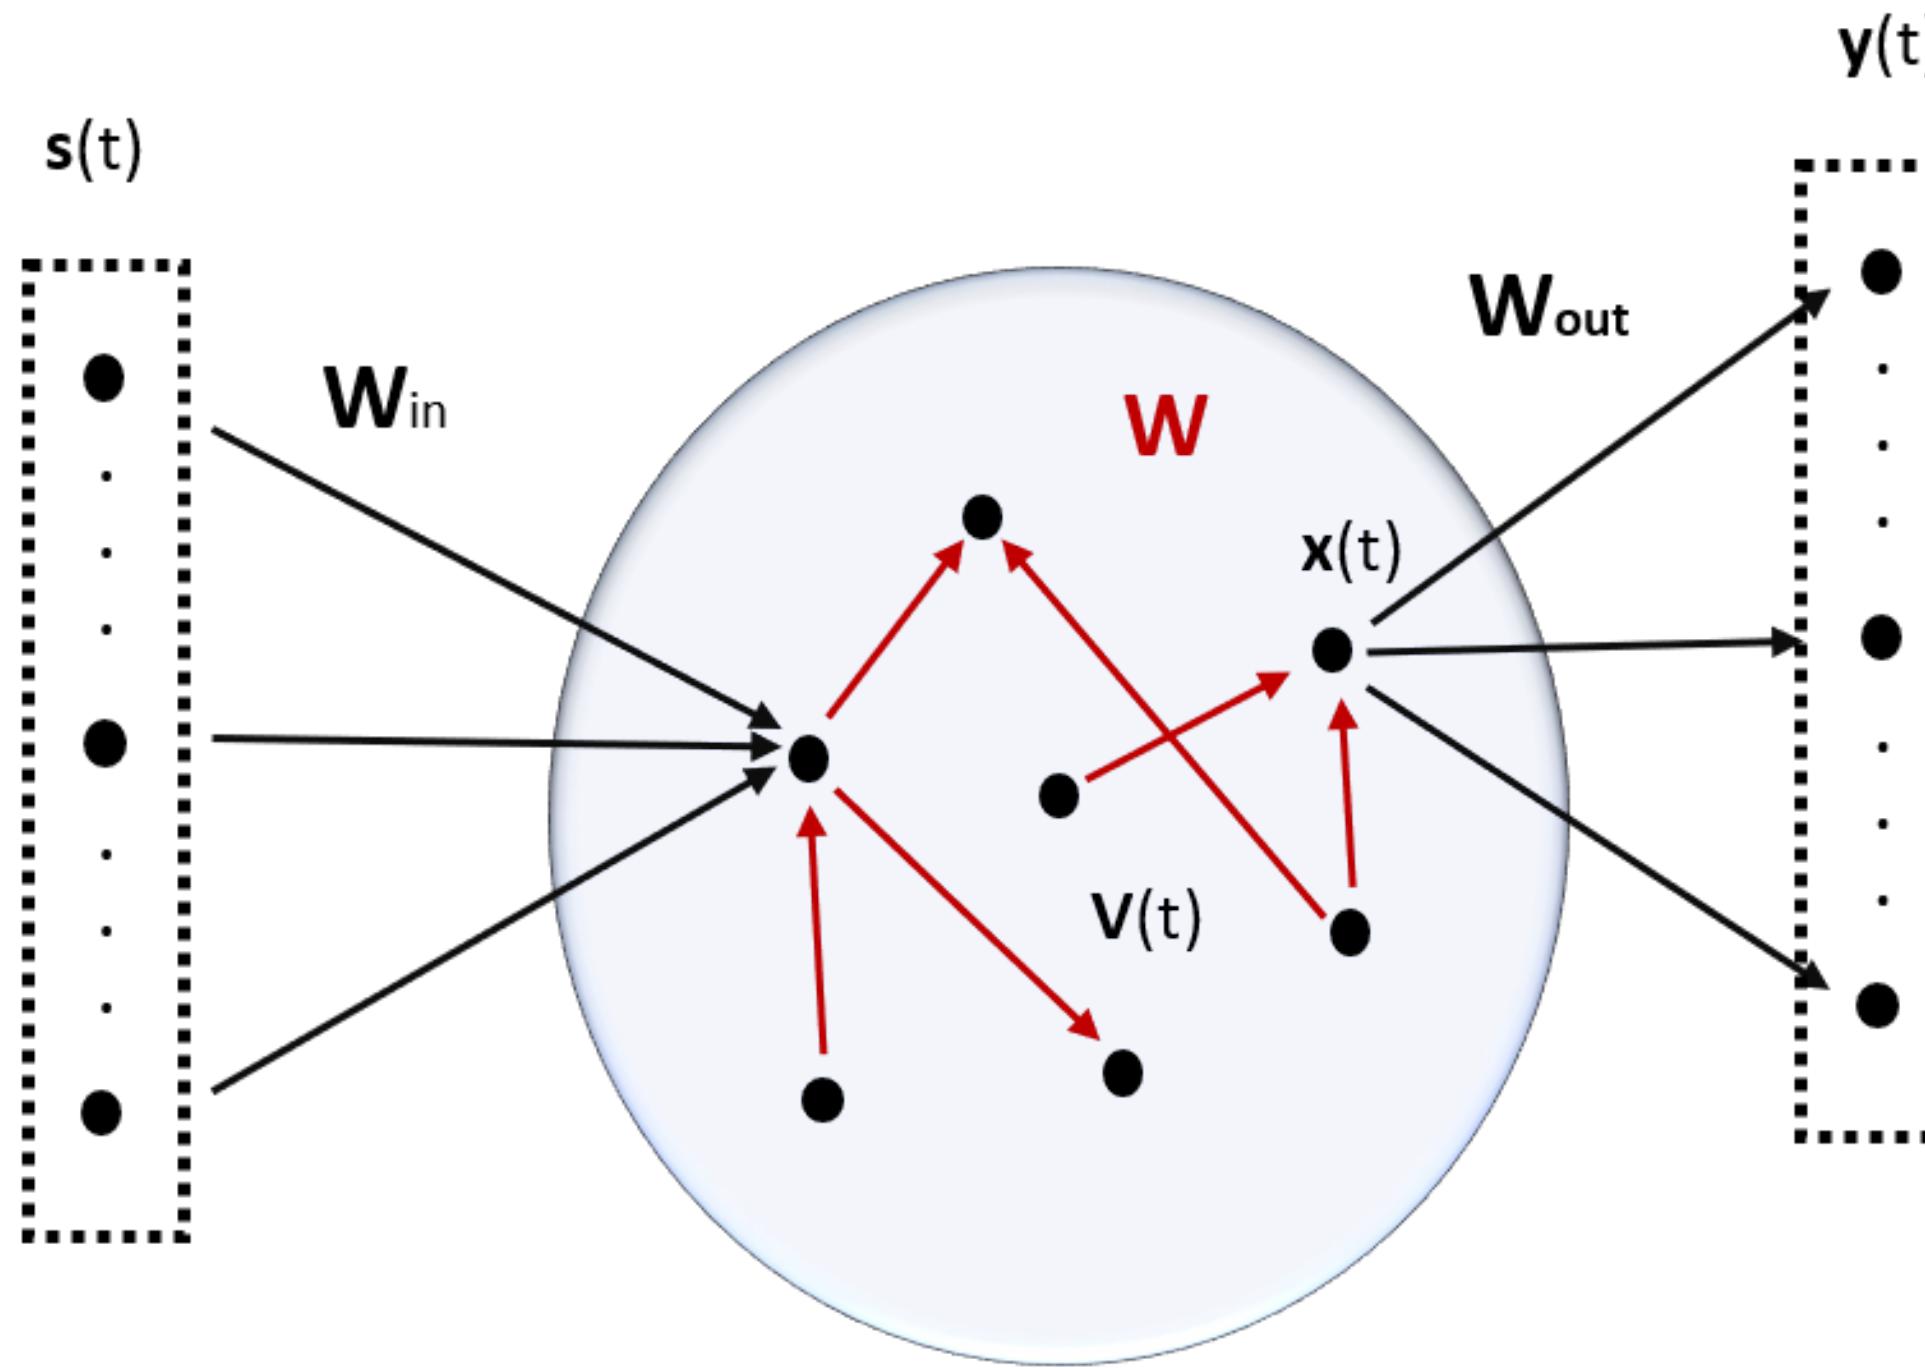
$$\alpha = \frac{\delta t}{\tau}$$

γ : gain factor

$$V(t + \delta t) = (1 - \alpha)V(t) + \alpha f(\gamma W_{in}s(t) + \rho WV(t)) \quad 0 < \rho < 1$$

Manneschi et al (2021) *Frontiers in Applied Mathematics and Statistics*

RESERVOIR COMPUTING



$$\alpha = \frac{\delta t}{\tau}$$

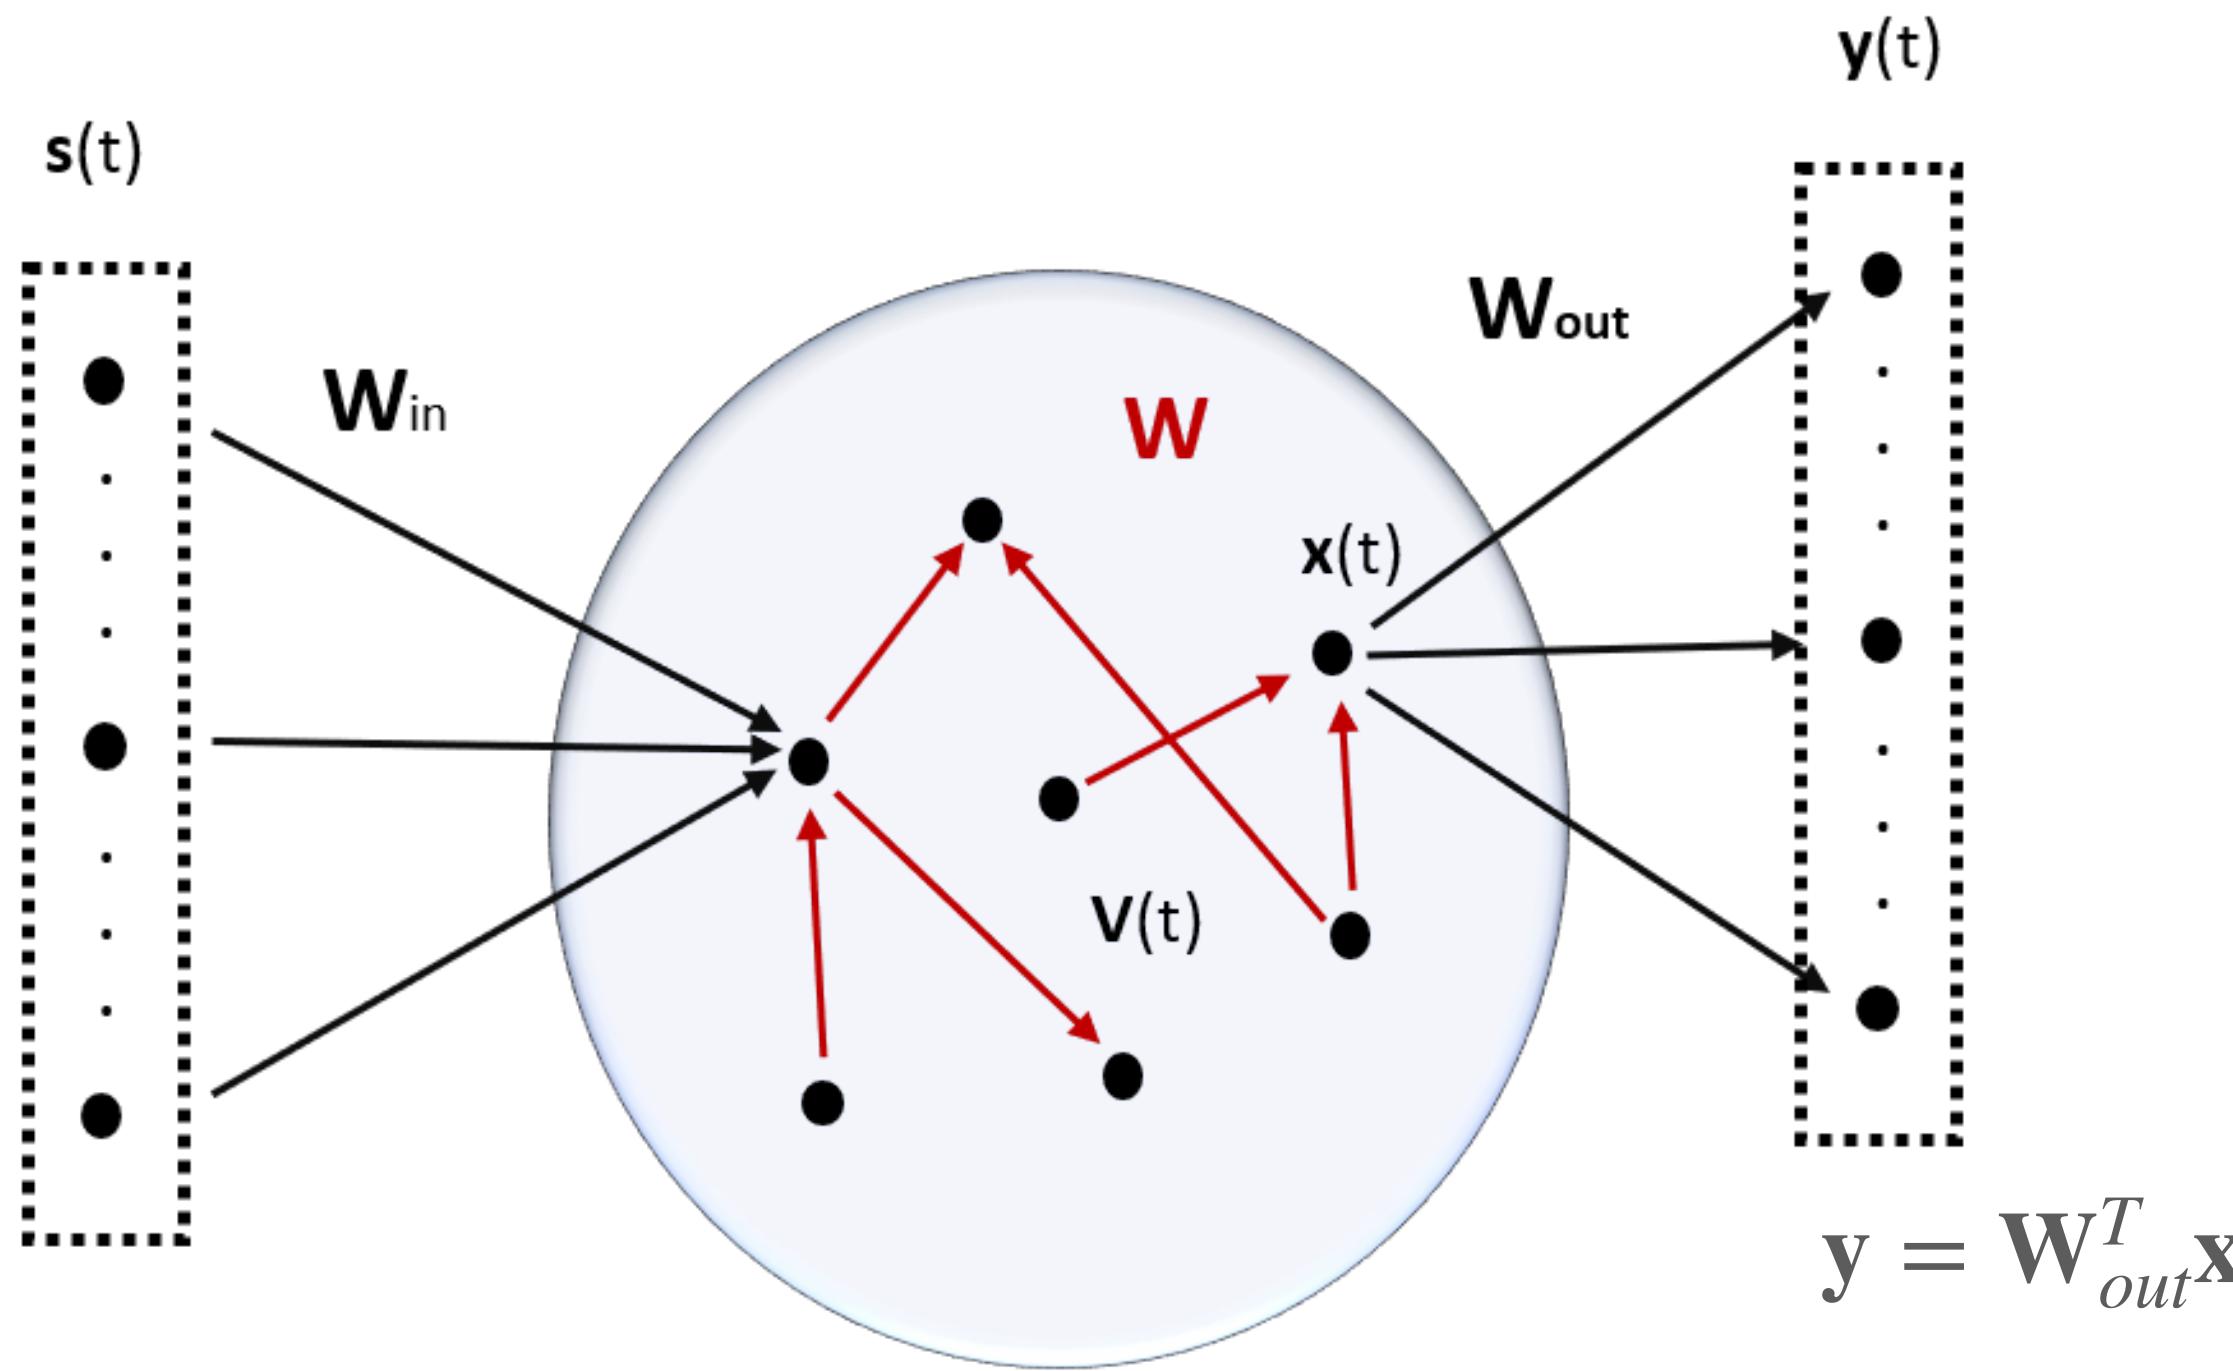
γ : gain factor

$$V(t + \delta t) = (1 - \alpha)V(t) + \alpha f(\gamma W_{in}s(t) + \rho WV(t)) \quad 0 < \rho < 1$$

$$x(t) = V(t)$$

Manneschi et al (2021) *Frontiers in Applied Mathematics and Statistics*

RESERVOIR COMPUTING



$$\alpha = \frac{\delta t}{\tau}$$

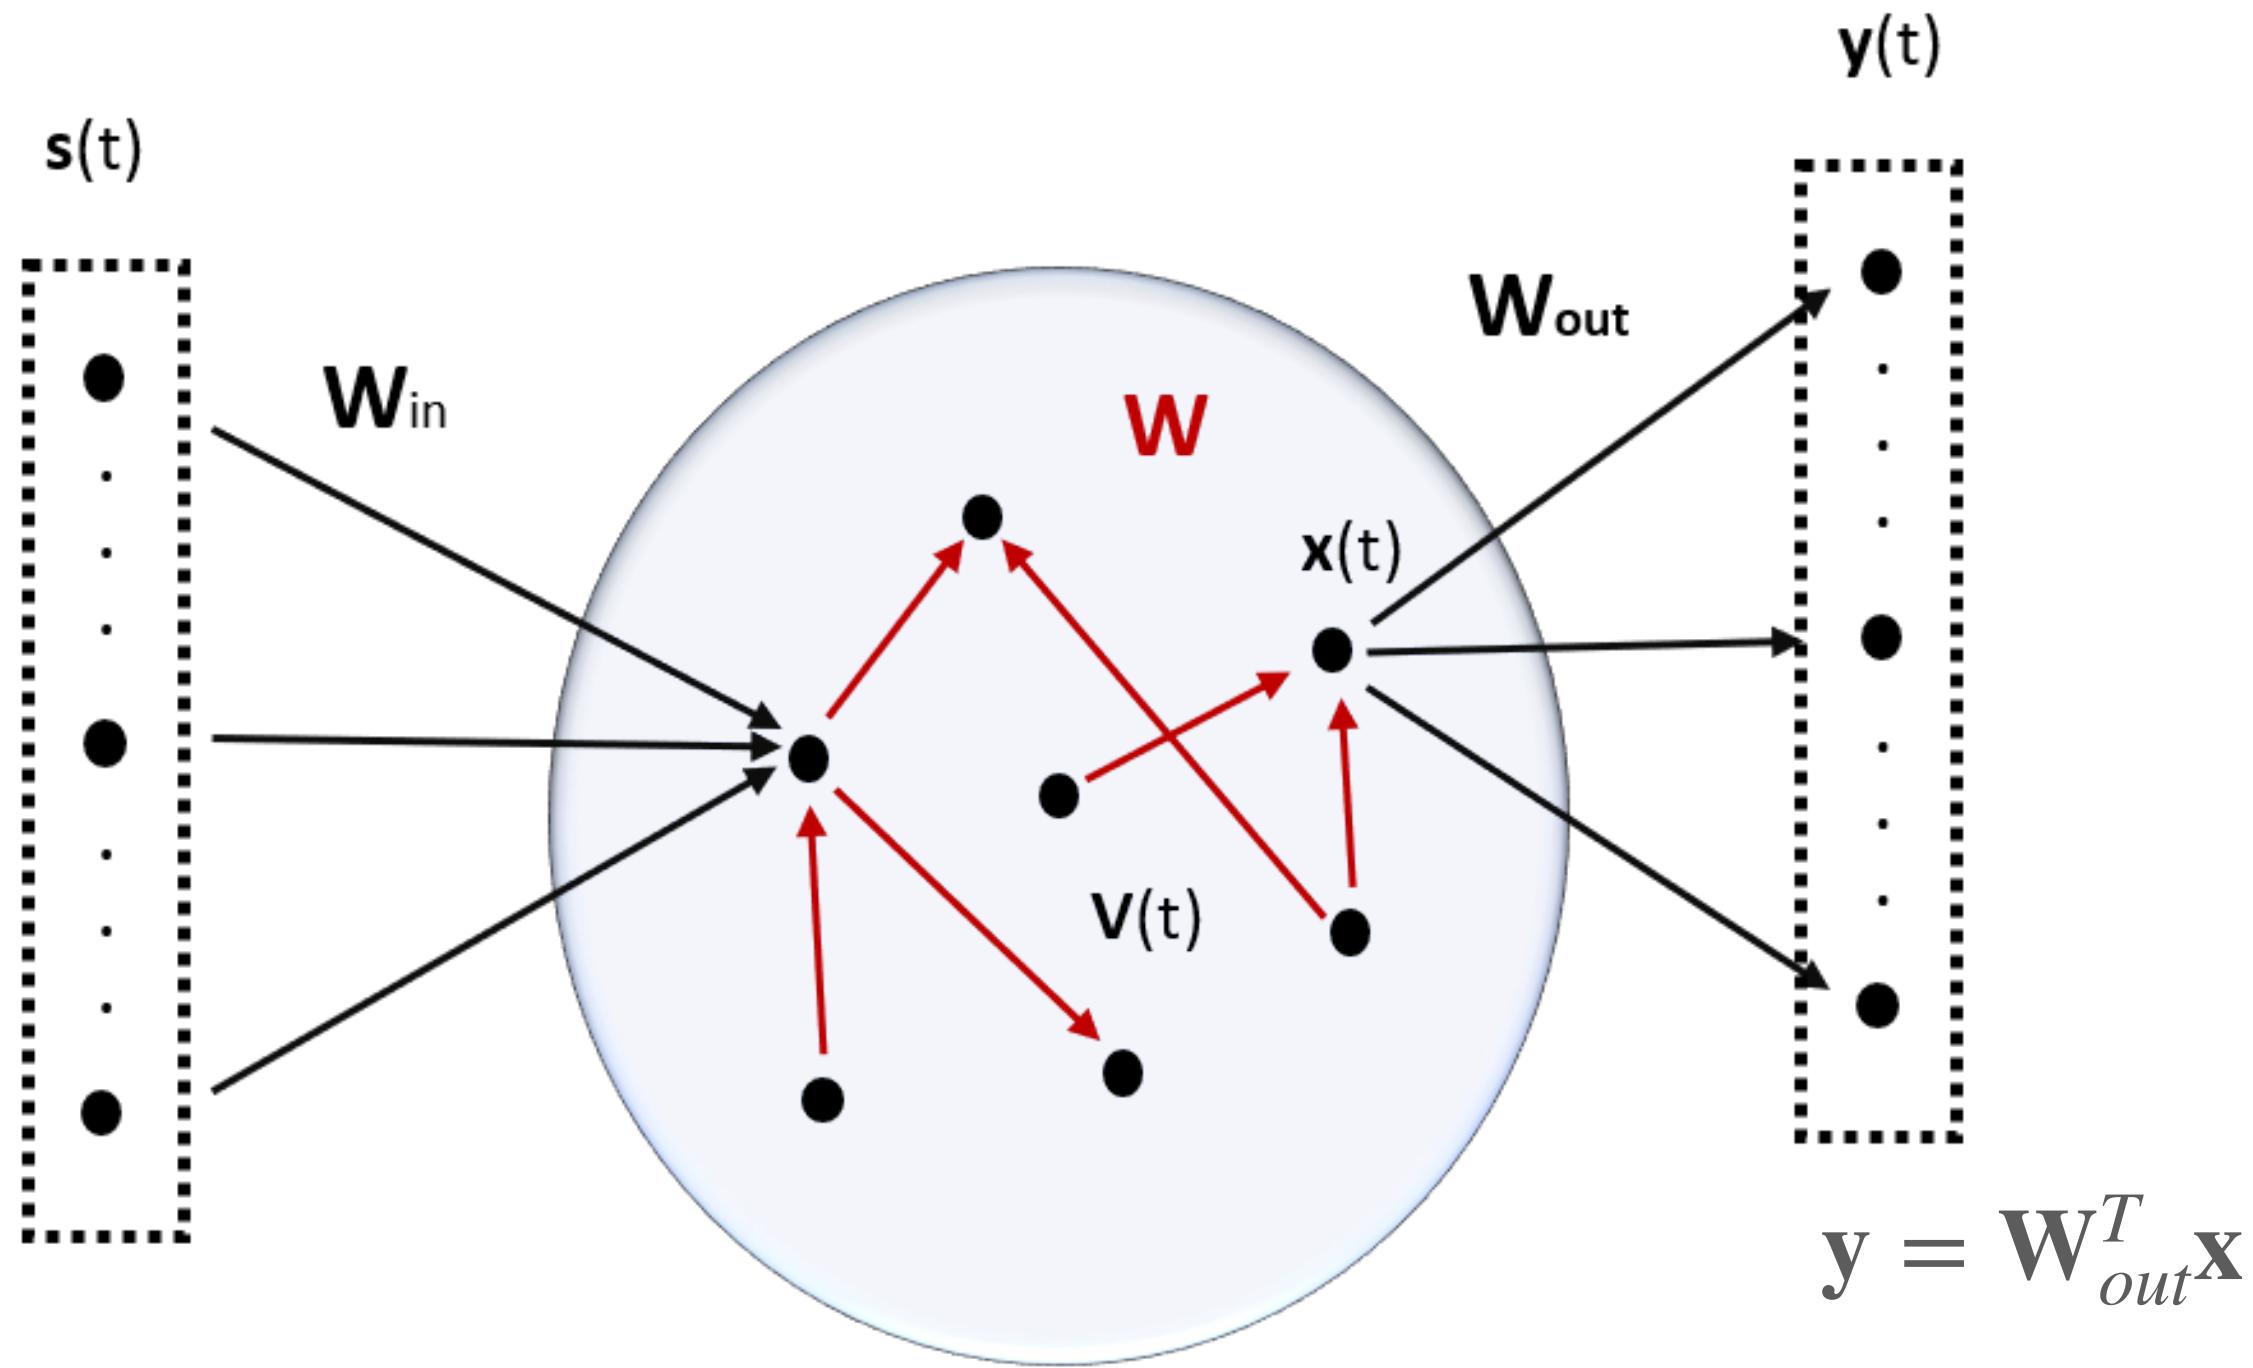
γ : gain factor

$$\mathbf{V}(t + \delta t) = (1 - \alpha)\mathbf{V}(t) + \alpha f(\gamma \mathbf{W}_{in} \mathbf{s}(t) + \rho \mathbf{W} \mathbf{V}(t)) \quad 0 < \rho < 1$$

$$\mathbf{x}(t) = \mathbf{V}(t)$$

Manneschi et al (2021) *Frontiers in Applied Mathematics and Statistics*

RESERVOIR COMPUTING



*Echo state property:
Eigenvalues of W
inside unit circle*

$$\alpha = \frac{\delta t}{\tau}$$

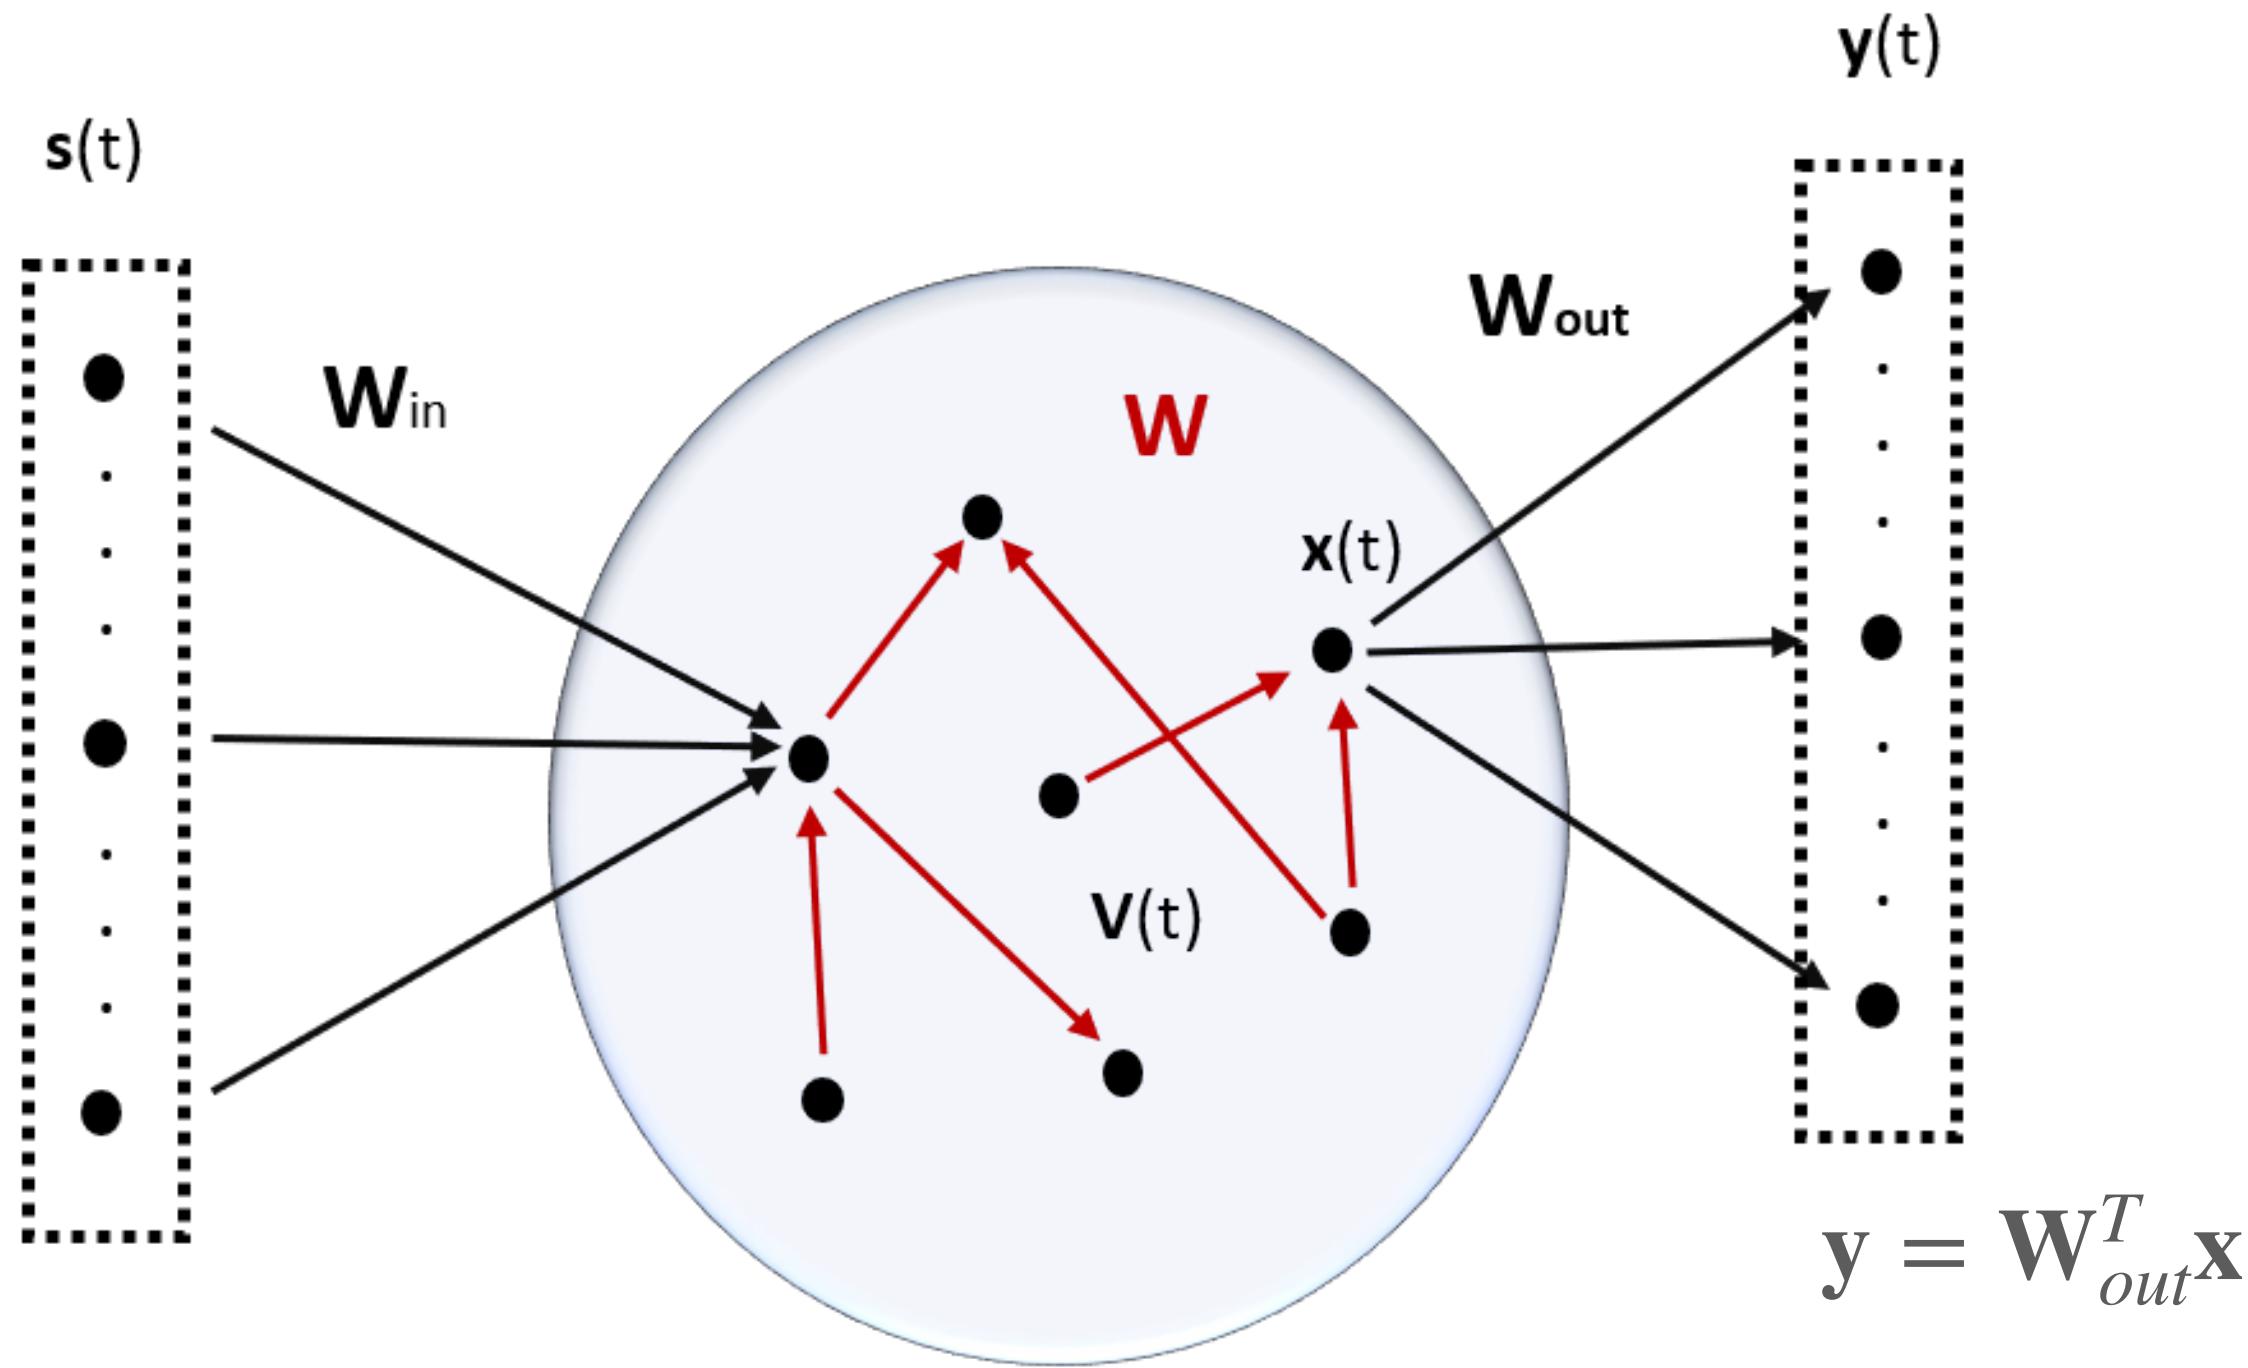
γ : gain factor

$$V(t + \delta t) = (1 - \alpha)V(t) + \alpha f(\gamma W_{in}s(t) + \rho WV(t)) \quad 0 < \rho < 1$$

$$x(t) = V(t)$$

*Manneschi et al (2021) Frontiers in
Applied Mathematics and Statistics*

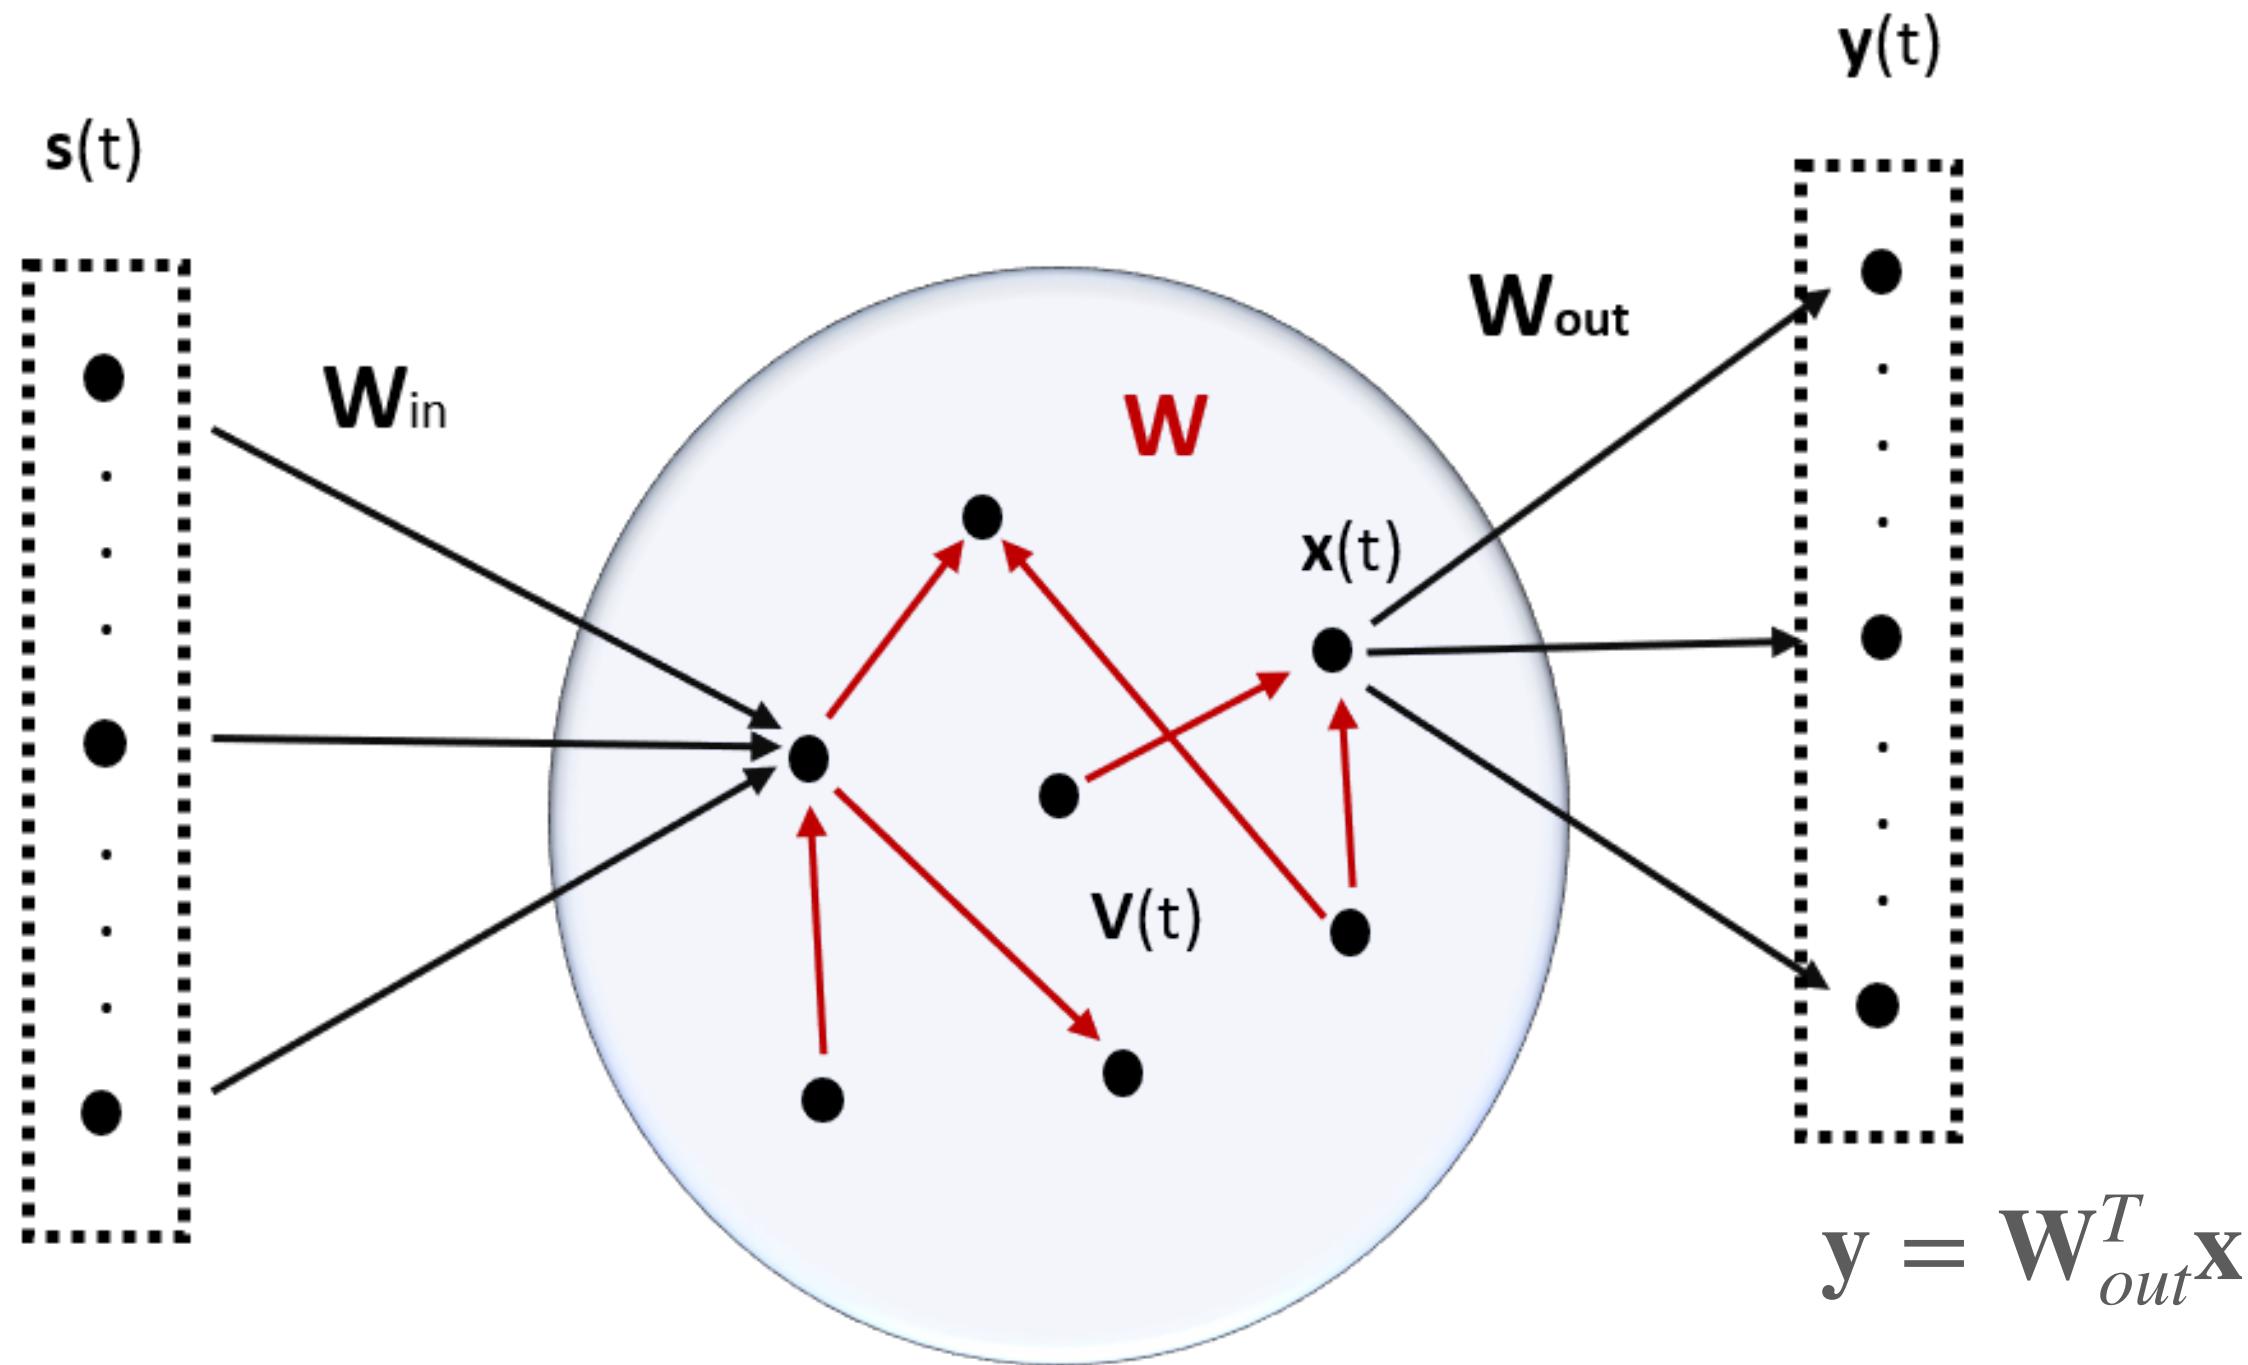
SPARSE RESERVOIR COMPUTING (SPARCE)



$$V(t + \delta t) = (1 - \alpha)V(t) + \alpha f(\gamma W_{in}s(t) + \rho WV(t))$$

Manneschi et al (2021) IEEE TNNLS
<https://arxiv.org/abs/1907.08600>

SPARSE RESERVOIR COMPUTING (SPARCE)

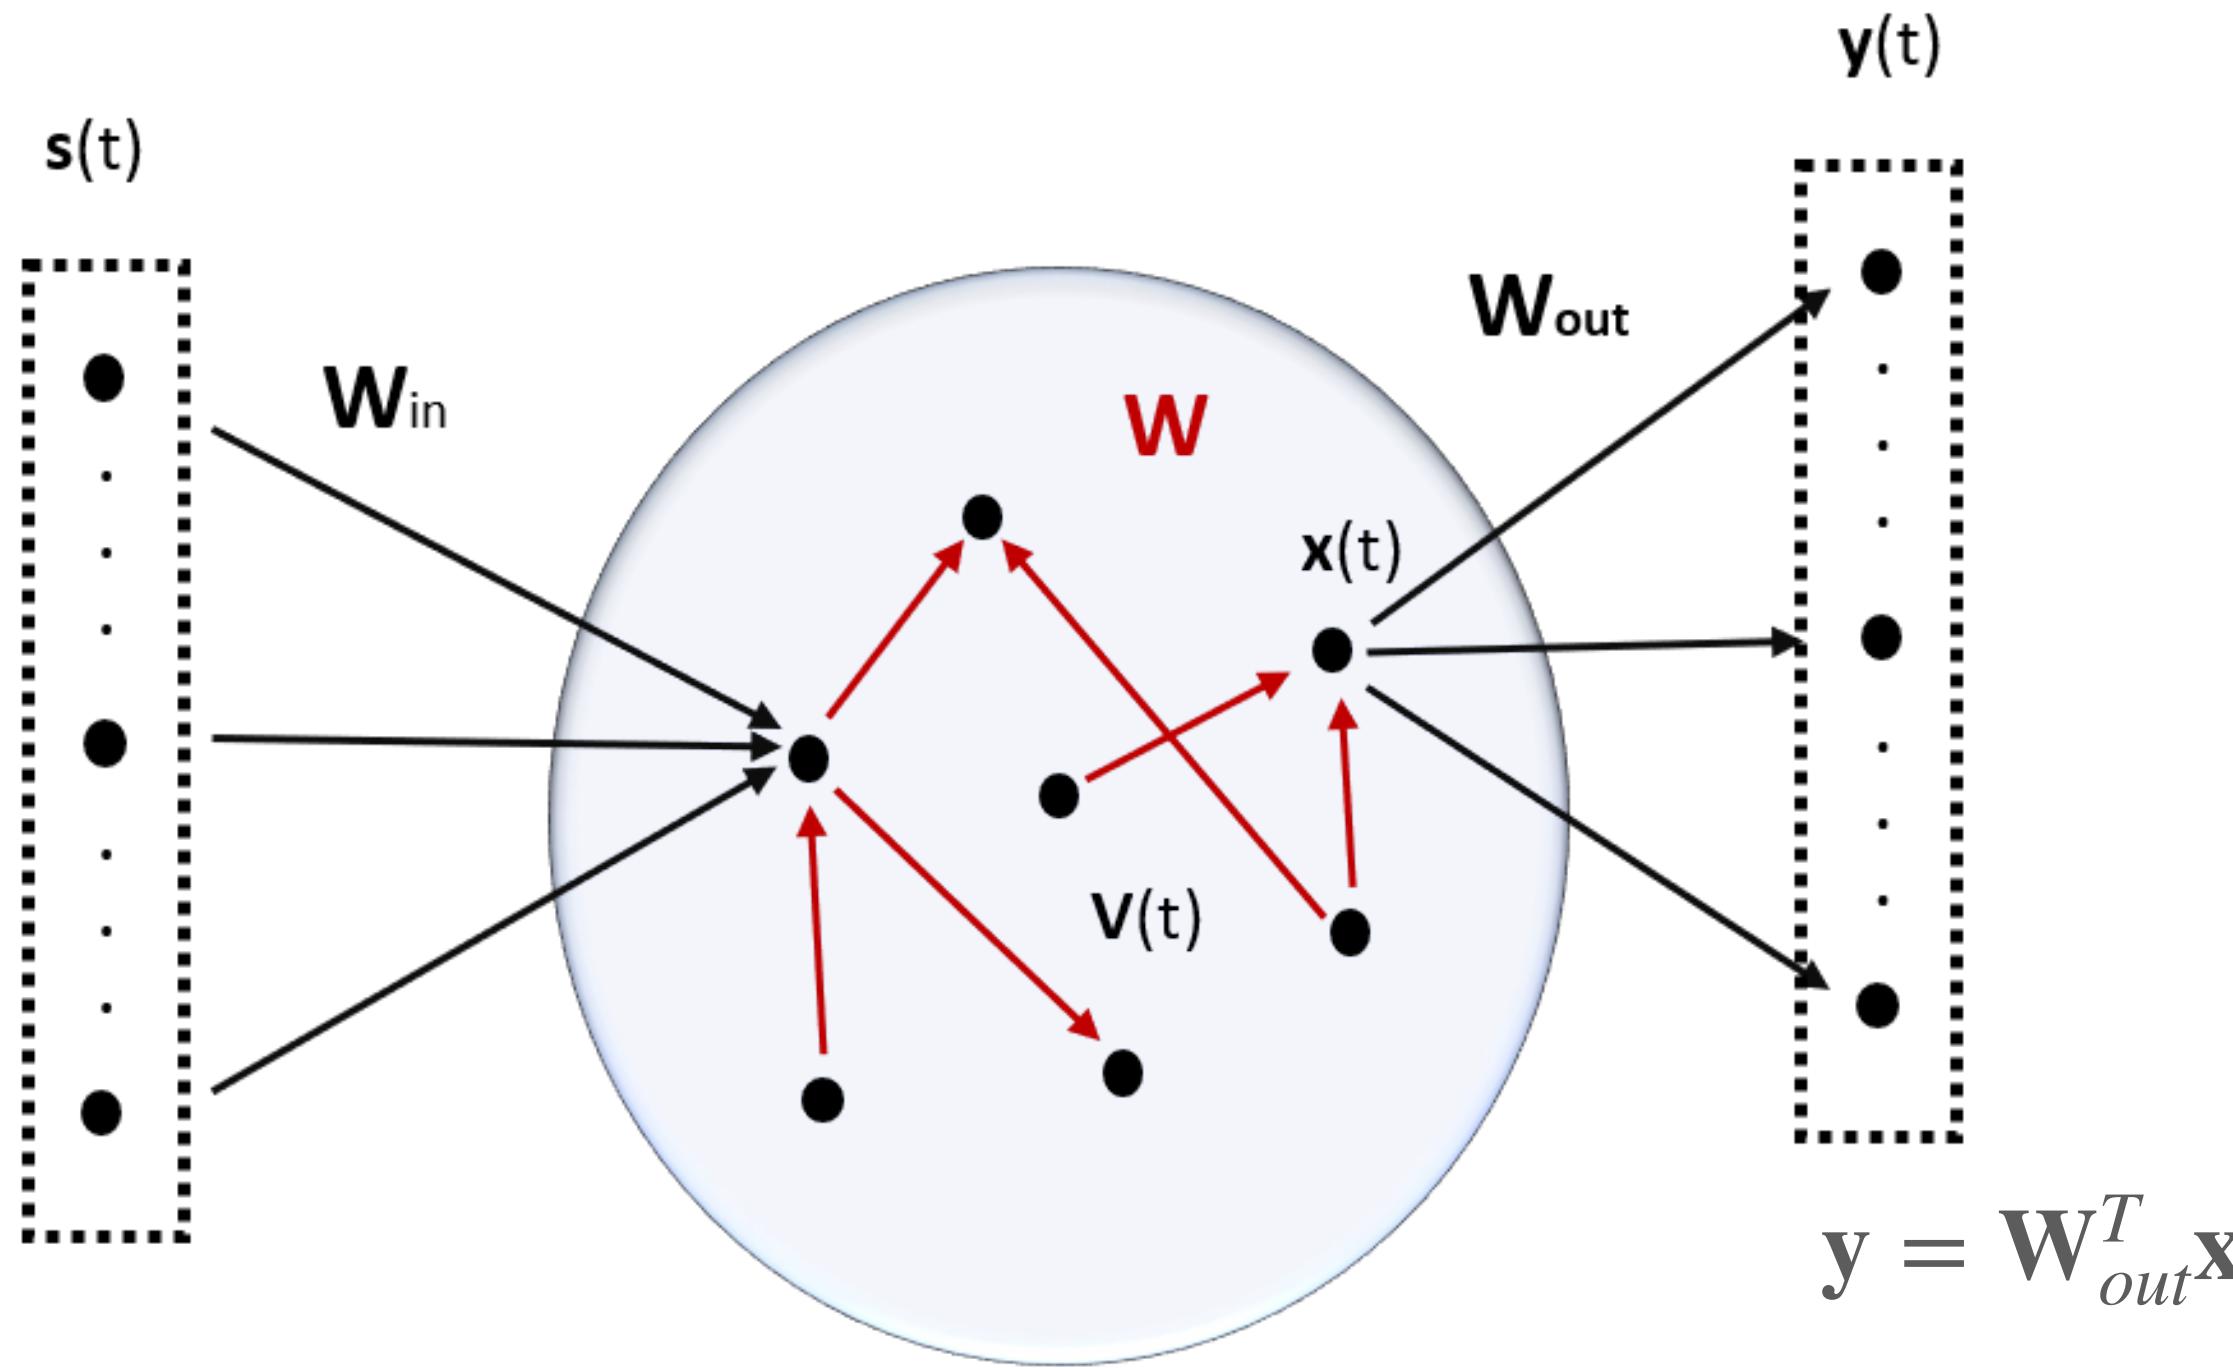


$$V(t + \delta t) = (1 - \alpha)V(t) + \alpha f(\gamma W_{in}s(t) + \rho WV(t))$$

$$x_i(t) = sign(V_i(t)) \operatorname{relu}\left[|V_i(t)| - \theta_i \right]$$

Manneschi et al (2021) IEEE TNNLS
<https://arxiv.org/abs/1907.08600>

SPARSE RESERVOIR COMPUTING (SPARCE)



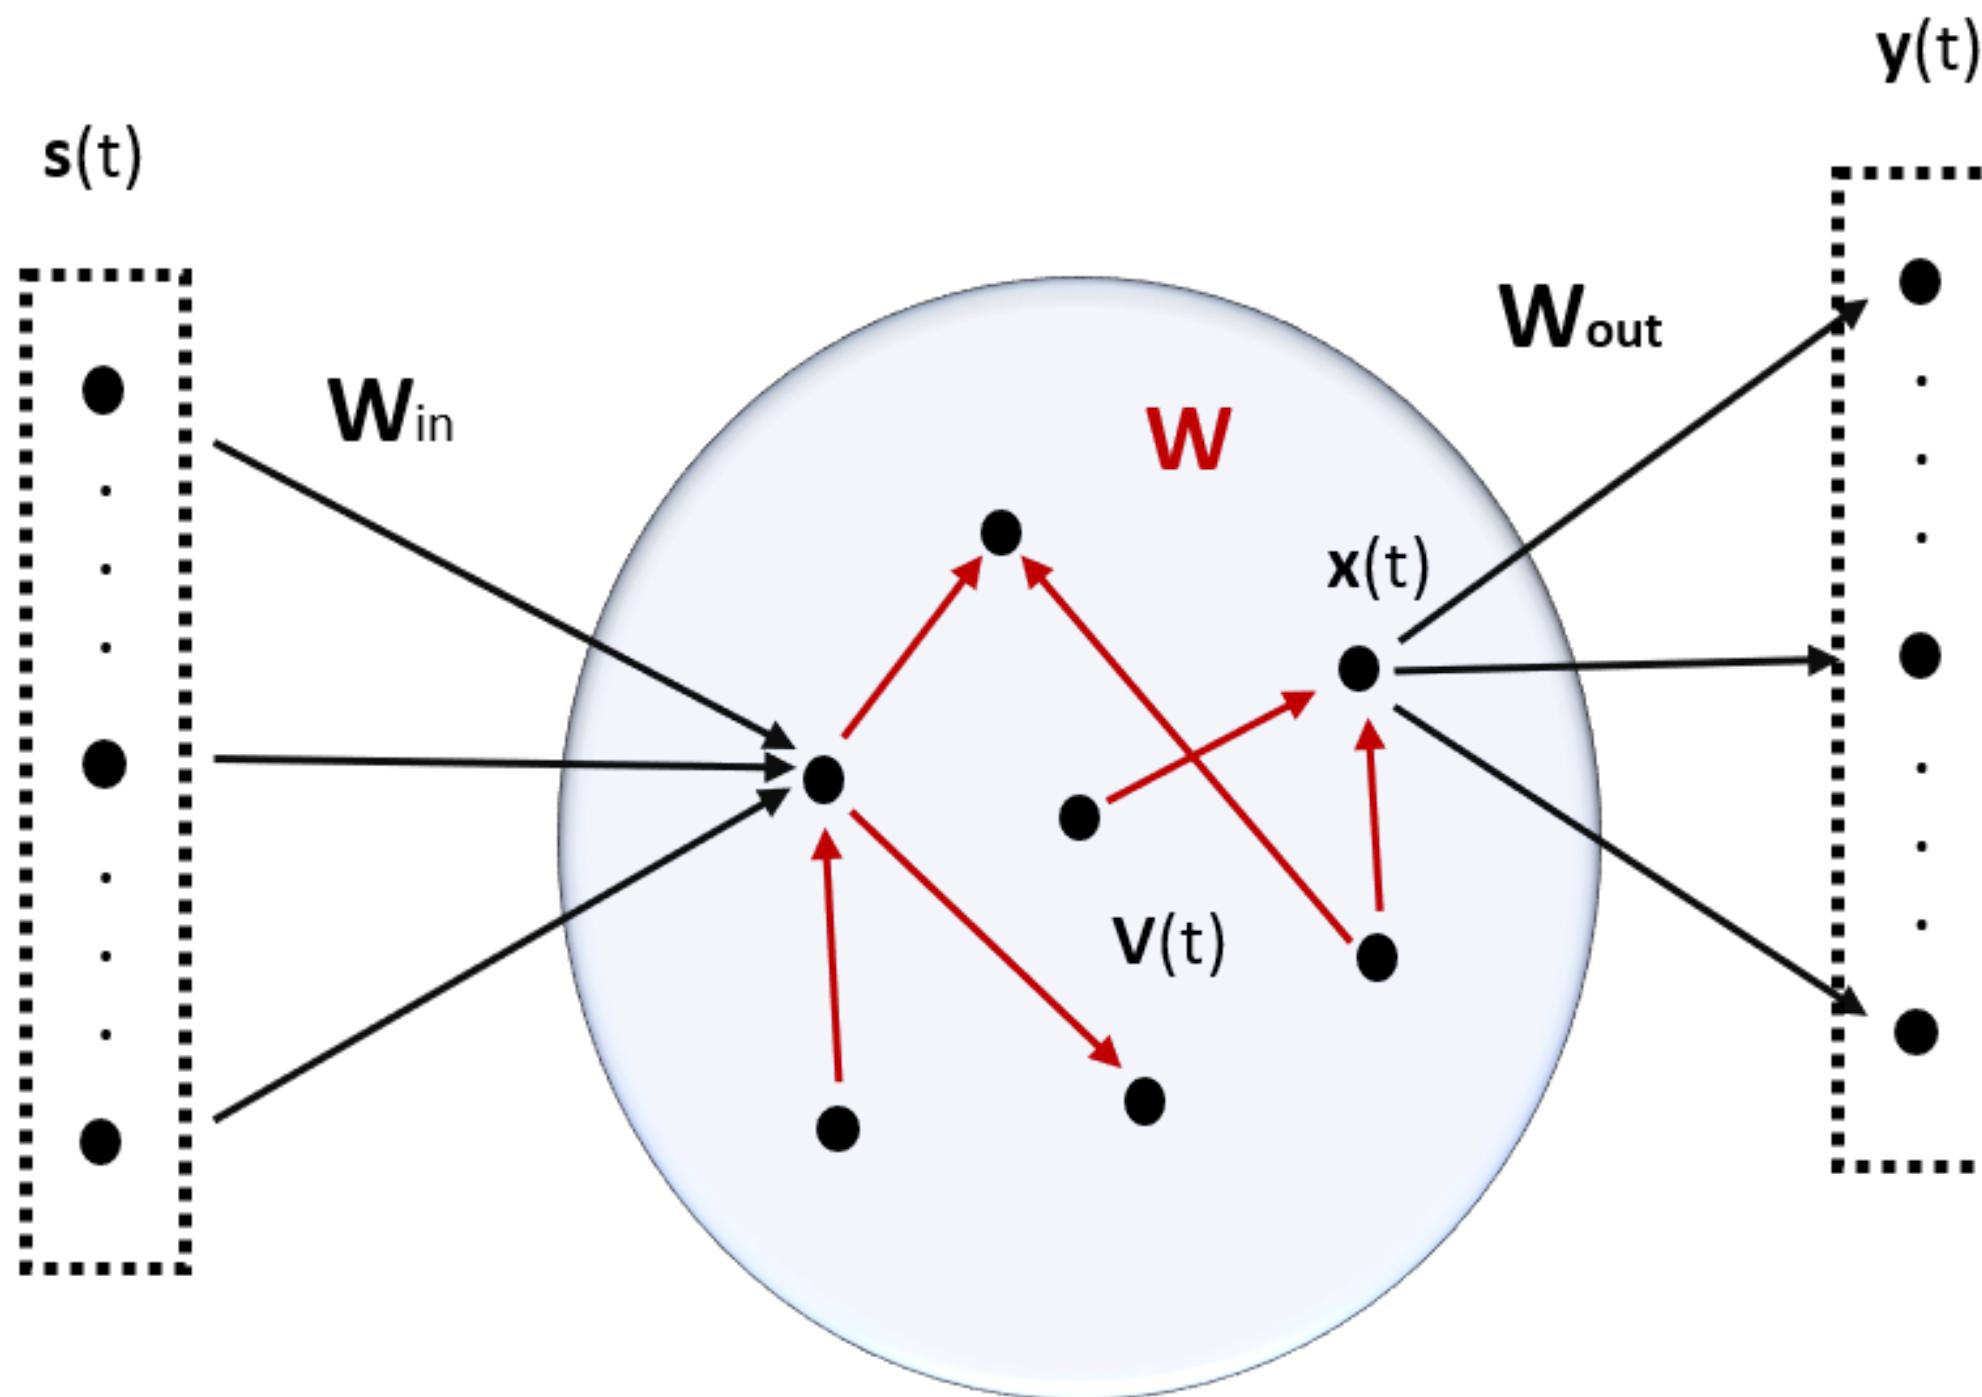
$$relu(x) = \max(0, x)$$

$$\mathbf{V}(t + \delta t) = (1 - \alpha)\mathbf{V}(t) + \alpha f(\gamma \mathbf{W}_{in} \mathbf{s}(t) + \rho \mathbf{W} \mathbf{V}(t))$$

$$x_i(t) = sign(V_i(t)) \, relu \left[|V_i(t)| - \theta_i \right]$$

Manneschi et al (2021) IEEE TNNLS
<https://arxiv.org/abs/1907.08600>

READ-OUT WEIGHTS UPDATE RULE



$$E = \frac{1}{2} \sum_j [y_j^t - y_j]^2$$
$$\Delta W_{ji}^{out} = -\eta_w \frac{\partial E}{\partial W_{ji}^{out}} = \eta_w [y_j^t - y_j] x_i(t)$$

On line setting: (more) biologically plausible, avoid inversion of large matrices.

LEARNING THE READ-OUT THRESHOLDS

$$x_i = \text{sign}(V_i) \text{relu} \left[|V_i| - \theta_i \right]$$

Gradient descent, local thresholds θ_i

LEARNING THE READ-OUT THRESHOLDS

$$x_i = \text{sign}(V_i) \text{relu}(|V_i| - \theta_i)$$

Gradient descent, local thresholds θ_i

$$E = \frac{1}{2} \sum_j \left[y_j^t - \sum_l W_{jl}^{out} \text{sign}(V_l) \text{relu}(|V_l| - \theta_l) \right]^2$$

LEARNING THE READ-OUT THRESHOLDS

$$x_i = \text{sign}(V_i) \text{relu}(|V_i| - \theta_i)$$

Gradient descent, local thresholds θ_i

$$E = \frac{1}{2} \sum_j \left[\underline{y_j^t} - \sum_l W_{jl}^{out} \text{sign}(V_l) \text{relu}(|V_l| - \theta_l) \right]^2$$

Target

LEARNING THE READ-OUT THRESHOLDS

$$x_i = \text{sign}(V_i) \text{relu}(|V_i| - \theta_i)$$

Gradient descent, local thresholds θ_i

$$E = \frac{1}{2} \sum_j \left[\underbrace{y_j^t}_{\text{Target}} - \underbrace{\sum_l W_{jl}^{\text{out}} \text{sign}(V_l) \text{relu}(|V_l| - \theta_l)}_{\text{Output } y_j} \right]^2$$

LEARNING THE READ-OUT THRESHOLDS

$$x_i = \text{sign}(V_i) \text{relu}(|V_i| - \theta_i)$$

Gradient descent, local thresholds θ_i

$$E = \frac{1}{2} \sum_j \left[\underbrace{y_j^t}_{\text{Target}} - \underbrace{\sum_l W_{jl}^{out} \text{sign}(V_l) \text{relu}(|V_l| - \theta_l)}_{\text{Output } y_j} \right]^2$$

$$\Delta \theta_i = -\eta \frac{\partial E}{\partial \theta_i}$$

LEARNING THE READ-OUT THRESHOLDS

$$x_i = \text{sign}(V_i) \operatorname{relu} \left[|V_i| - \theta_i \right] \quad \text{Gradient descent, local thresholds } \theta_i$$

$$E = \frac{1}{2} \sum_j \left[y_j^t - \sum_l W_{jl}^{out} sign(V_l) relu[|V_l| - \theta_l] \right]^2$$

Target

Output y_j

$$\Delta\theta_i = -\eta \frac{\partial E}{\partial \theta_i} = -\eta \sum_j [y_j^t - y_j] W_{ji}^{out} sign(V_i) H(|V_i| - \theta_i)$$

LEARNING THE READ-OUT THRESHOLDS

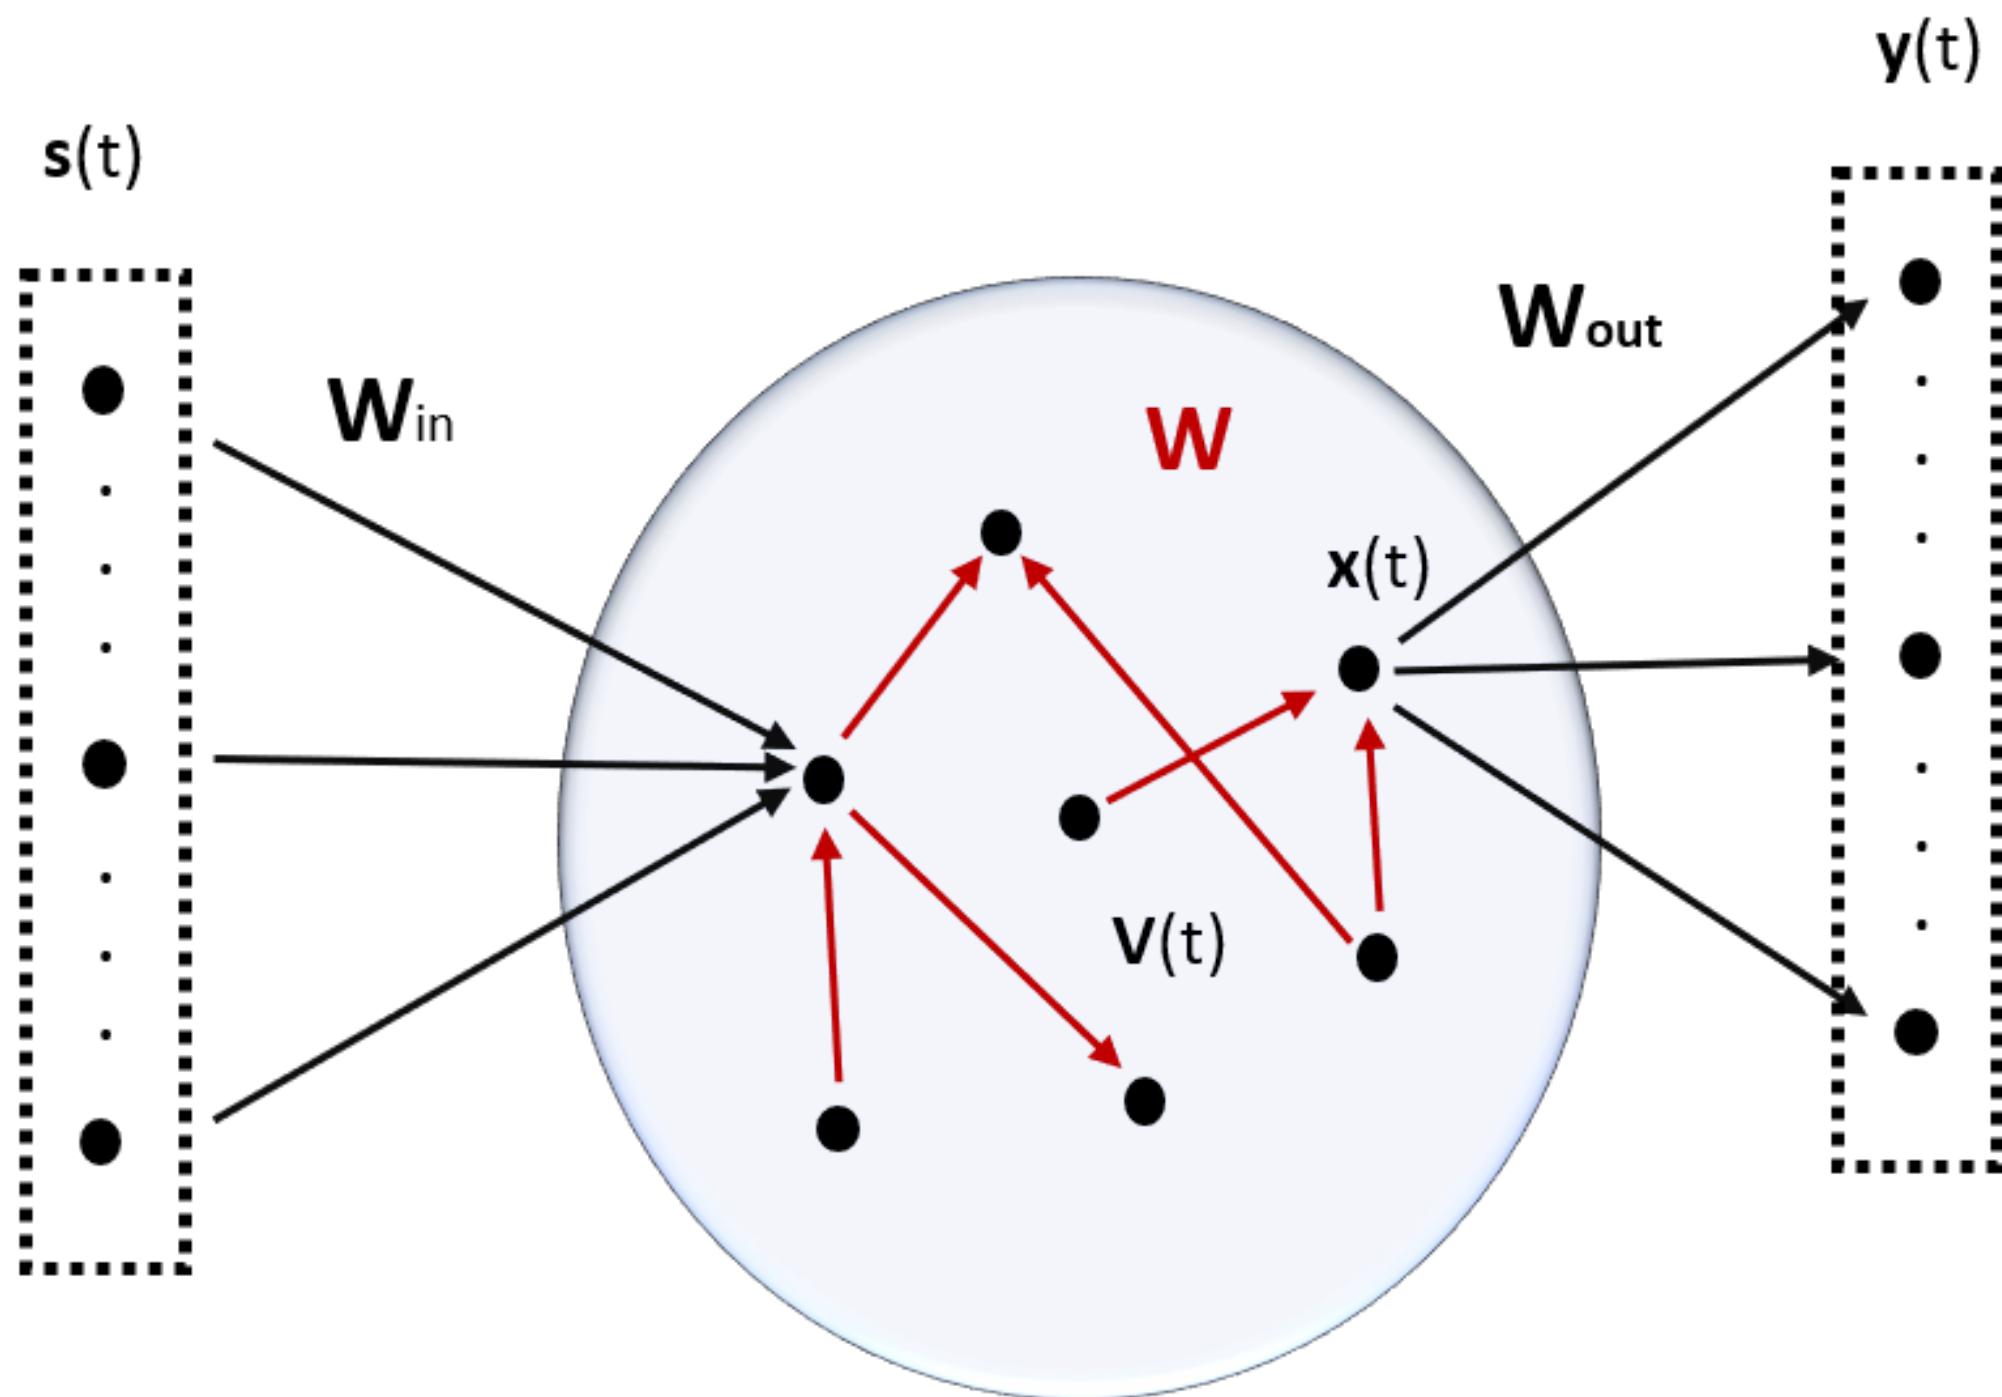
$$x_i = \text{sign}(V_i) \text{relu} \left[|V_i| - \theta_i \right] \quad \text{Gradient descent, local thresholds } \theta_i$$

$$E = \frac{1}{2} \sum_j \left[y_j^t - \sum_l W_{jl}^{out} sign(V_l) relu[|V_l| - \theta_l] \right]^2$$

— Target
 — Output y_j

$$\begin{aligned}\Delta\theta_i &= -\eta \frac{\partial E}{\partial \theta_i} = -\eta \sum_j [y_j^t - y_j] W_{ji}^{out} sign(V_i) H(|V_i| - \theta_i) \\ &= -\eta \sum_j [y_j^t - y_j] W_{ji}^{out} sign(x_i)\end{aligned}$$

THRESHOLDS UPDATE RULE

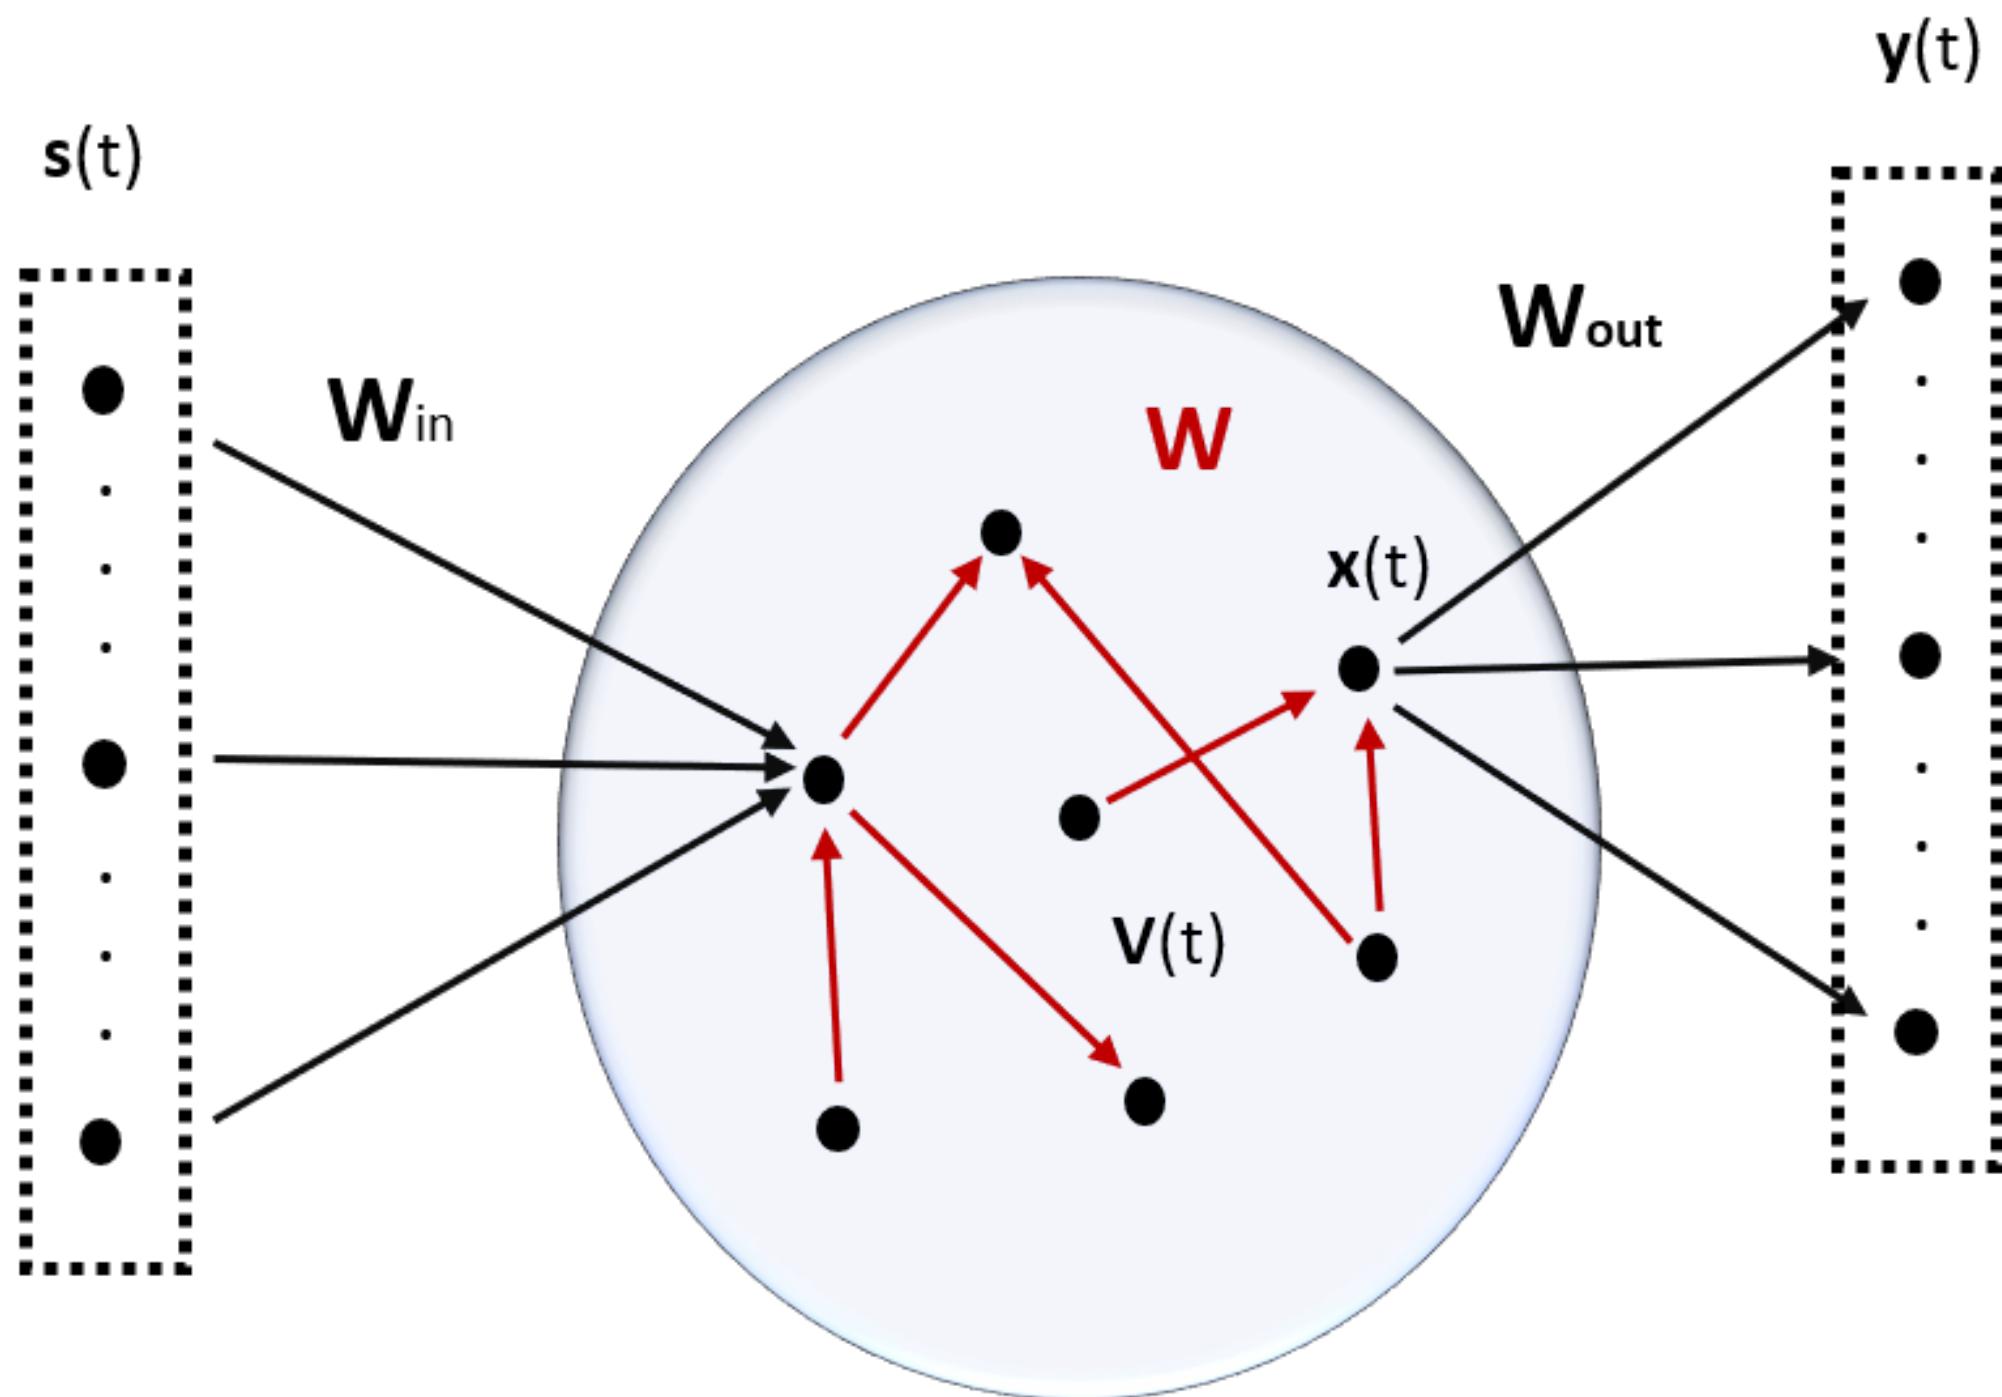


$$E = \frac{1}{2} \sum_j [y_j^t - y_j]^2$$

$$\Delta\theta_i = -\eta_\theta \sum_j [y_j^t - y_j] W_{ji}^{out} sign(x_i)$$

On line setting: (more) biologically plausible, avoid inversion of large matrices.

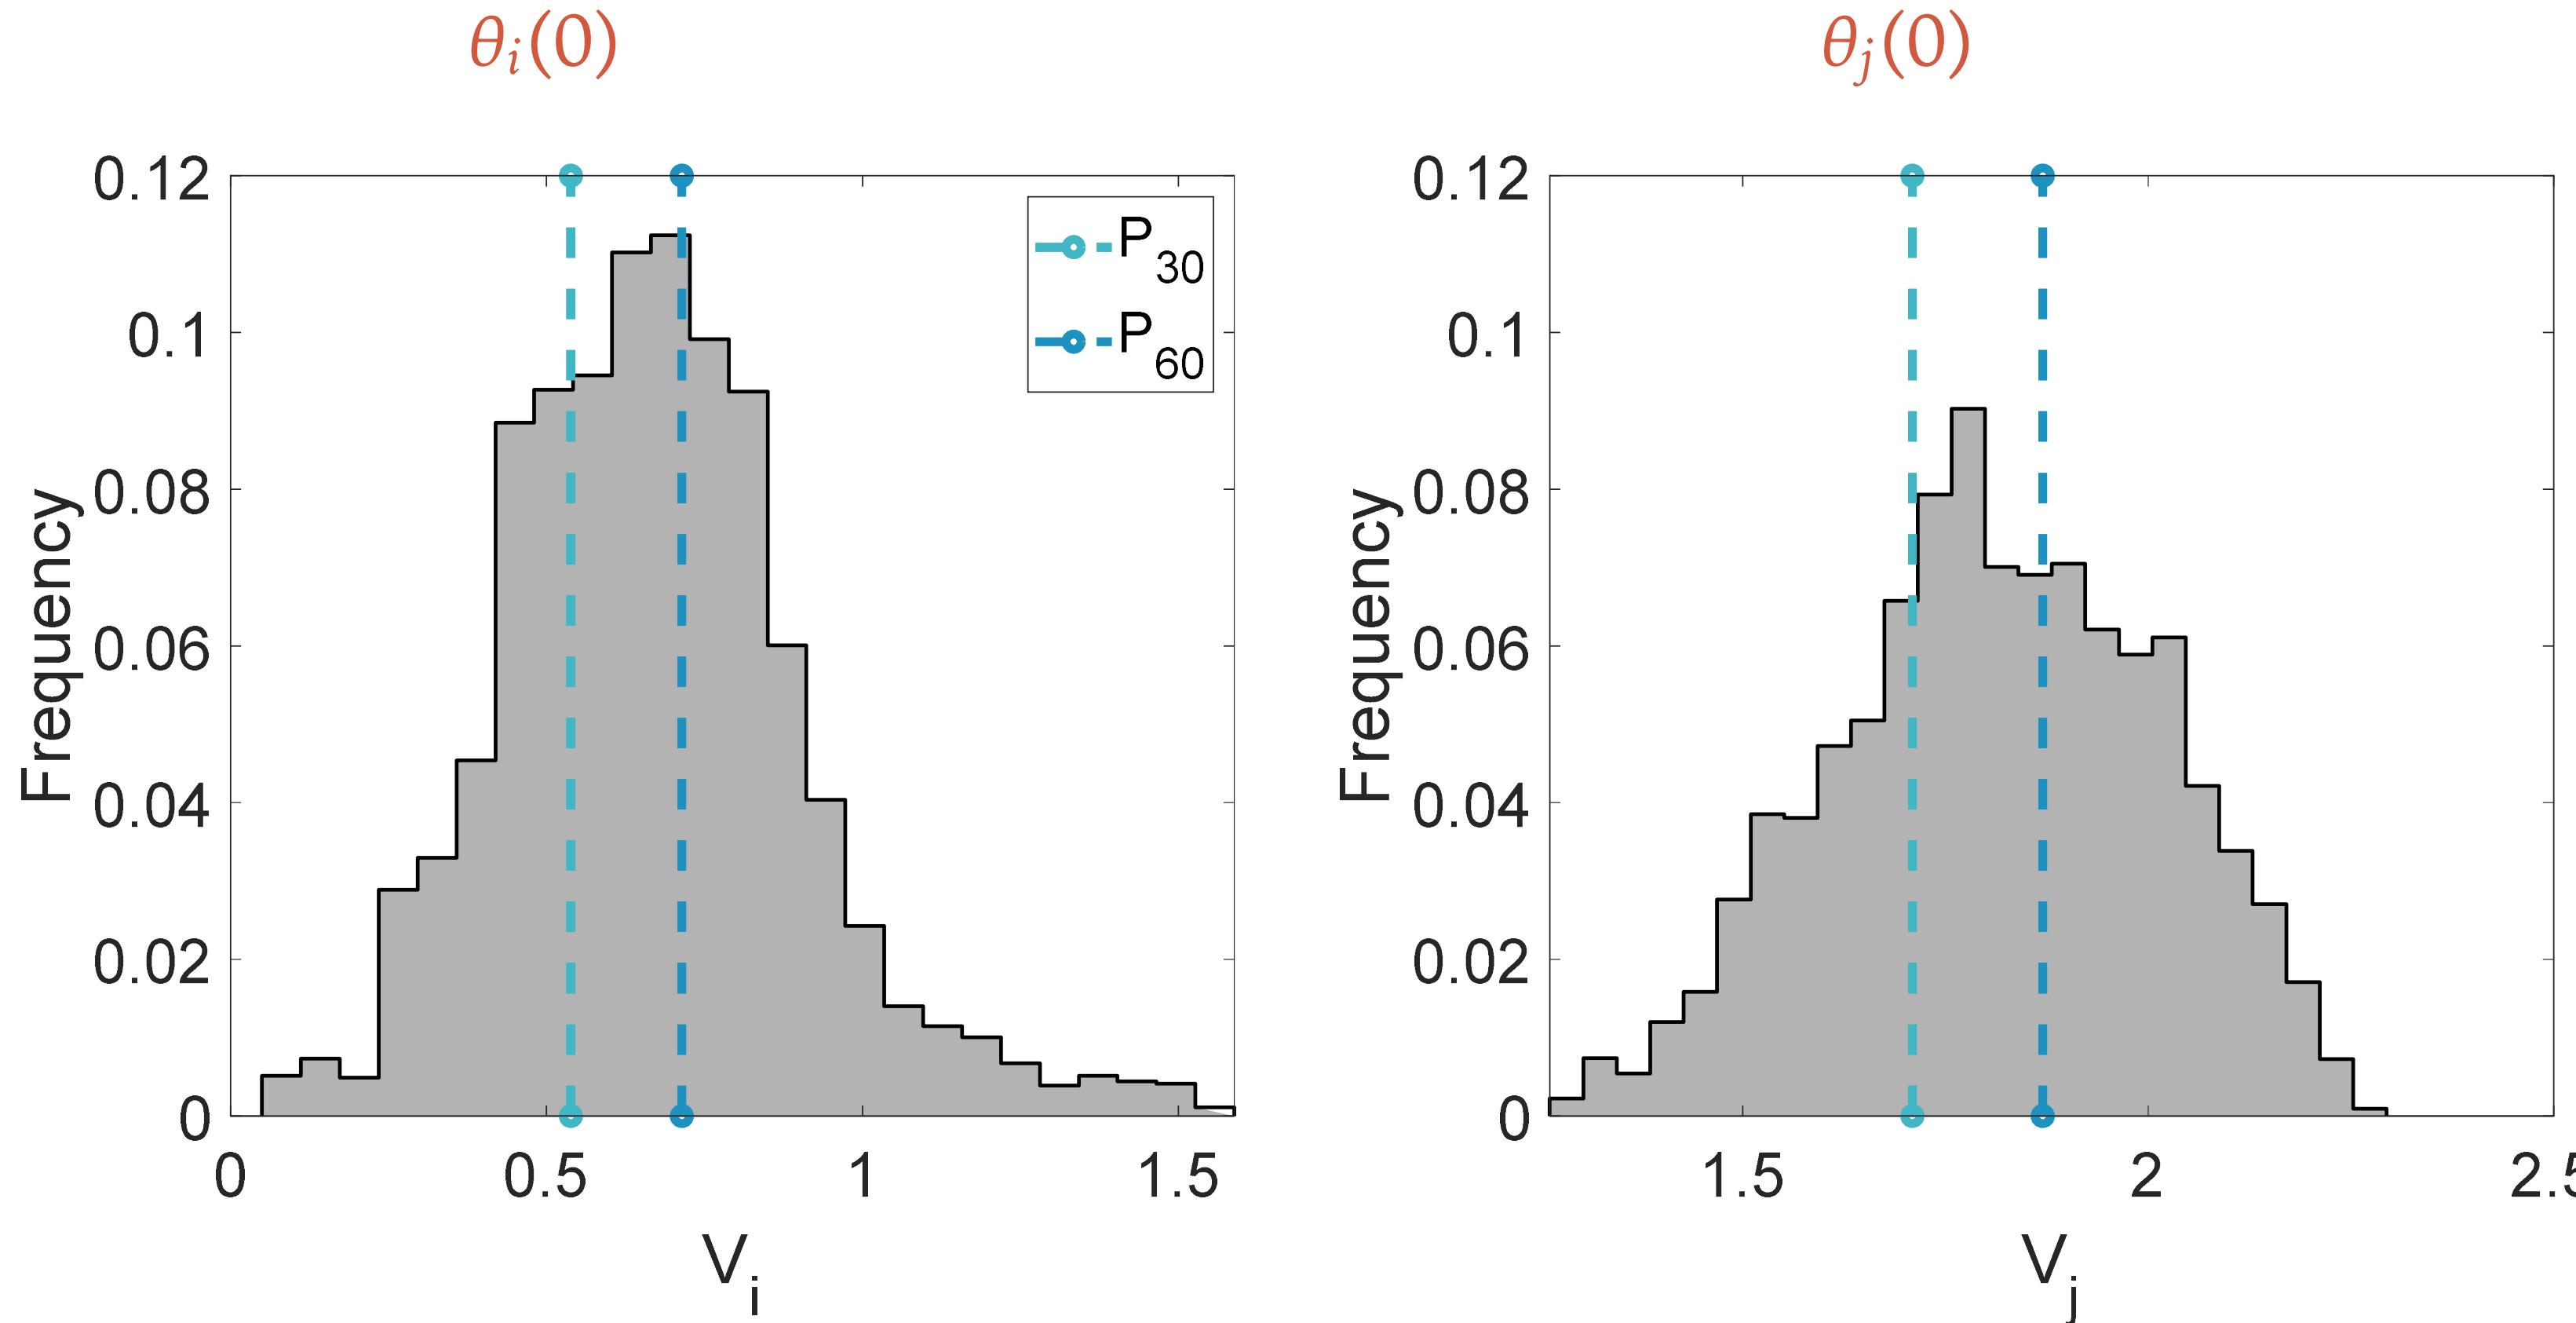
THRESHOLDS UPDATE RULE



$$E = \frac{1}{2} \sum_j [y_j^t - y_j]^2$$
$$\Delta\theta_i = -\eta_\theta \sum_j [y_j^t - y_j] W_{ji}^{out} \text{sign}(x_i)$$

On line setting: (more) biologically plausible, avoid inversion of large matrices.

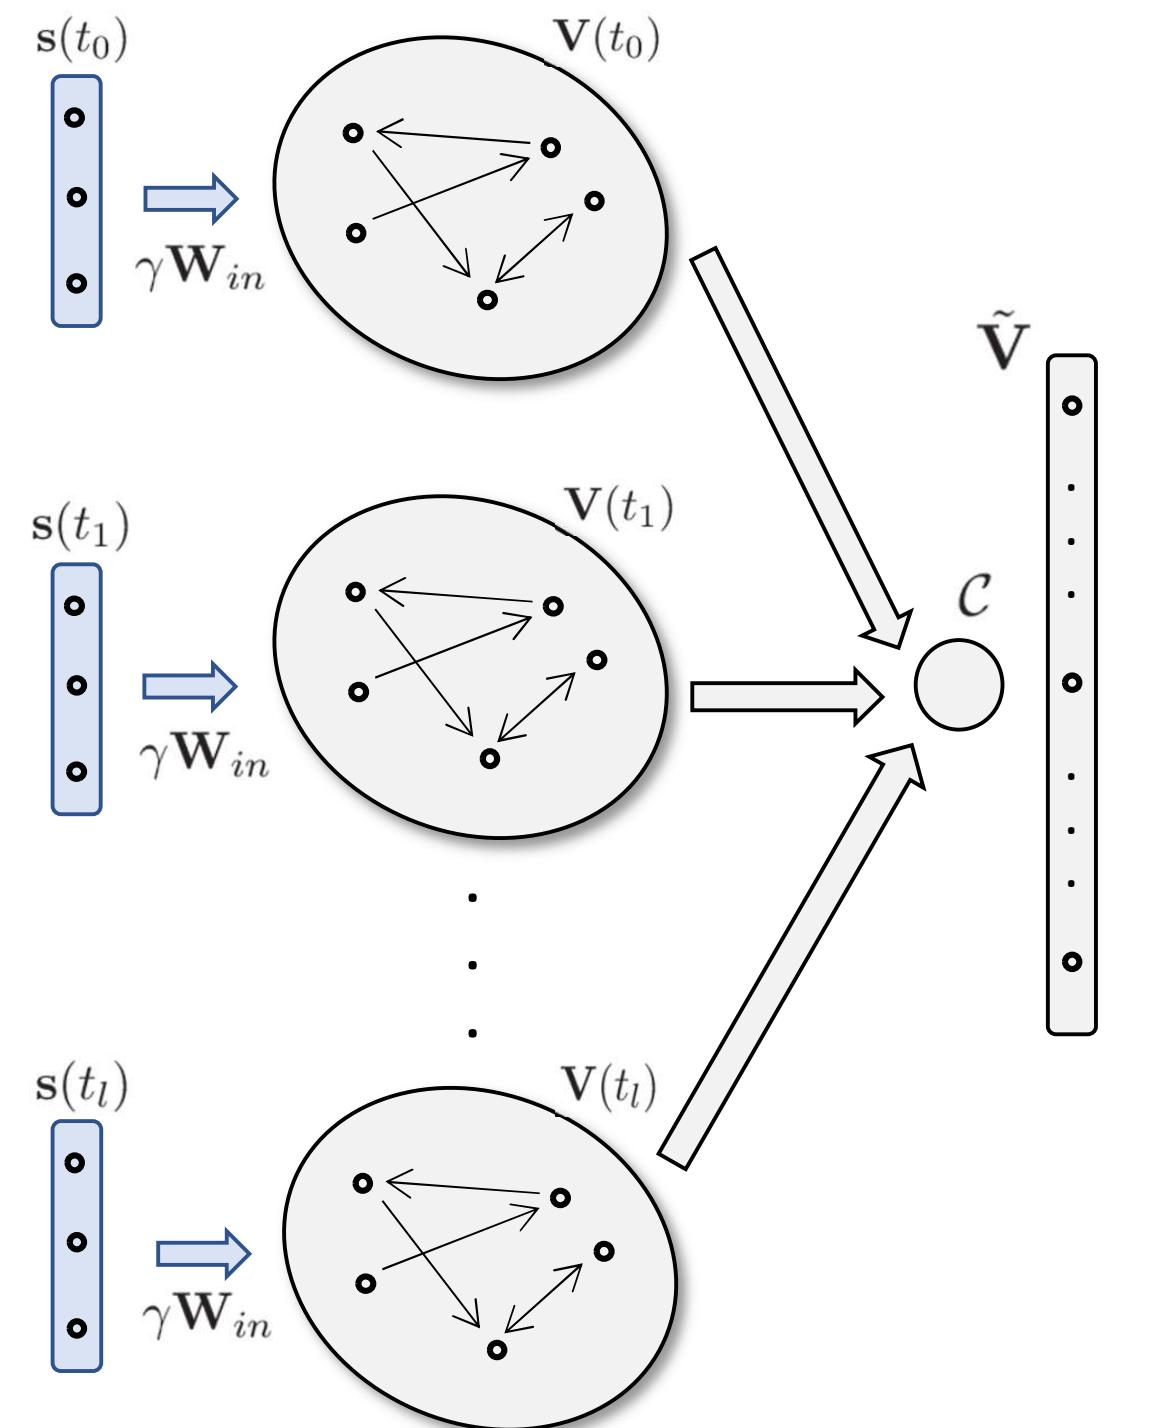
“DEMOCRATIC” INITIALISATION OF THRESHOLDS



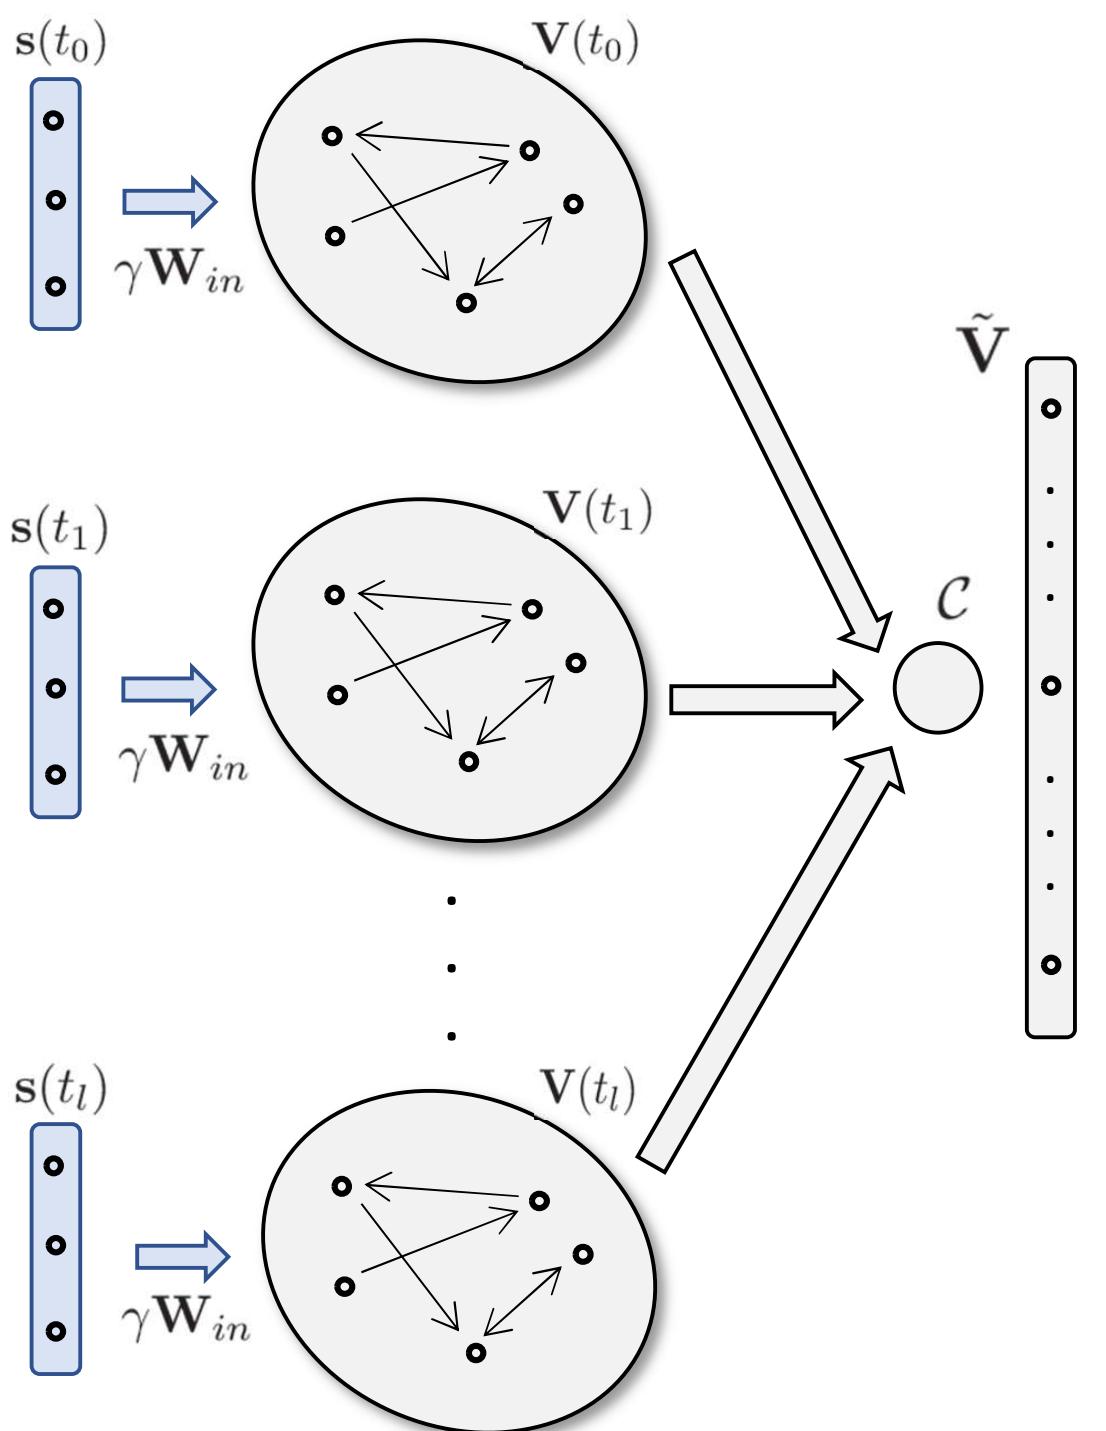
Define the initial sparsity level by the distribution of the neuron activity.

Can be learned on-line.

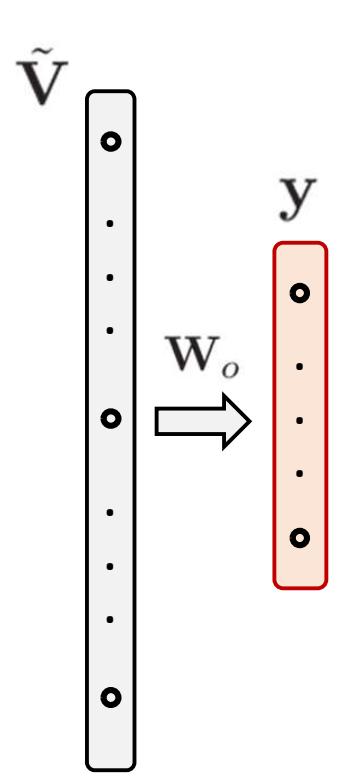
SPARCE SCHEME



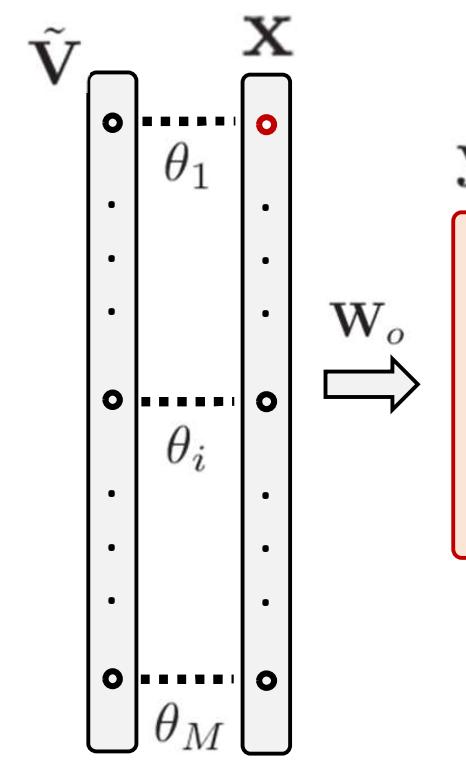
SPARCE SCHEME



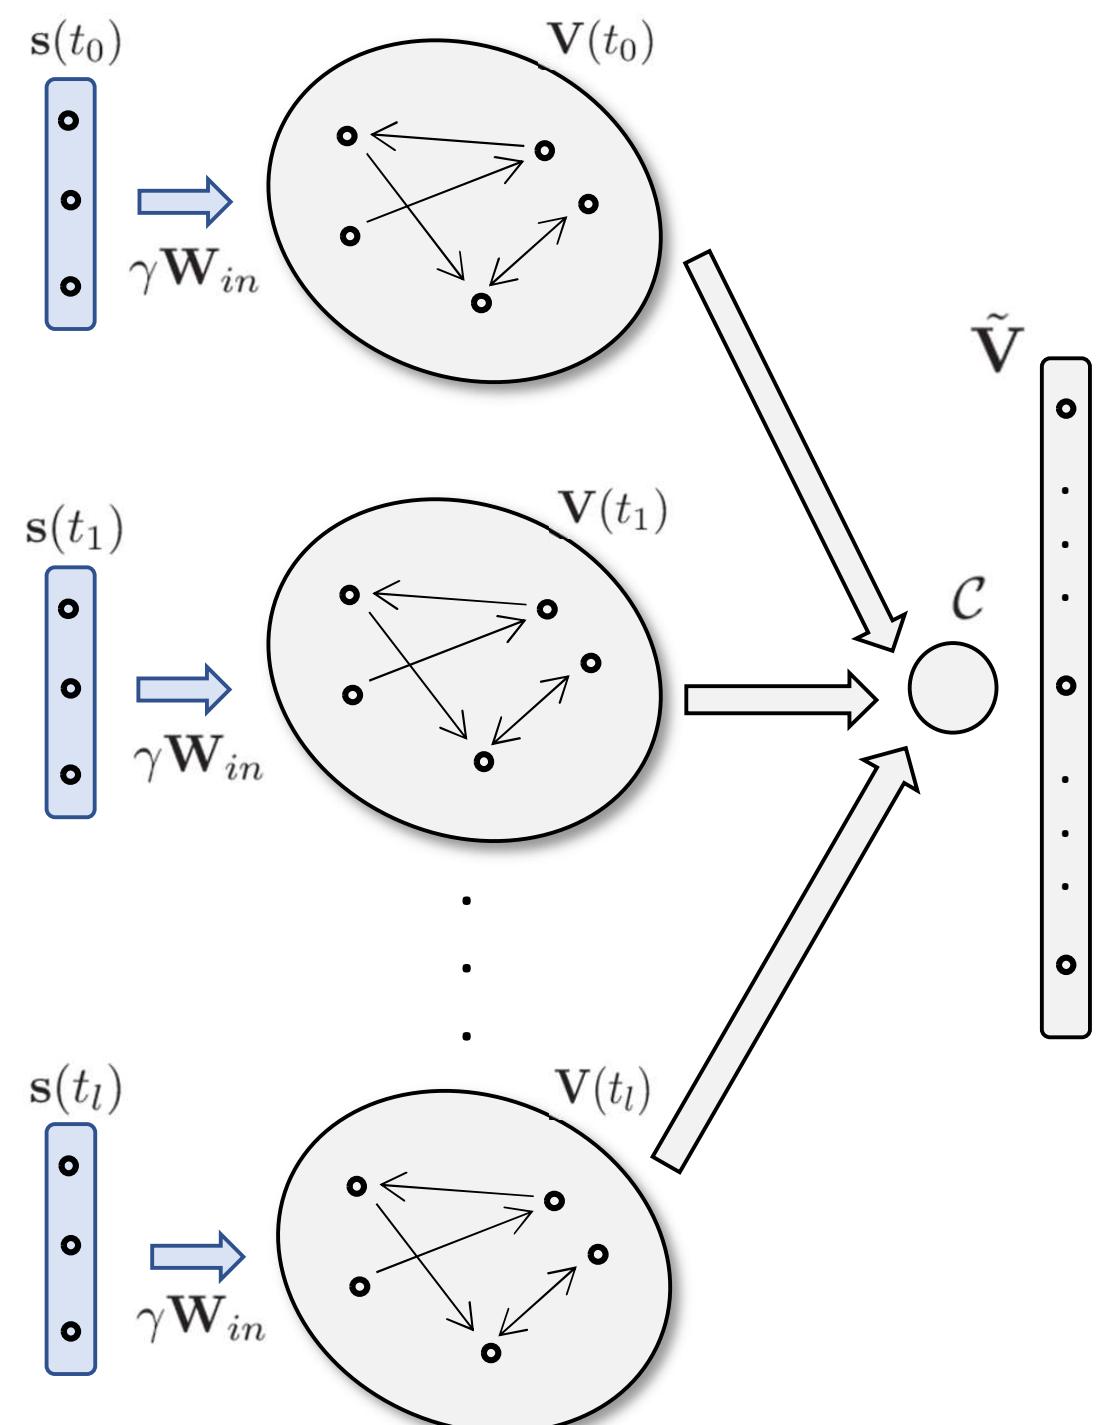
ESN



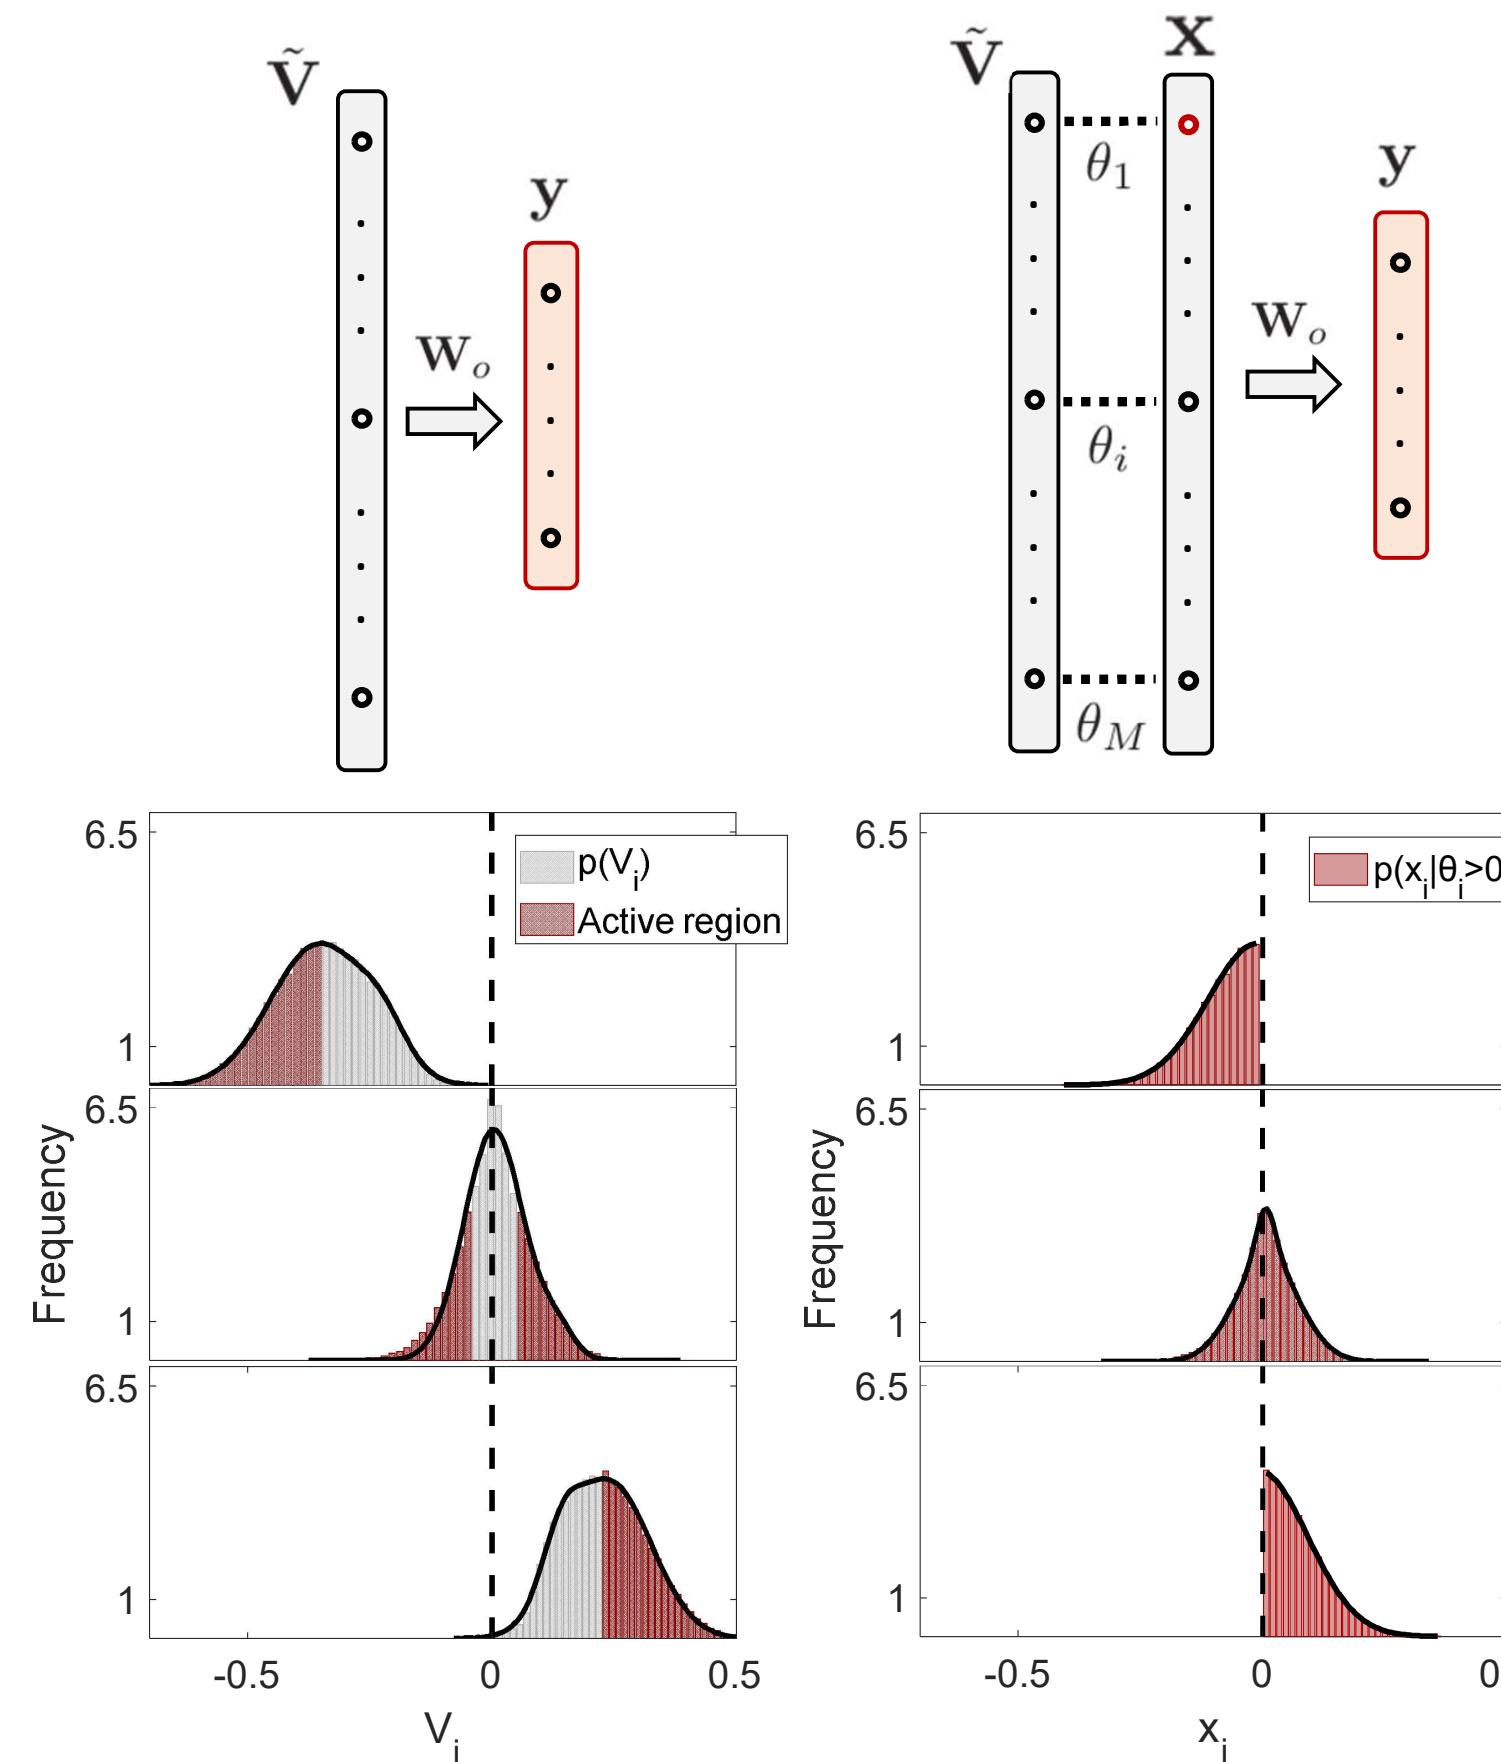
SpaRCe-ESN



SPARCE SCHEME

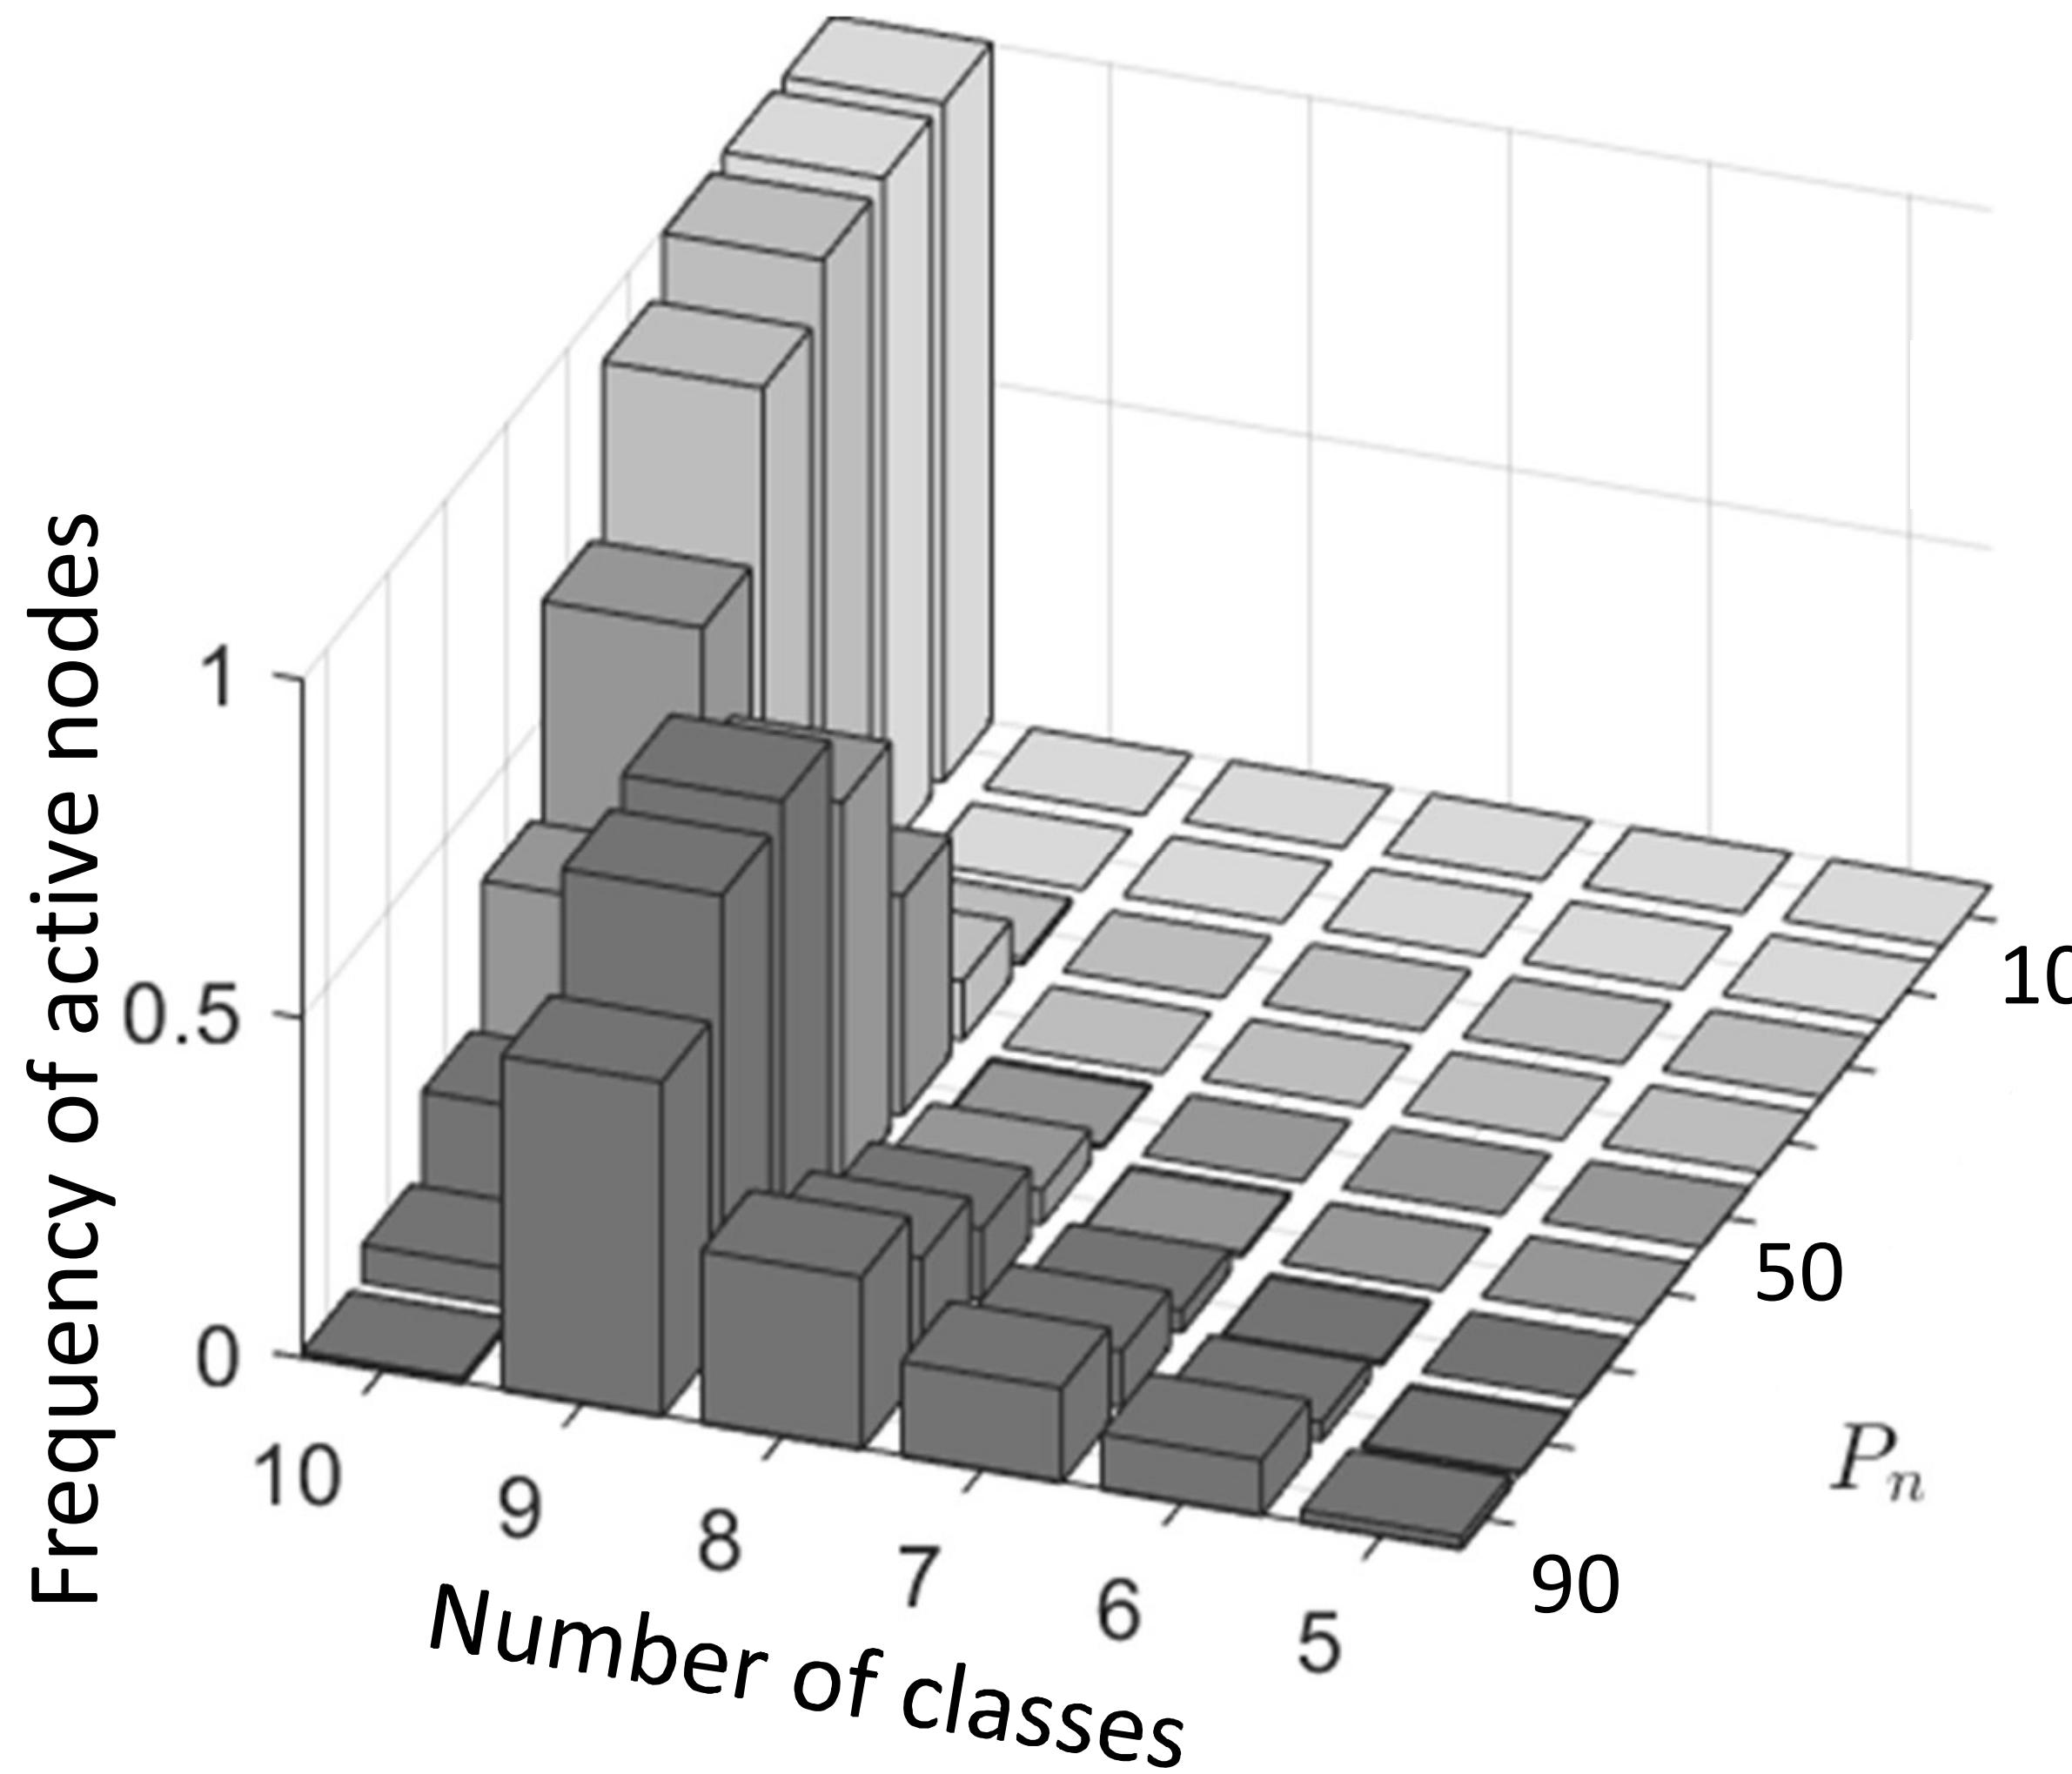


ESN



$$x_i(t) = \text{sign}(V_i(t)) \text{relu}\left[|V_i(t)| - \theta_i \right]$$

SPARSITY INTRODUCES NEURONAL SPECIALISATION



Example: column by column MNIST

7 9

10

50

P_n

ANALYSIS OF THRESHOLD LEARNING

$$\Delta\theta_i = -\eta \frac{\partial E}{\partial \theta_i} = -\eta \sum_j [y_j^t - y_j] W_{ji}^{out} sign(x_i)$$

ANALYSIS OF THRESHOLD LEARNING

$$\Delta\theta_i = -\eta \frac{\partial E}{\partial \theta_i} = -\eta \sum_j [y_j^t - y_j] W_{ji}^{out} sign(x_i)$$

$$= -\eta \sum_j [y_j^t - \sum_l W_{jl}^{out} x_l] W_{ji}^{out} sign(x_i)$$

ANALYSIS OF THRESHOLD LEARNING

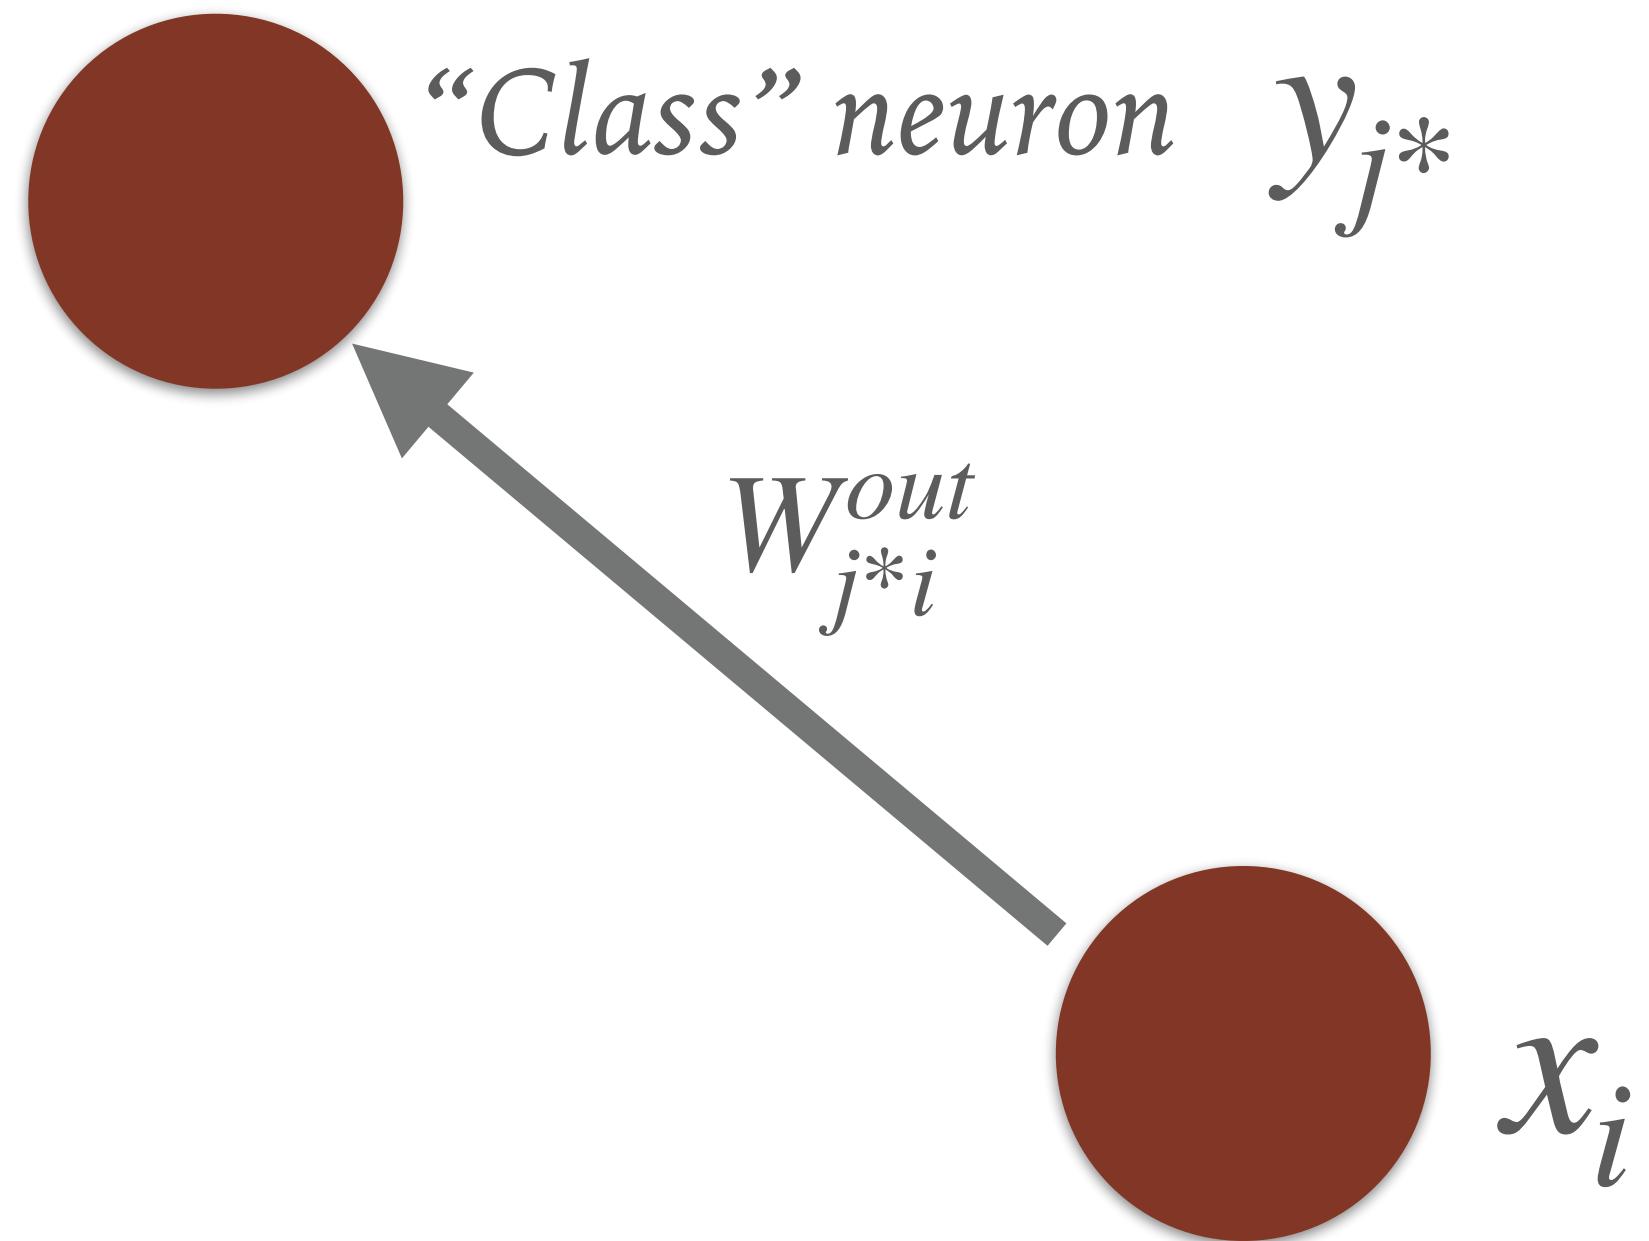
$$\begin{aligned}\Delta\theta_i &= -\eta \frac{\partial E}{\partial \theta_i} = -\eta \sum_j [y_j^t - y_j] W_{ji}^{out} \text{sign}(x_i) \\ &= -\eta \sum_j [y_j^t - \sum_l W_{jl}^{out} x_l] W_{ji}^{out} \text{sign}(x_i) \\ &= -\eta \sum_j y_j^t W_{ji}^{out} \text{sign}(x_i) + \eta \sum_j \sum_l W_{jl}^{out} x_l W_{ji}^{out} \text{sign}(x_i)\end{aligned}$$

ANALYSIS OF THRESHOLD LEARNING

$$\begin{aligned}\Delta\theta_i &= -\eta \frac{\partial E}{\partial \theta_i} = -\eta \sum_j [y_j^t - y_j] W_{ji}^{out} \text{sign}(x_i) \\ &= -\eta \sum_j [y_j^t - \sum_l W_{jl}^{out} x_l] W_{ji}^{out} \text{sign}(x_i) \\ &= -\eta \sum_j y_j^t W_{ji}^{out} \text{sign}(x_i) + \eta \sum_j \sum_l W_{jl}^{out} x_l W_{ji}^{out} \text{sign}(x_i) \\ \text{one-hot encoding} &\quad = -\eta W_{j^*i}^{out} \text{sign}(x_i) + \eta \sum_j \sum_l W_{jl}^{out} x_l W_{ji}^{out} \text{sign}(x_i)\end{aligned}$$

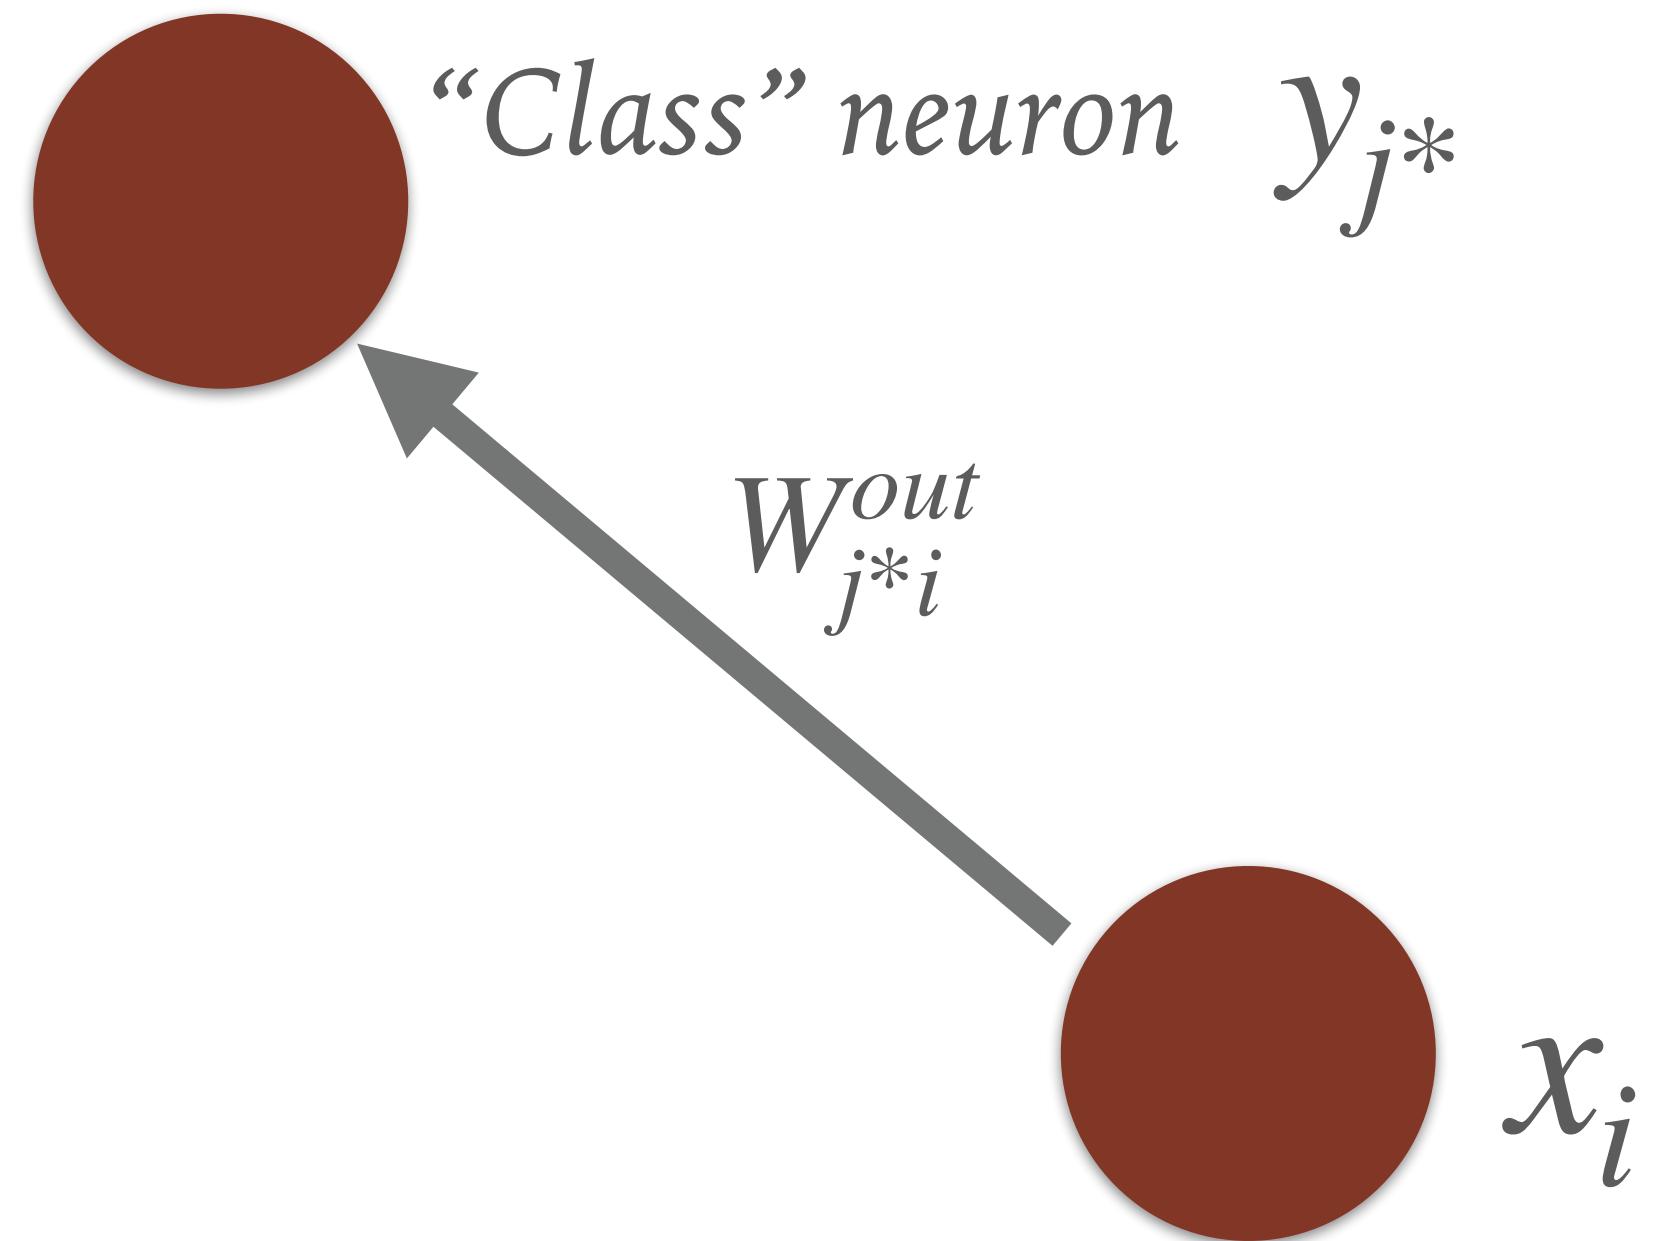
TWO FORCES DRIVE THRESHOLD LEARNING

$$\Delta\theta_i = -\eta_\theta \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \eta_\theta \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$



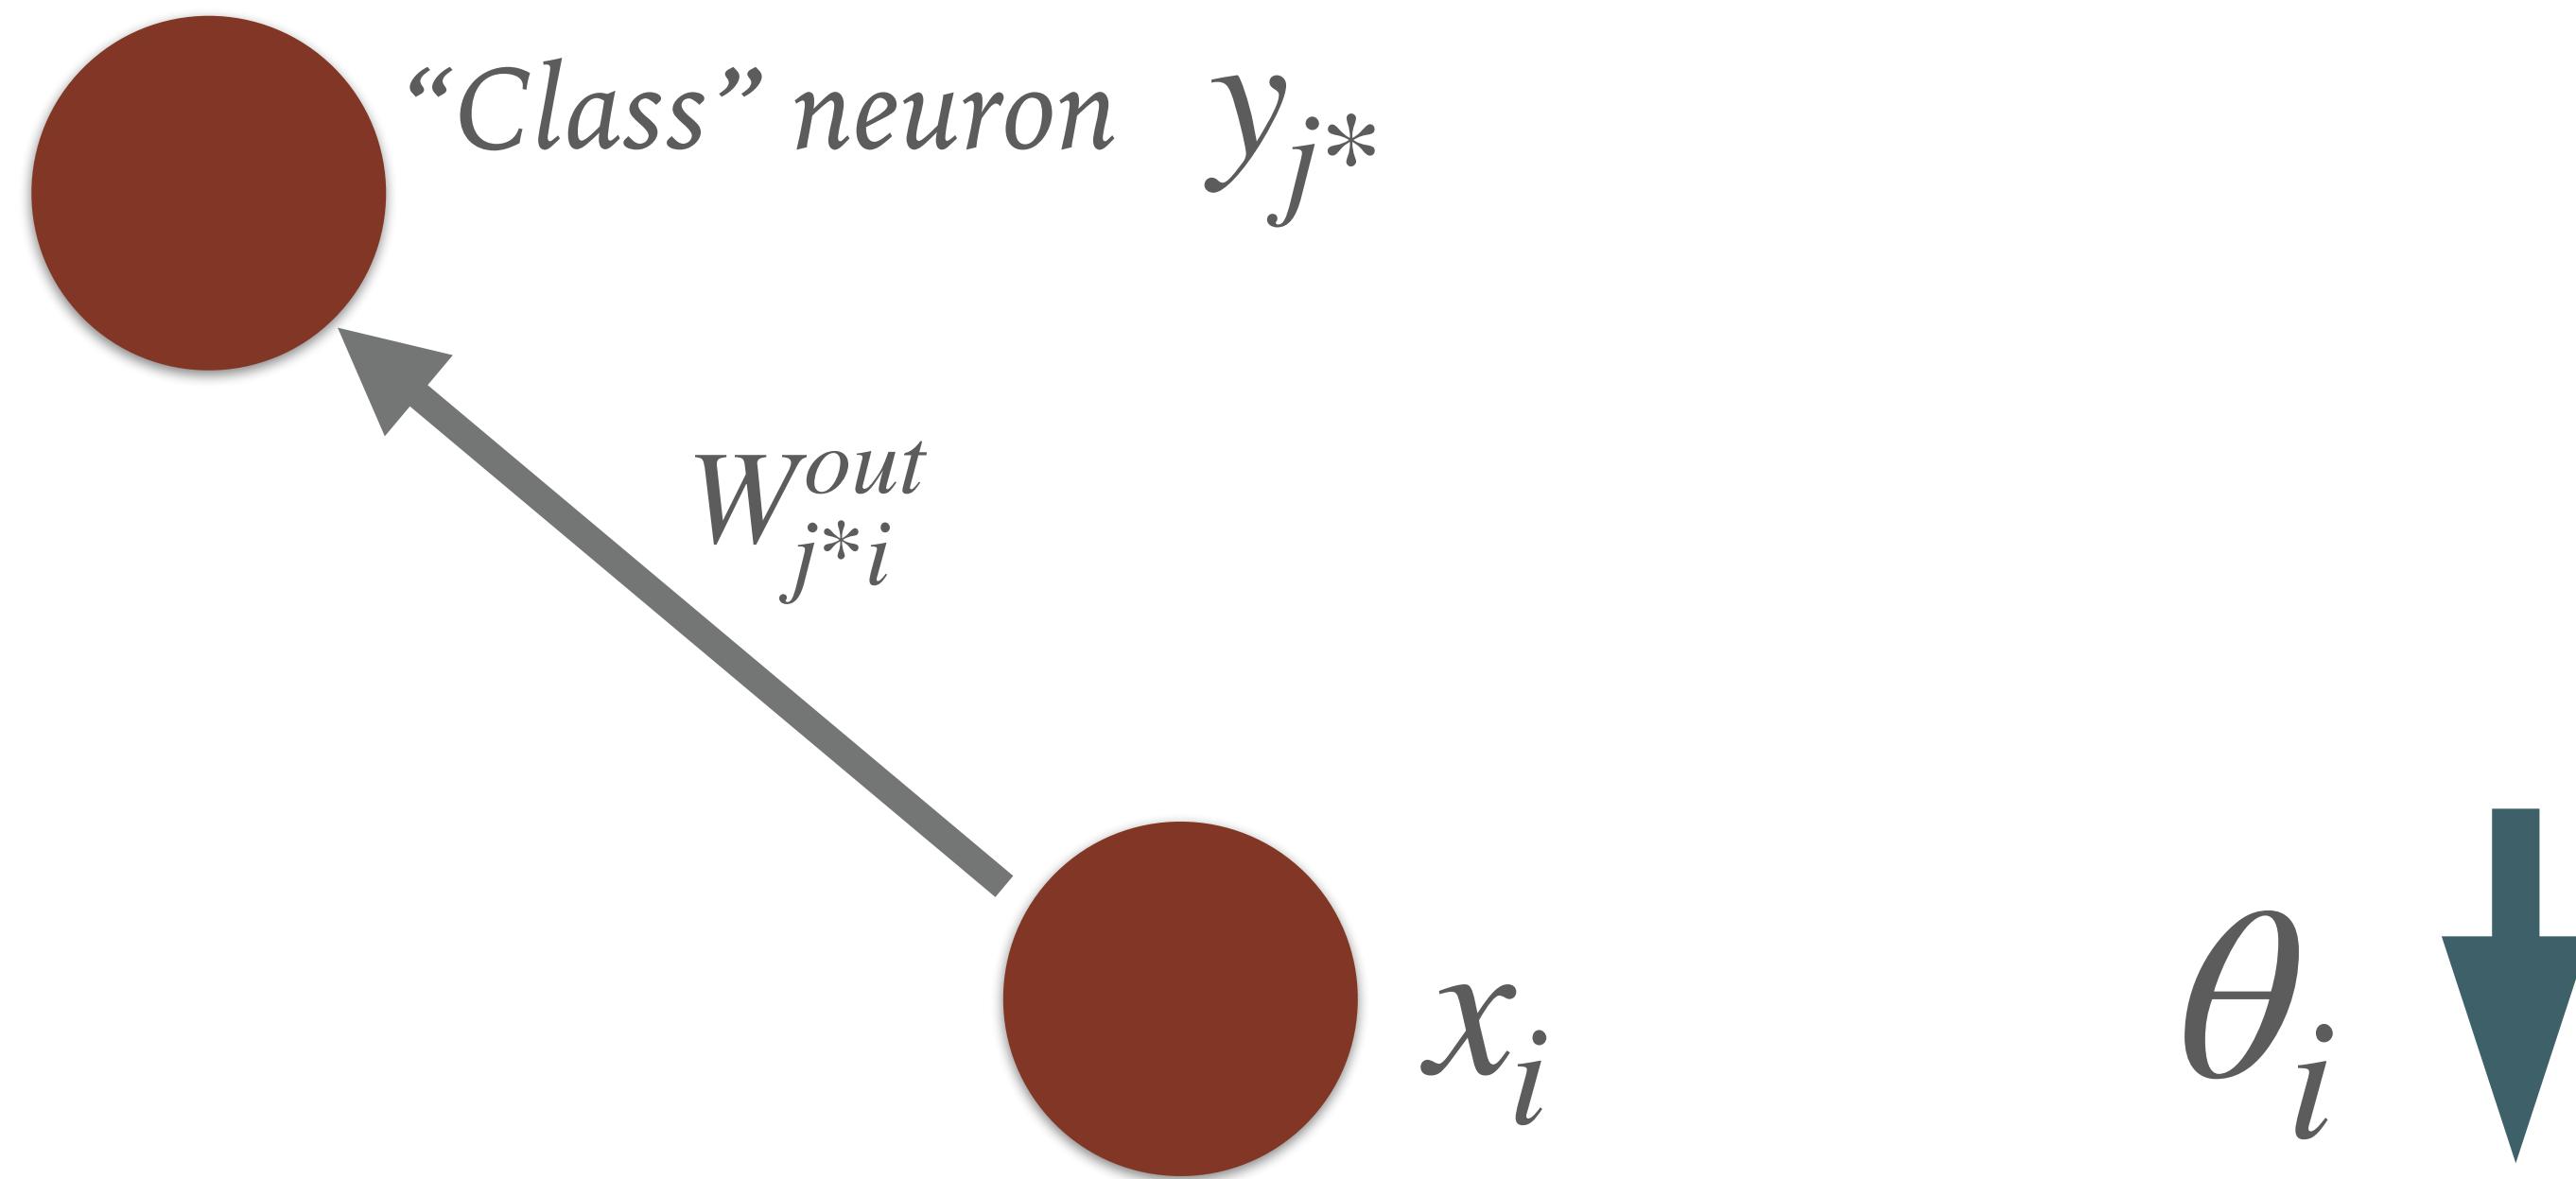
TWO FORCES DRIVE THRESHOLD LEARNING

$$\Delta\theta_i = -\eta_\theta \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \eta_\theta \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$



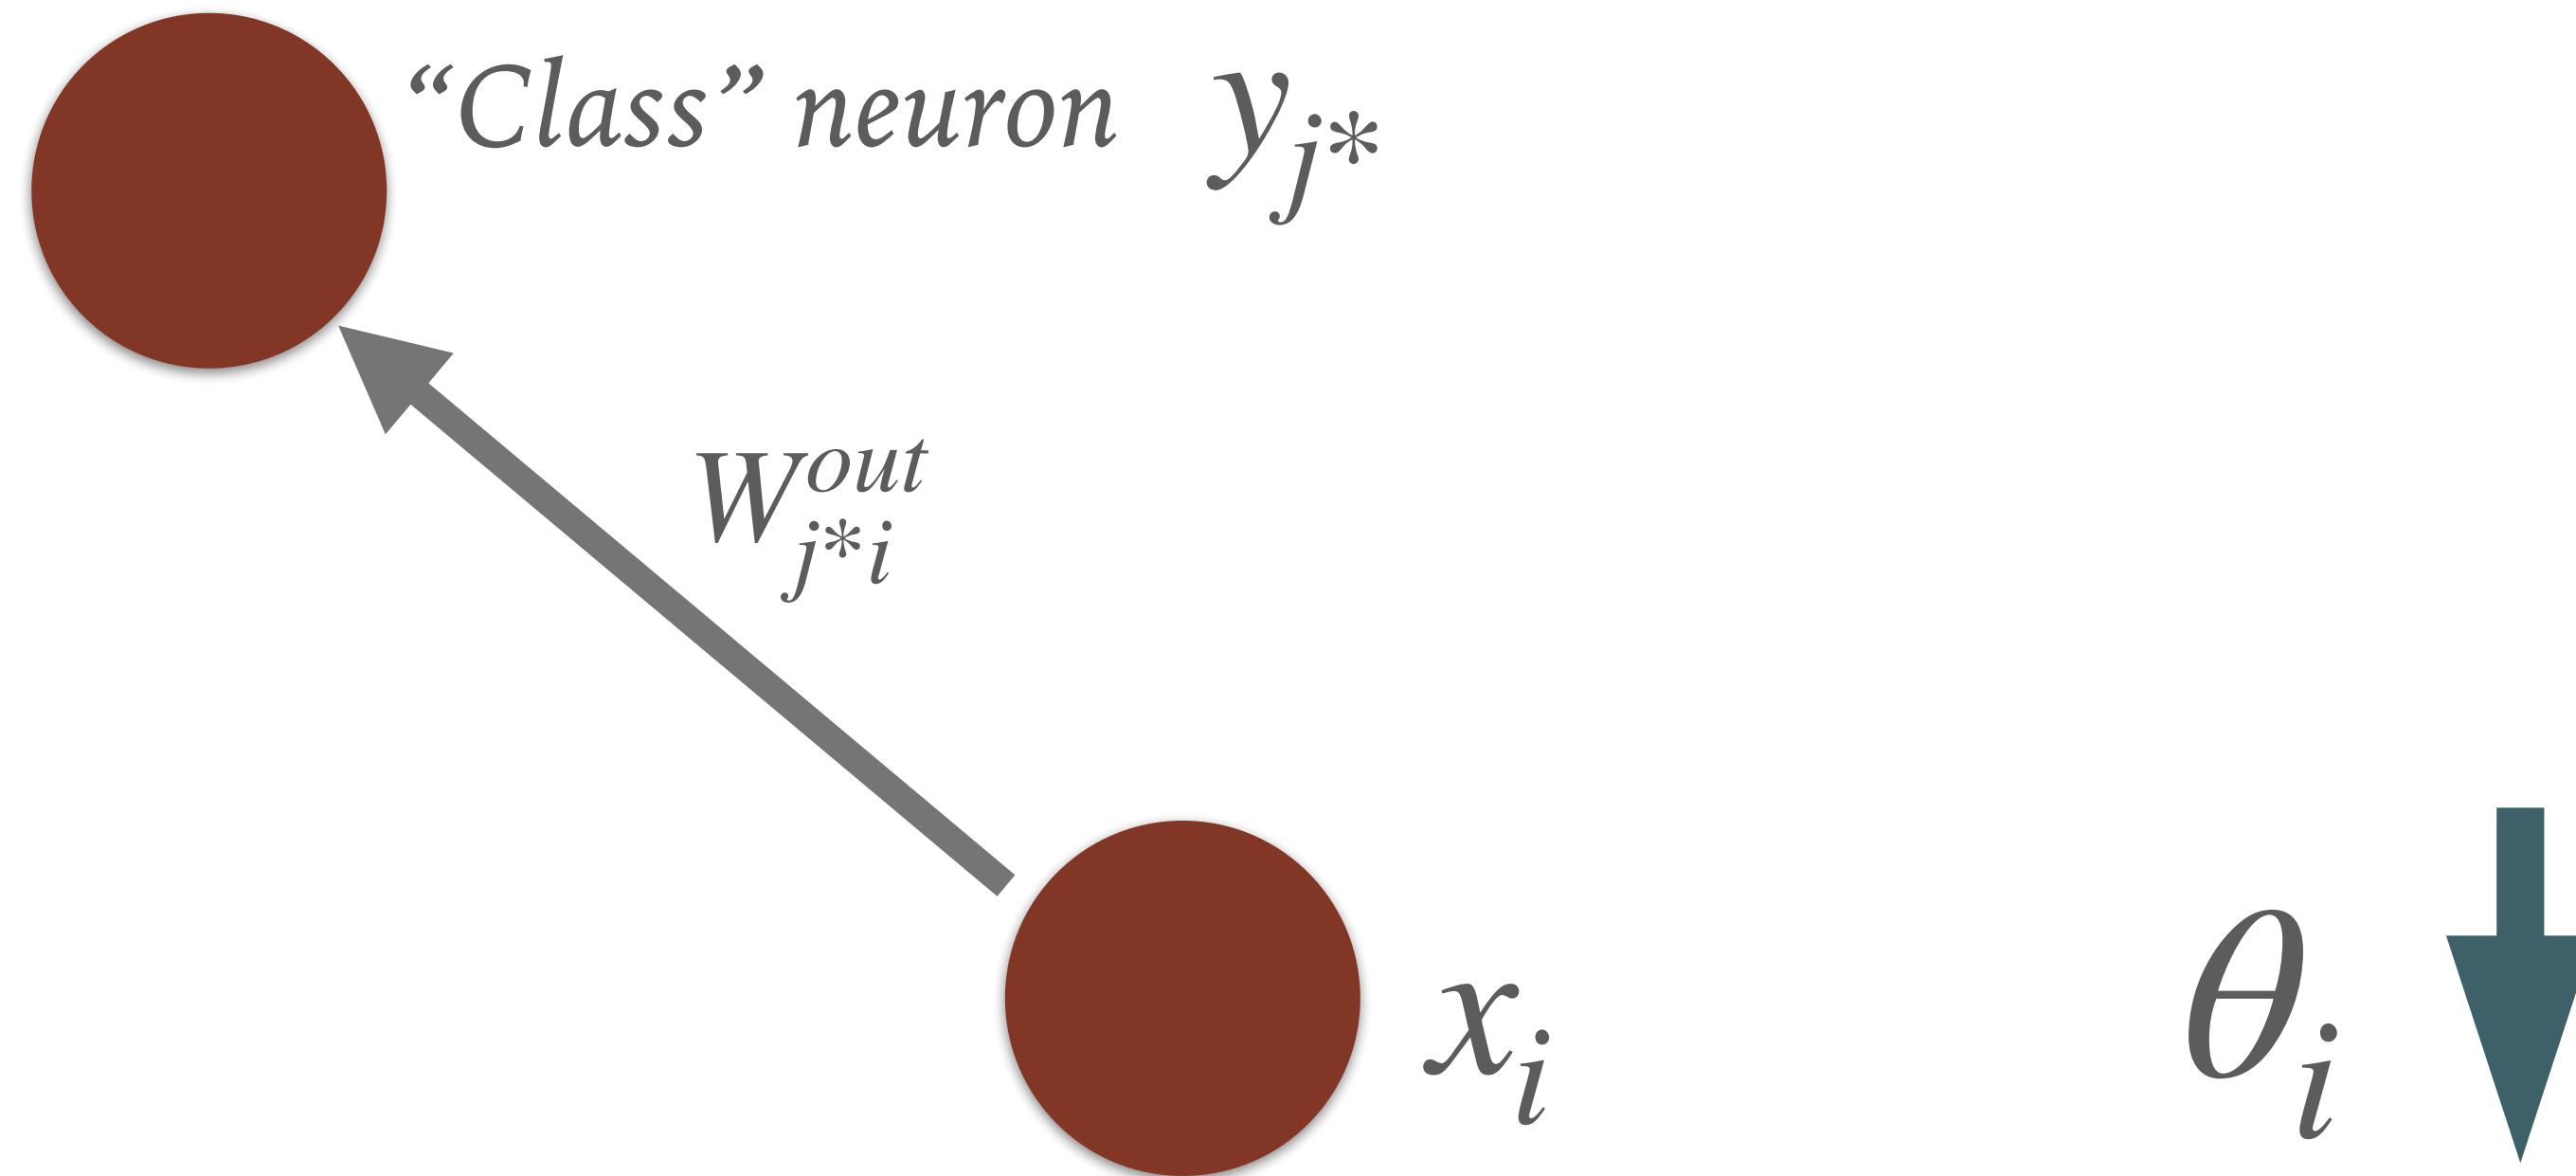
TWO FORCES DRIVE THRESHOLD LEARNING

$$\Delta\theta_i = -\eta_\theta \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \eta_\theta \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$



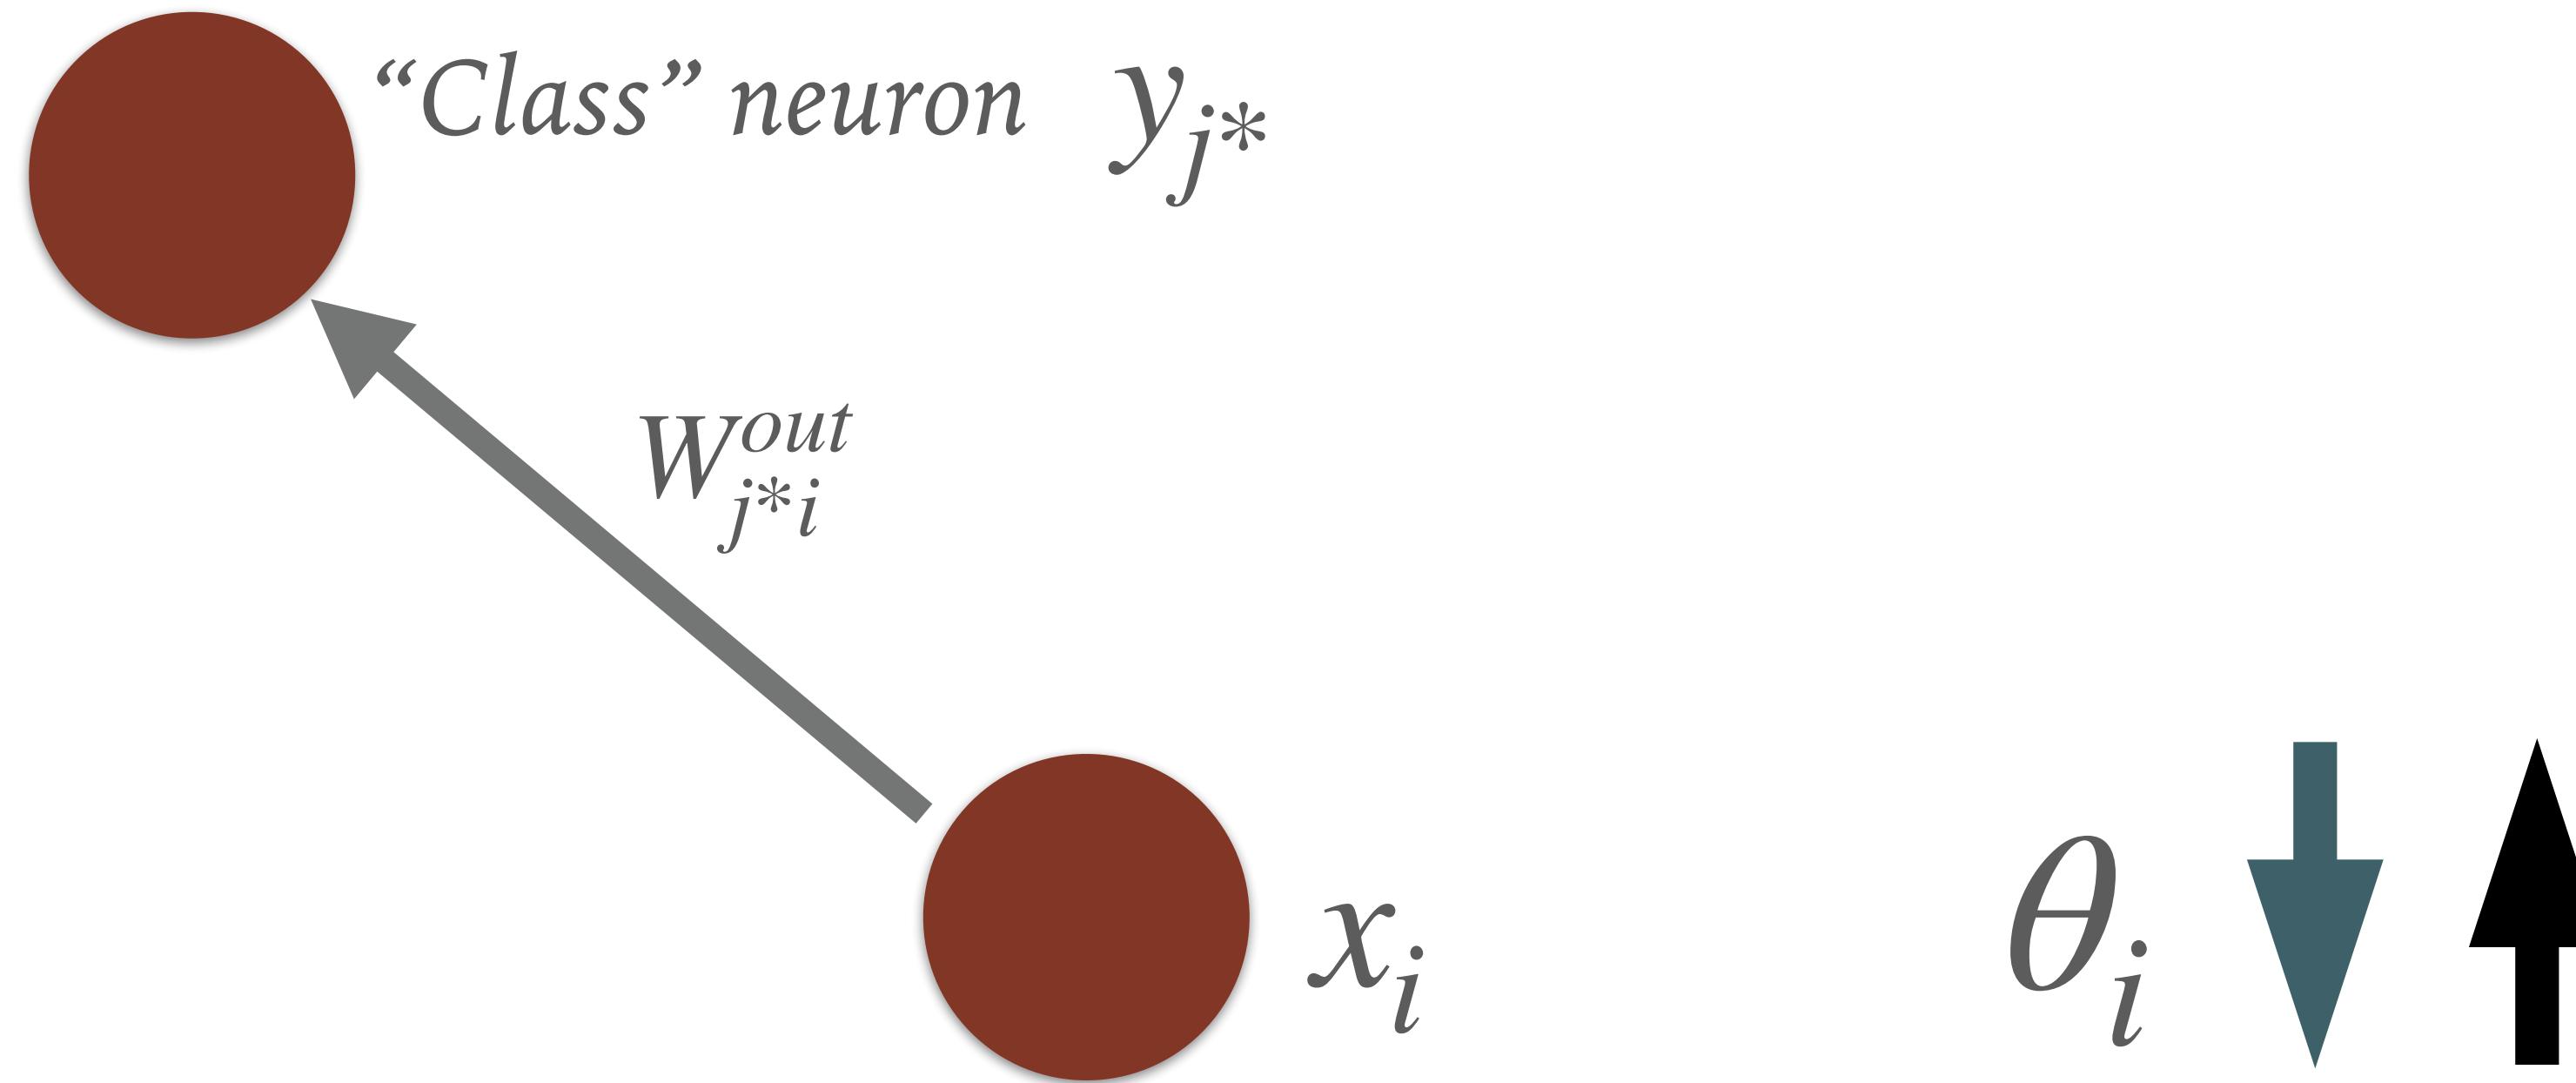
TWO FORCES DRIVE THRESHOLD LEARNING

$$\Delta\theta_i = -\eta_\theta \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \eta_\theta \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$



TWO FORCES DRIVE THRESHOLD LEARNING

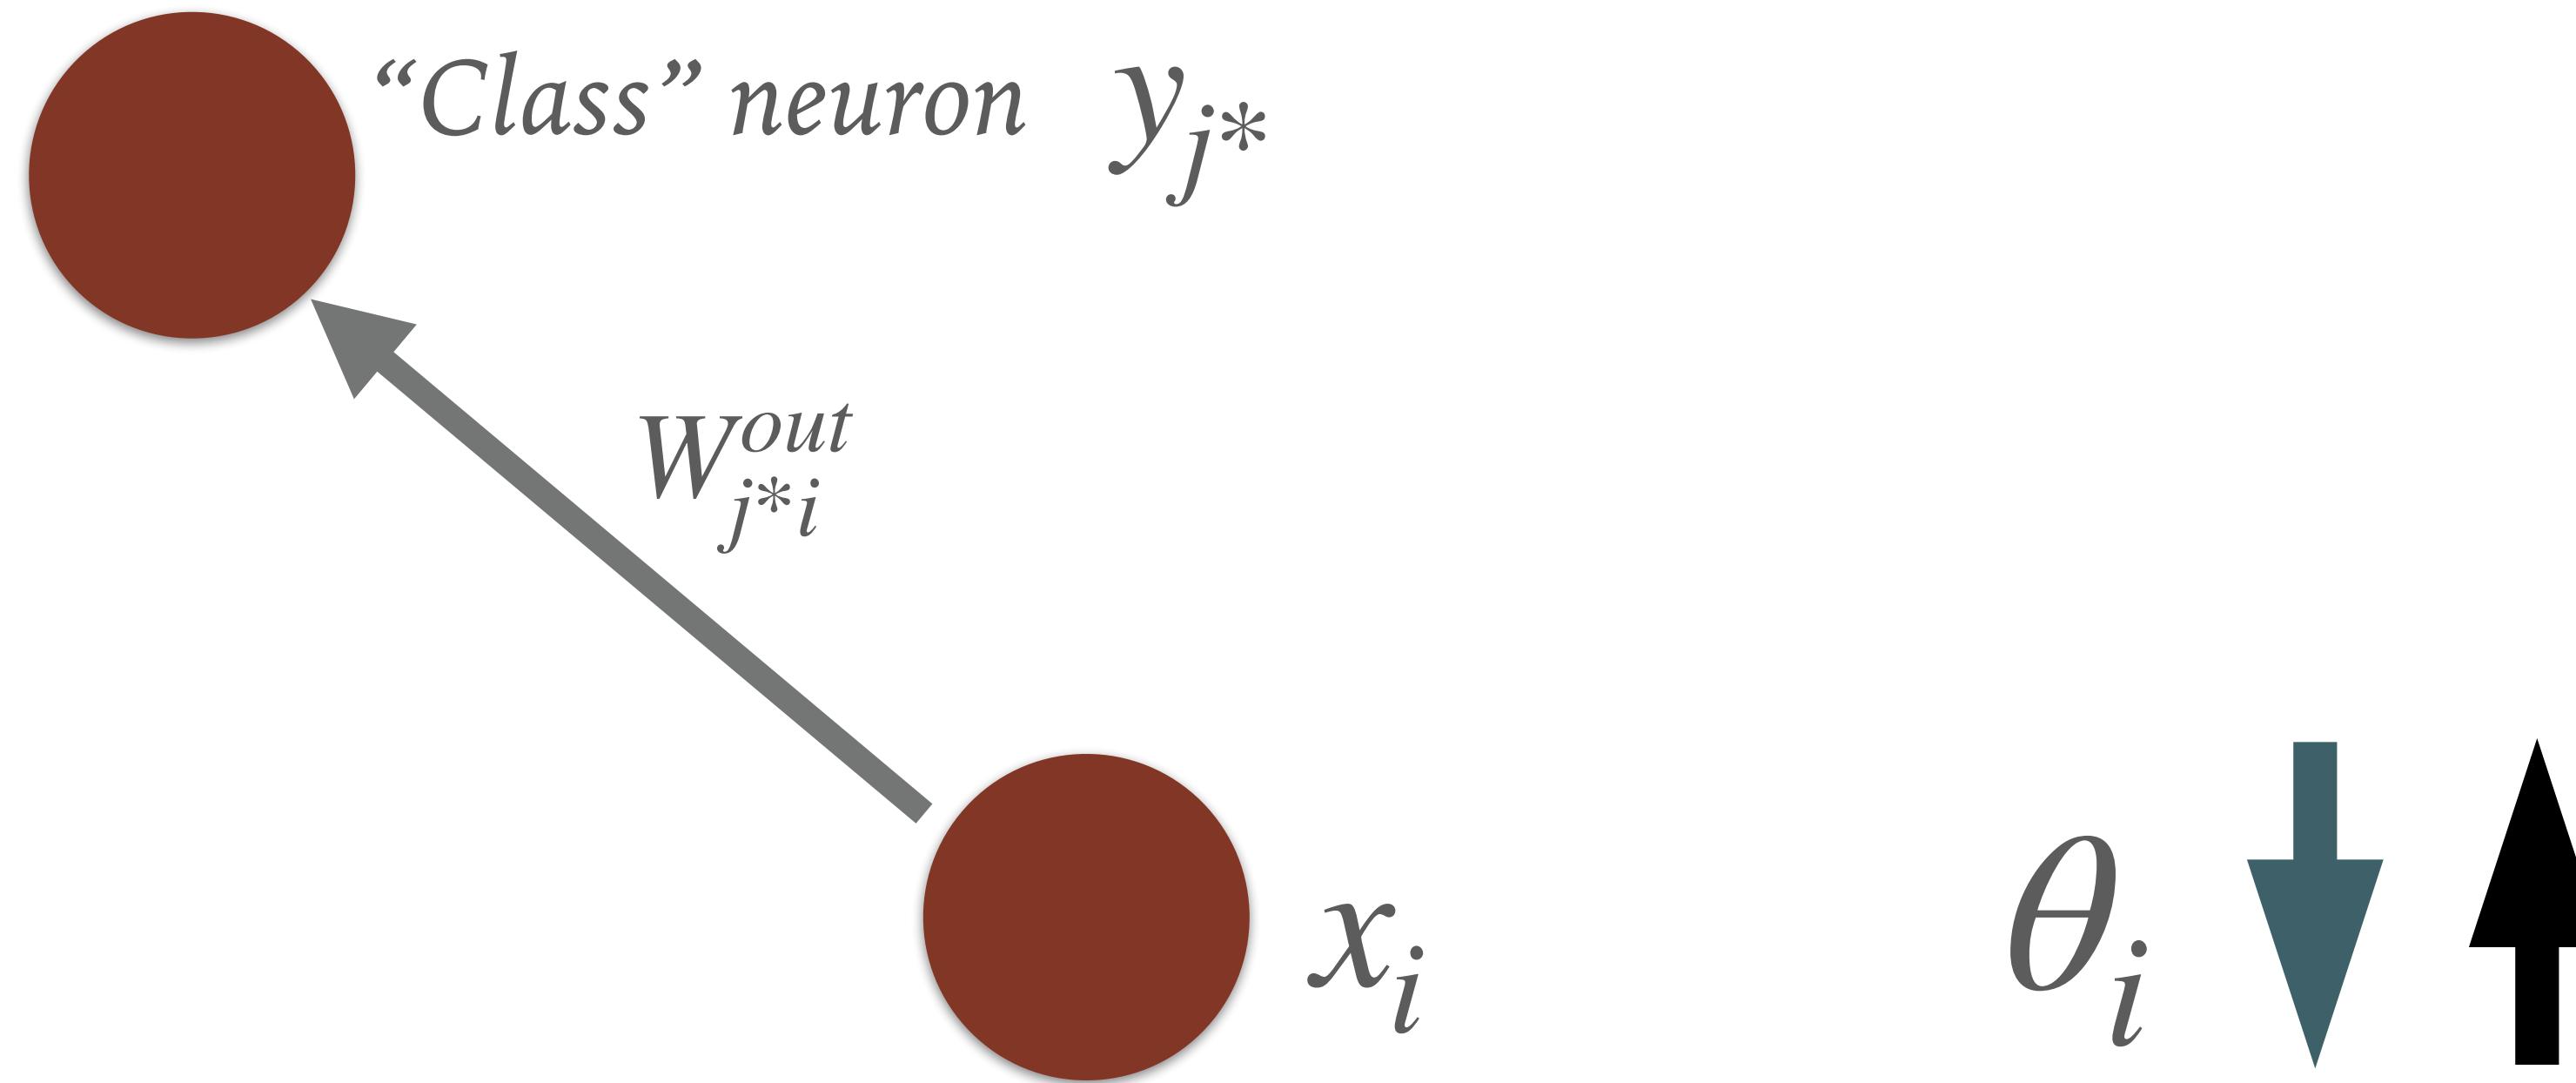
$$\Delta\theta_i = -\eta_\theta \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \eta_\theta \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$



TWO FORCES DRIVE THRESHOLD LEARNING

$$\Delta\theta_i = -\eta_\theta \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \eta_\theta \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$

reduces correlations



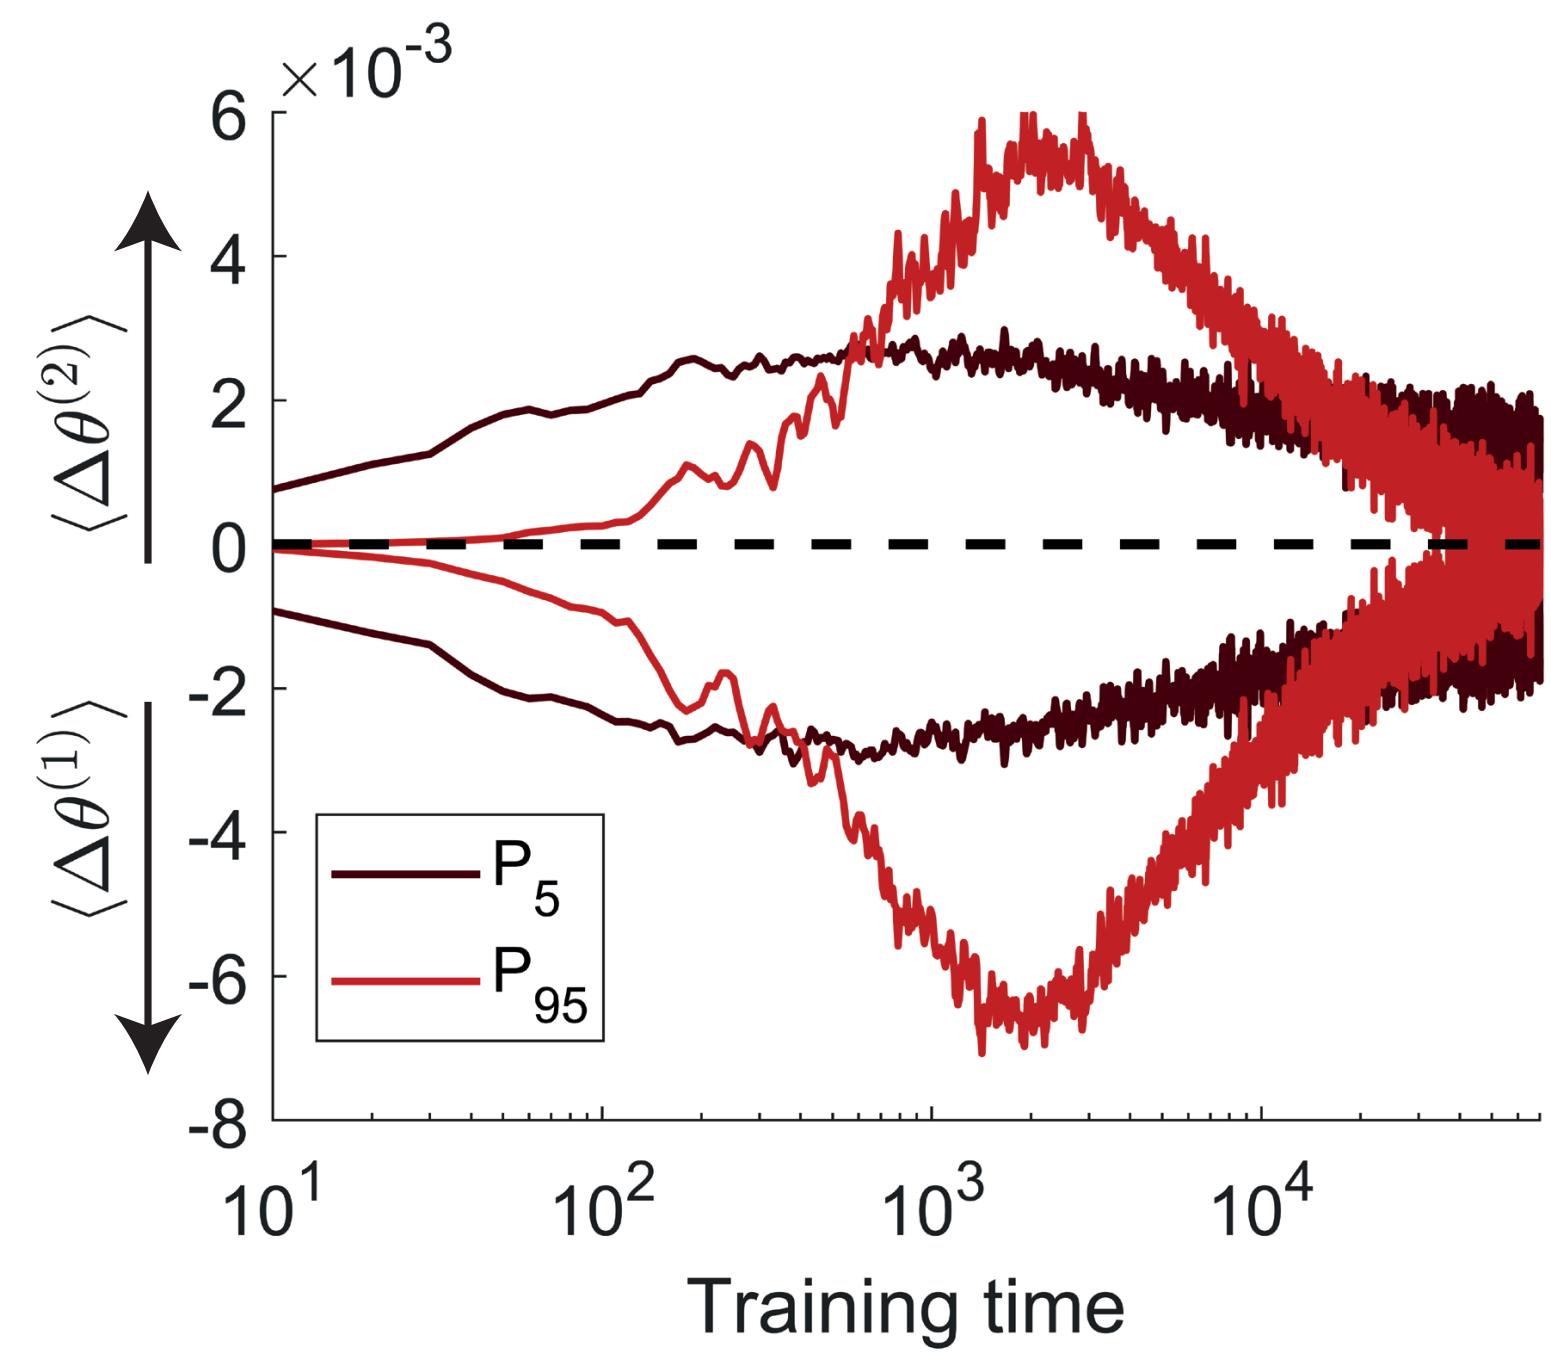
TWO FORCES DRIVE THRESHOLD LEARNING

$$\Delta\theta_i/\eta_\theta = -y_{j^*}^t \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$

$\Delta\theta^{(2)}$

$\Delta\theta^{(1)}$

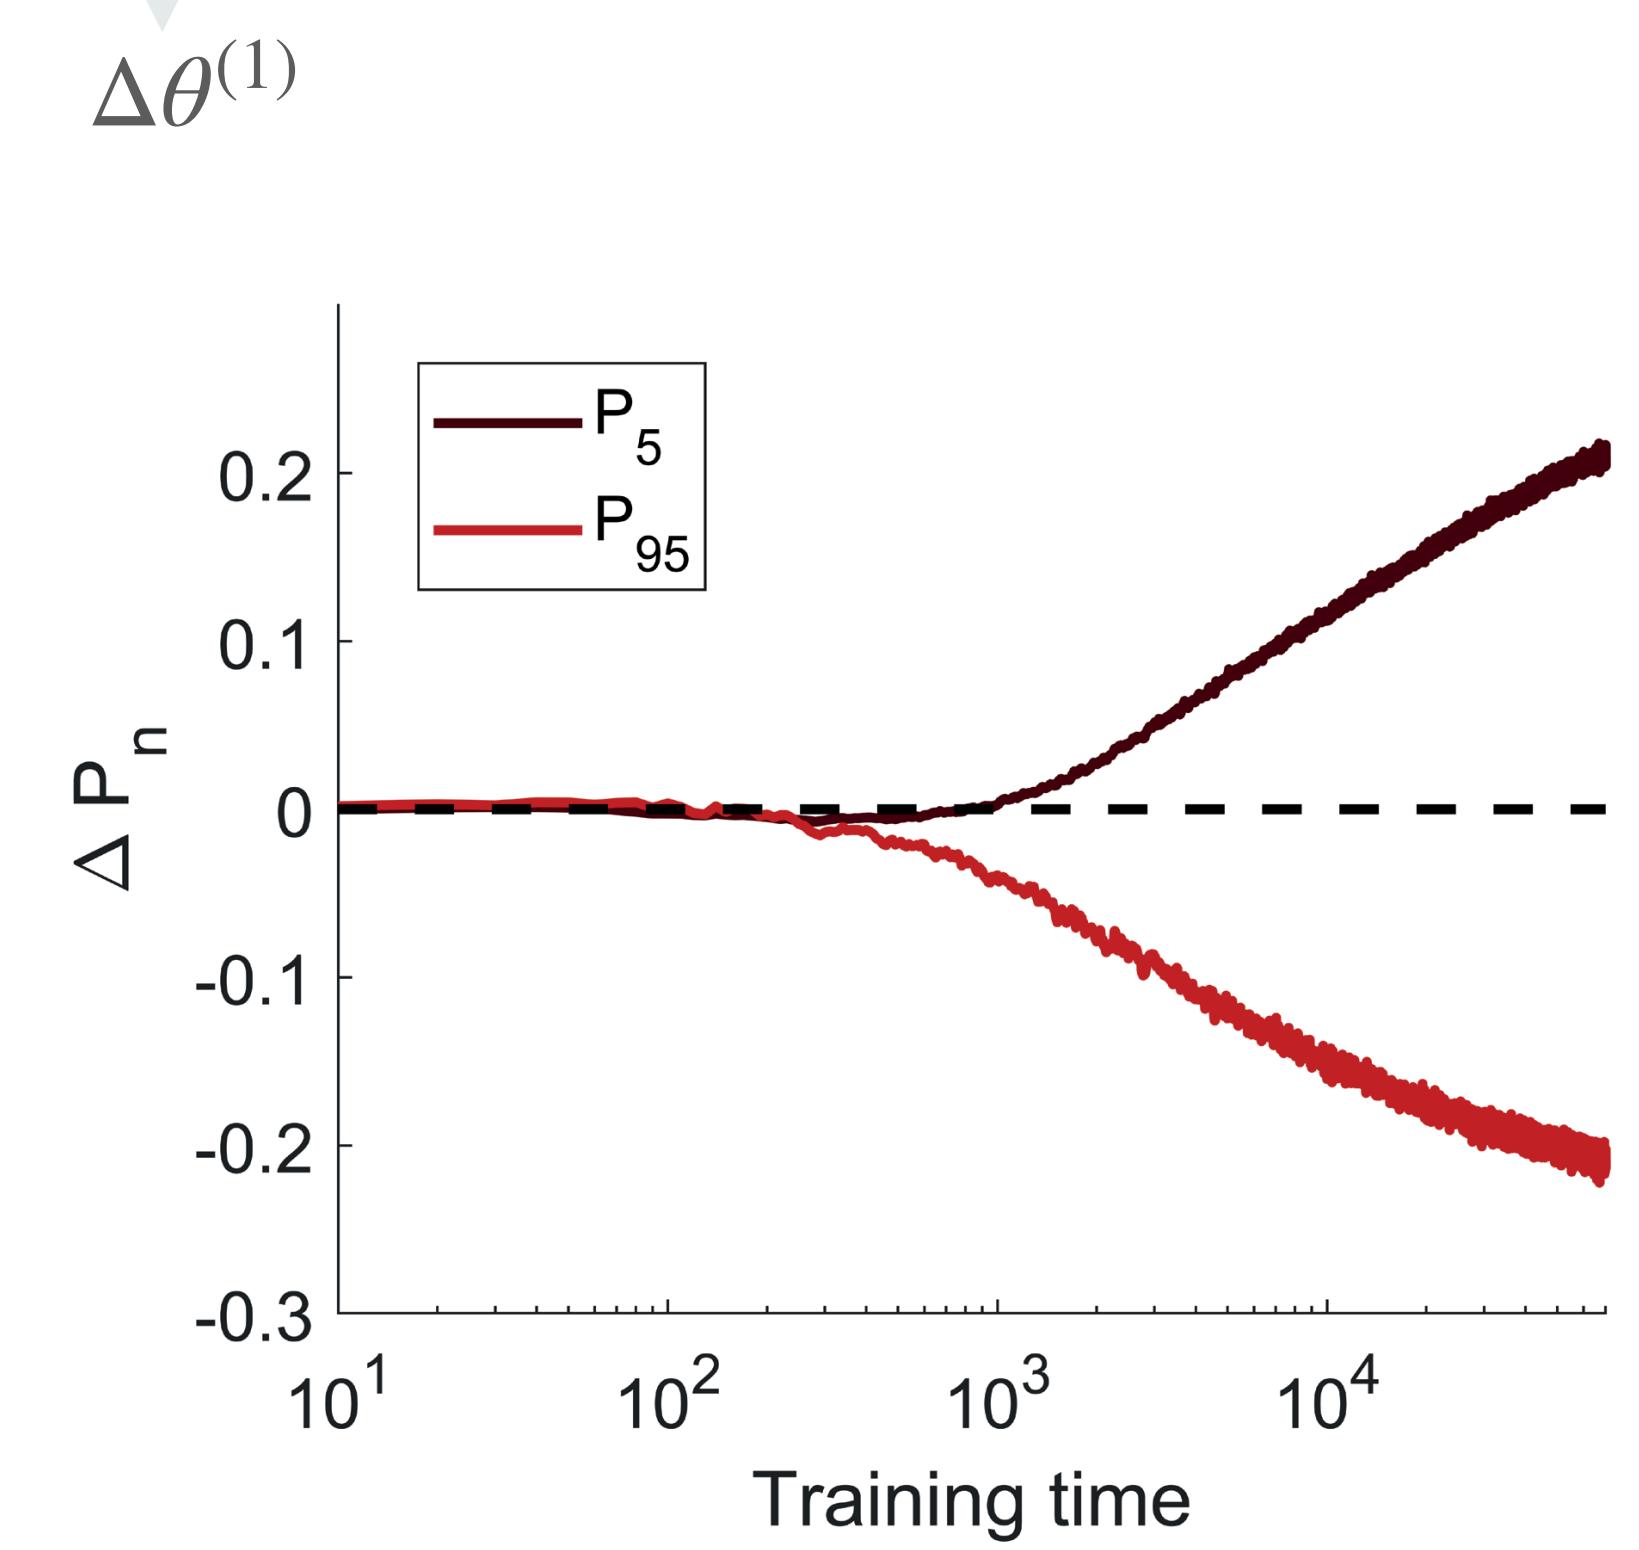
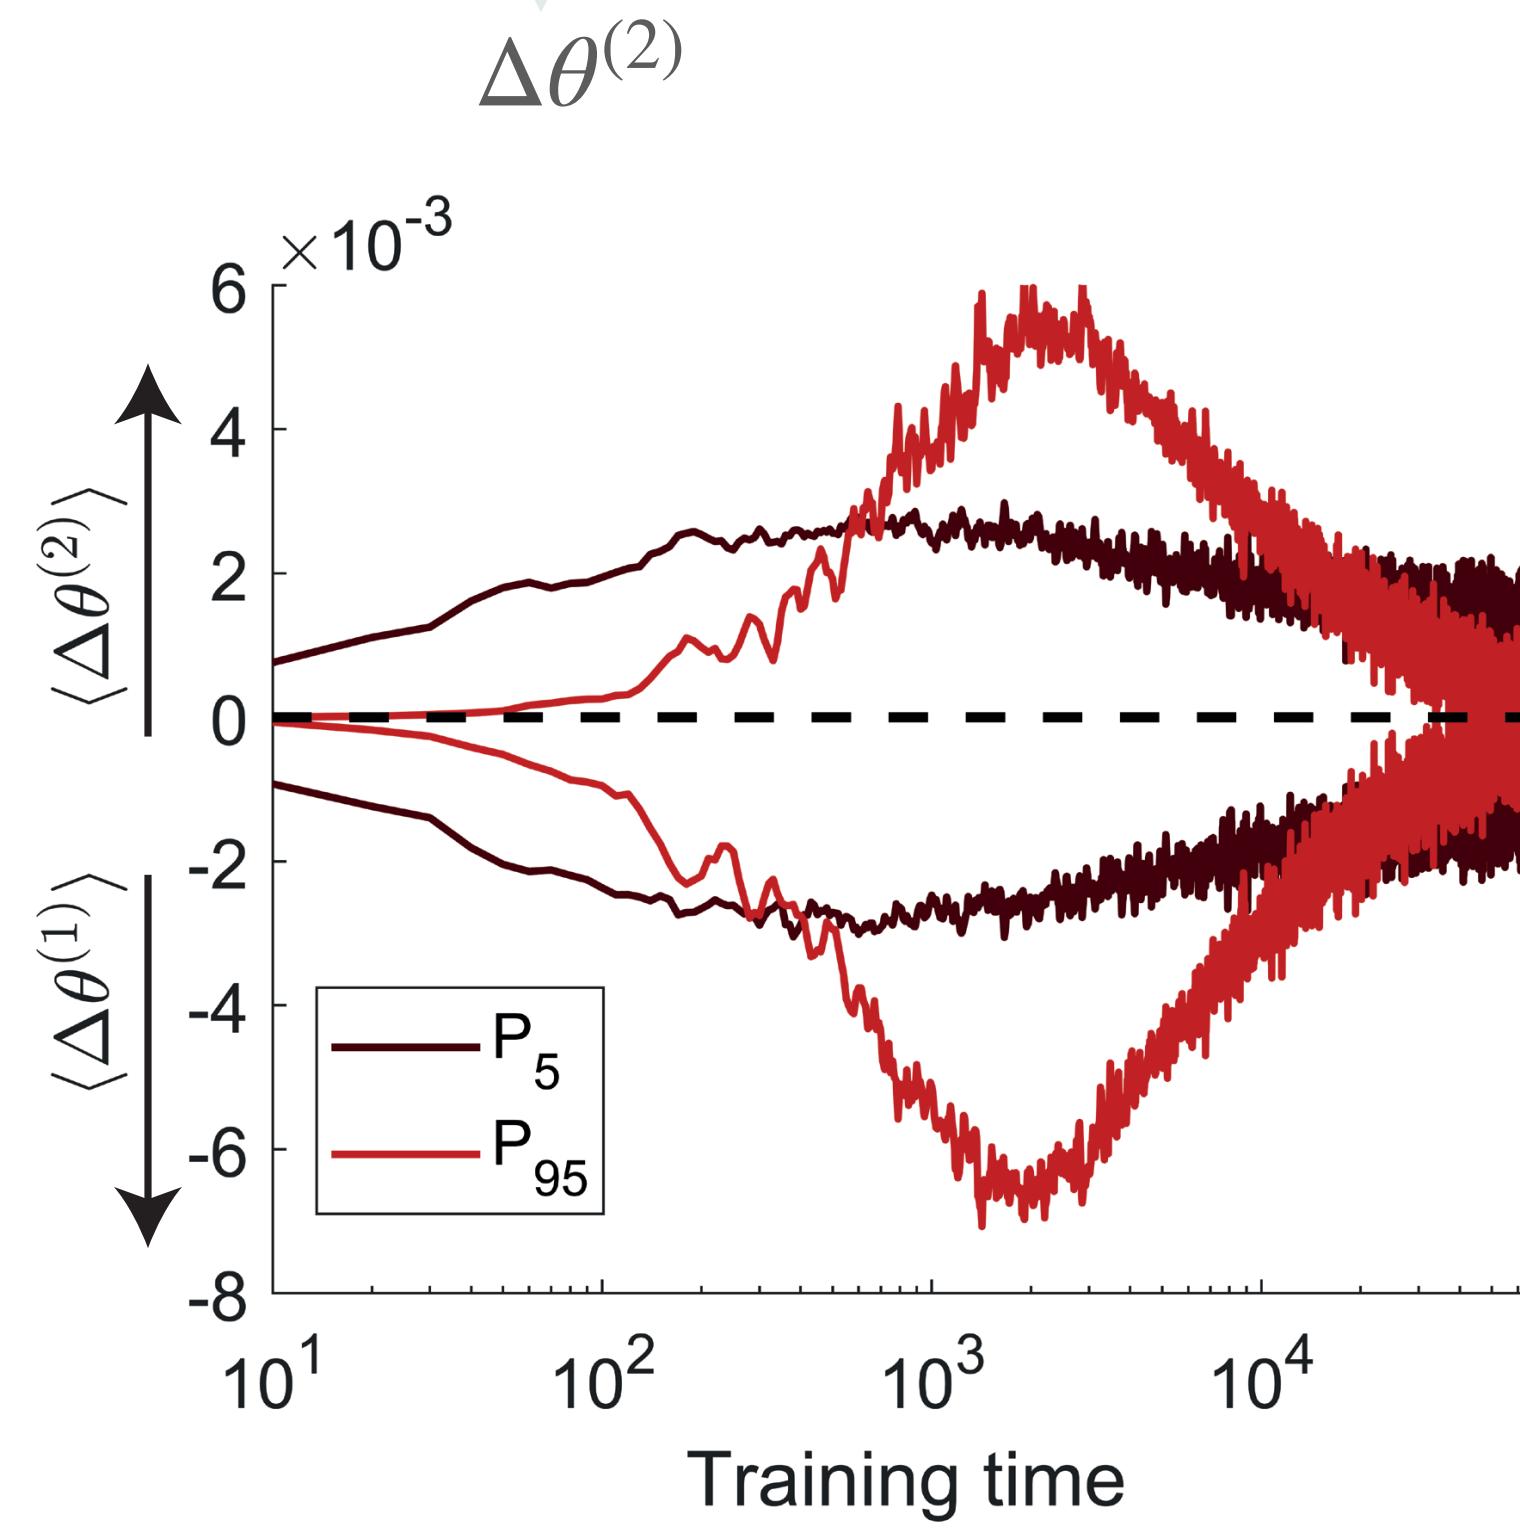
Example: column by column MNIST



7 9

TWO FORCES DRIVE THRESHOLD LEARNING

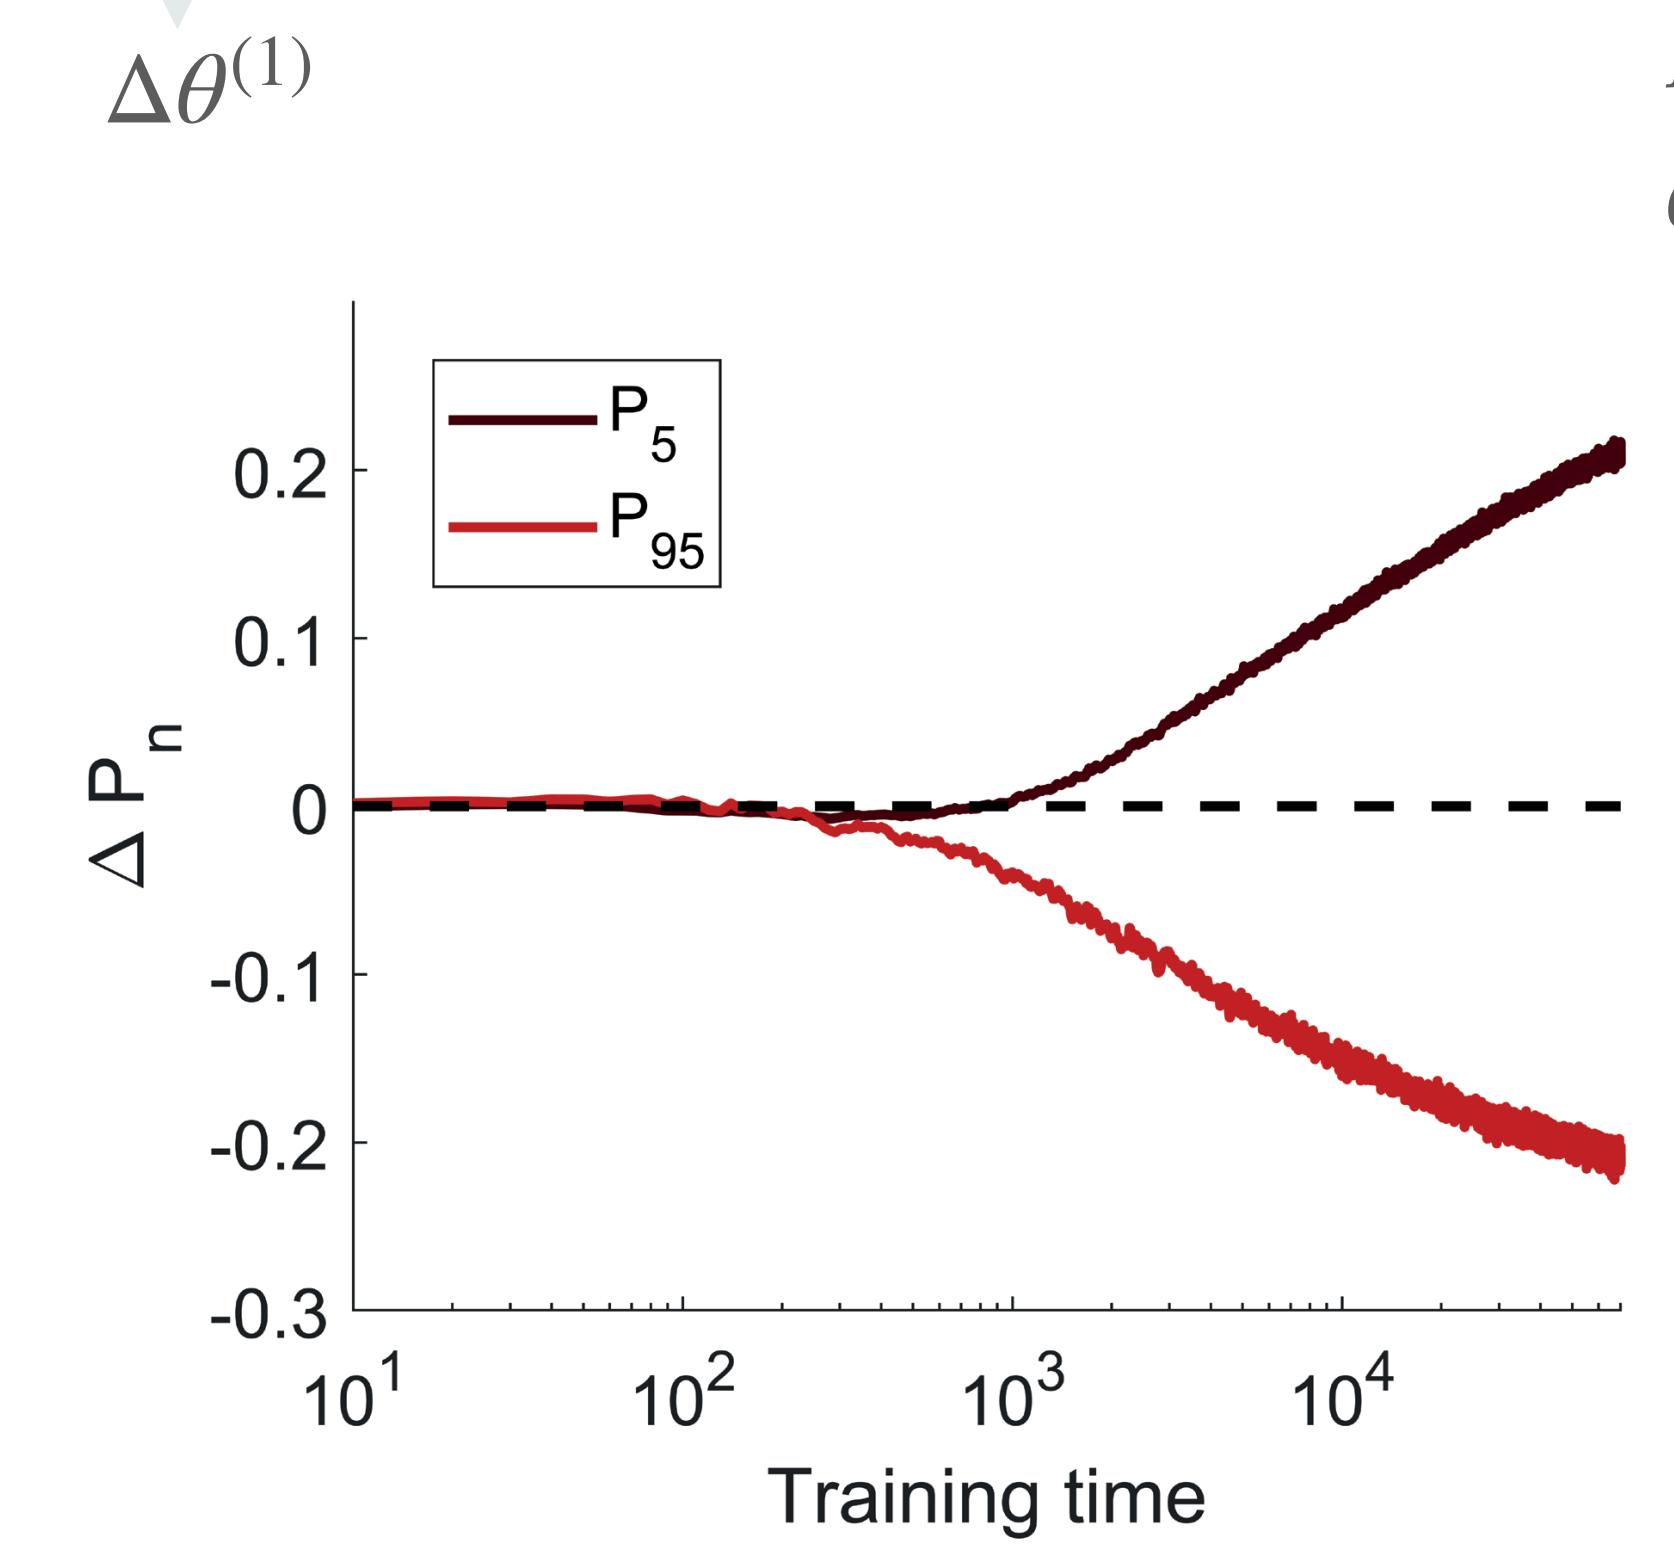
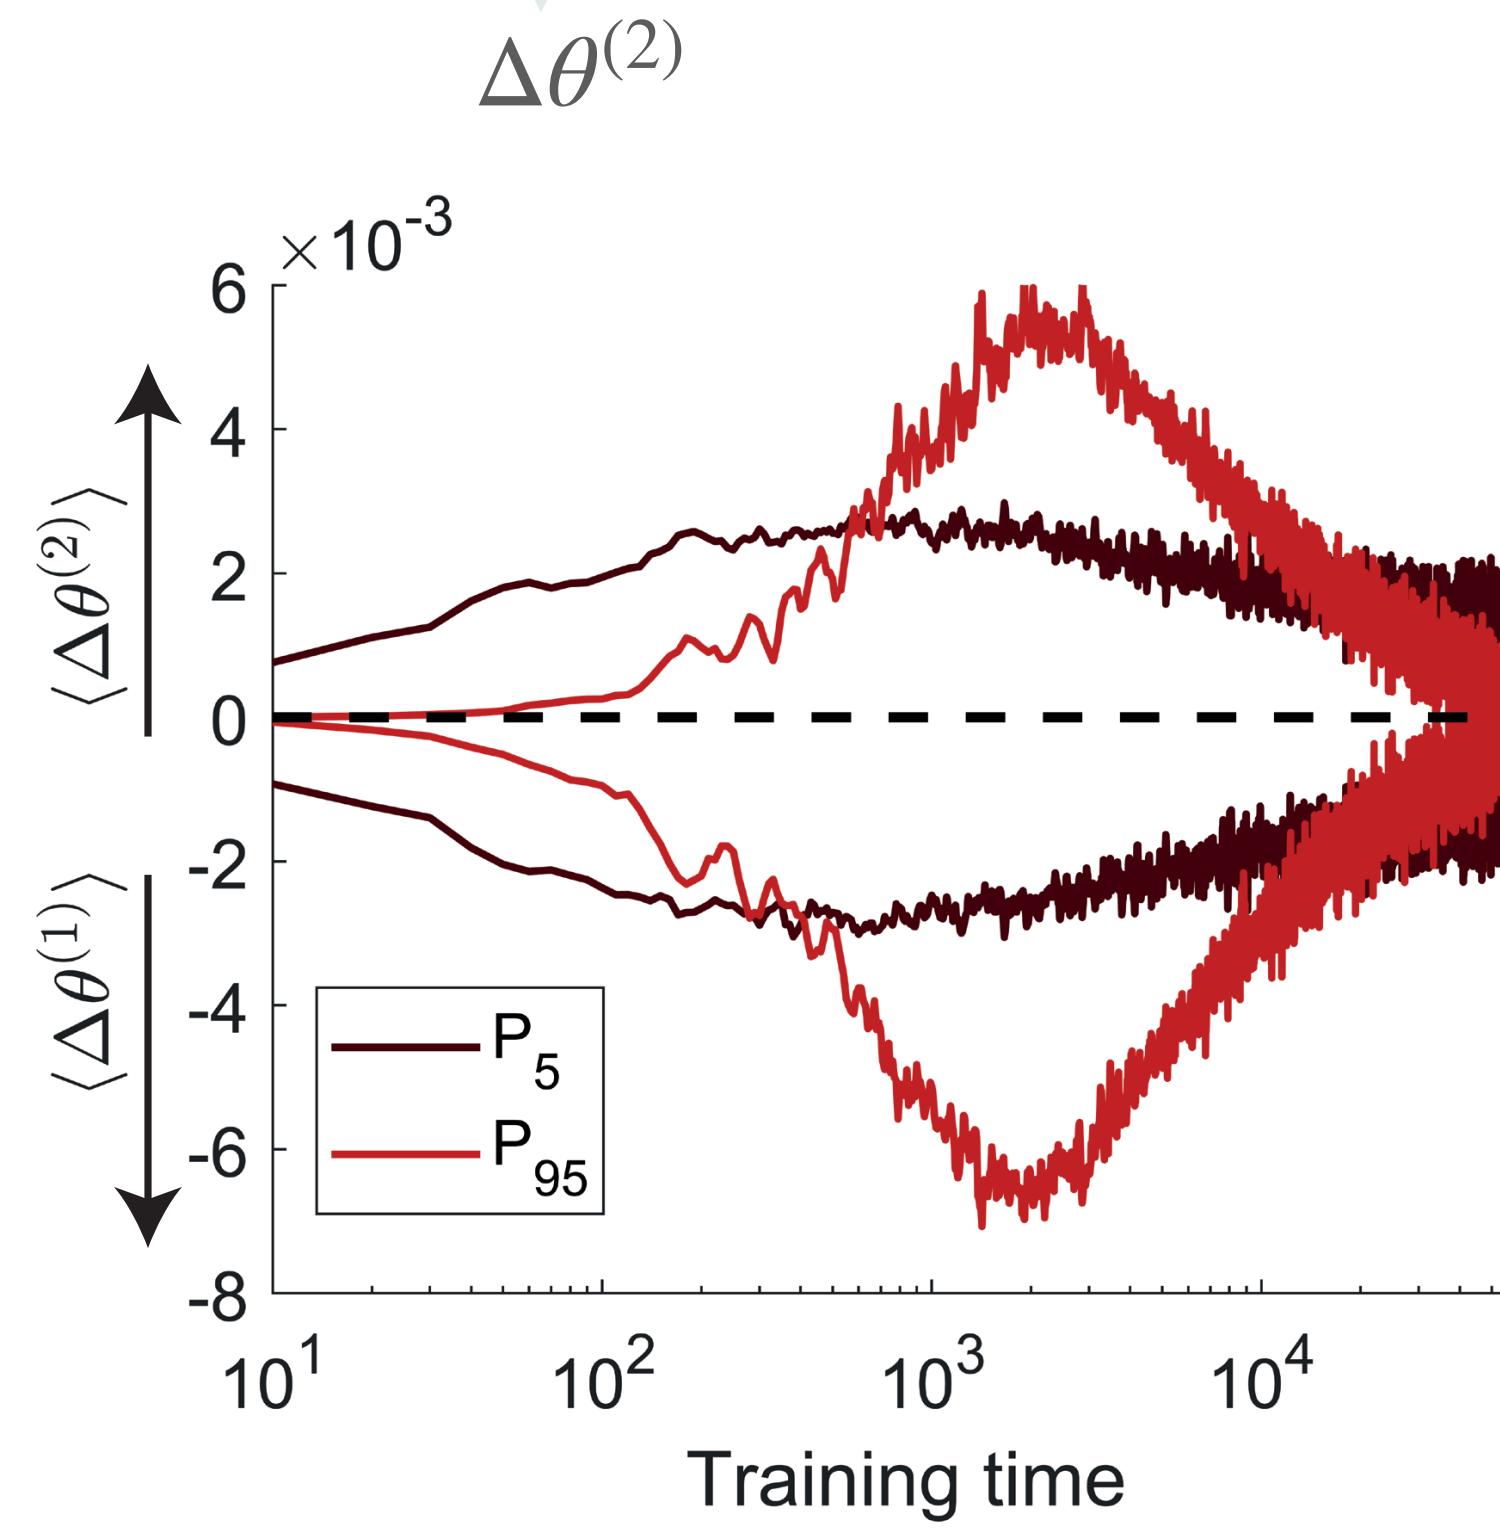
$$\Delta\theta_i/\eta_\theta = -y_{j^*}^t \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$



Example: column by column MNIST

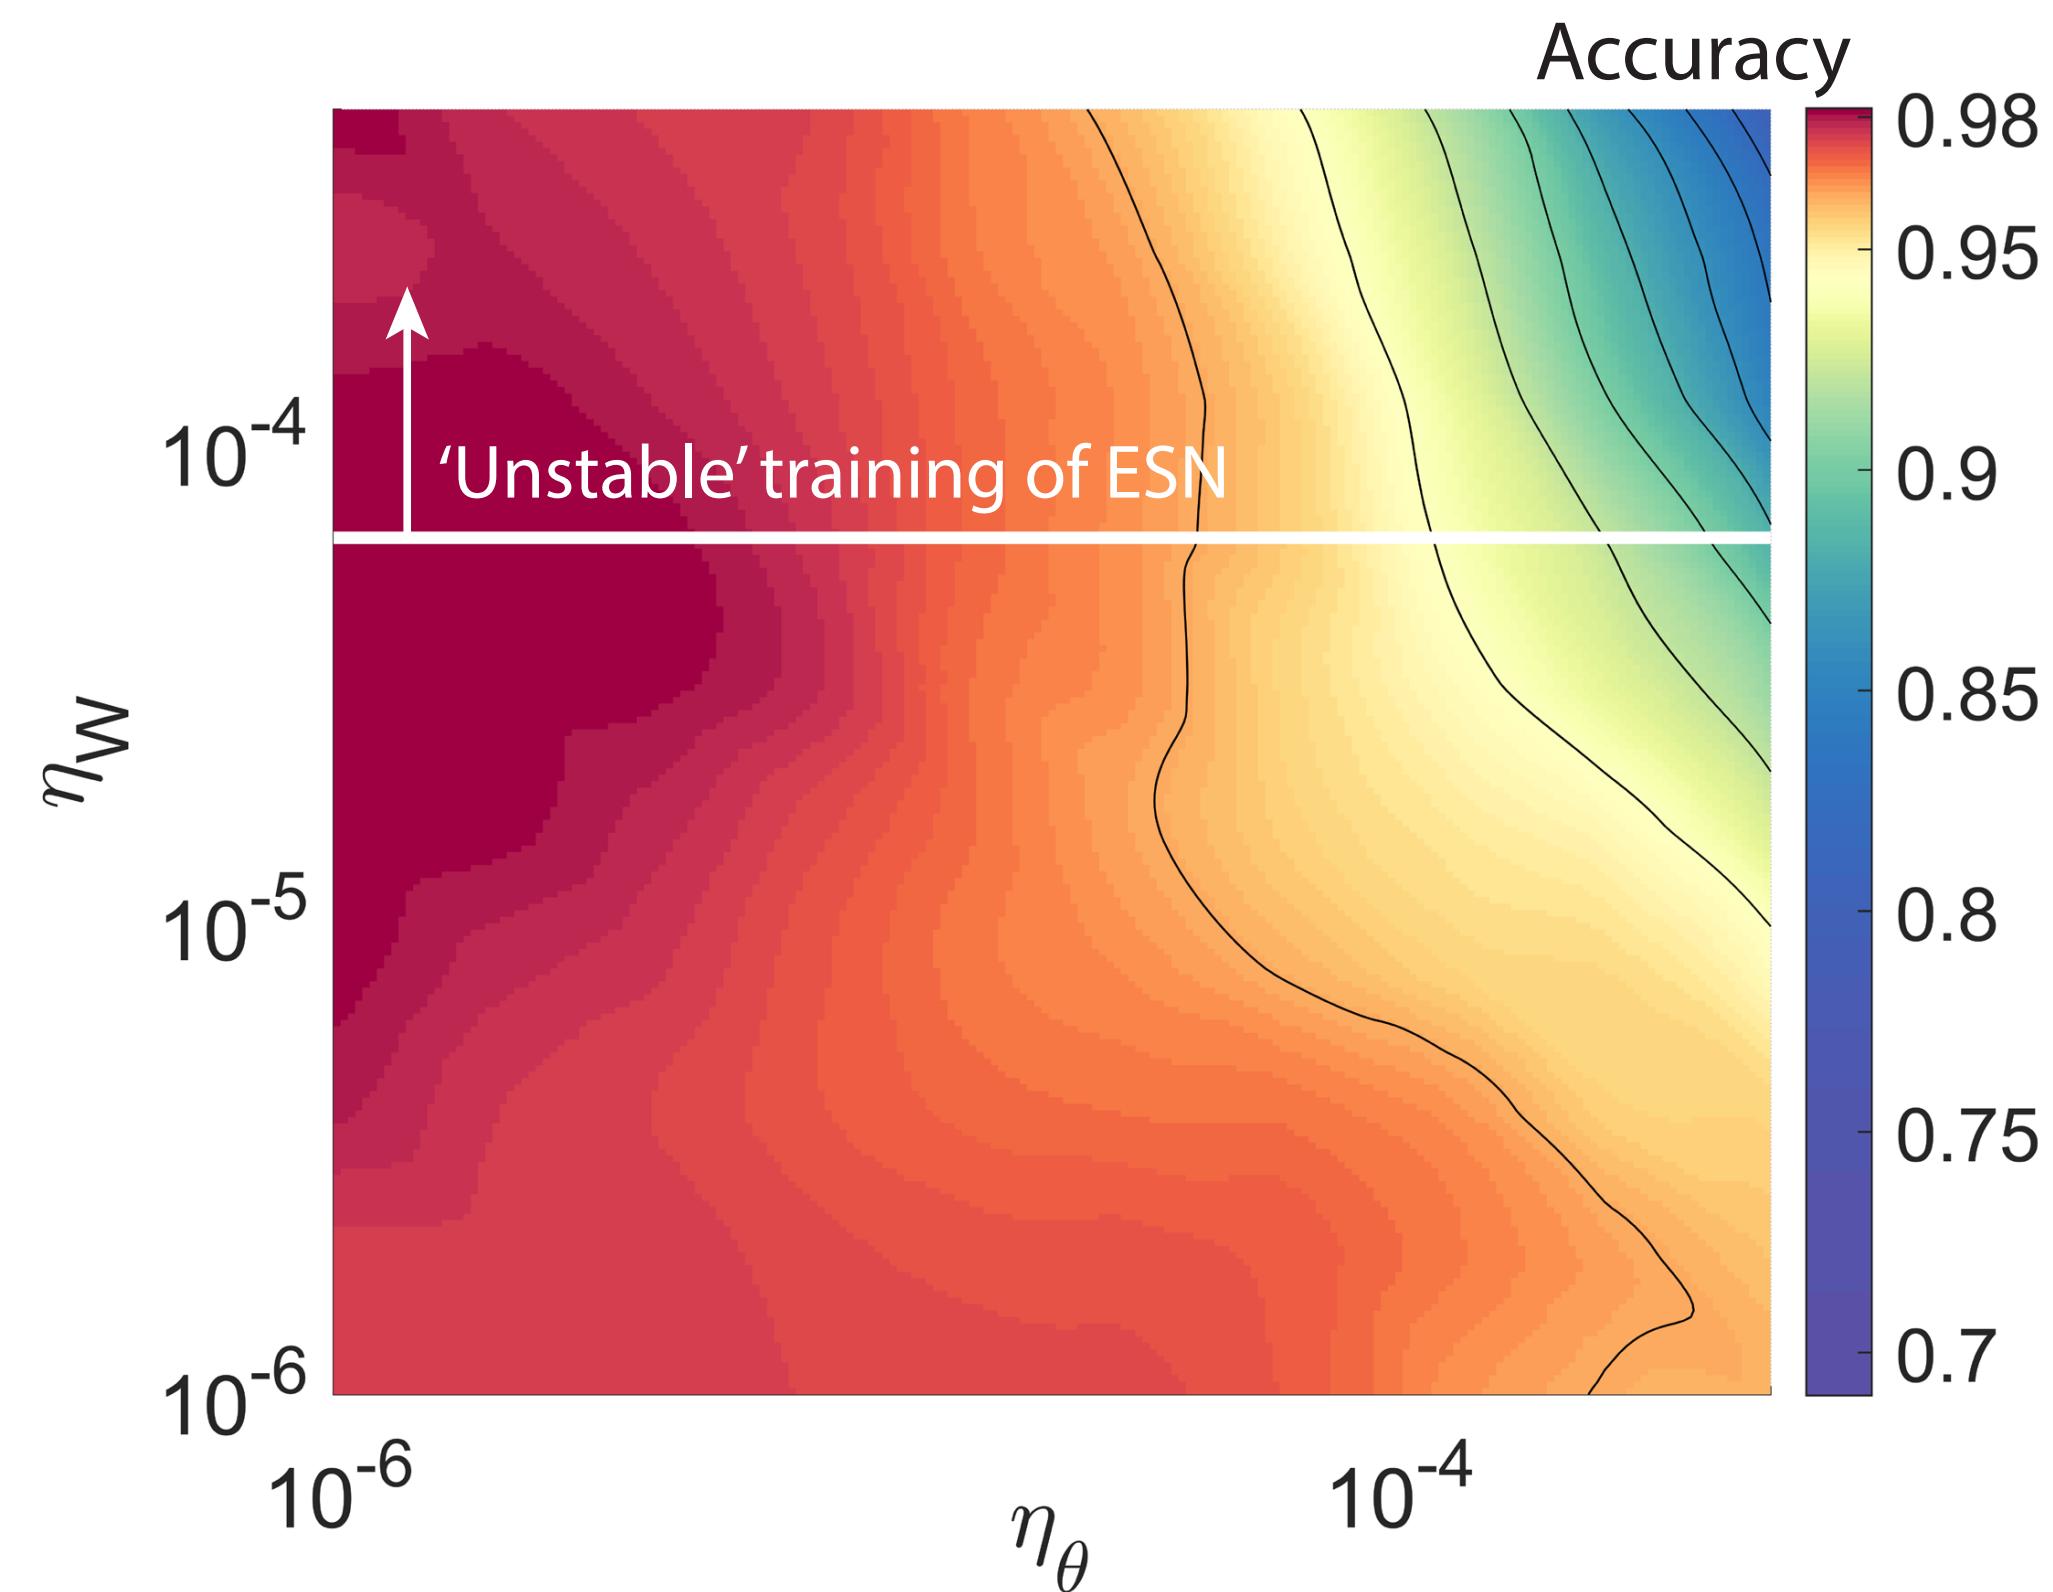
TWO FORCES DRIVE THRESHOLD LEARNING

$$\Delta\theta_i/\eta_\theta = -y_{j^*}^t \left(W_{j^*i}^{out} \text{sign}(x_i) \right) + \sum_j \sum_l \left(W_{jl}^{out} x_l \right) \left(W_{ji}^{out} \text{sign}(x_i) \right)$$

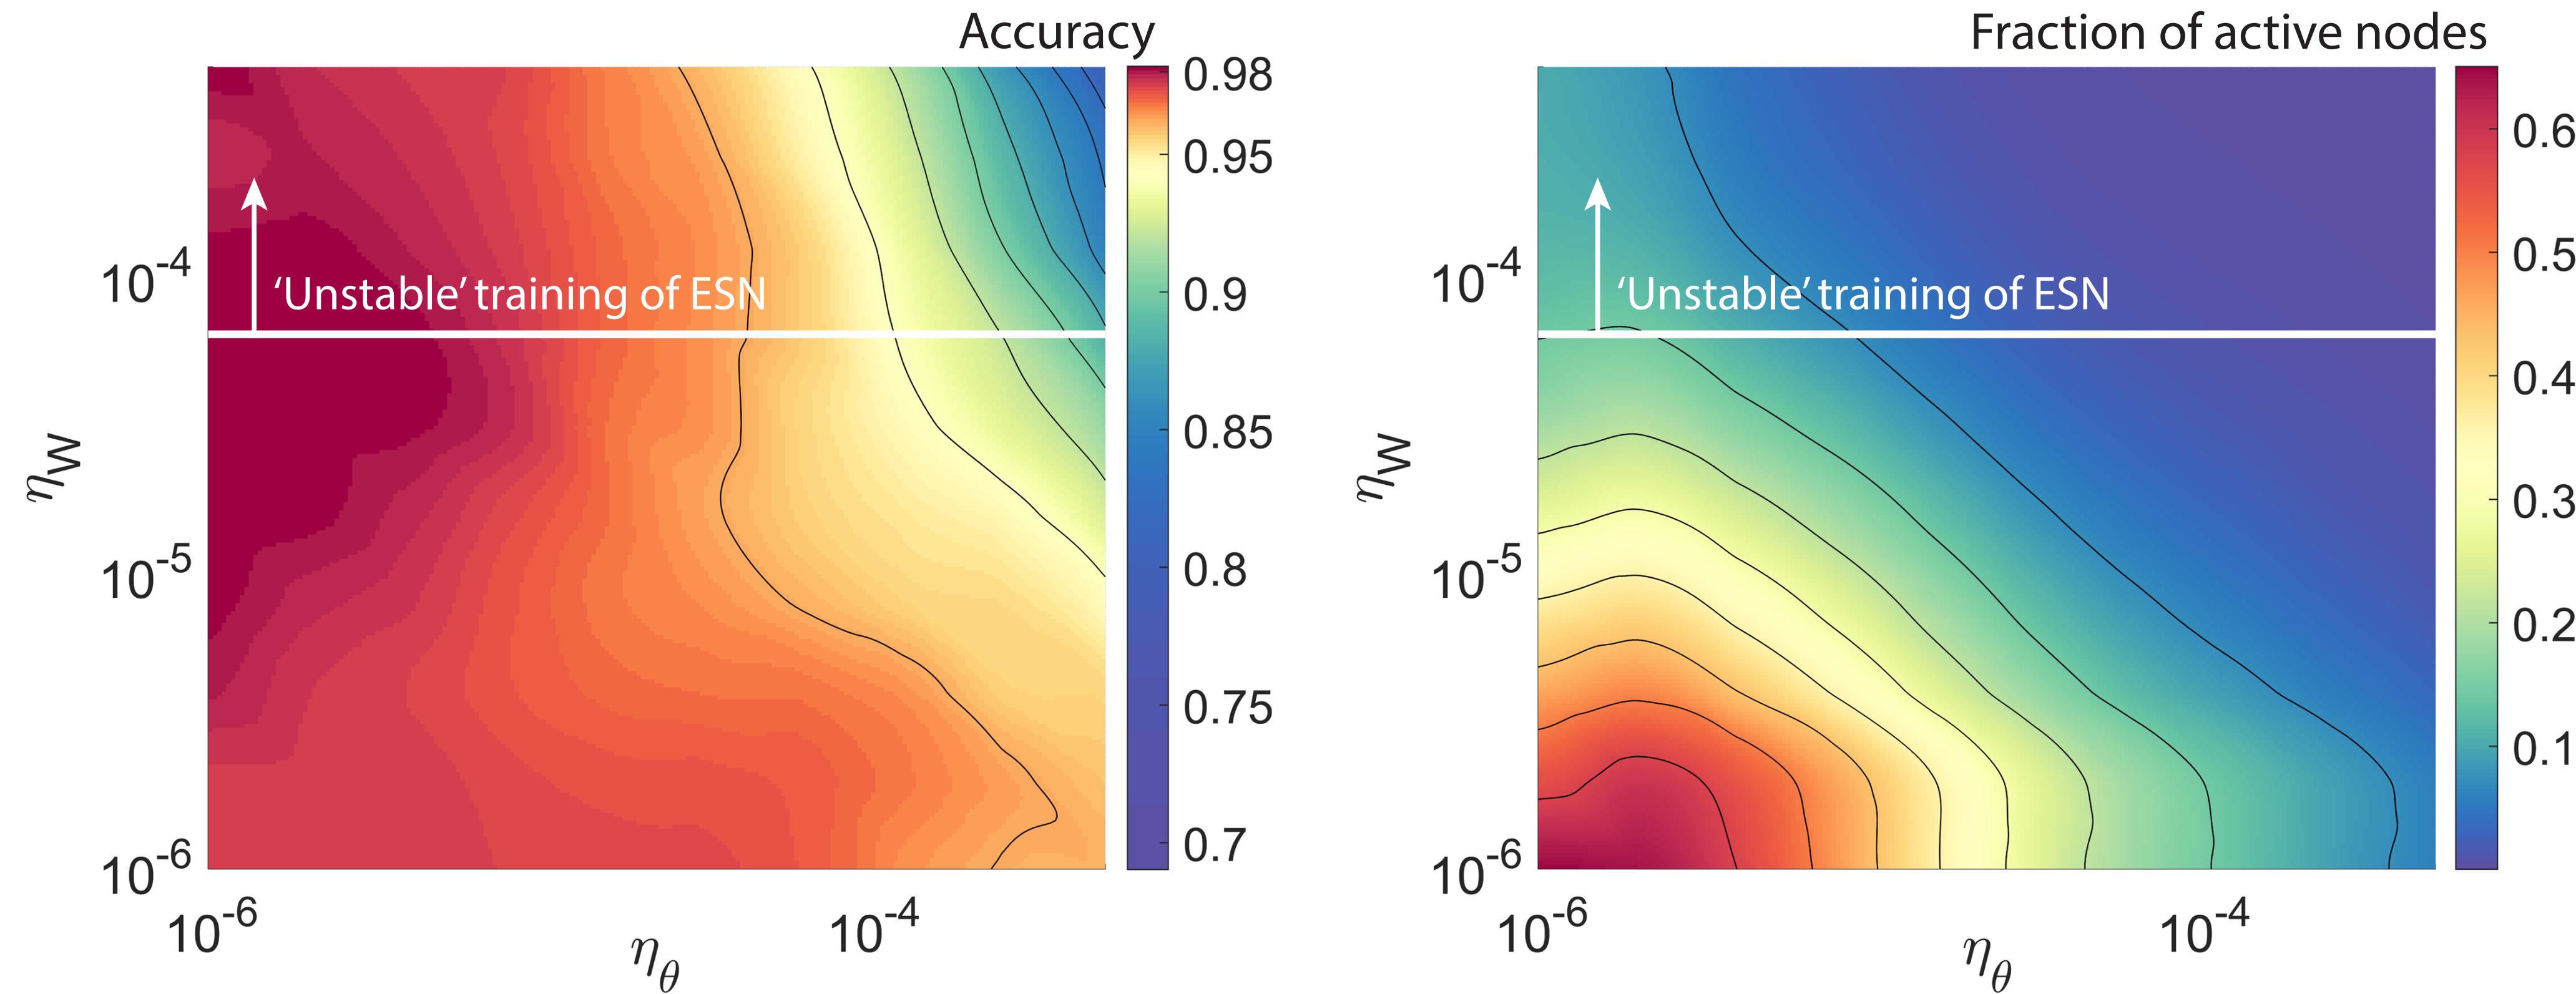


Conclusions also hold for cross-entropy error function

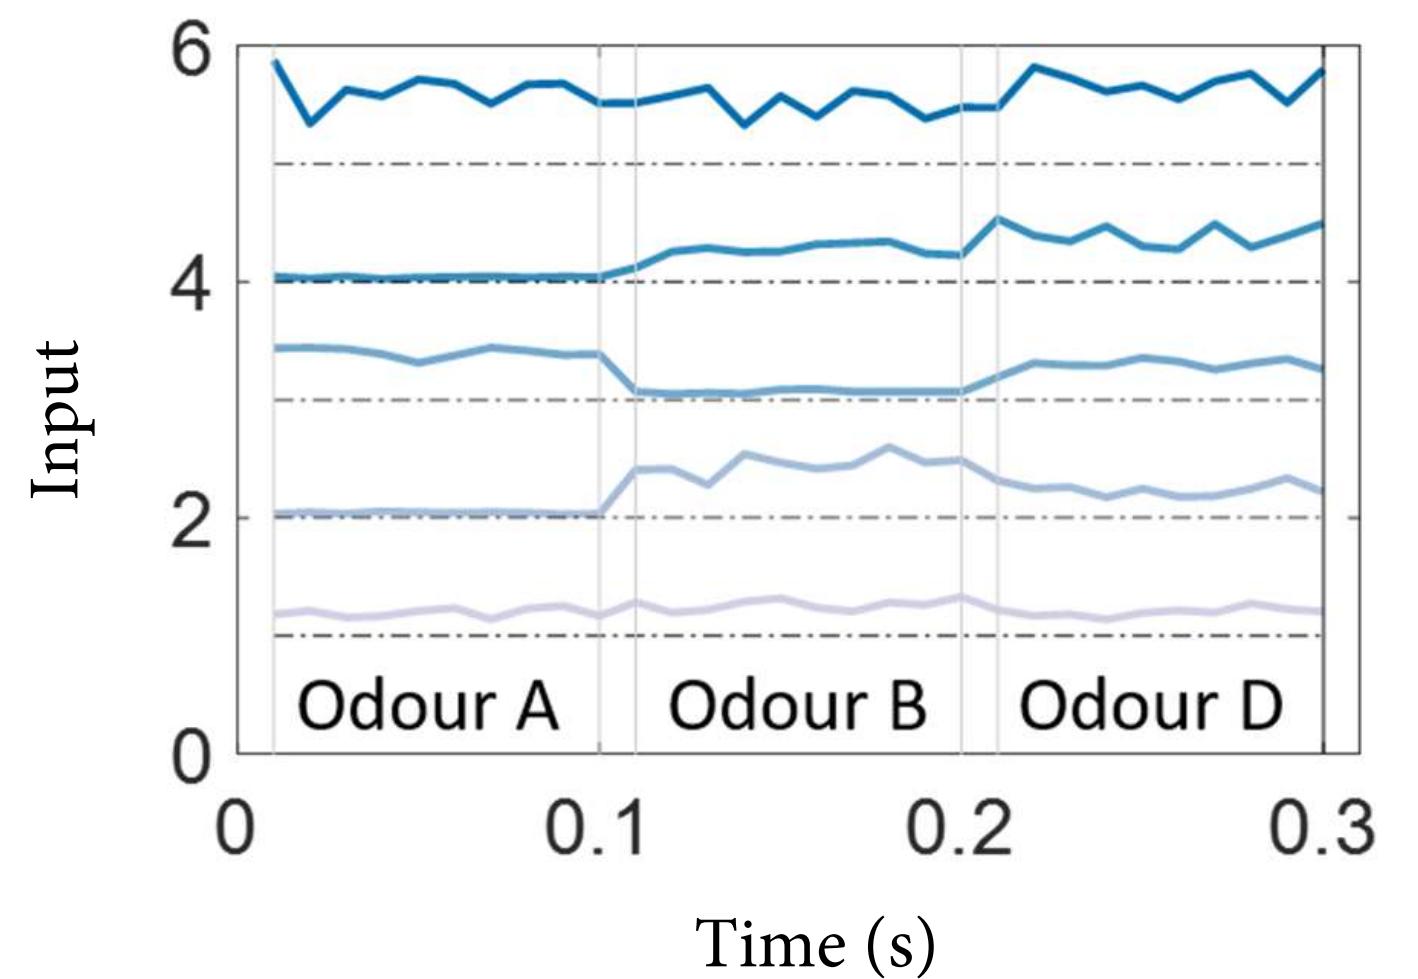
THRESHOLDS ALLOW FOR HIGHER LEARNING RATES



THRESHOLDS ALLOW FOR HIGHER LEARNING RATES

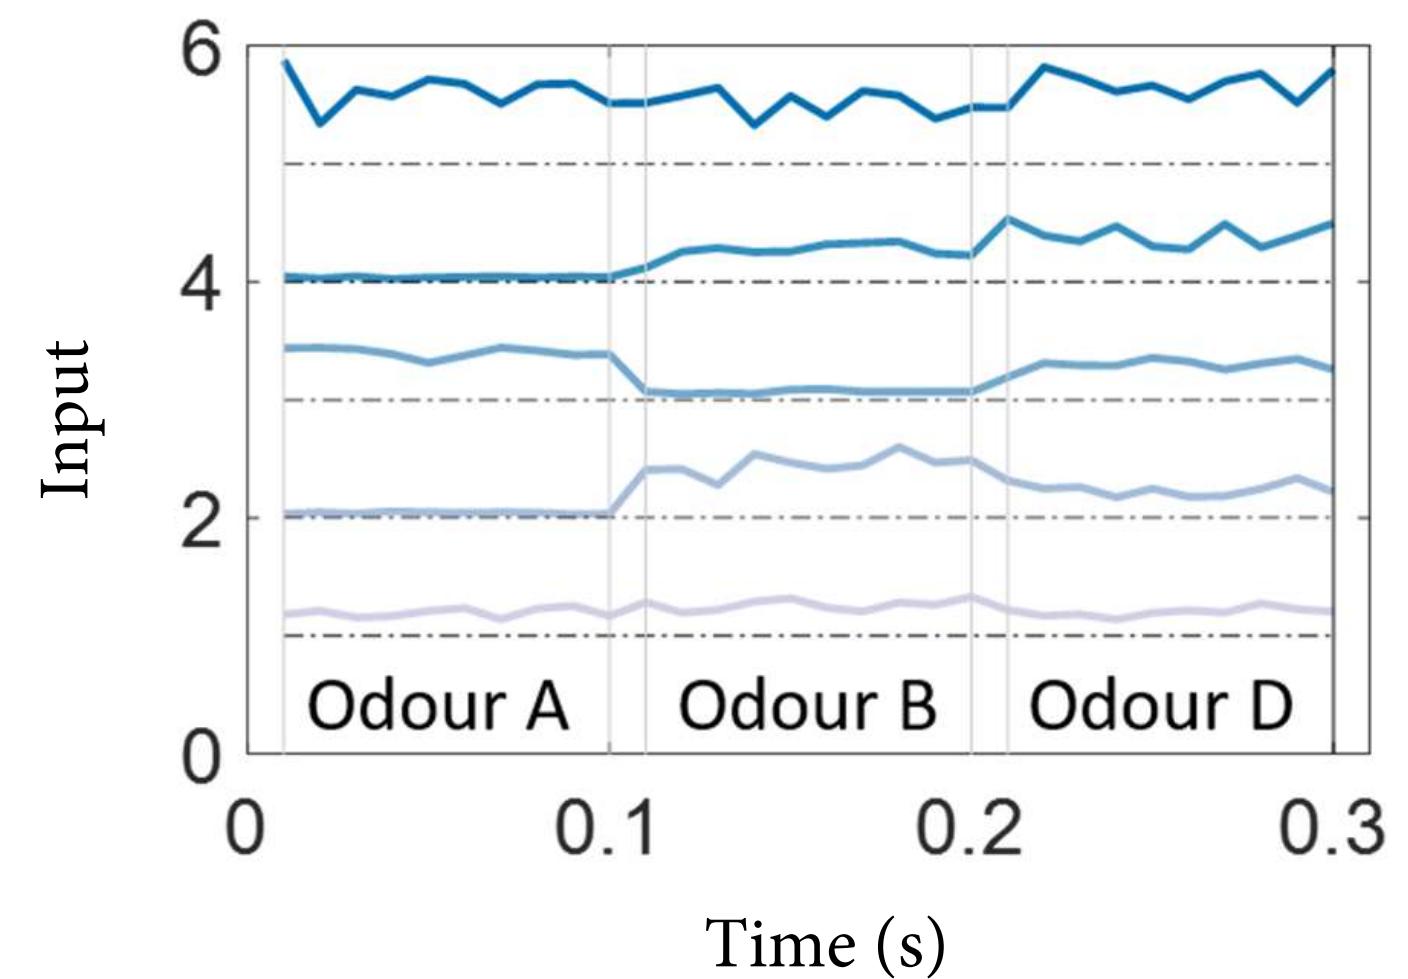


BENCHMARK: ODOURS SEQUENCE



Elissa A Hallem and John R Carlson. Coding of odors by a receptor repertoire. *Cell*, 125(1):143–160, 2006.

BENCHMARK: ODOURS SEQUENCE

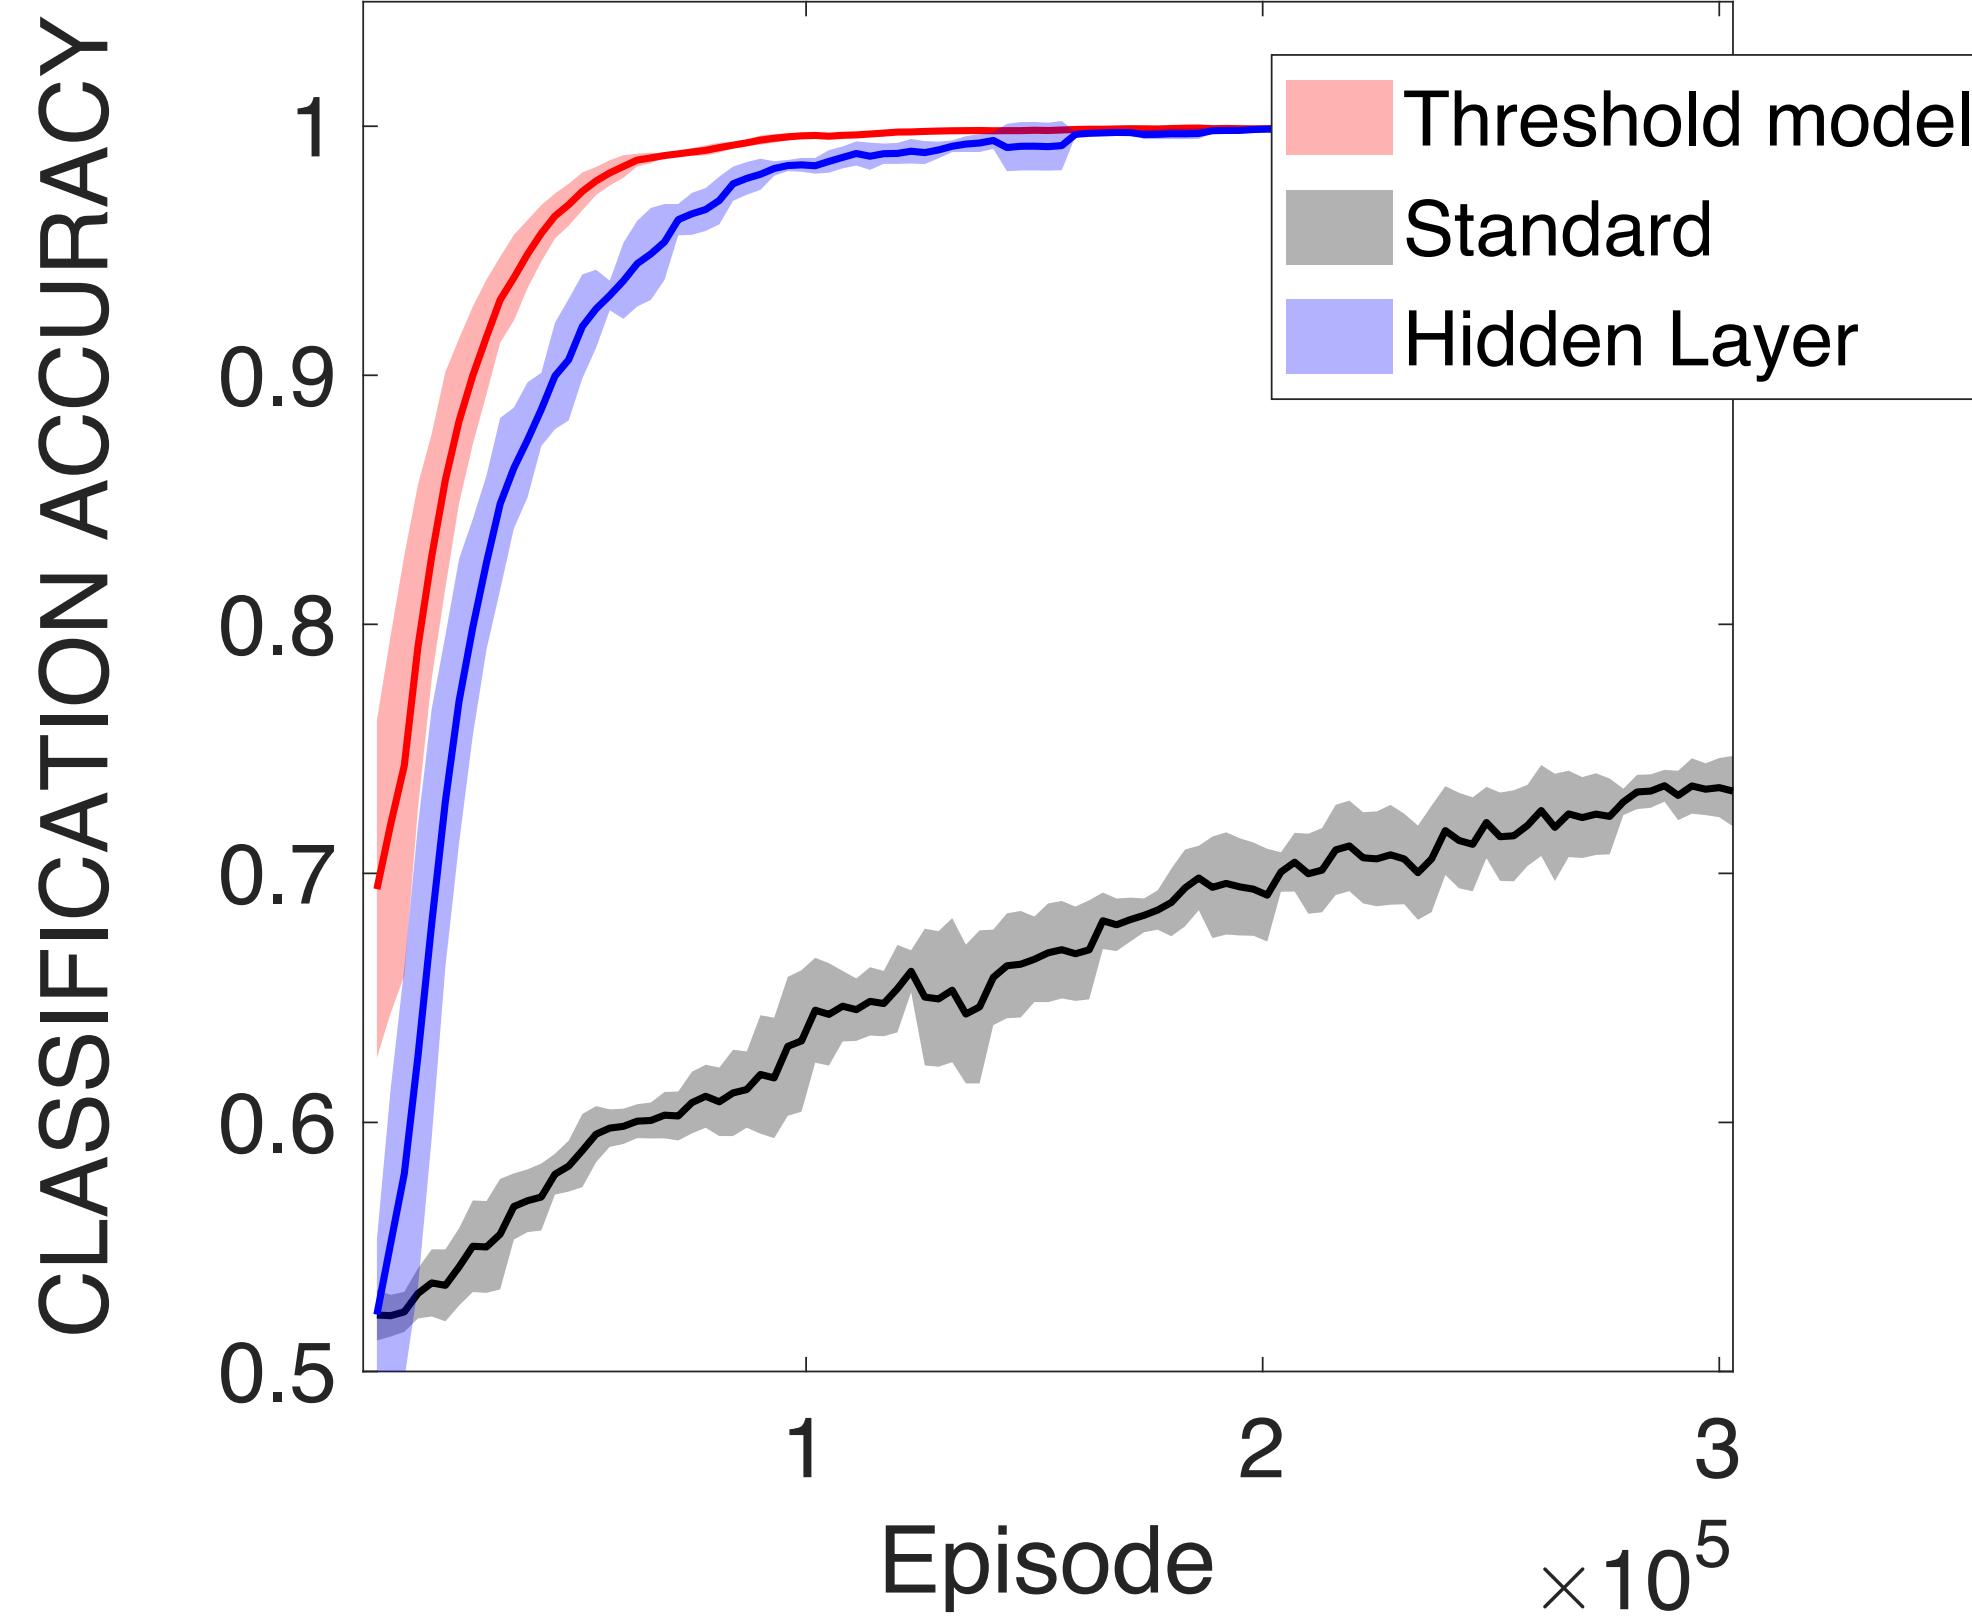


Class			
A	B	D	1
A	B	E	2
A	F	C	1
A	G	C	2
H	B	C	1
I	B	C	2

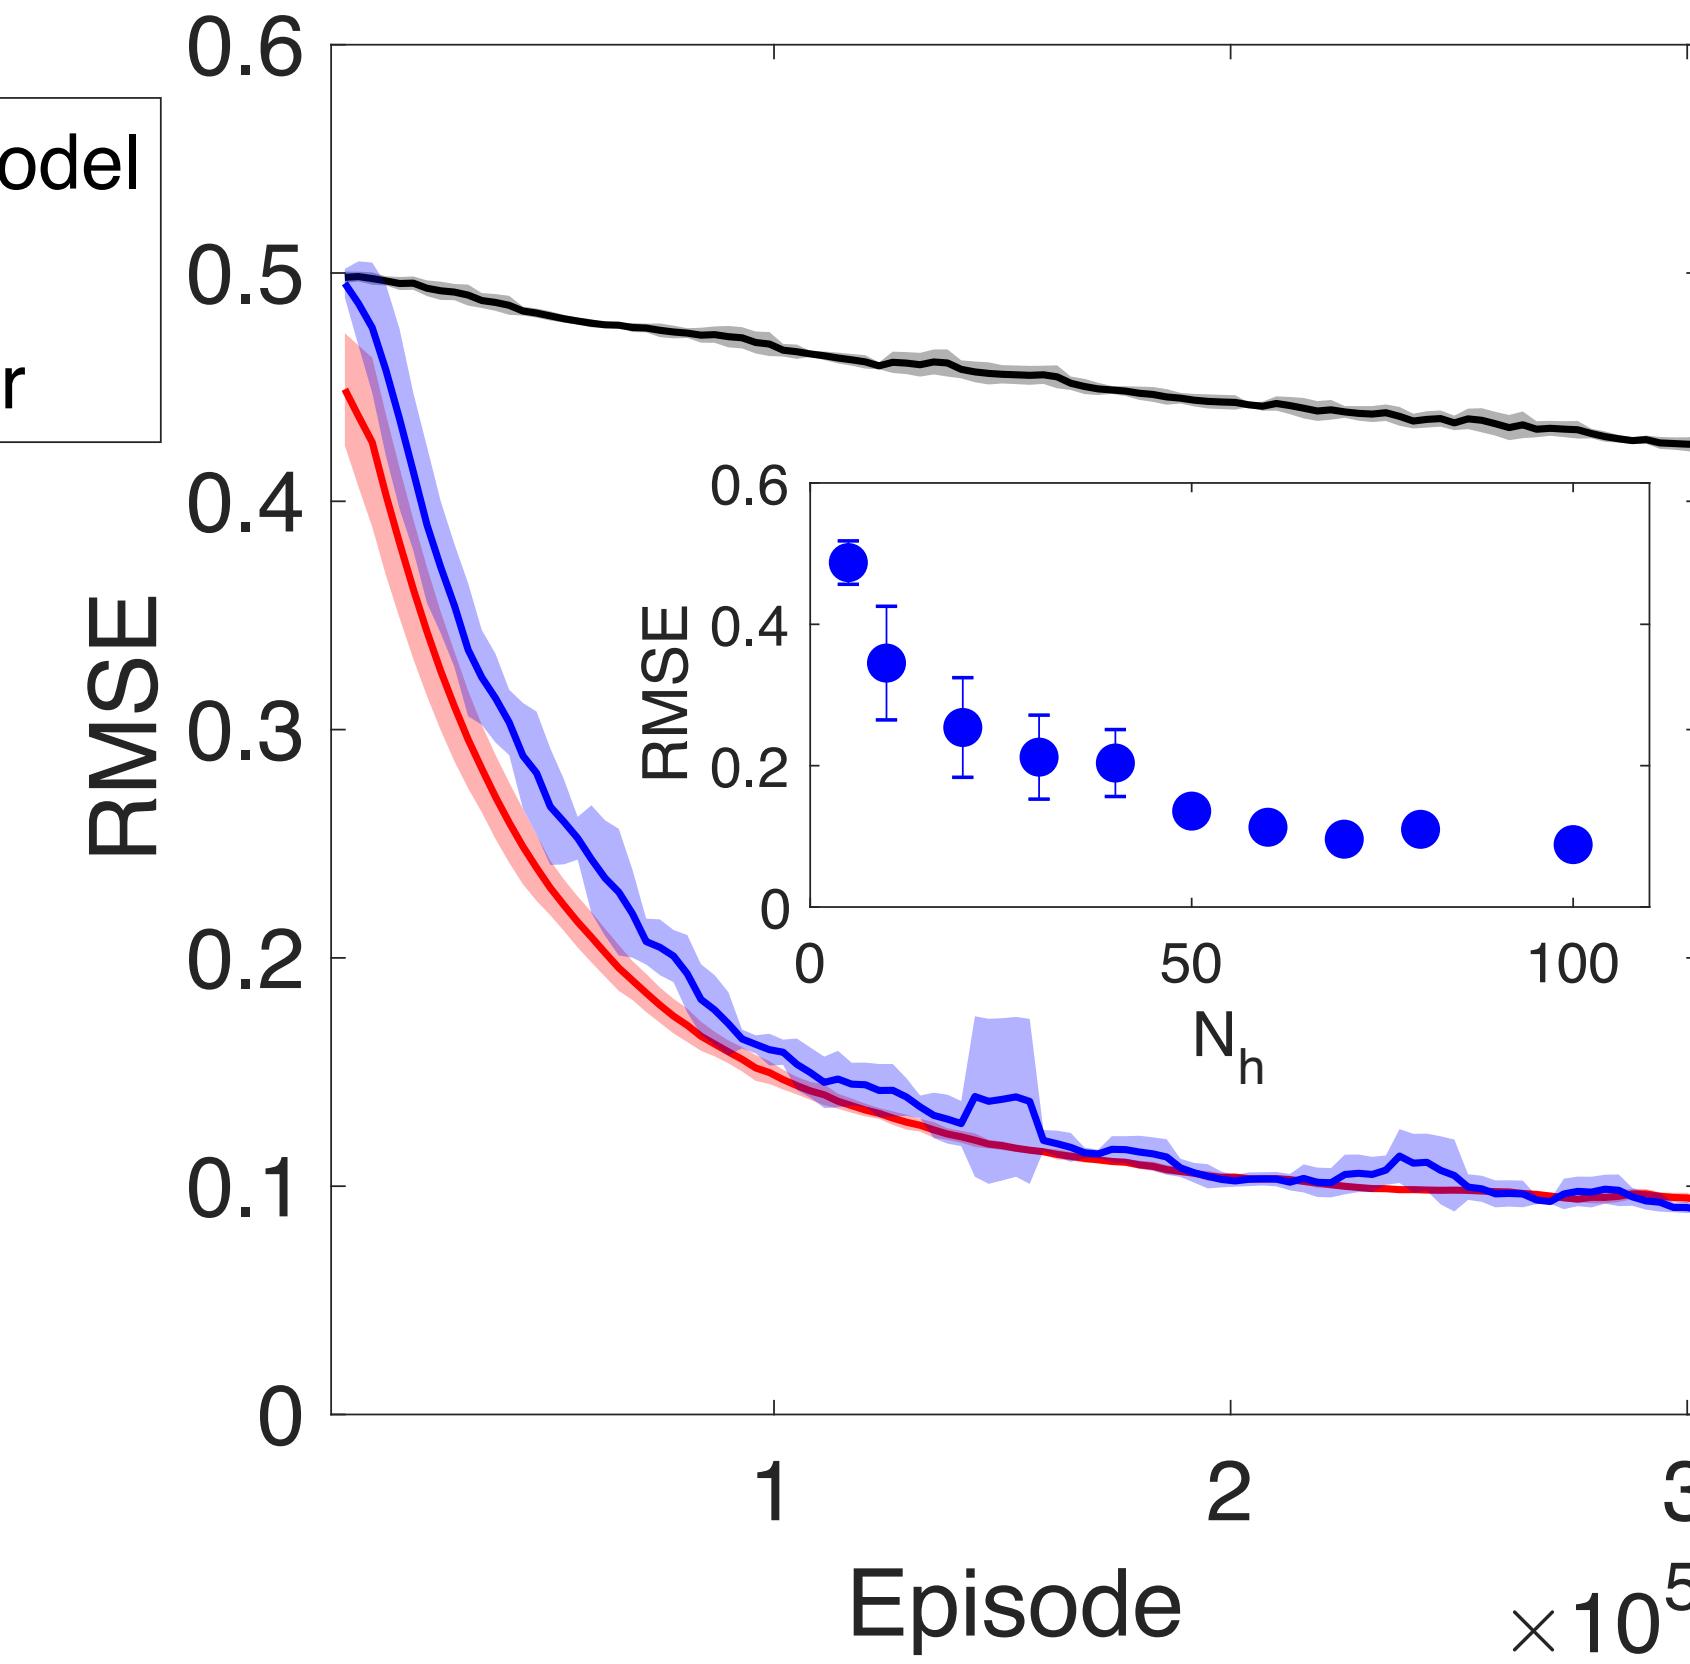
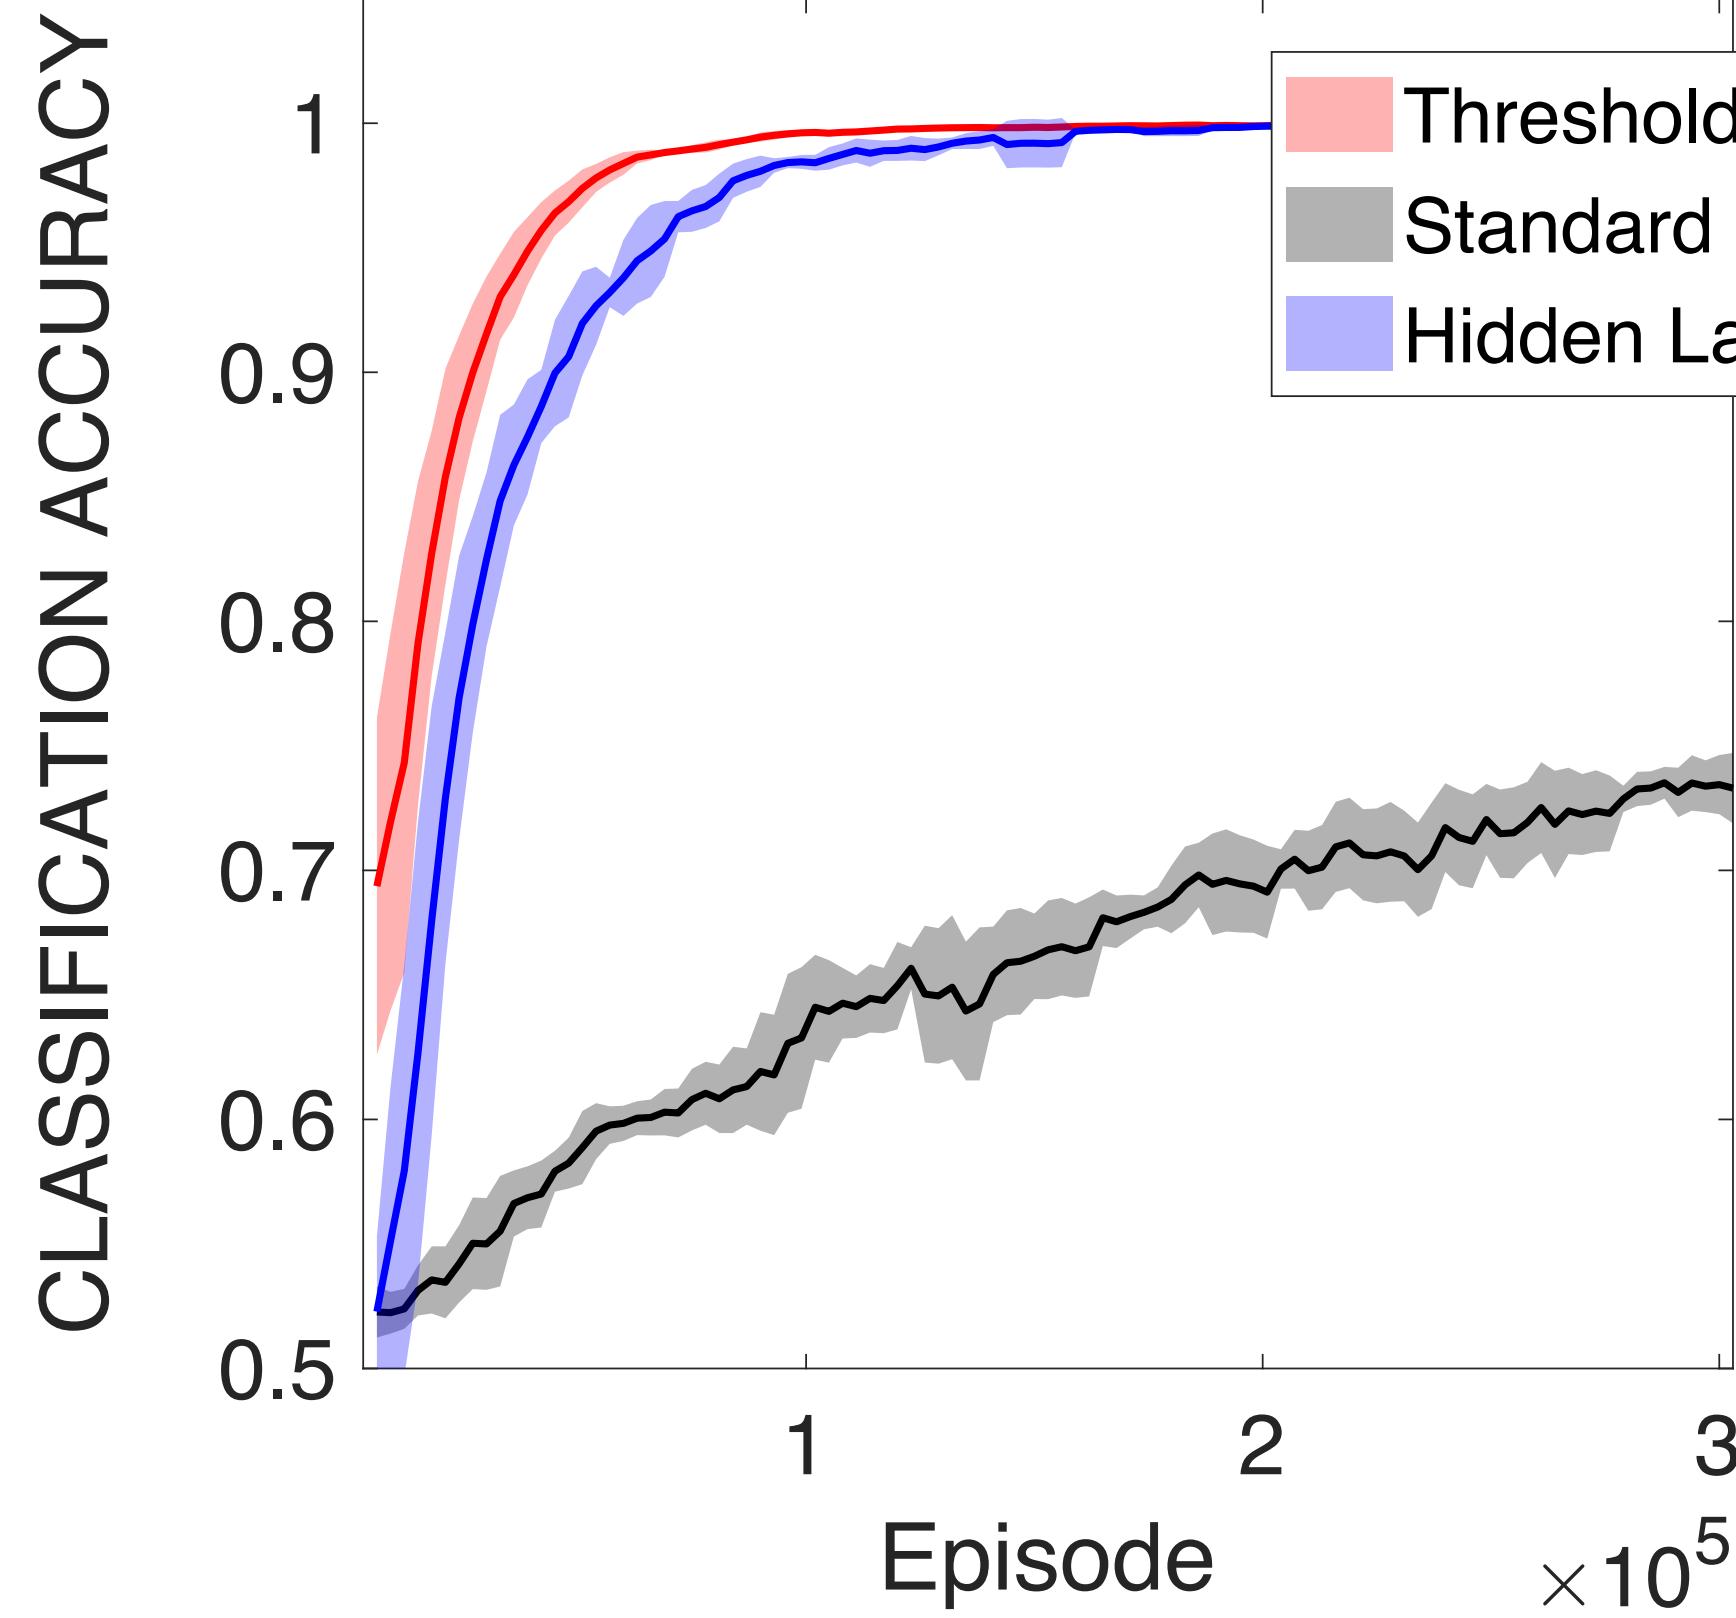
Class			
L	M	D	2
L	M	E	1
L	F	N	2
L	G	N	1
H	M	N	2
I	M	N	1

Elissa A Hallem and John R Carlson. Coding of odors by a receptor repertoire. *Cell*, 125(1):143–160, 2006.

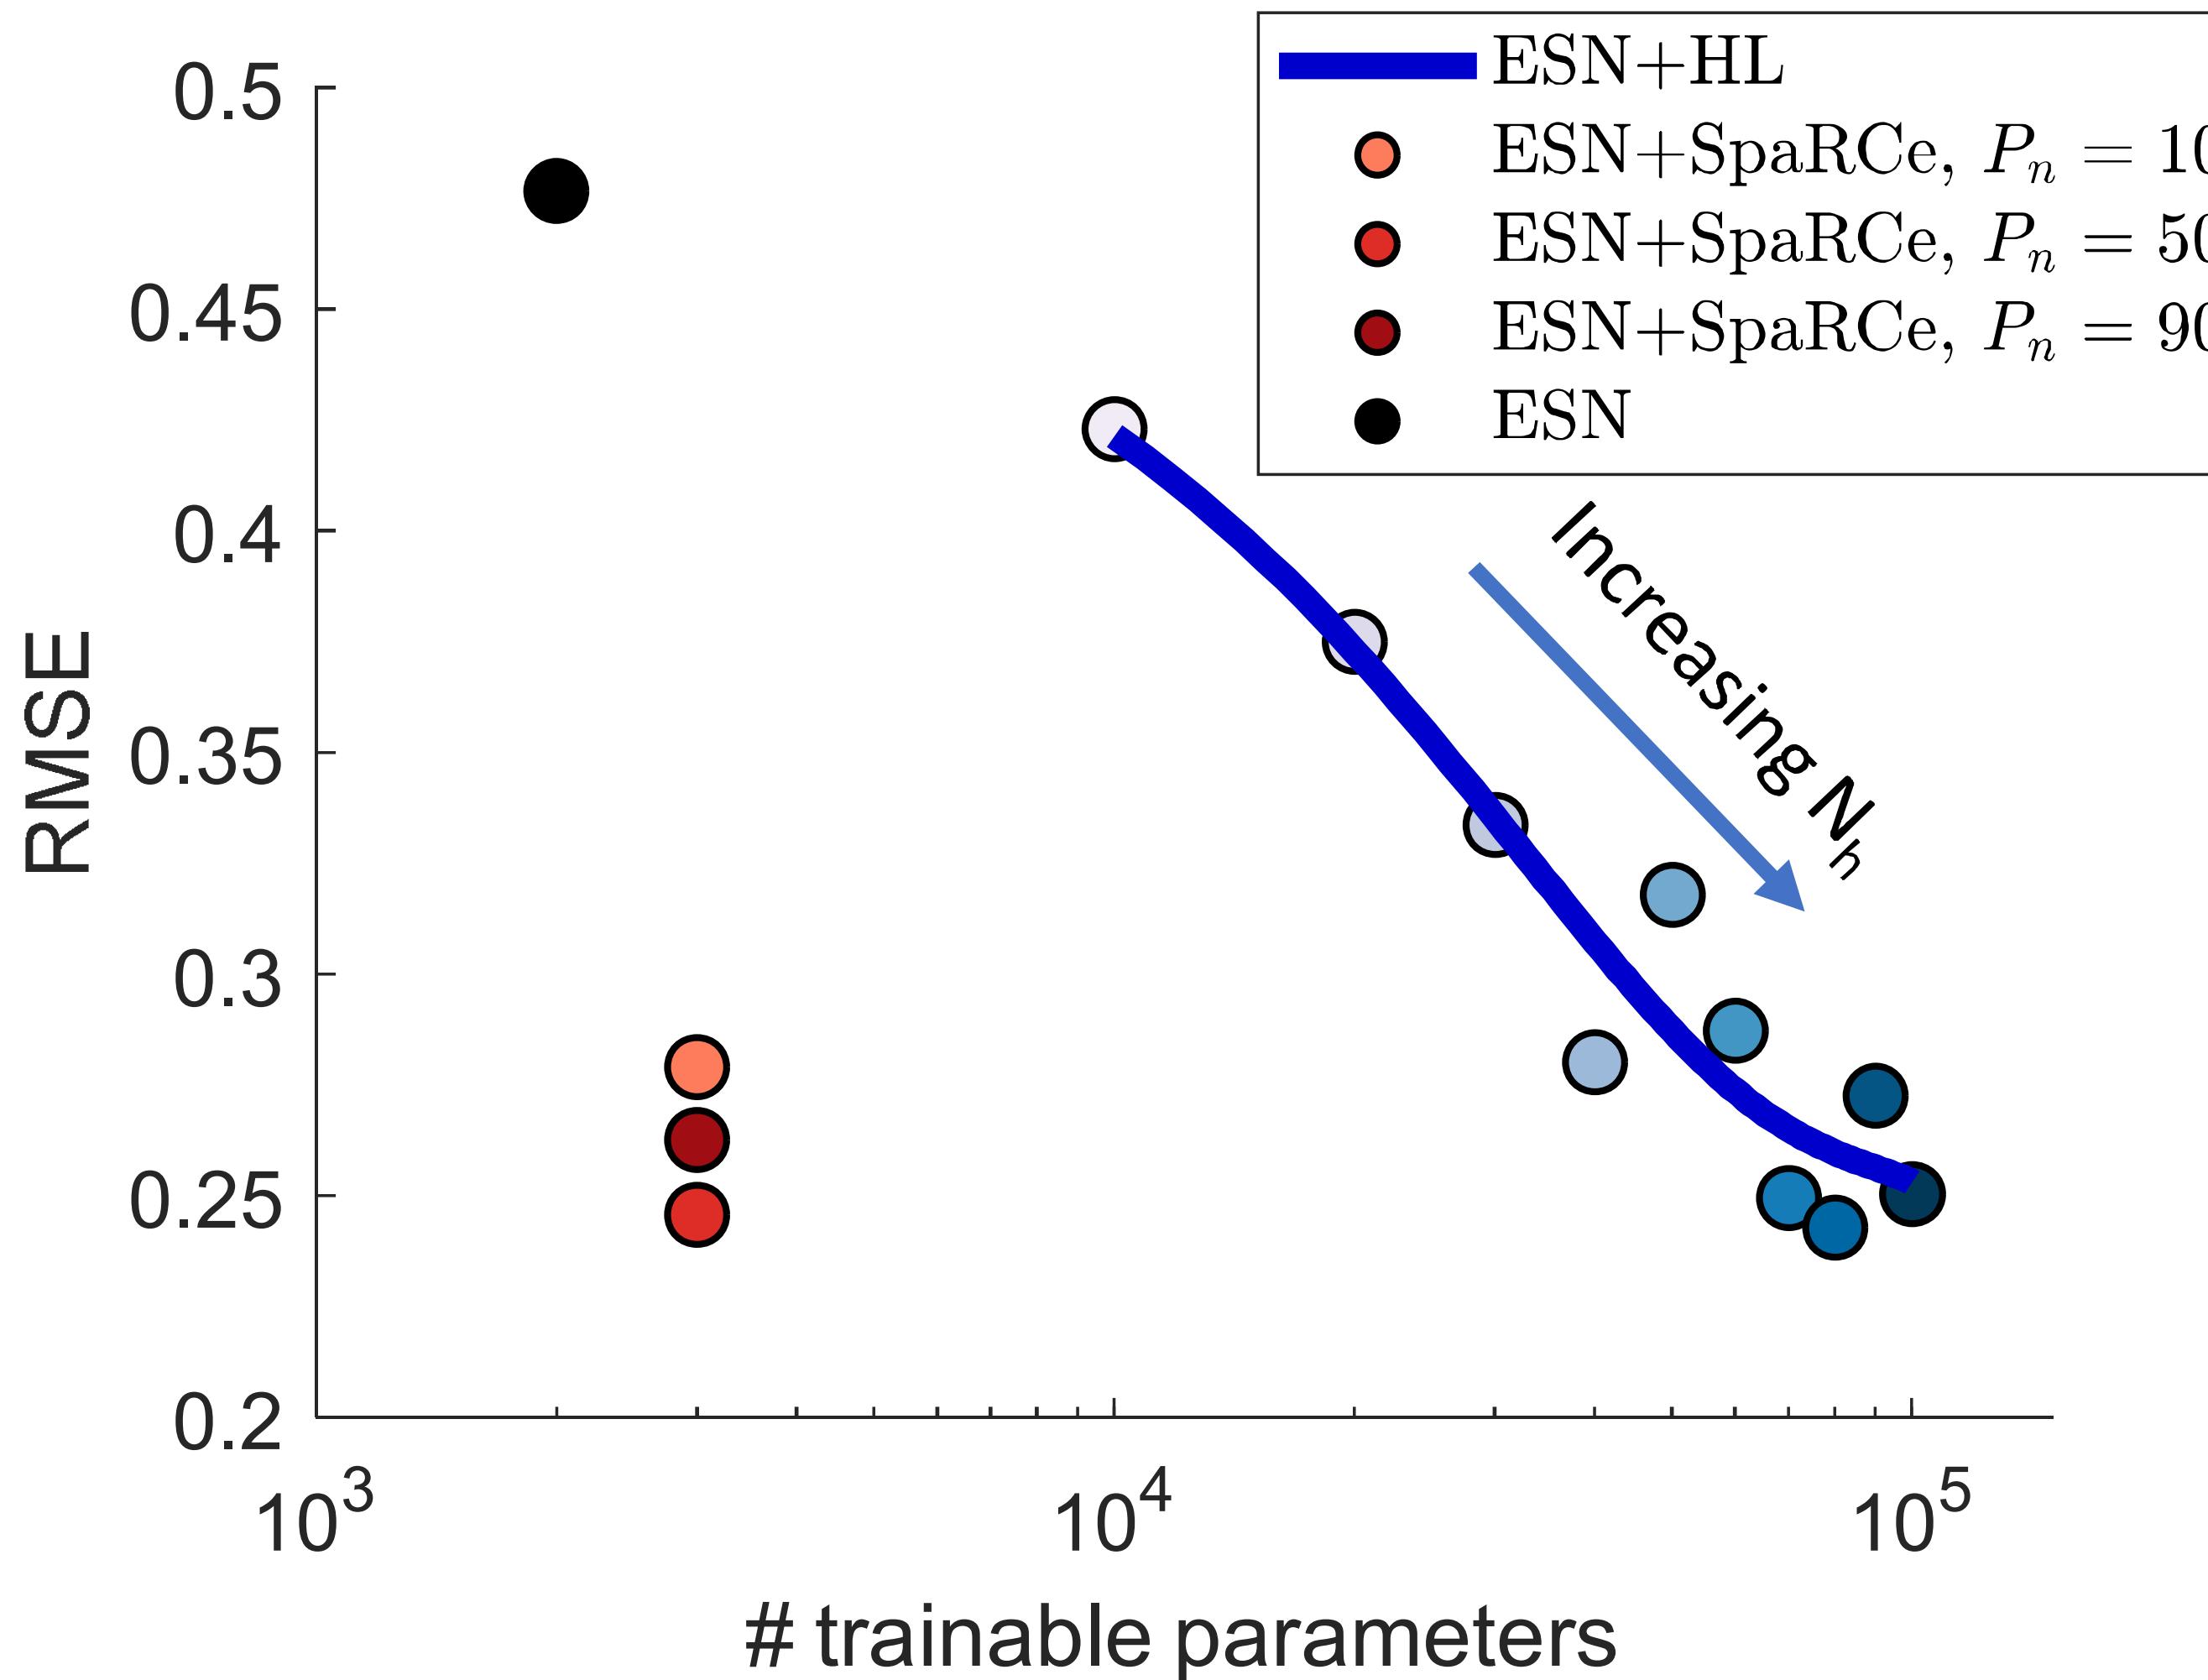
PERFORMANCE: ODOURS SEQUENCE



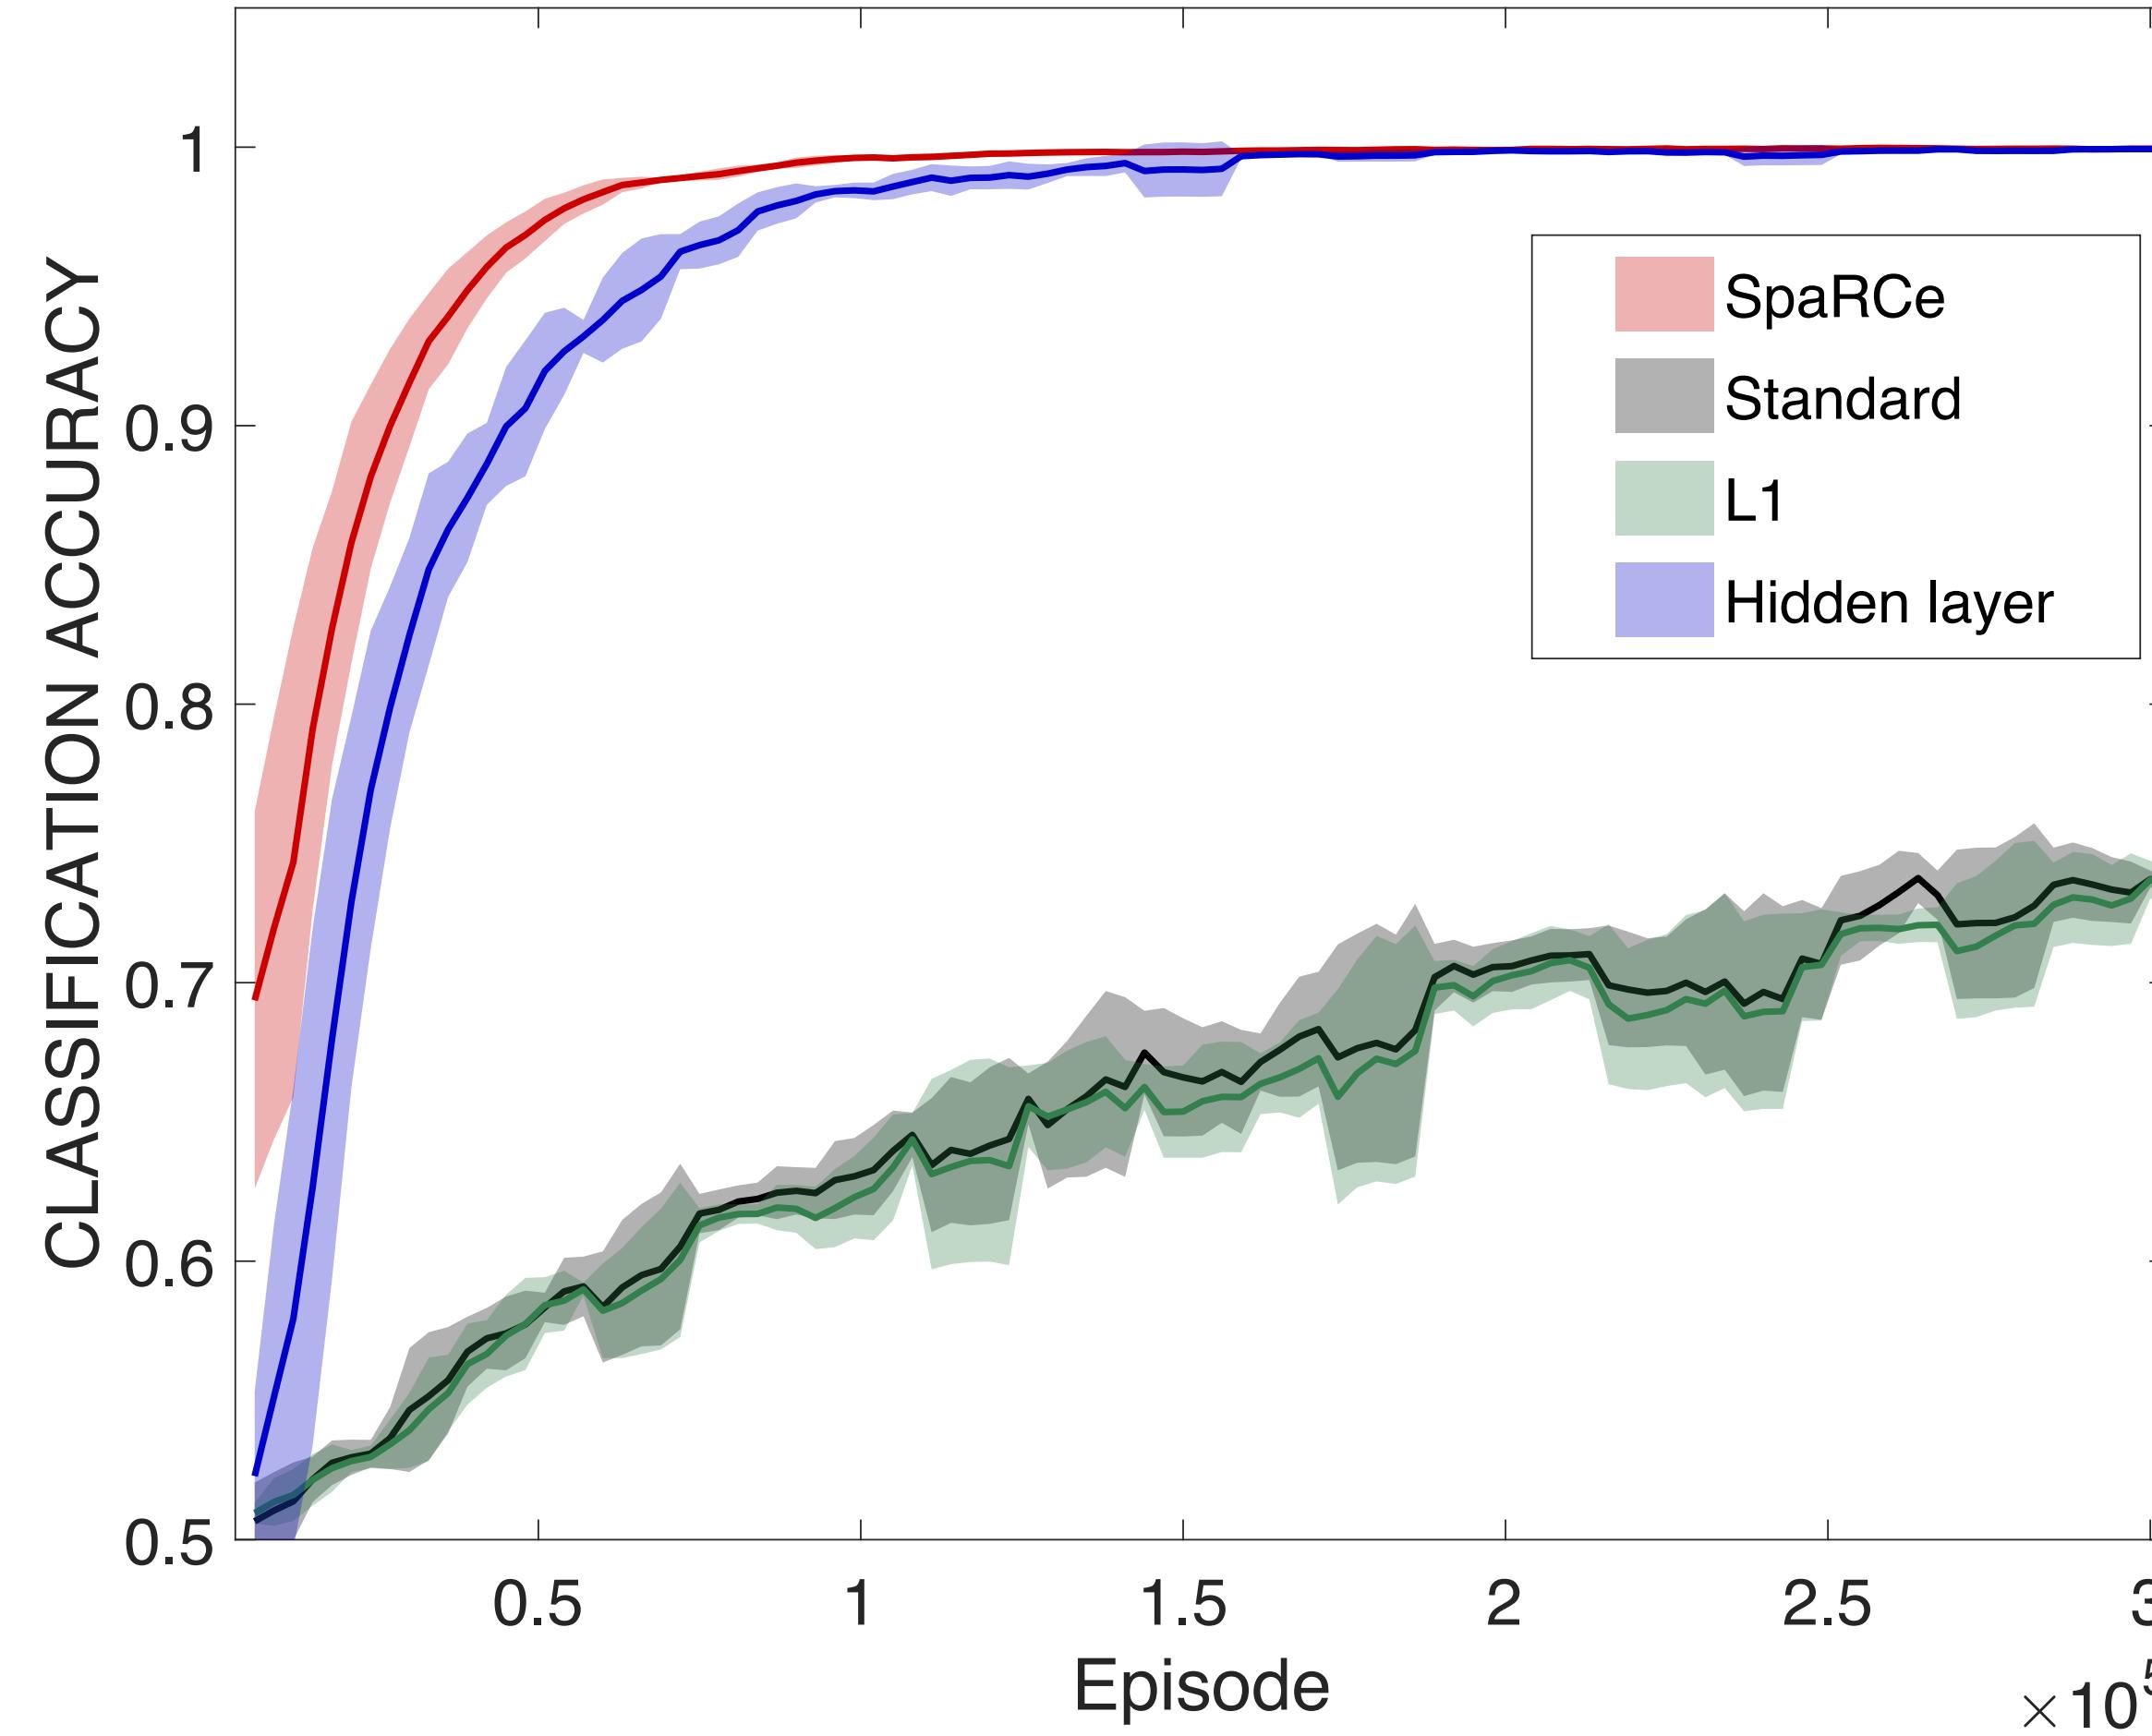
PERFORMANCE: ODOURS SEQUENCE



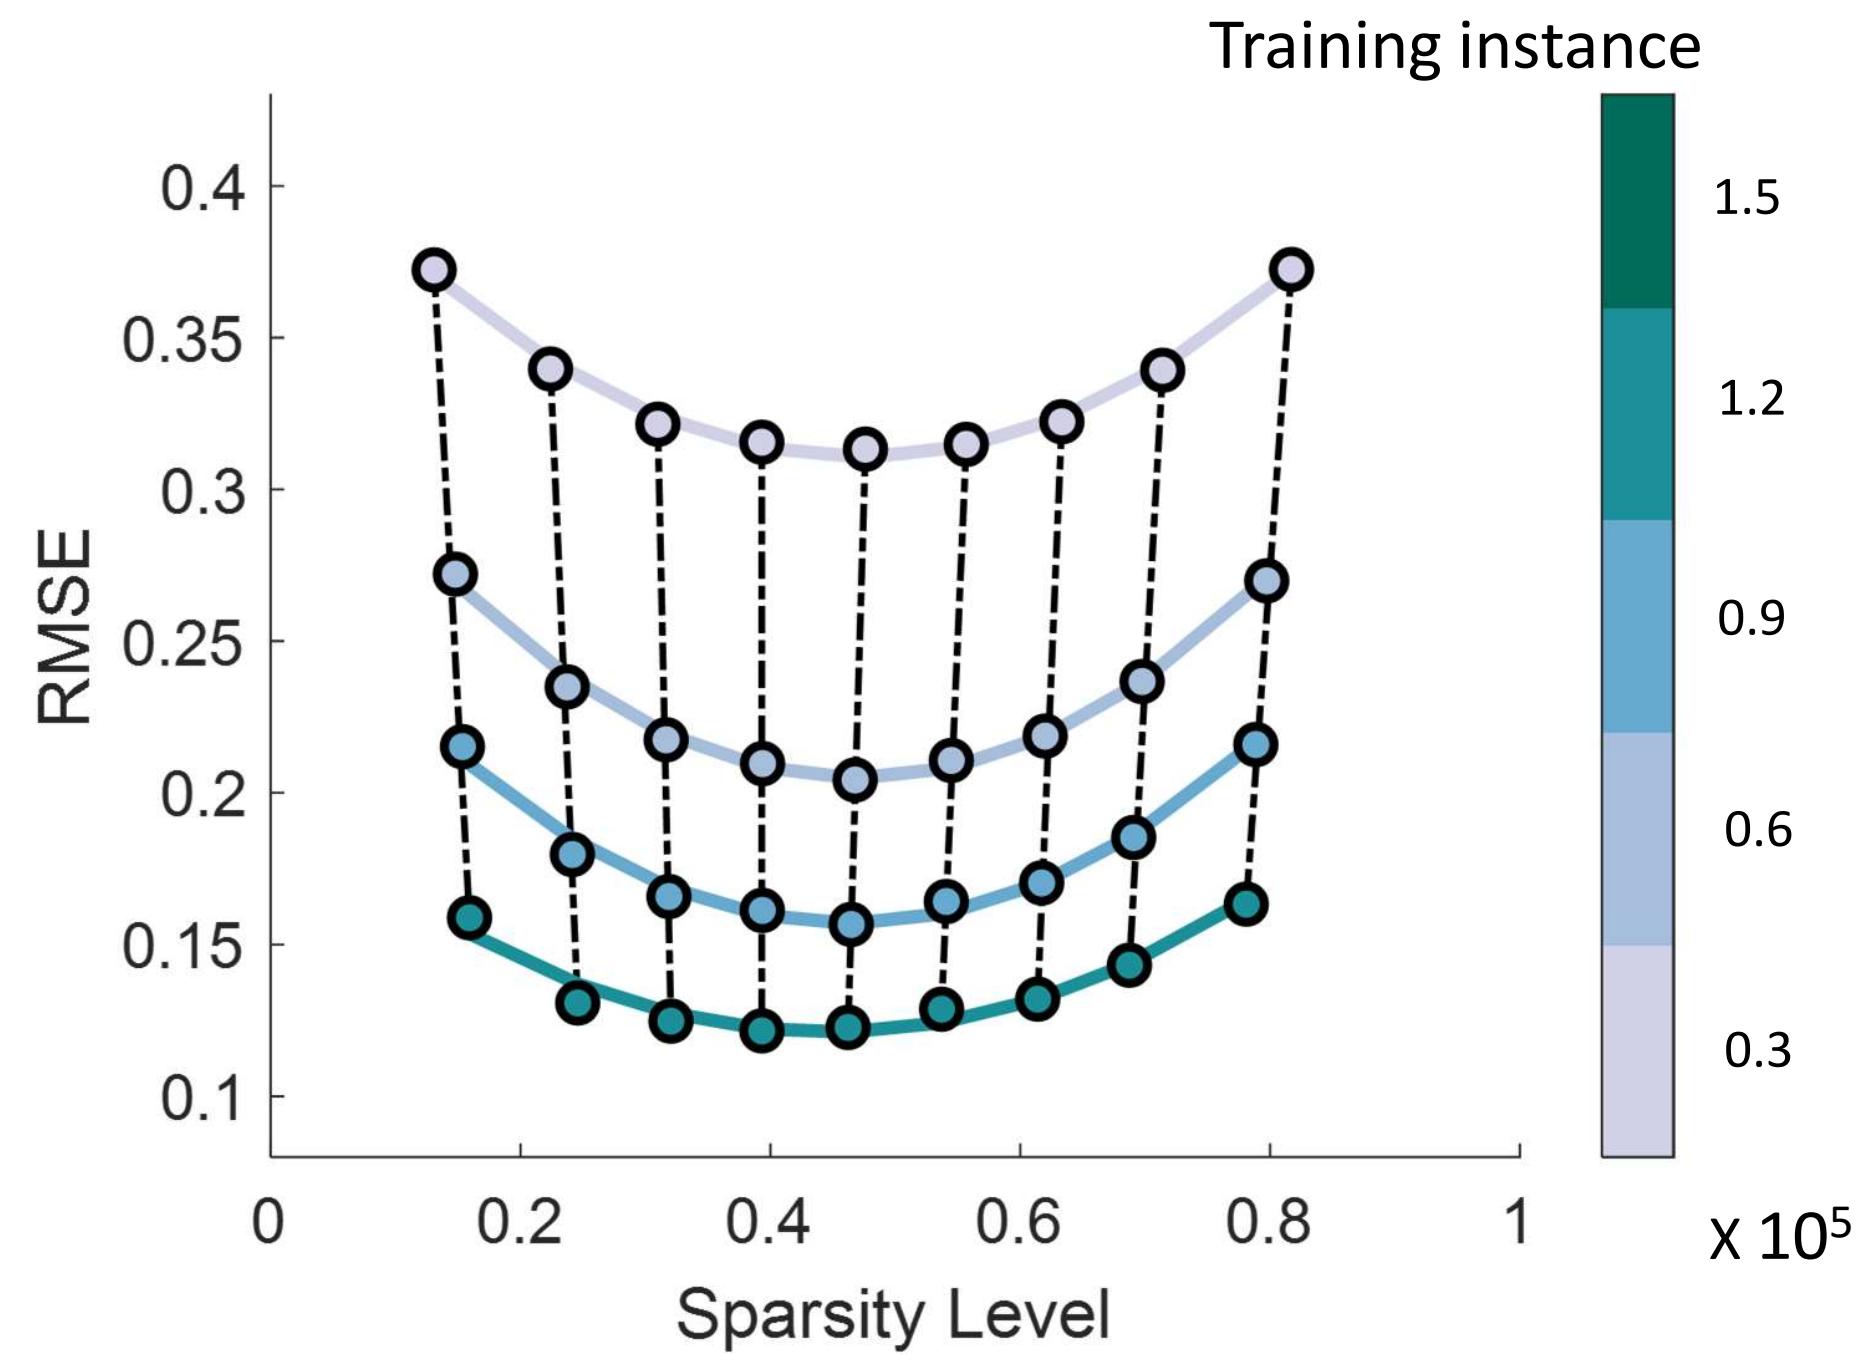
NUMBER OF LEARNABLE PARAMETERS



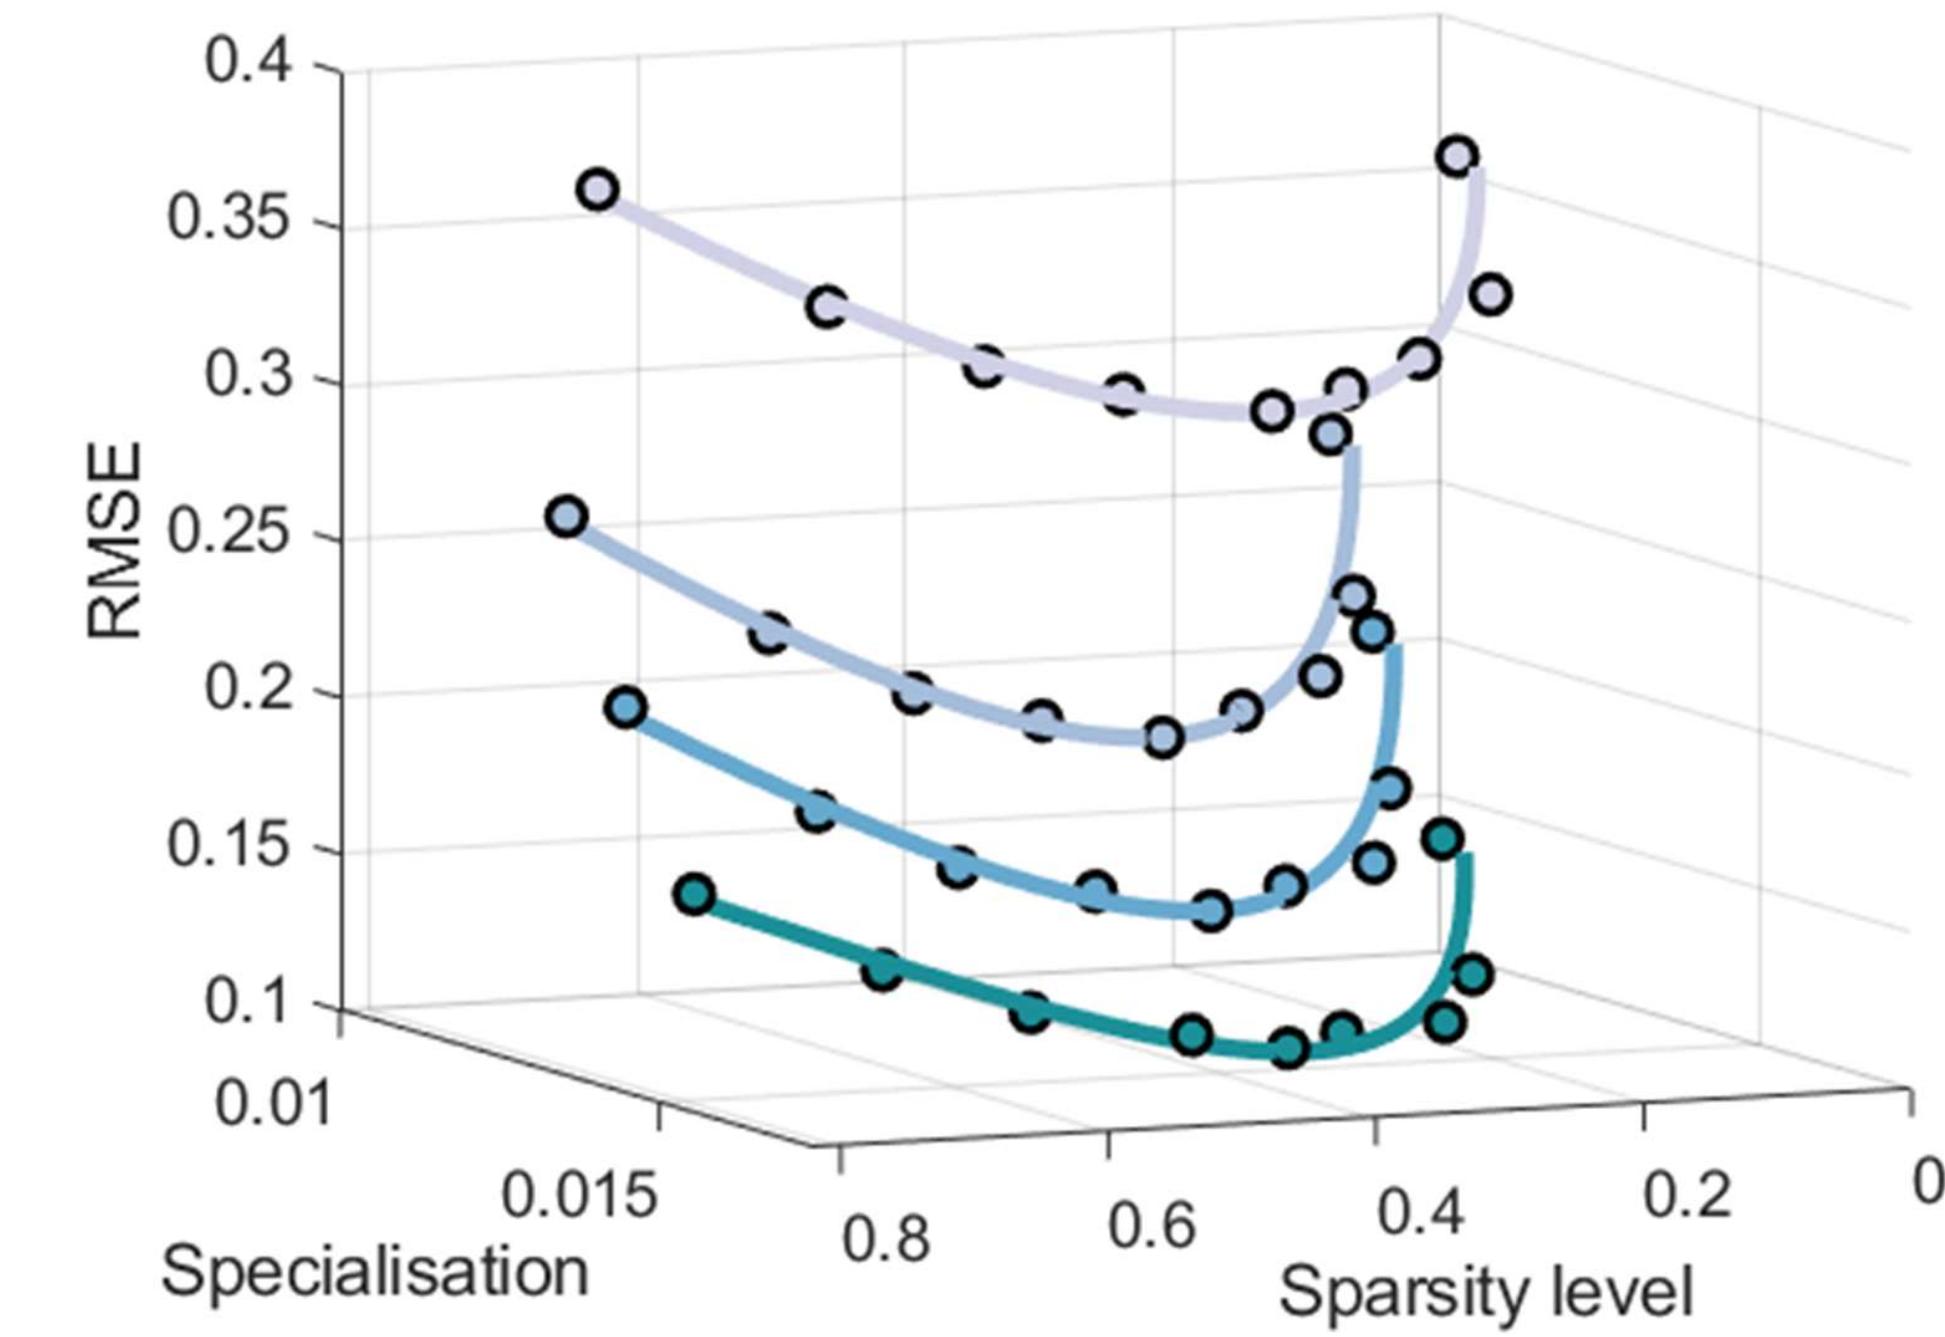
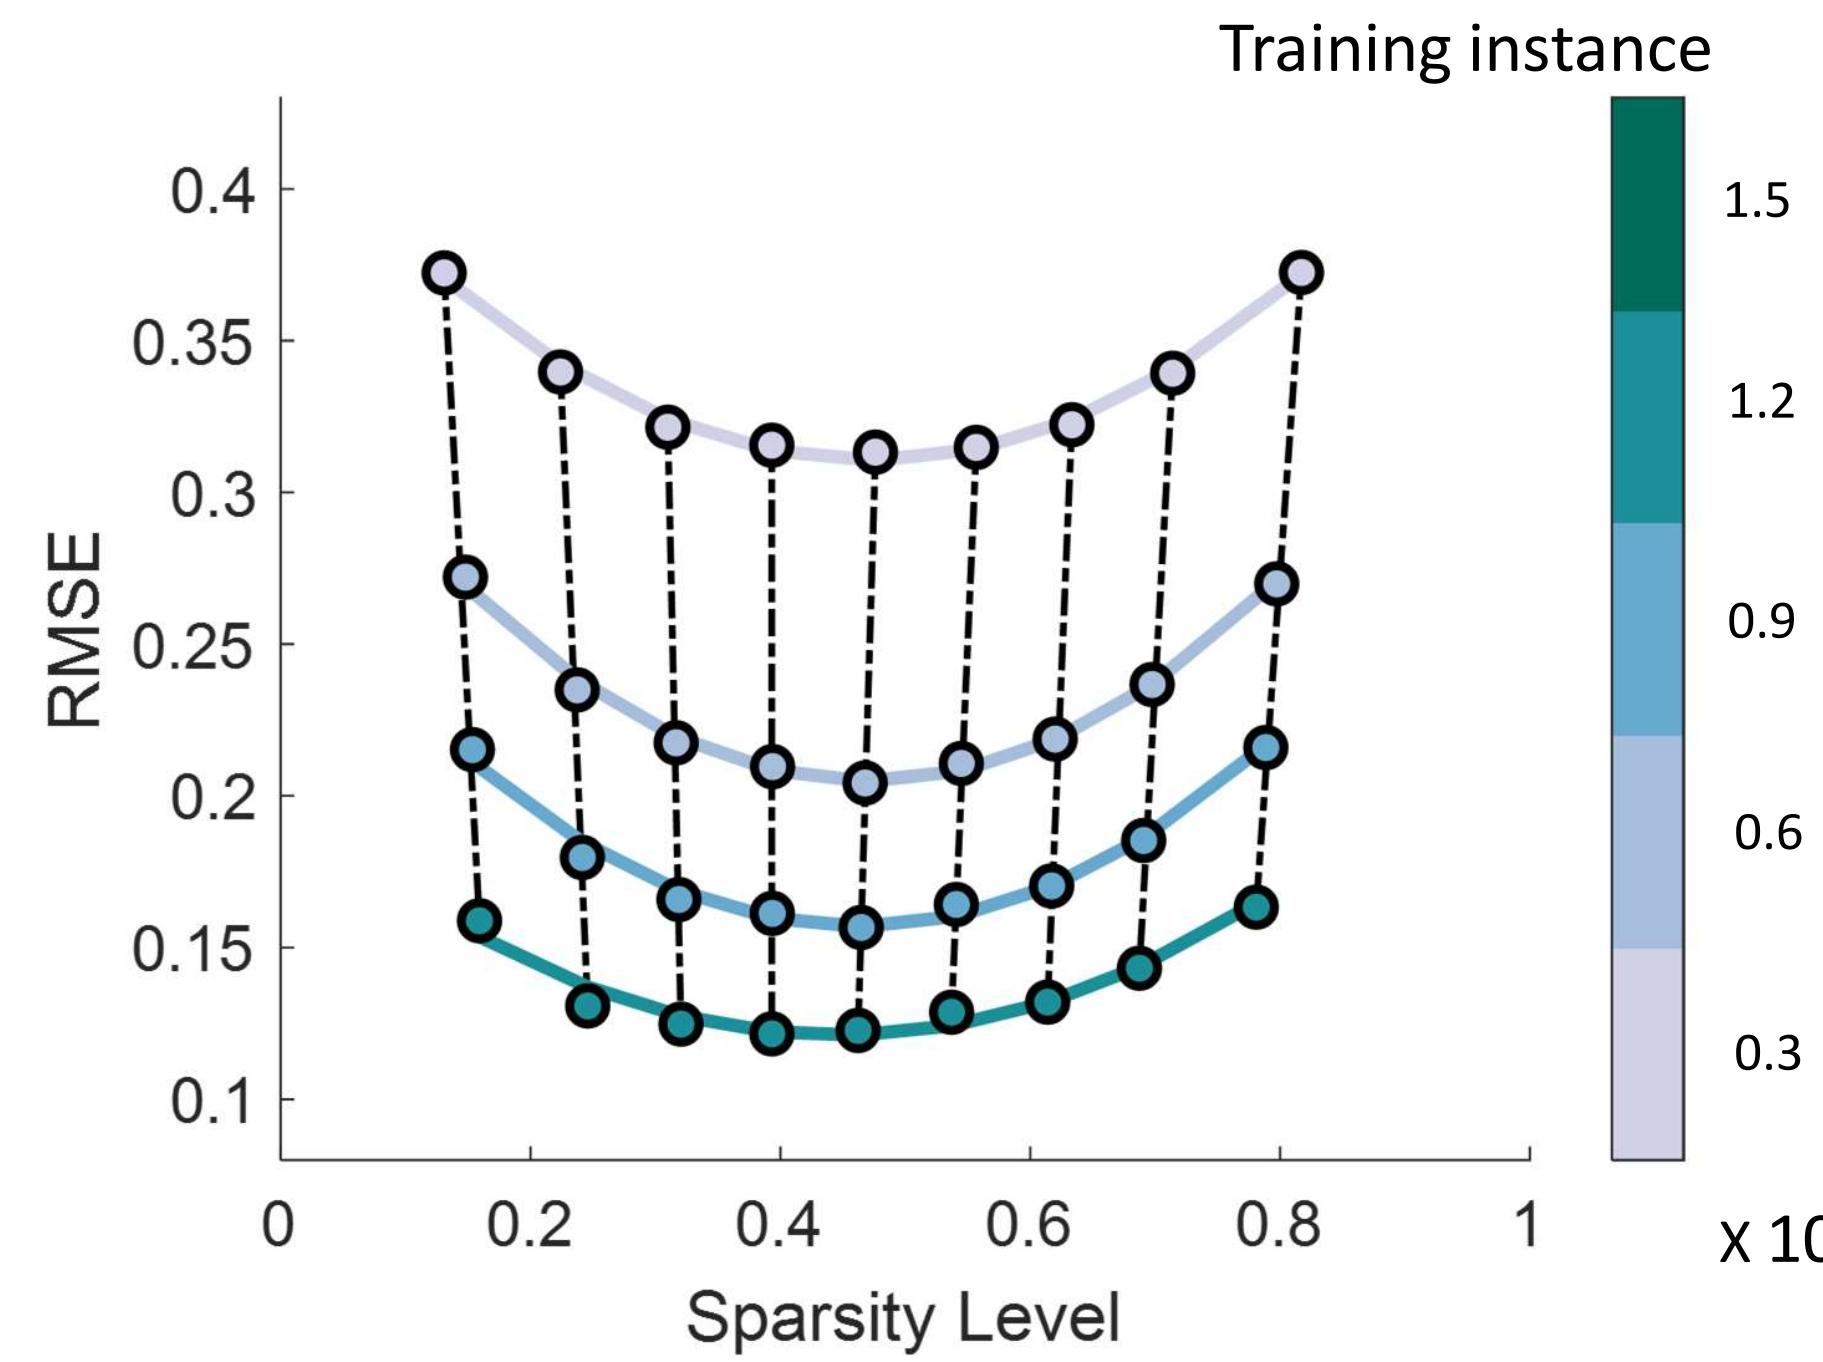
L1 DOES NOT IMPROVE PERFORMANCE



MINIMISING RMSE MODULATES THE NETWORK SPARSITY



MINIMISING RMSE MODULATES THE NETWORK SPARSITY

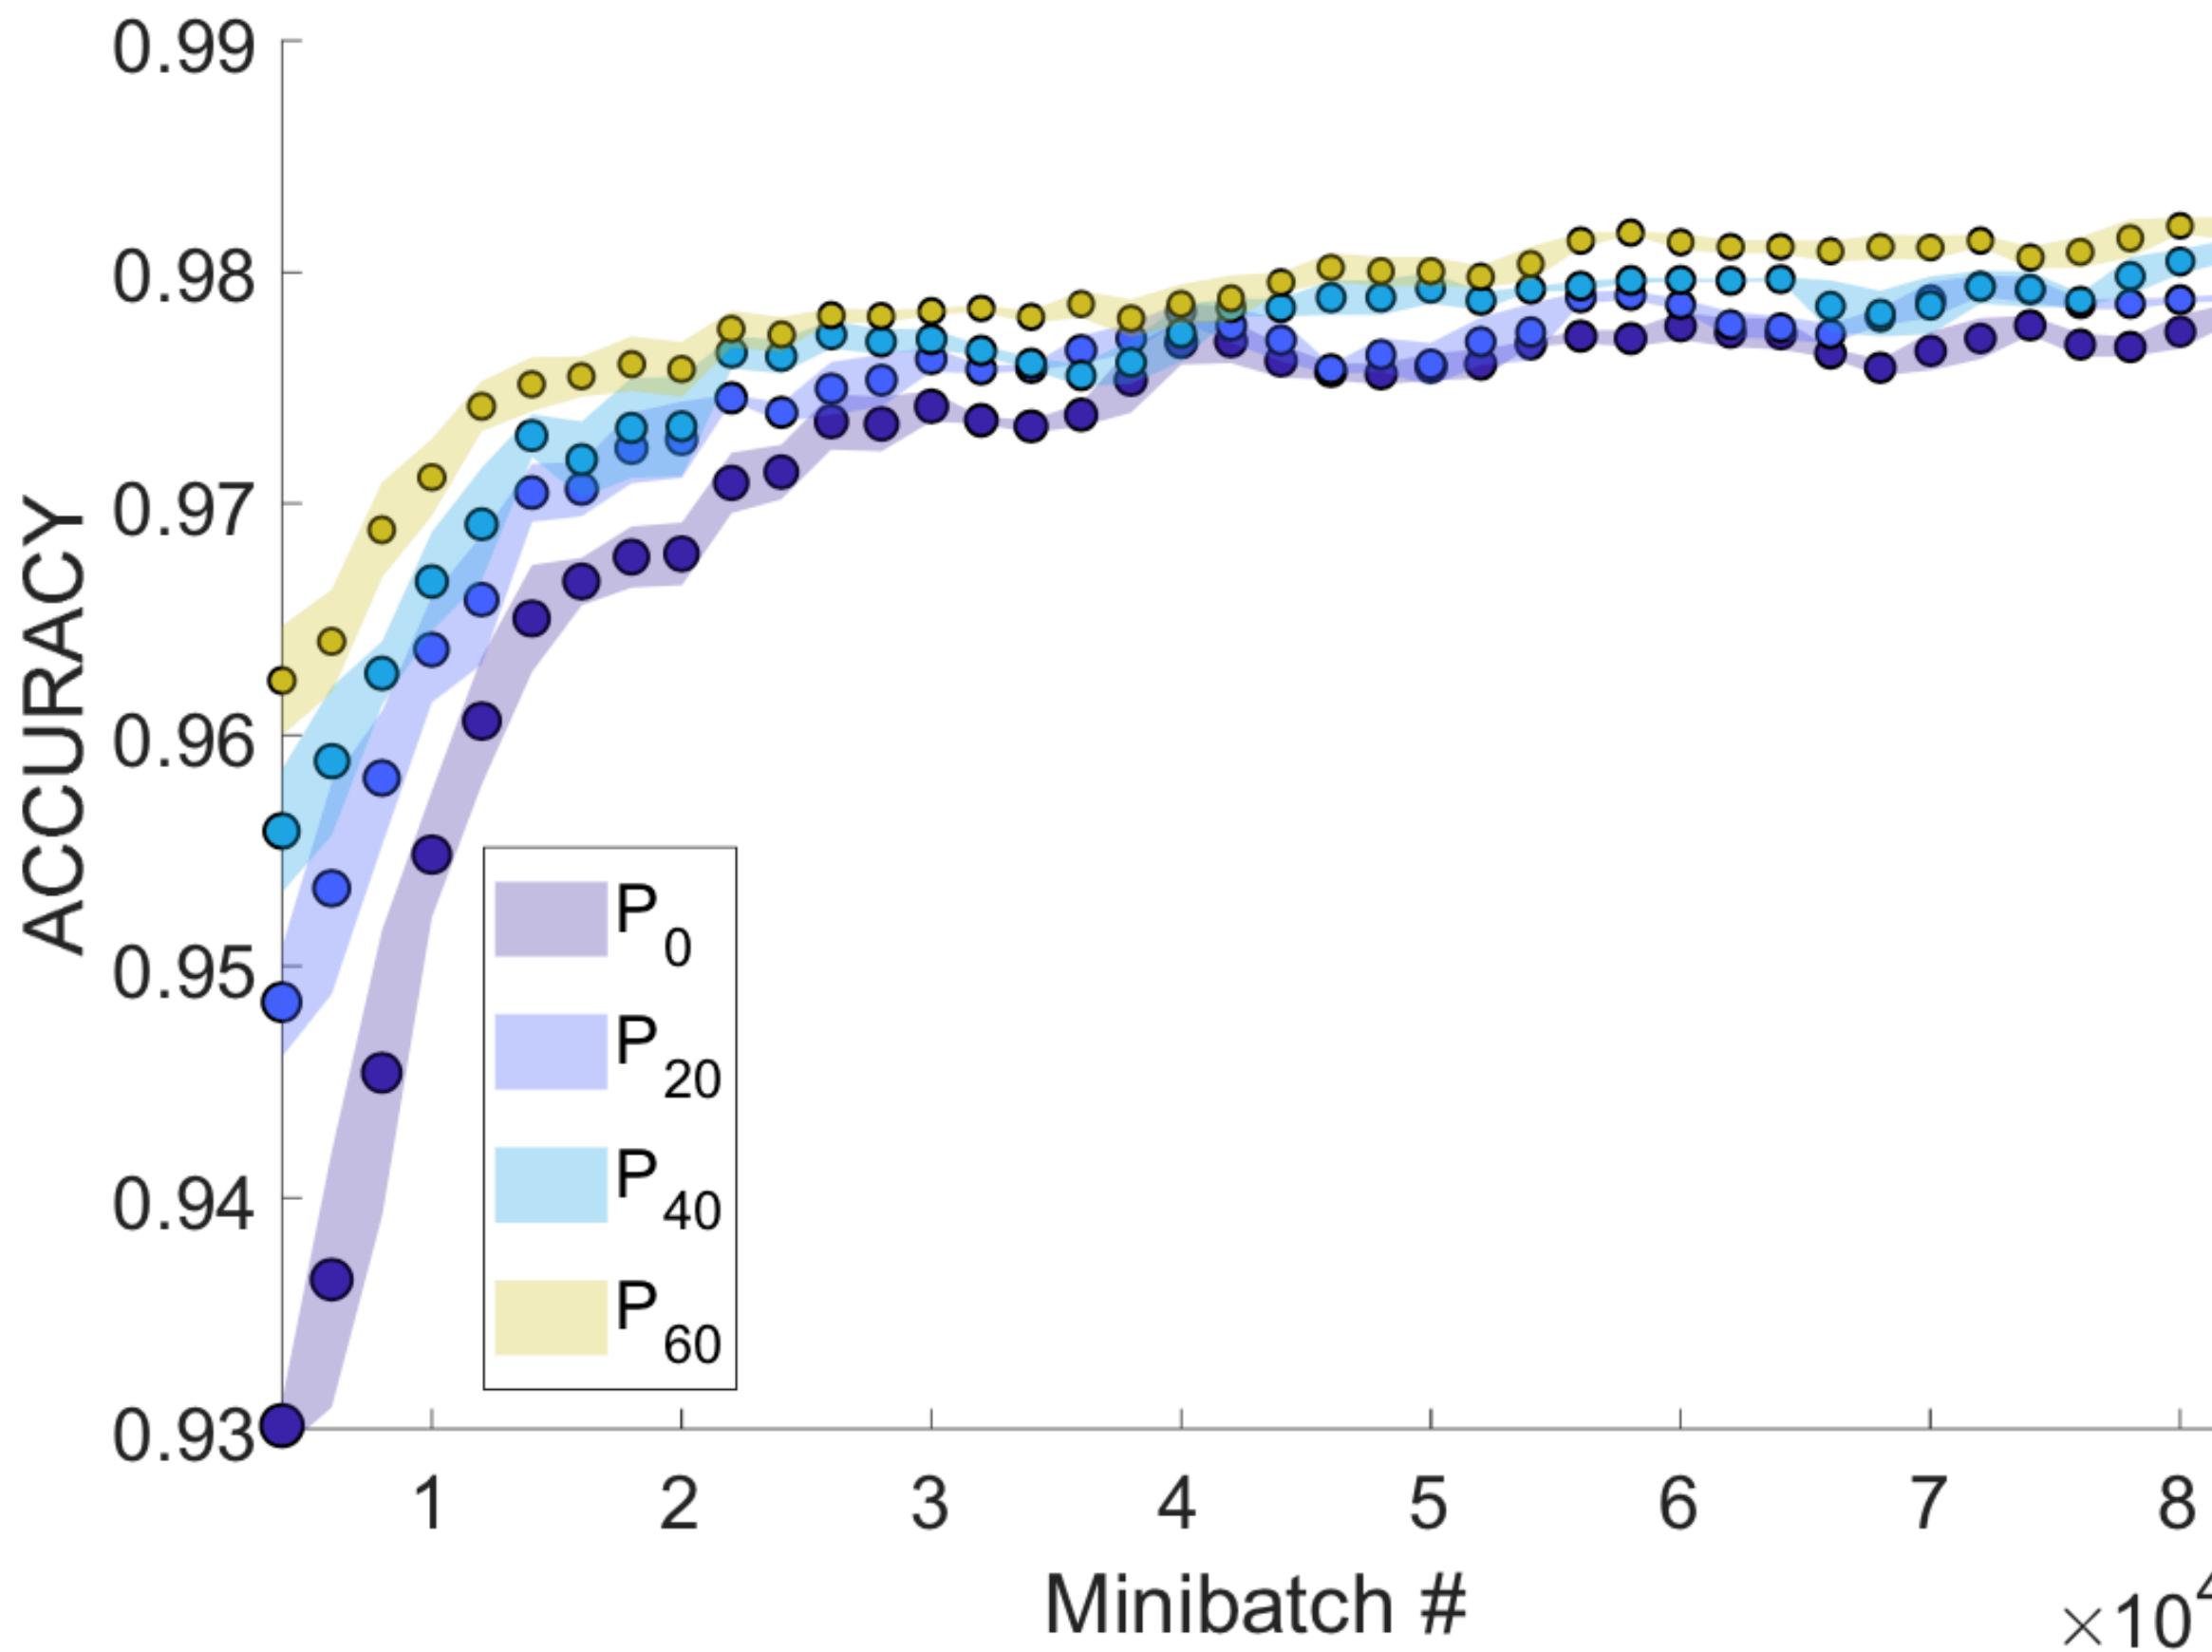


BENCHMARK: SEQUENTIAL (COLUMN BY COLUMN) MNIST

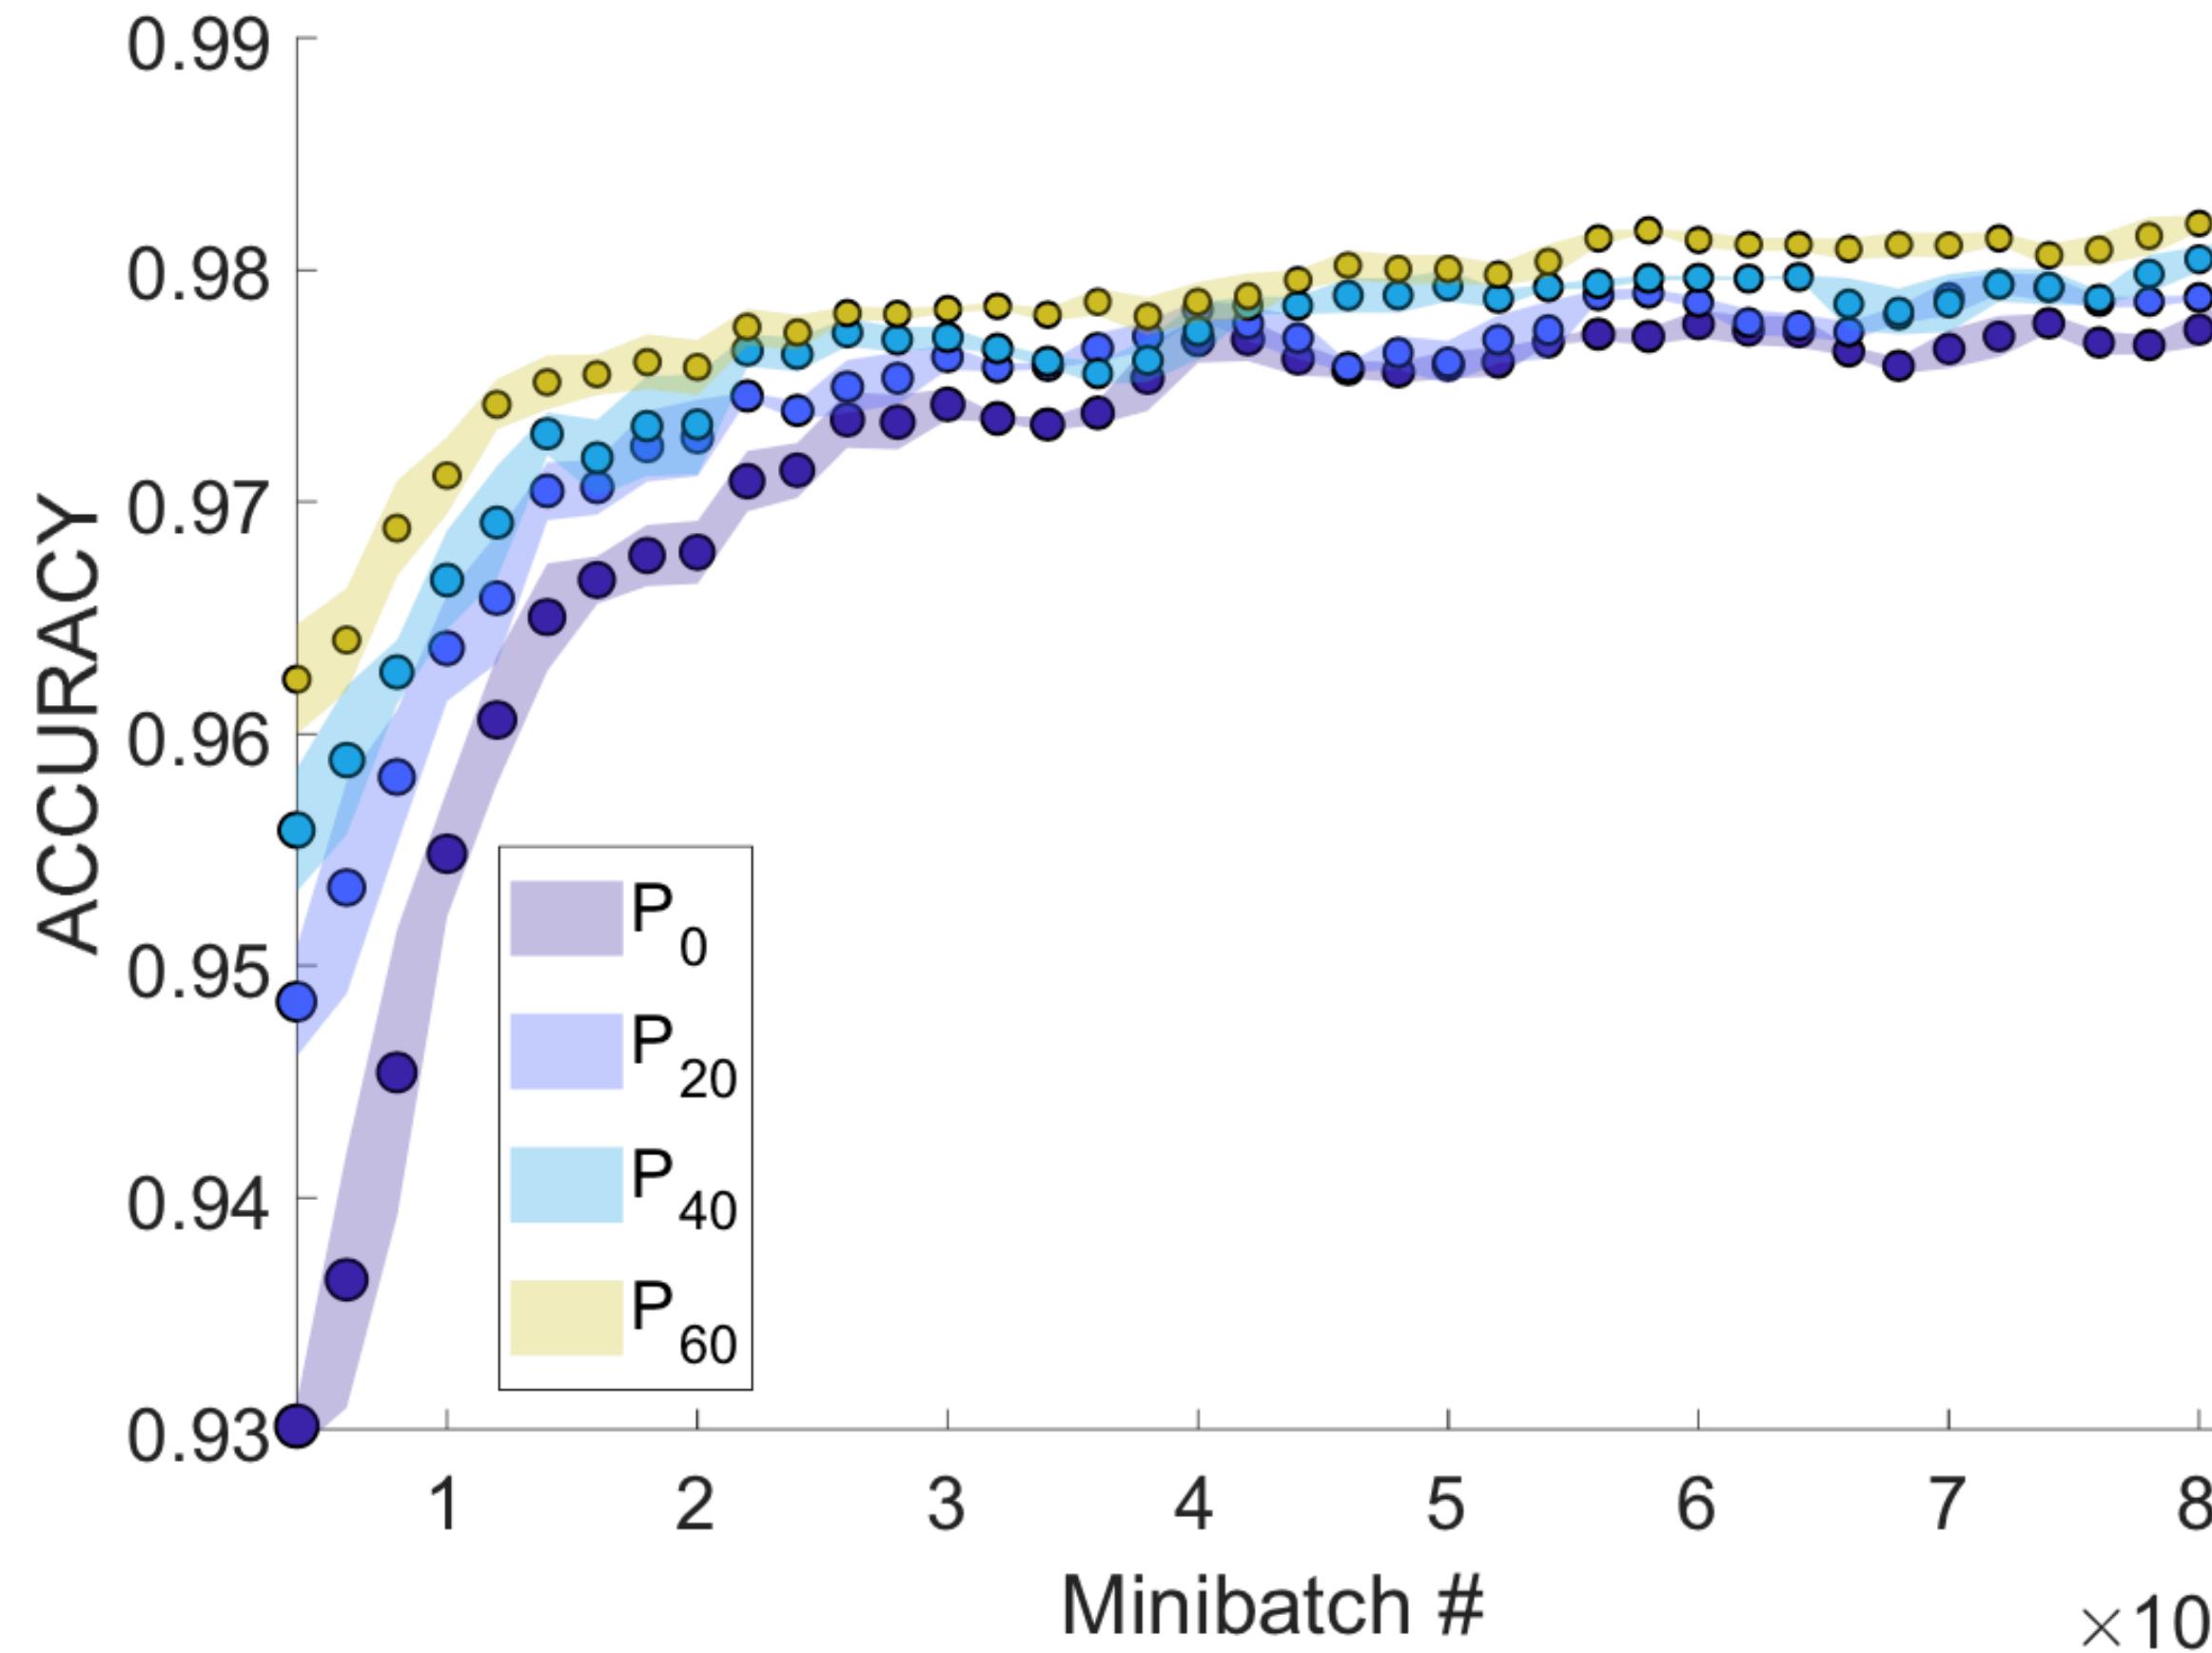
7	9	7	5	7
5	8	6	0	7
9	8	8	9	1
7	8	3	2	8
4	6	6	3	6
6	0	2	3	7
9	2	7	3	5
6	7	6	9	2

MNIST database - LeCun et al., Proceedings of the IEEE, 1998

PERFORMANCE: SEQUENTIAL (COLUMN BY COLUMN) MNIST

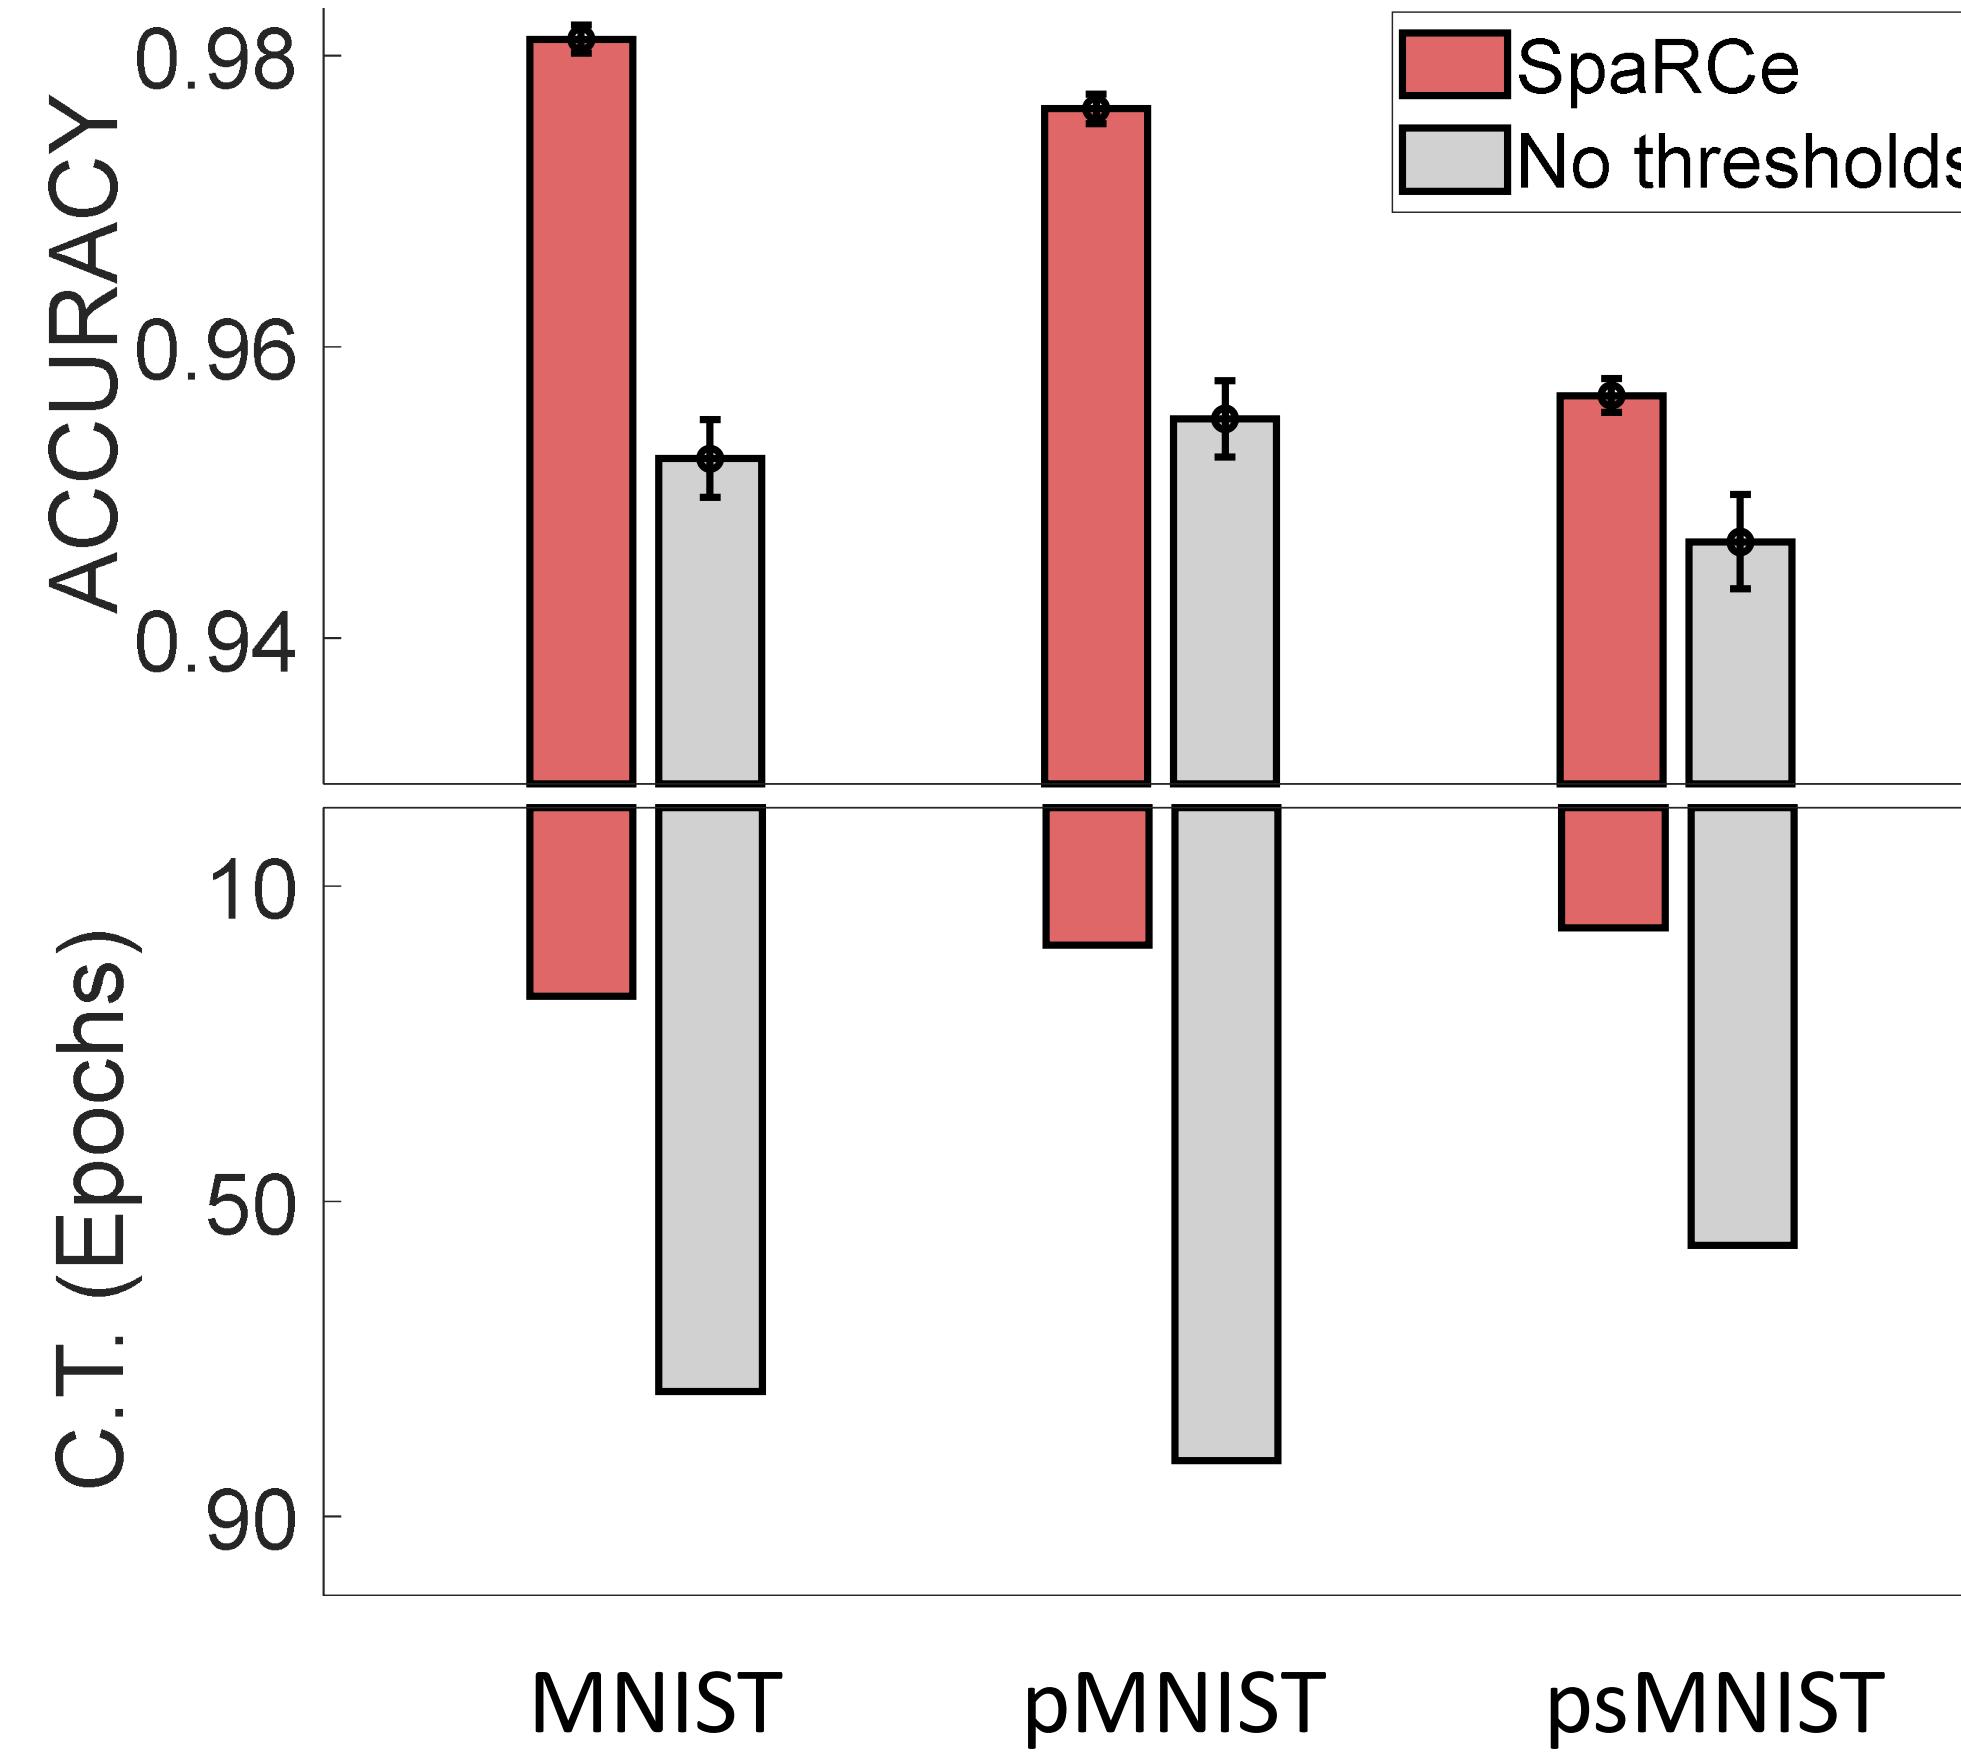


PERFORMANCE: SEQUENTIAL (COLUMN BY COLUMN) MNIST

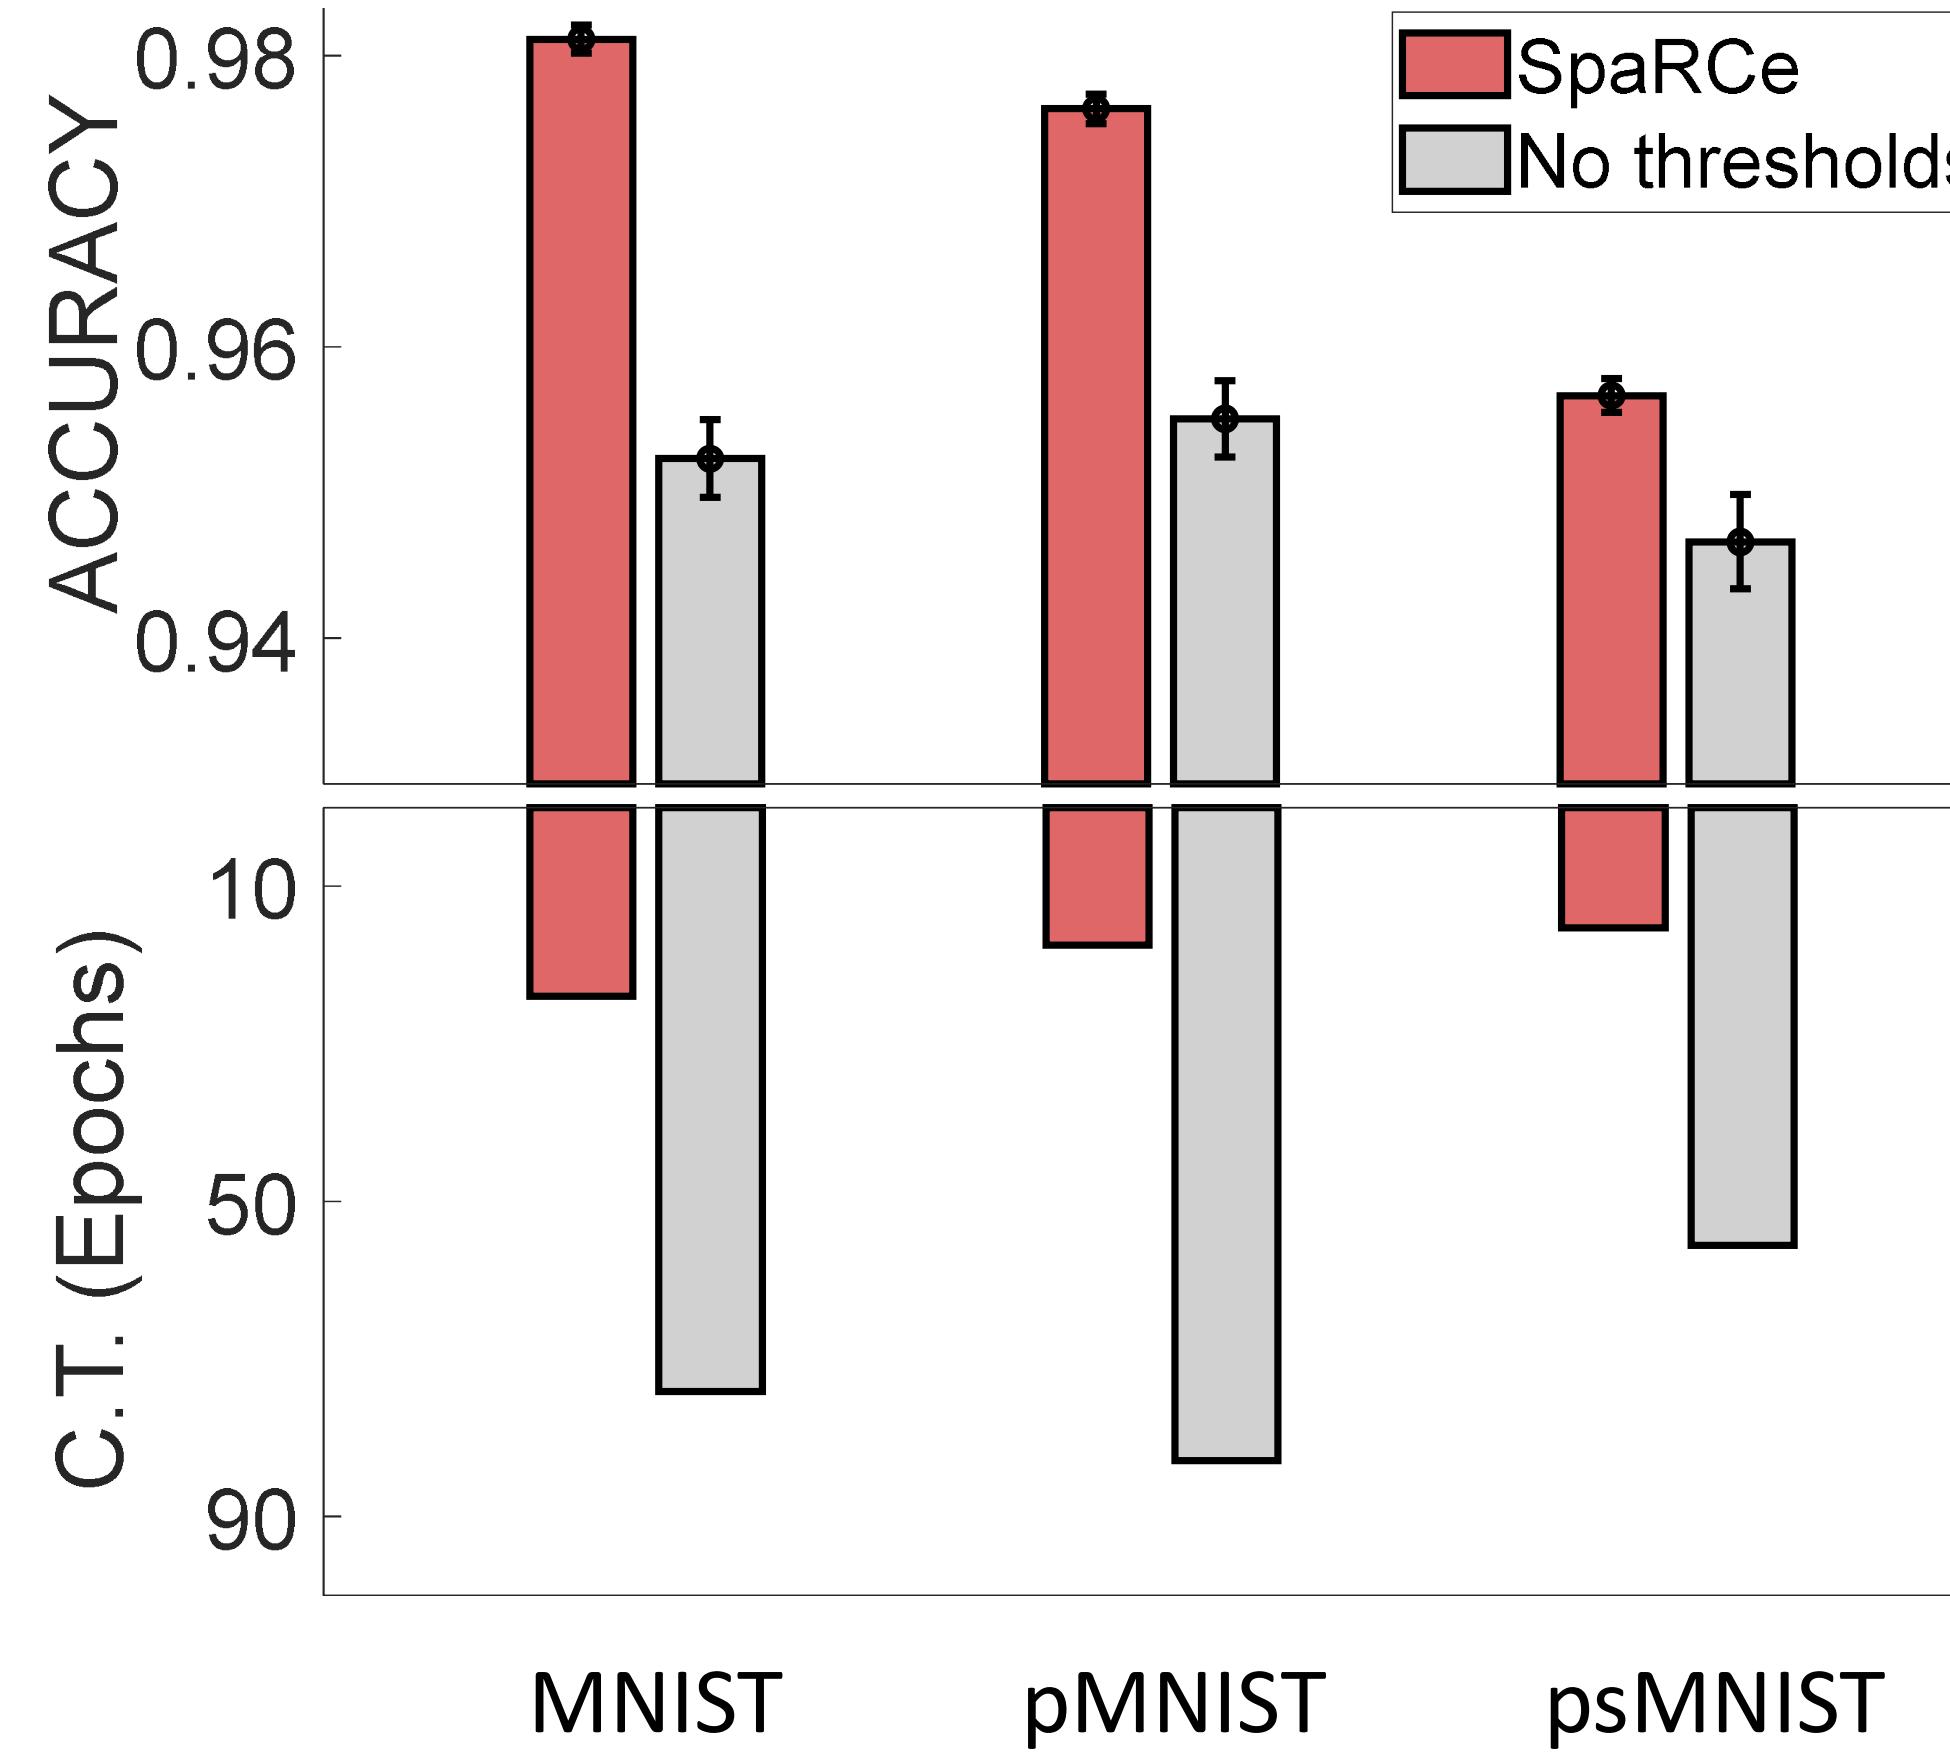


In practice we can start at P_{50} and adjust learning rate to reach maximum performance.

PERFORMANCE COMPARISON ON MNIST VARIADS



PERFORMANCE COMPARISON ON MNIST VARIADS

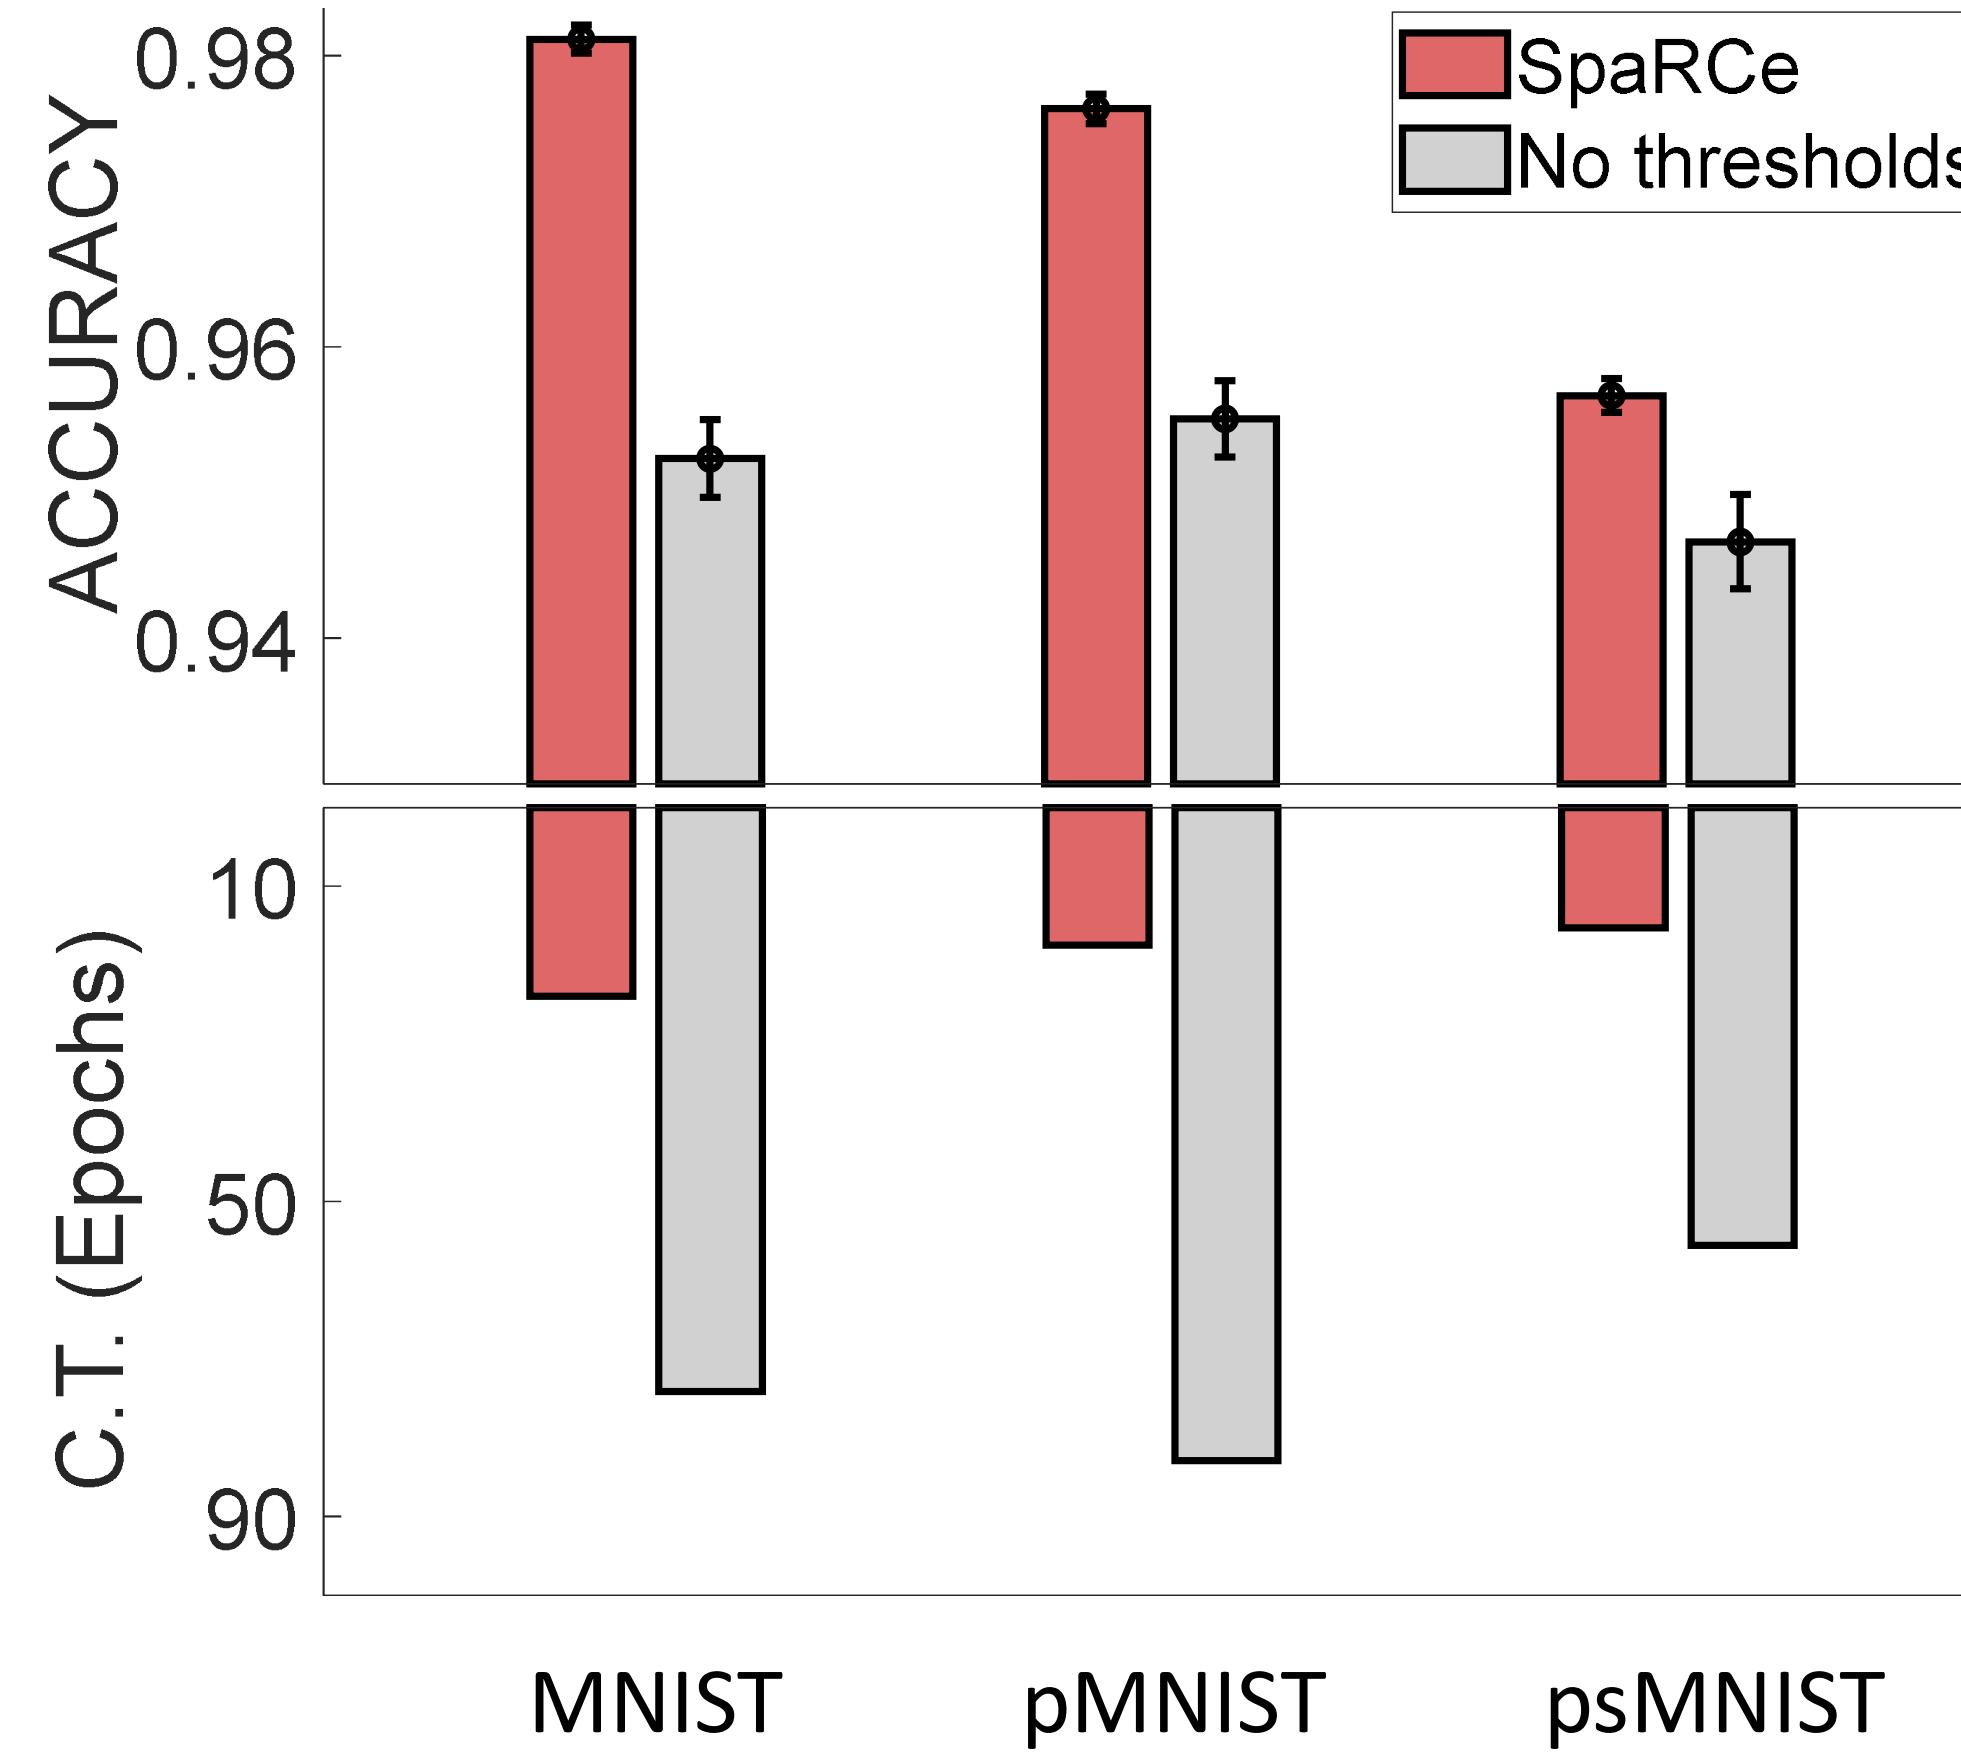


MNIST	
<i>ESN</i>	95.2
<i>SpaRCe</i>	98.1
<i>MLP</i>	97.0 98.5*
<i>Conv.</i>	98.3 99.6*

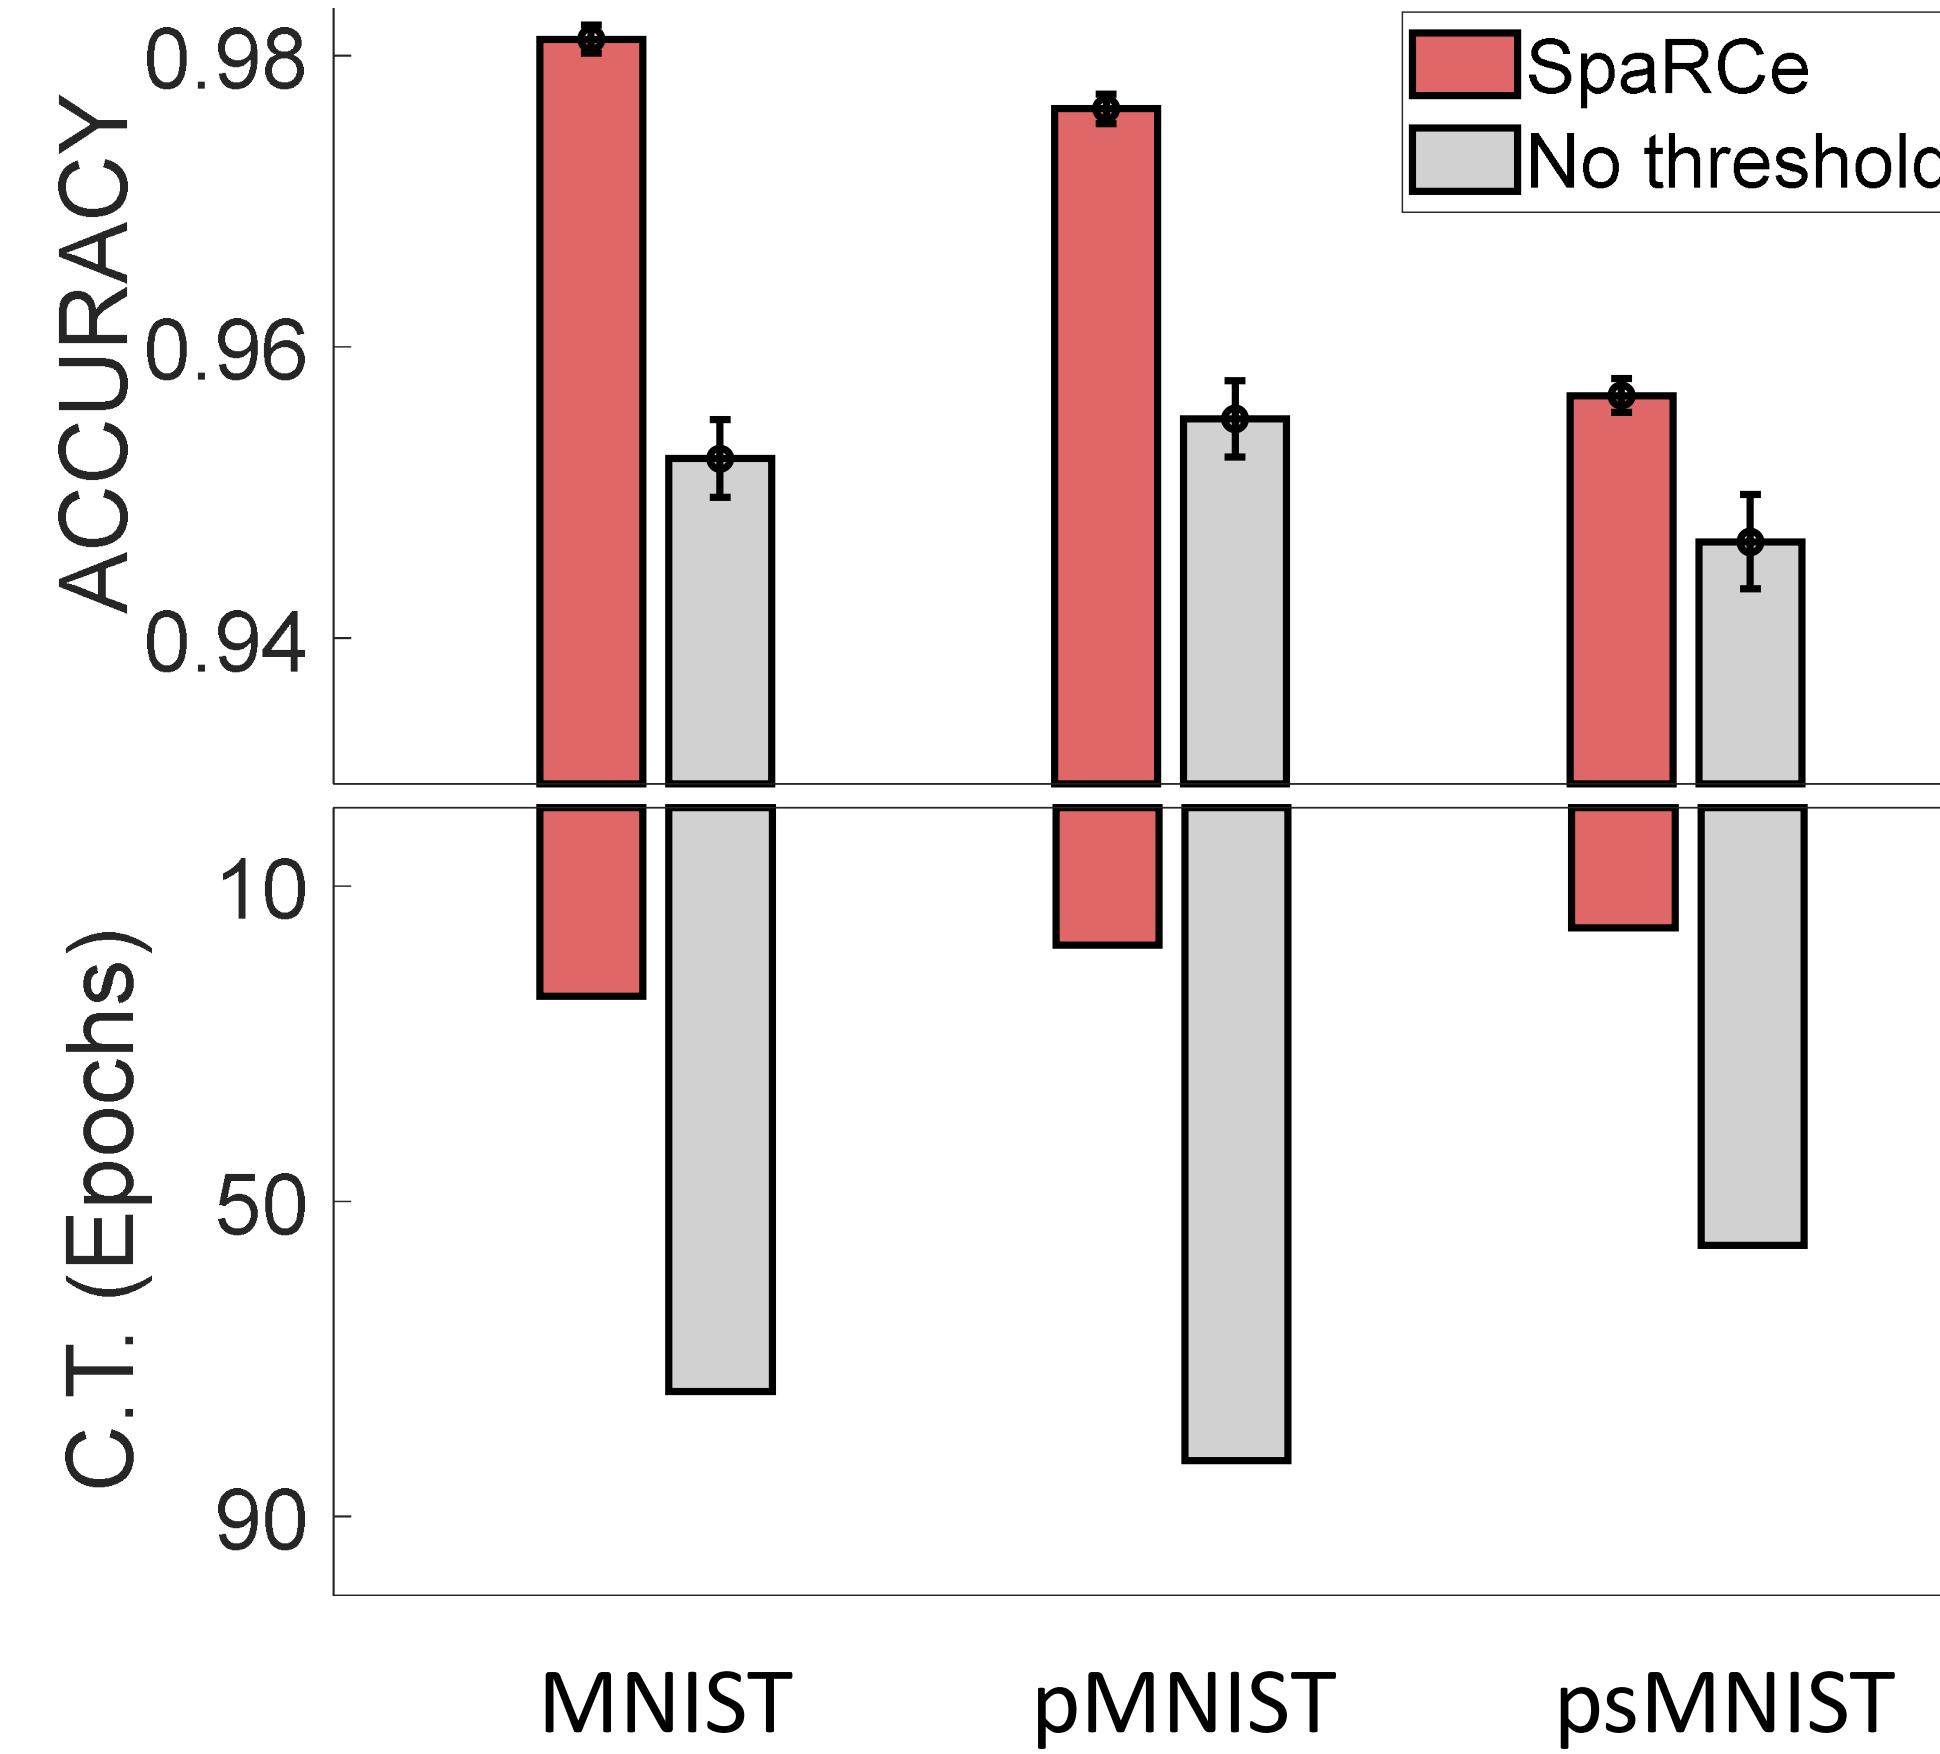
*Deng (2012) IEEE Signal Processing Magazine

** Chandar et al (2019) AAAI

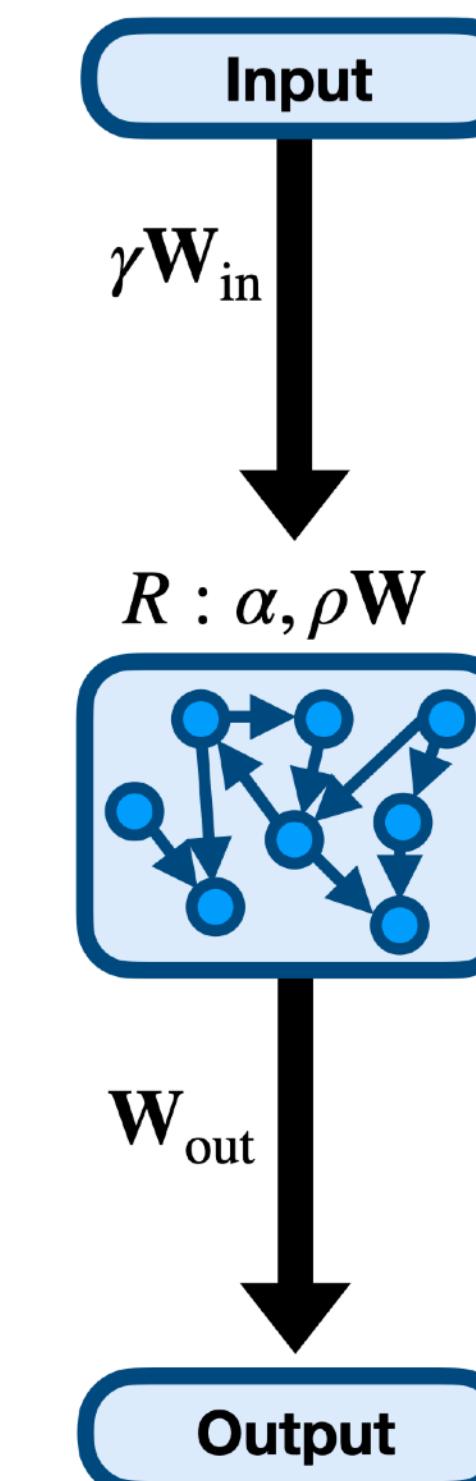
PERFORMANCE COMPARISON ON MNIST VARIADS



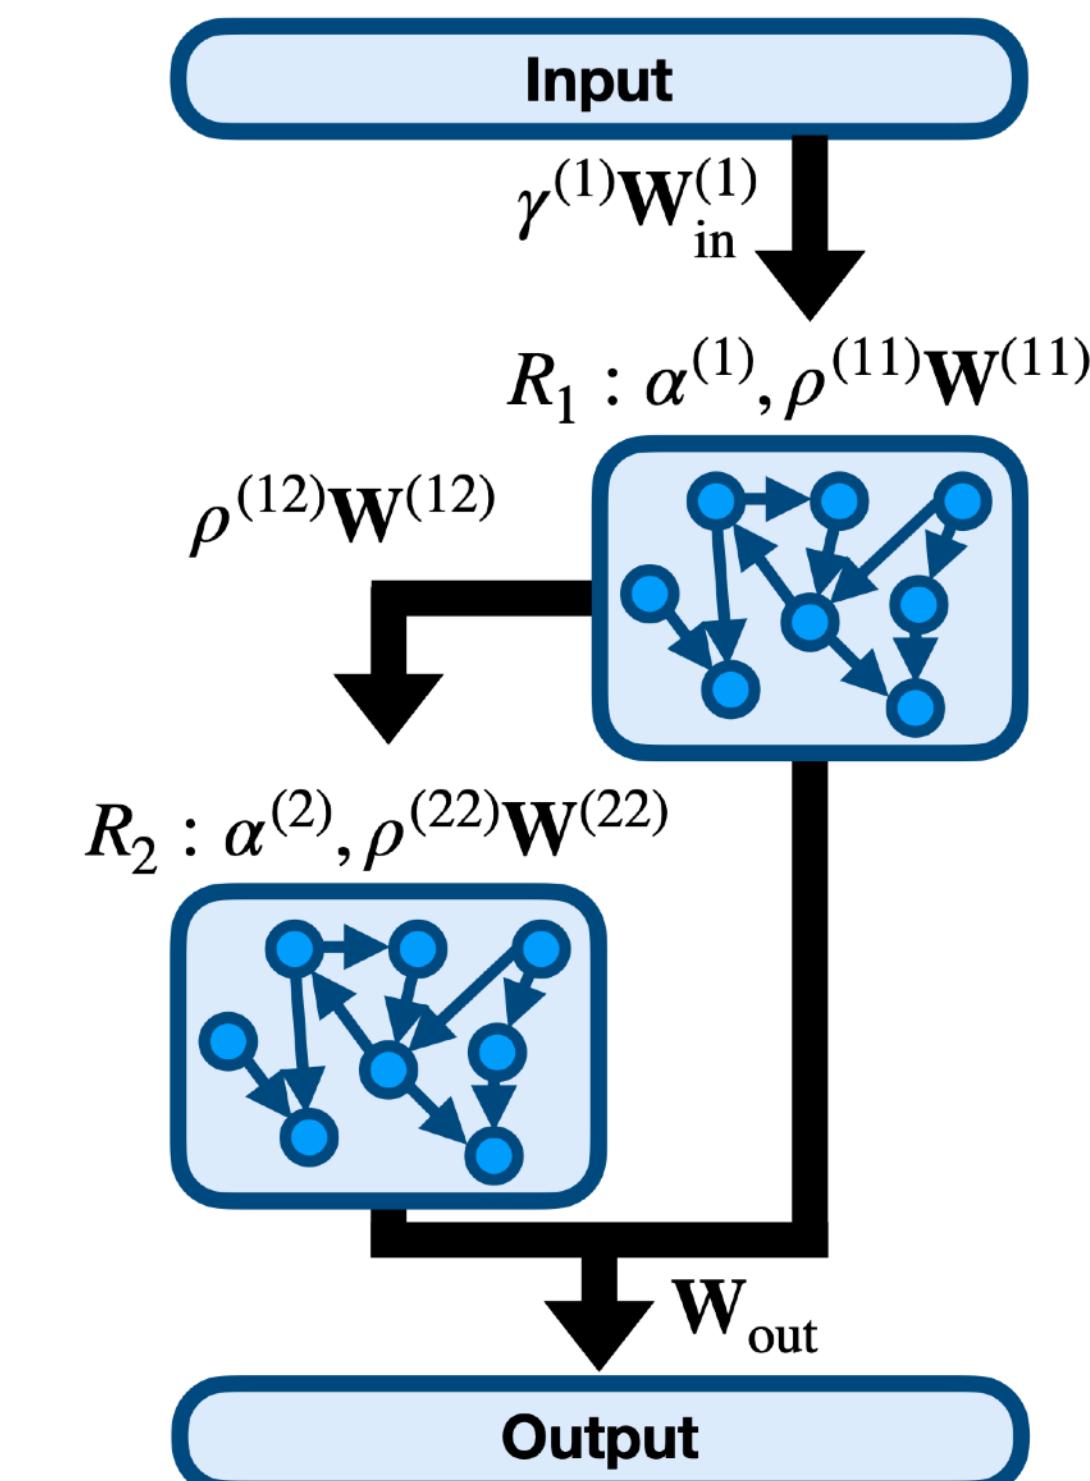
PERFORMANCE COMPARISON ON MNIST VARIADS



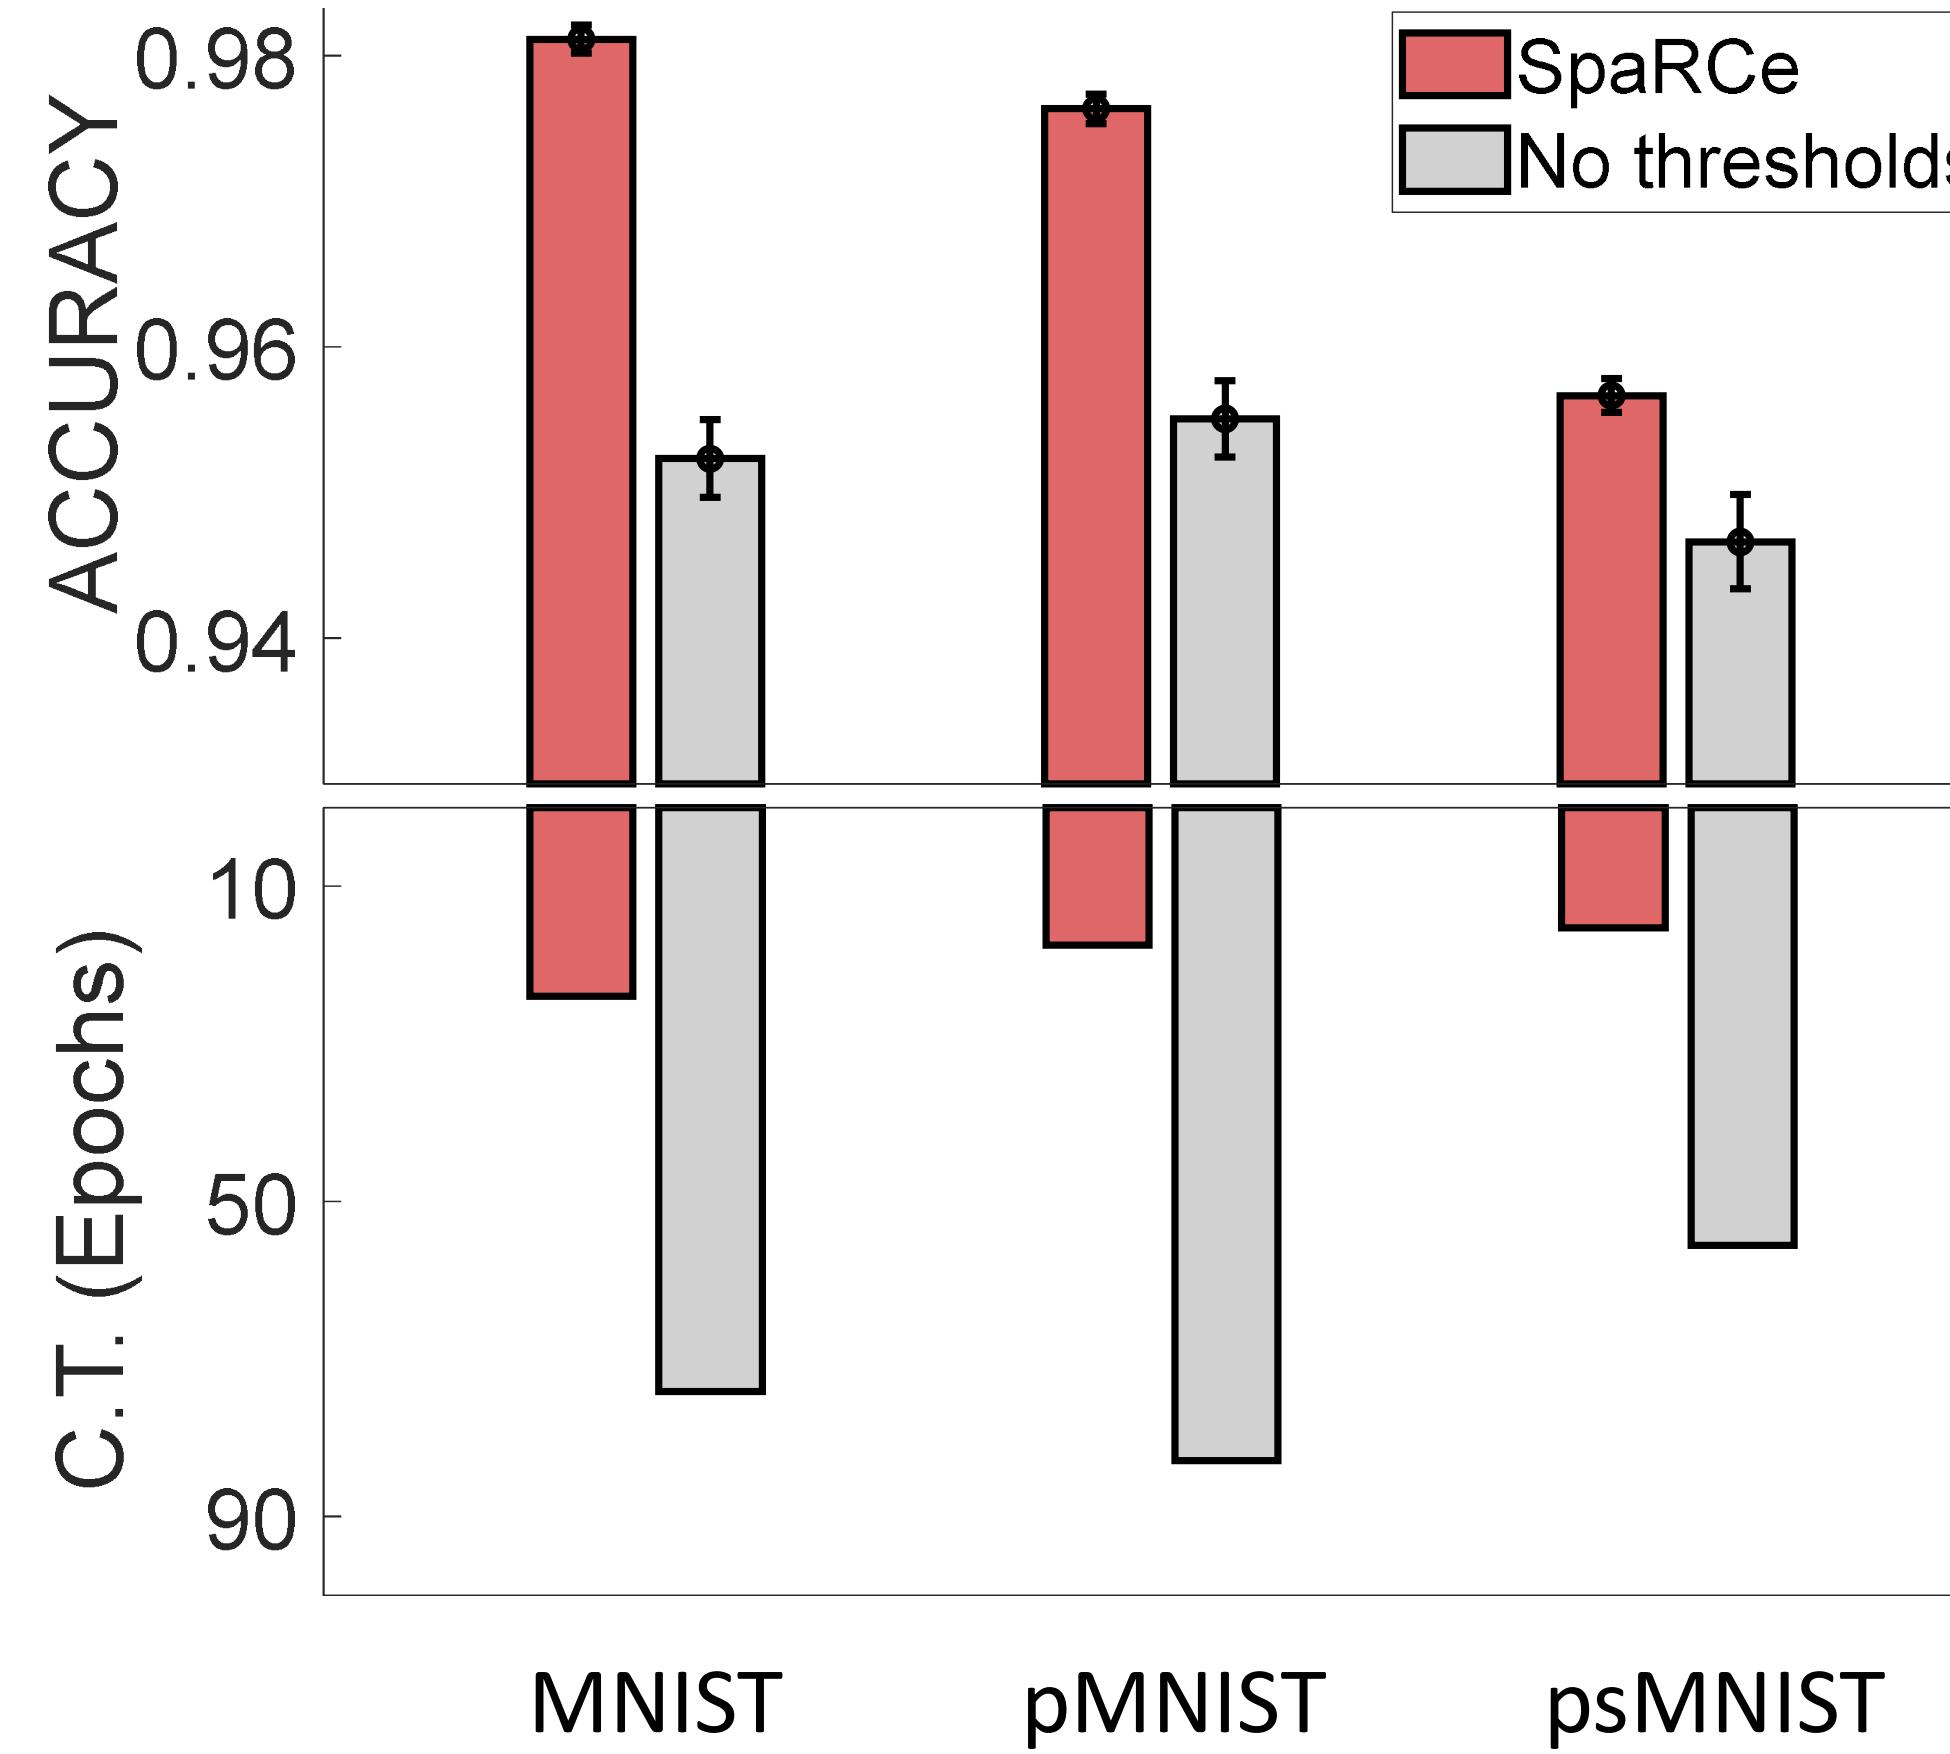
A. ESN



B. Hierarchical ESN



PERFORMANCE COMPARISON ON MNIST VARIADS

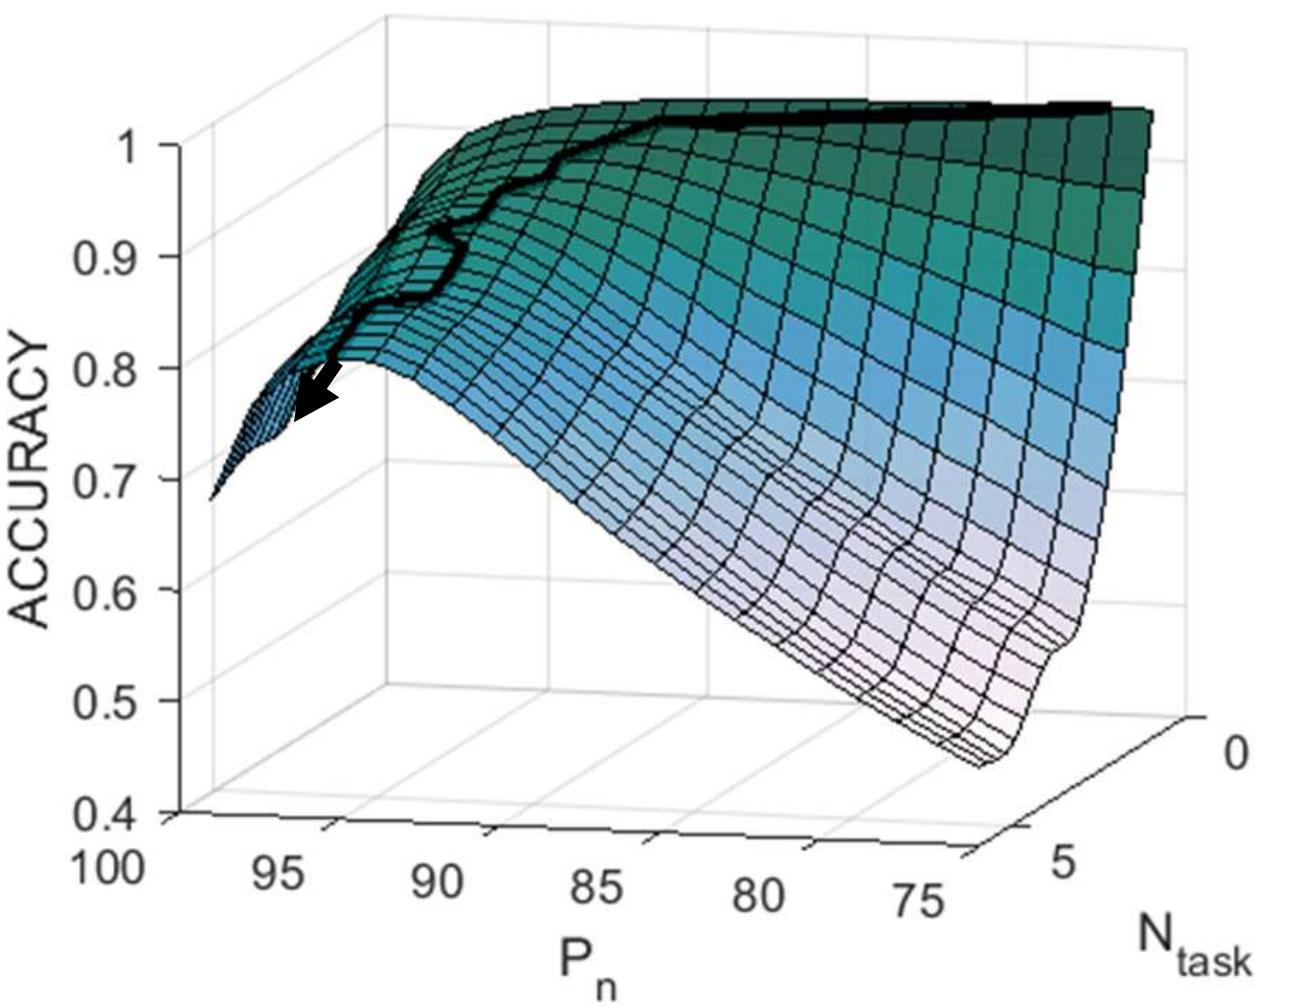
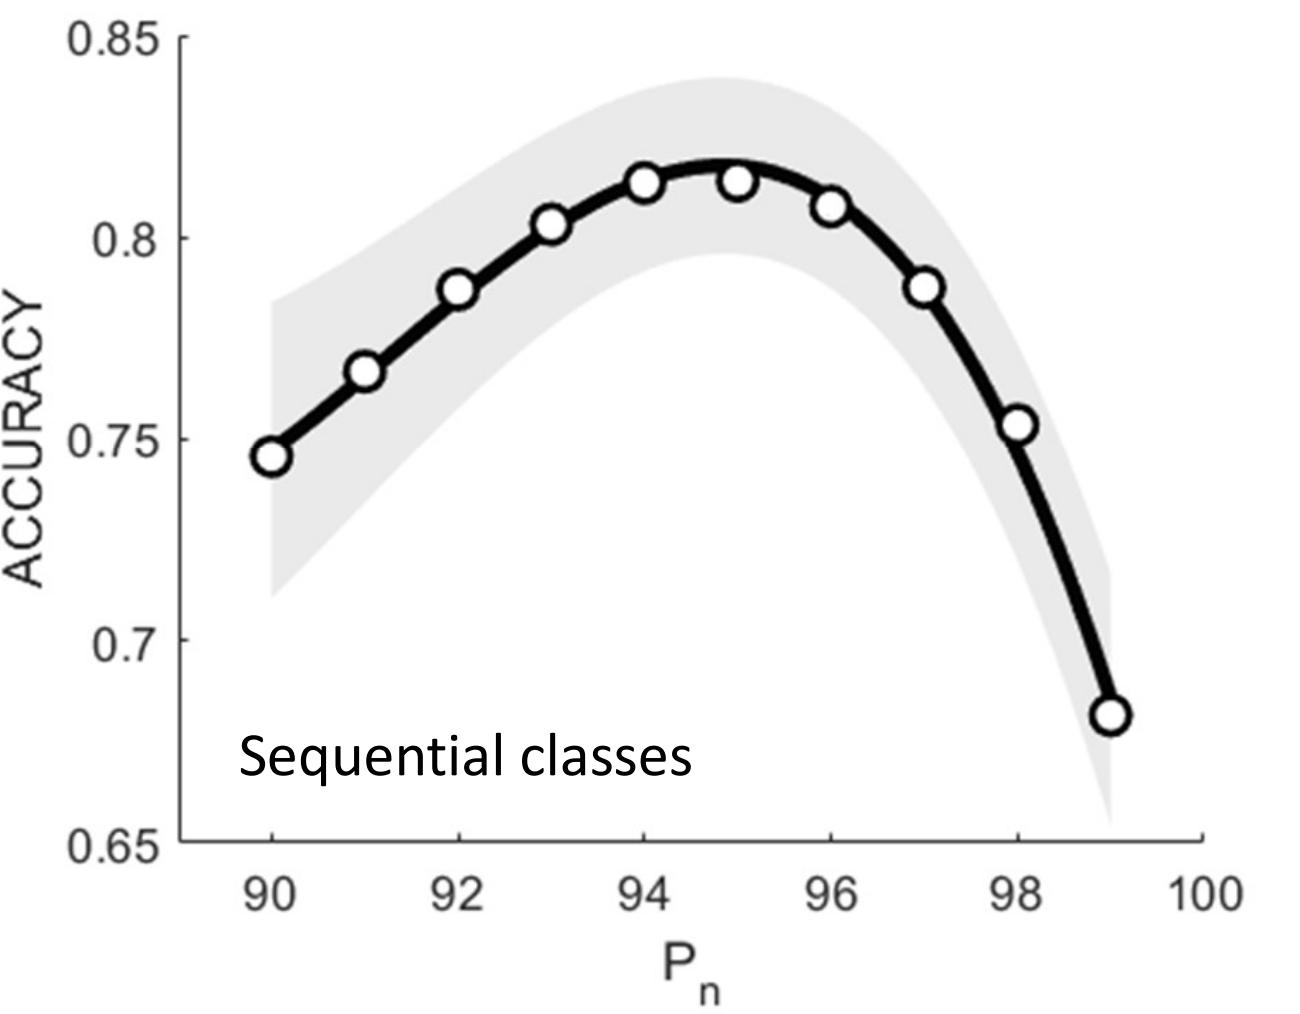


psMNIST			
Reservoir size	500	1000	1500
ESN^2	95.6	95.9	96.3
$SpaRCe^2$	96.2	96.6	96.9
$LSTM$	89.9**		
NRU	95.4**		

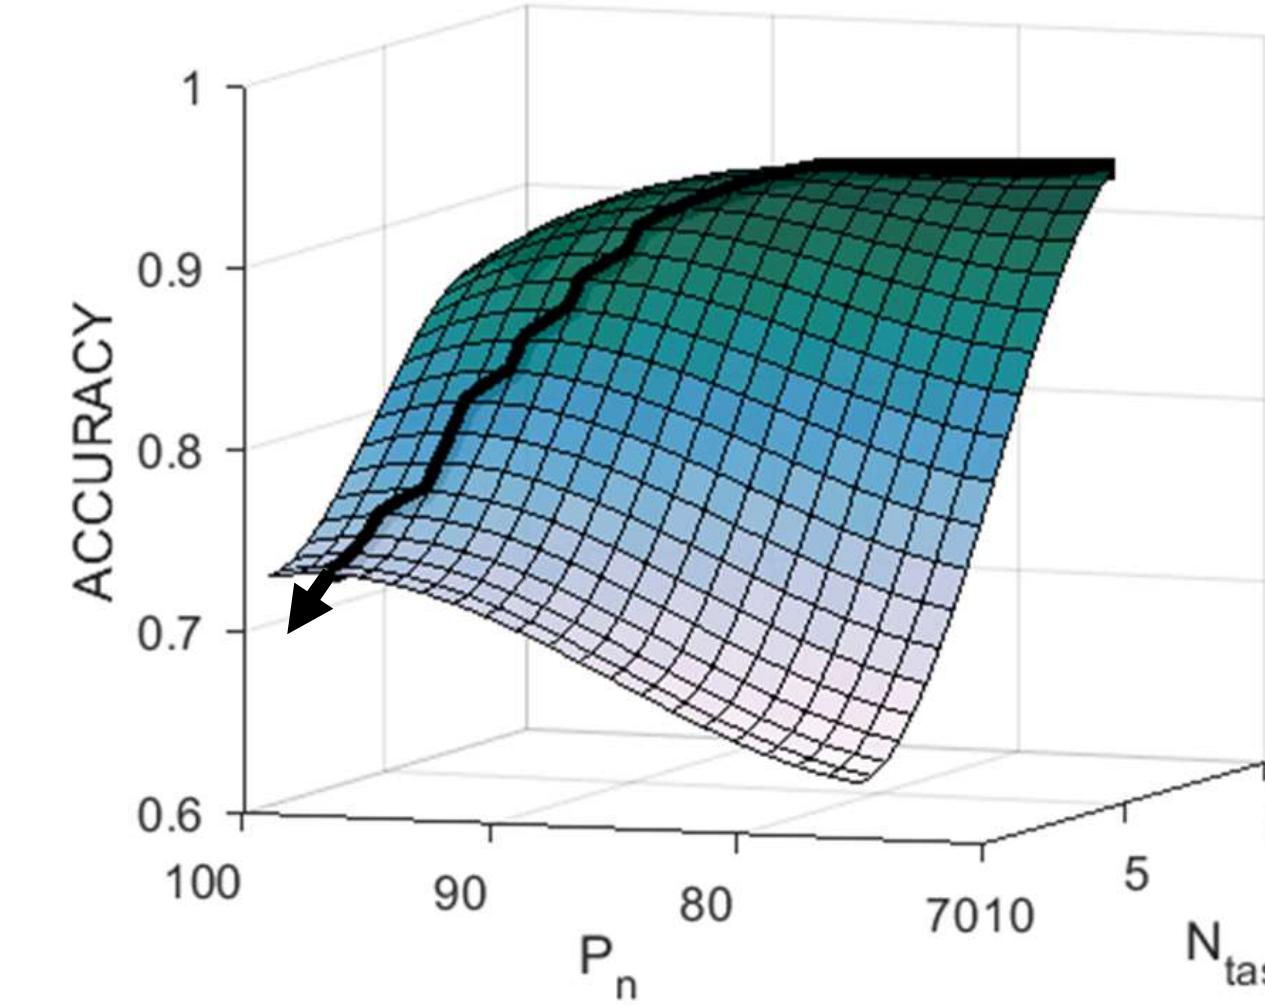
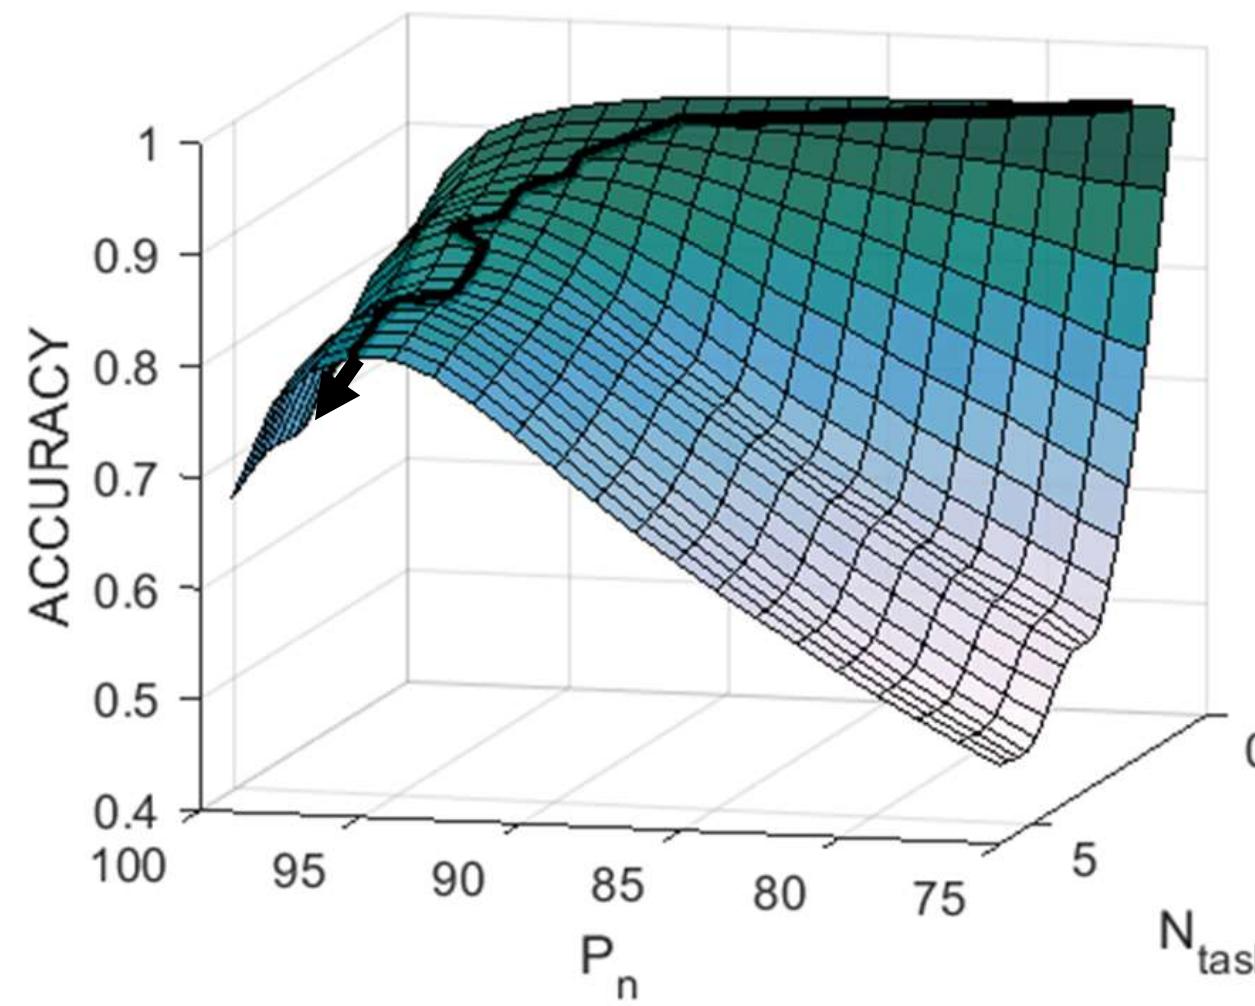
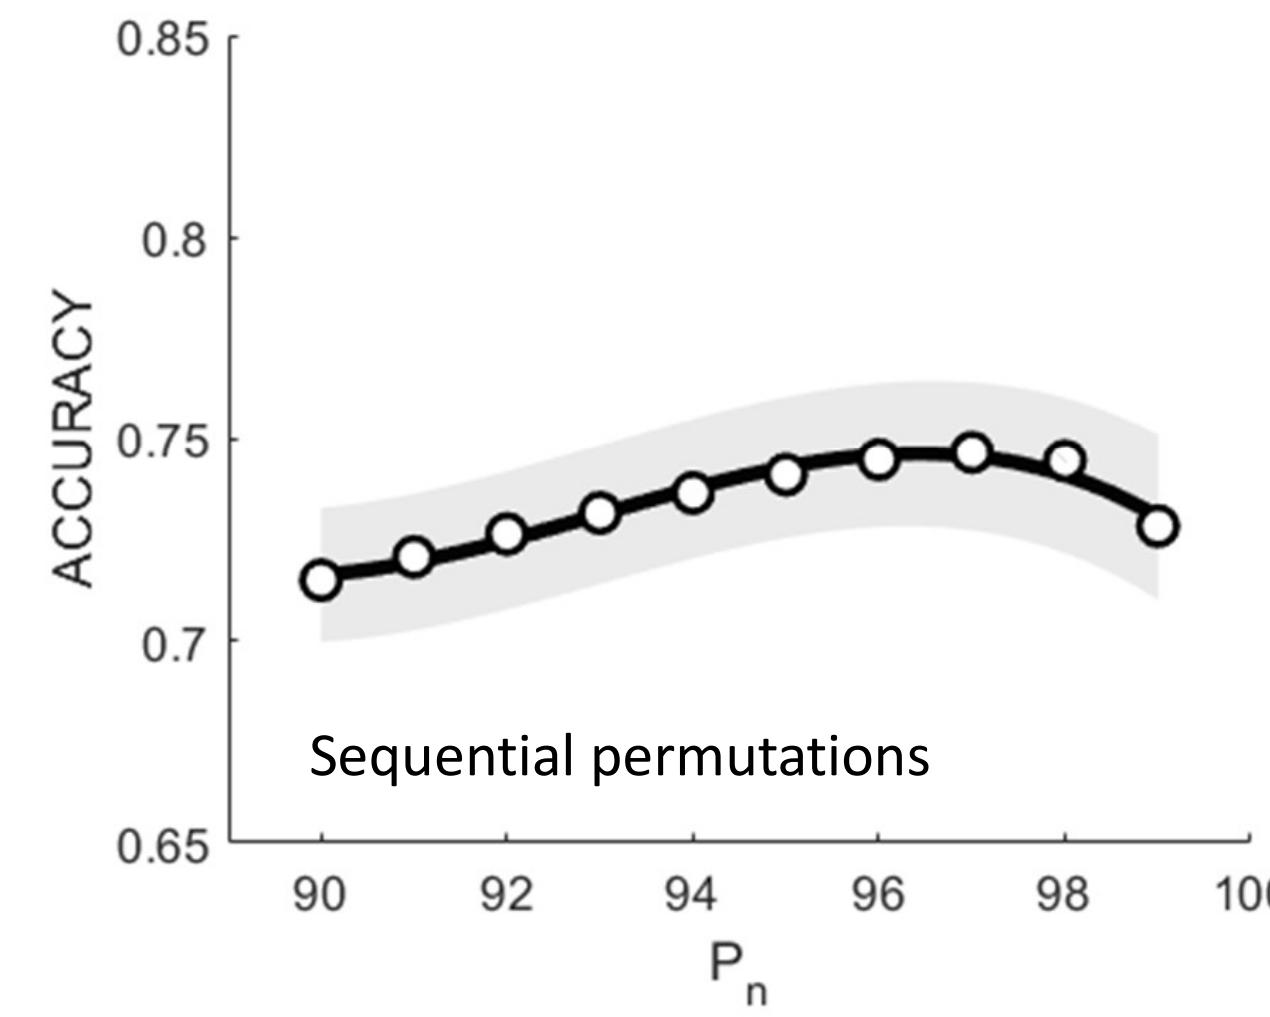
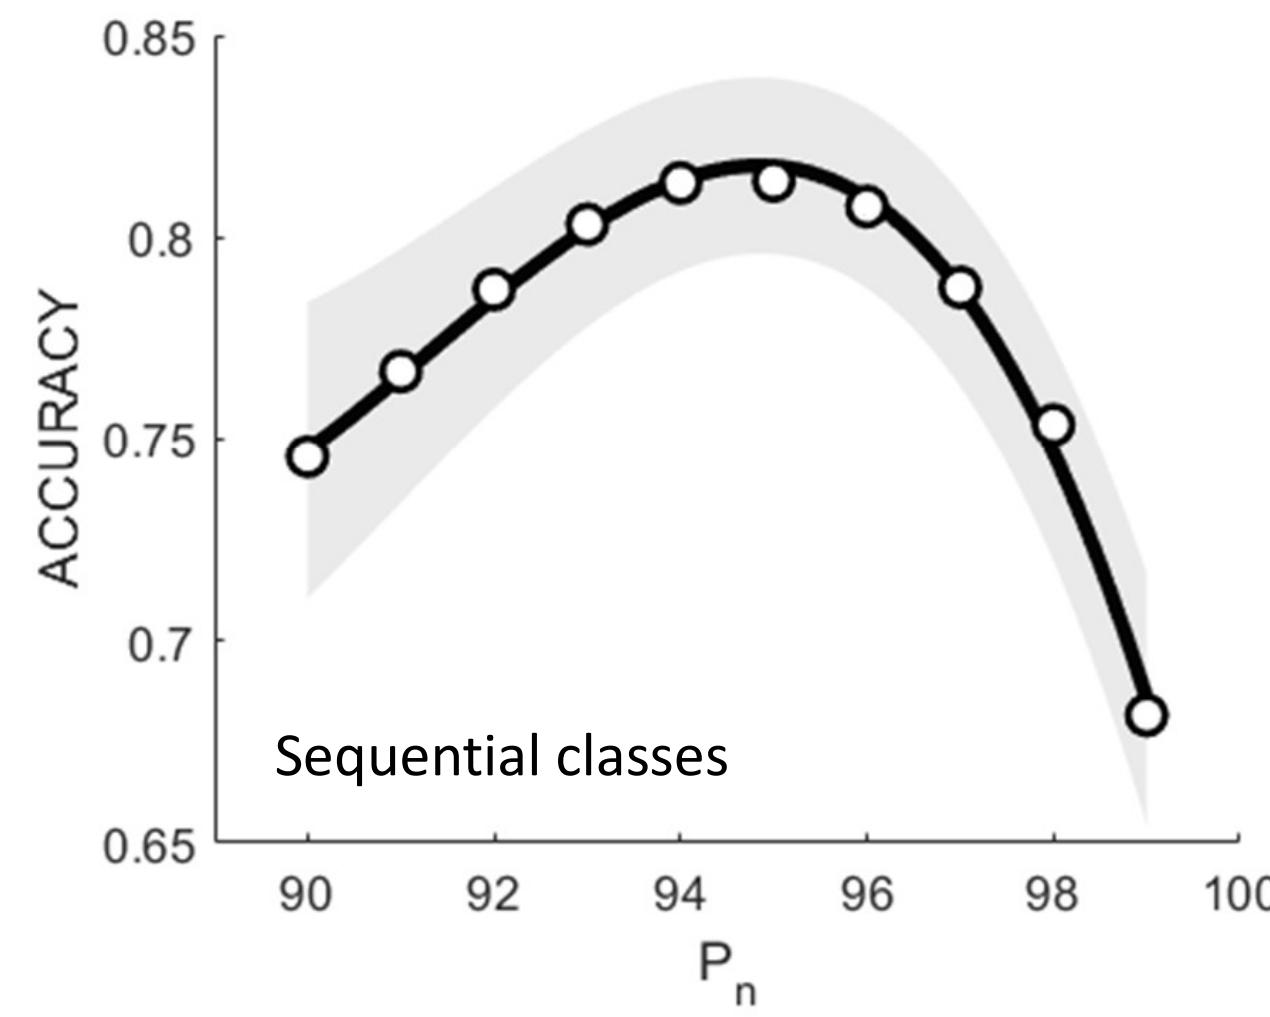
*Deng (2012) IEEE Signal Processing Magazine

**Chandar et al (2019) AAAI

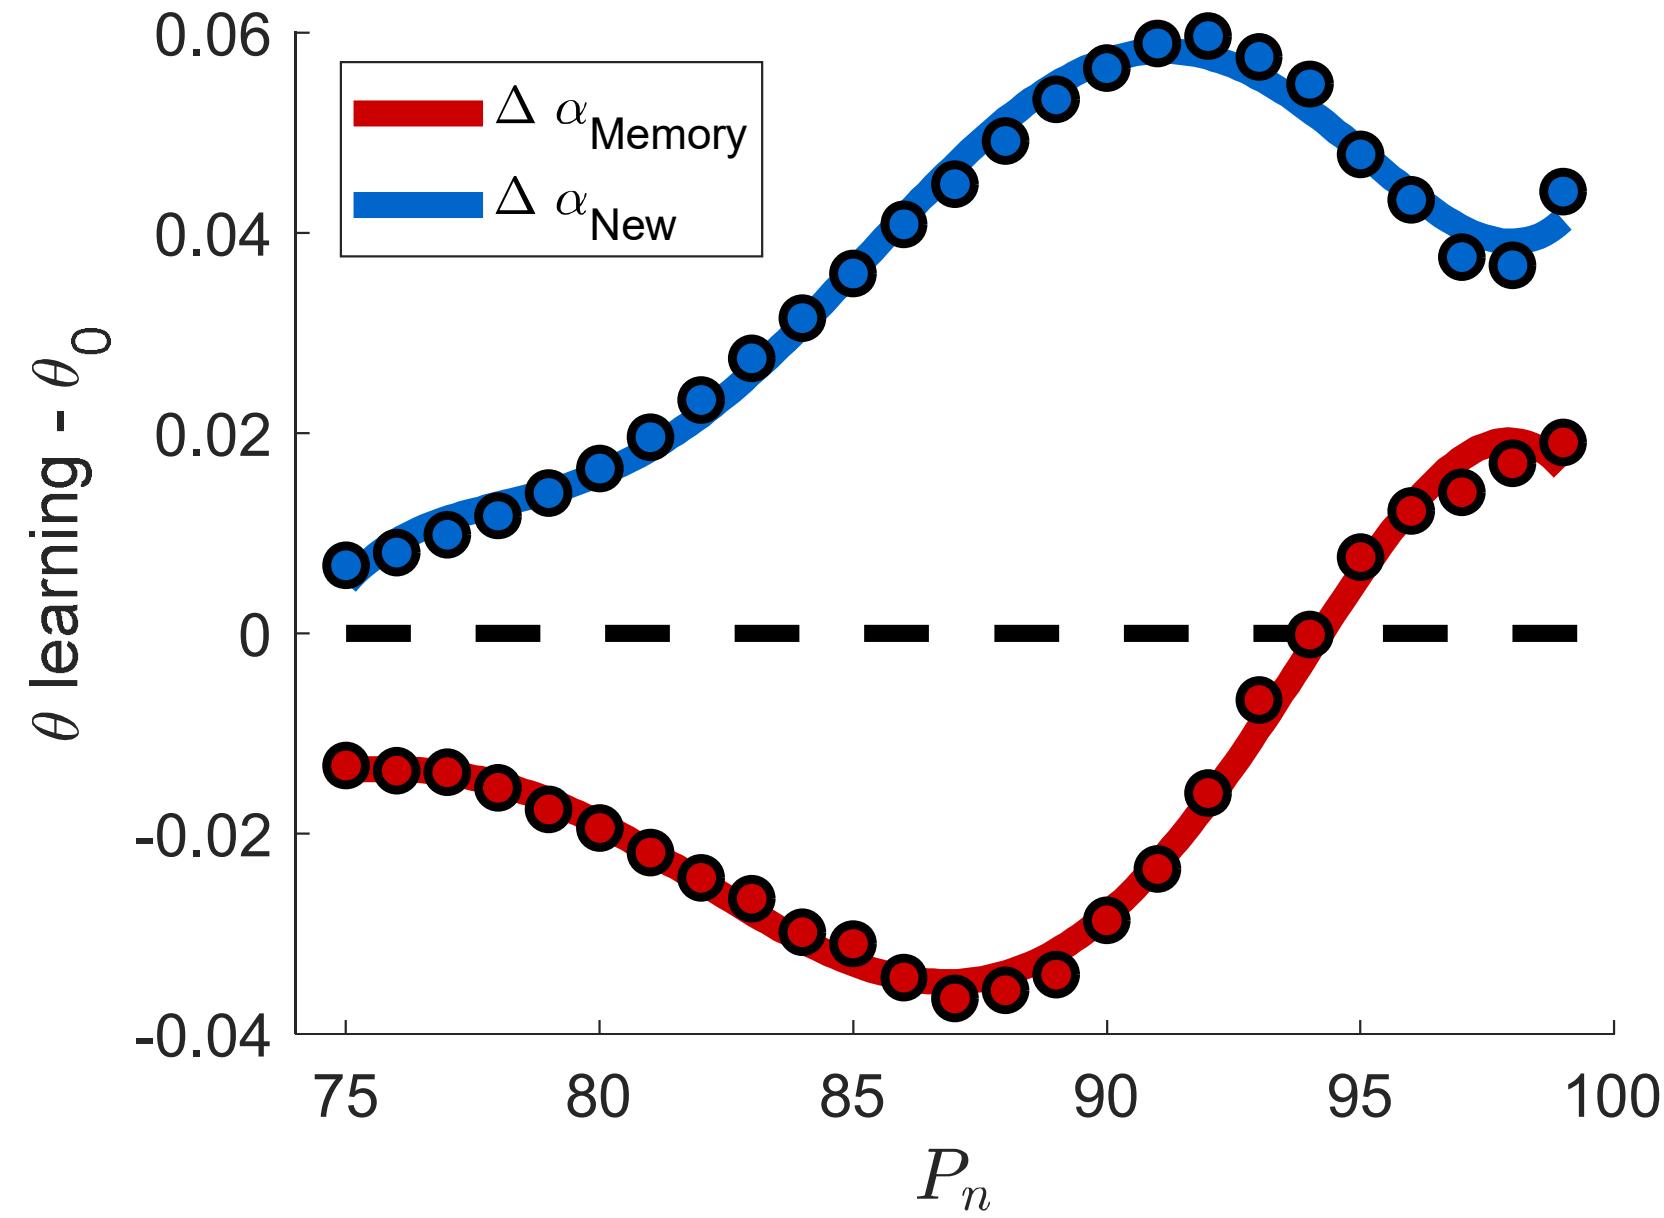
CATASTROPHIC FORGETTING



CATASTROPHIC FORGETTING



THRESHOLD INITIALISATION VS THRESHOLD LEARNING

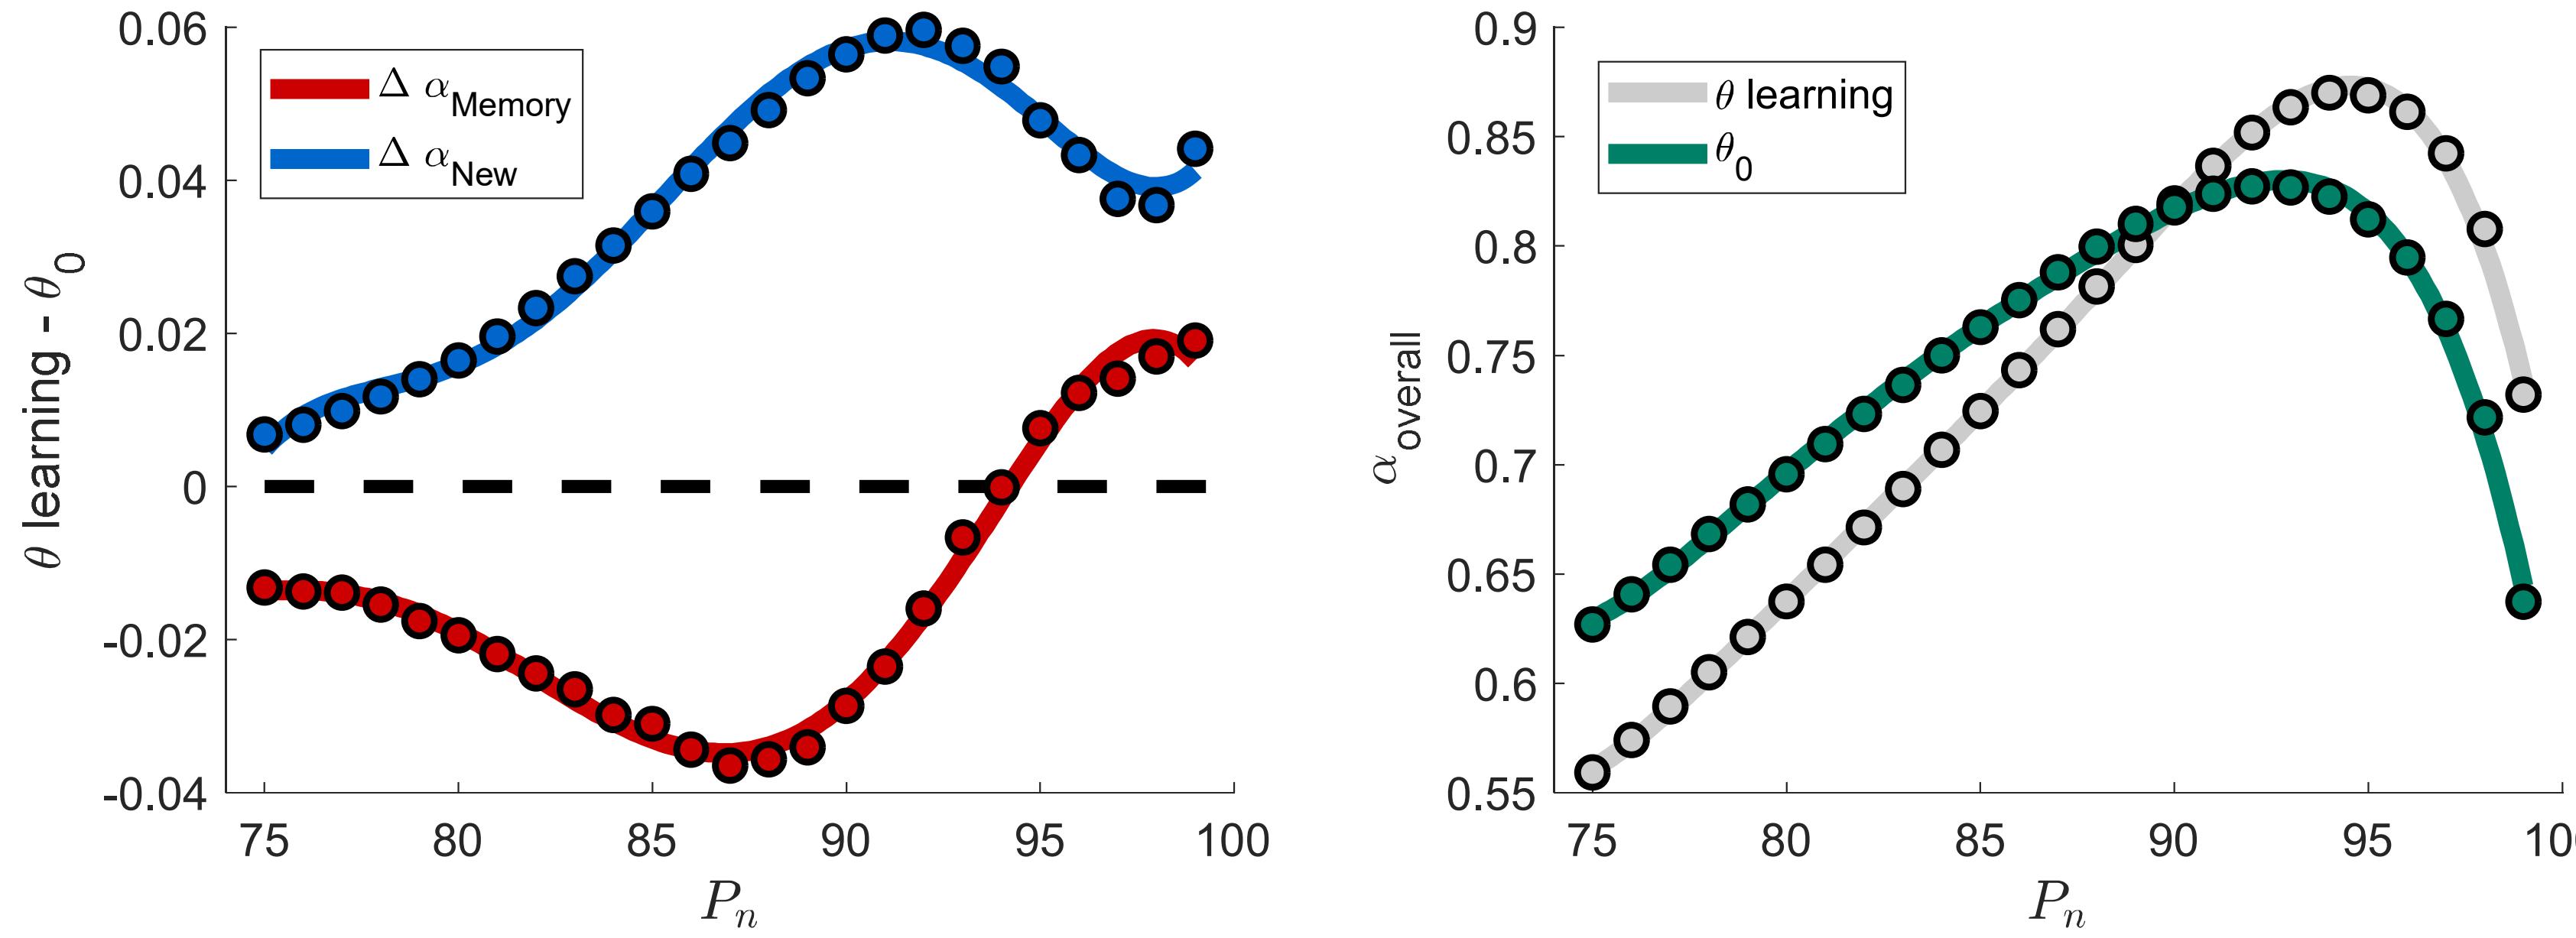


$$\alpha_{\text{Memory}} = \frac{1}{N_{\text{task}}} \sum_{n=1}^{N_{\text{task}}} [\text{acc}(n, N_{\text{task}}) - \text{acc}(n, n)]$$

$$\alpha_{\text{New}} = \frac{1}{N_{\text{task}}} \sum_{n=1}^{N_{\text{task}}} \text{acc}(n, n)$$

Metrics from: R. Kemker, M. McClure, A. Abitino, T. L. Hayes, and C. Kanan, Thirty-second AAAI conference on artificial intelligence, 2018.

THRESHOLD INITIALISATION VS THRESHOLD LEARNING



$$\alpha_{\text{Memory}} = \frac{1}{N_{\text{task}}} \sum_{n=1}^{N_{\text{task}}} [\text{acc}(n, N_{\text{task}}) - \text{acc}(n, n)]$$

$$\alpha_{\text{New}} = \frac{1}{N_{\text{task}}} \sum_{n=1}^{N_{\text{task}}} \text{acc}(n, n)$$

$$\alpha_{\text{Overall}} = \frac{1}{N_{\text{task}} - 1} \sum_{n=2}^{N_{\text{task}}} \frac{\text{acc}_n}{\text{acc}(1,1)}$$

Metrics from: R. Kemker, M. McClure, A. Abitino, T. L. Hayes, and C. Kanan, Thirty-second AAAI conference on artificial intelligence, 2018.

CATASTROPHIC FORGETTING

Results, $\alpha_{overall}$			
Sequential classes		Sequential permutations	
<i>ESN</i>	0.1	<i>ESN</i>	0.1
<i>SpaRCe</i>	0.870	<i>SpaRCe</i>	0.897
<i>EWC</i>	0.133 ⁺	<i>EWC</i>	0.746 ⁺
<i>FEL</i>	0.439 ⁺	<i>FEL</i>	0.279 ⁺
<i>GeppNet</i>	0.922 ⁺	<i>GeppNet</i>	0.364 ⁺

+R. Kemker, M. McClure, A. Abitino, T. L. Hayes, and C. Kanan,
Thirty- second AAAI conference on artificial intelligence, 2018.

<https://arxiv.org/abs/2110.05572v2>

EchoVPR: Echo State Networks for Visual Place Recognition

Anıl Özdemir^{1,†}, Mark Scerri^{1,†}, Andrew B. Barron², Andrew Philippides³,
Michael Mangan^{1,*}, Eleni Vasilaki^{1,4,*}, and Luca Manneschi^{1,*}

Abstract— Recognising previously visited locations is an important, but unsolved, task in autonomous navigation. Current visual place recognition (VPR) benchmarks typically challenge models to recover the position of a query image (or images) from sequential datasets that include both spatial and temporal components. Recently, Echo State Network (ESN) varieties have proven particularly powerful at solving machine learning tasks that require spatio-temporal modelling. These networks are simple, yet powerful neural architectures that—exhibiting memory over multiple time-scales and non-linear high-dimensional representations—can discover temporal relations in the data while still maintaining linearity in the learning. In this paper, we present a series of ESNs and analyse their applicability to the VPR problem. We report that the addition of ESNs to pre-processed convolutional neural networks led to a dramatic boost in performance in comparison to non-recurrent networks in five out of six standard benchmarks (GardensPoint, SPEDTest, ESSEX3IN1, Oxford RobotCar, and Nordland) demonstrating that ESNs are able to capture the temporal structure inherent in VPR problems. Moreover, we show that models that include ESNs can outperform class-leading VPR models which also exploit the sequential dynamics of the data. Finally, our results demonstrate that ESNs also improve generalisation abilities, robustness, and accuracy further supporting their suitability to VPR applications.

change (e.g. DenseVLAD [4], NetVLAD [5], AMOSNet [6], SuperGlue [7], DELG [8], Patch-NetVLAD [9], and HEAPUtil [10]). While matching single images is successful in many scenarios, it can suffer from the effects of aliasing, individual image corruption, or sampling mismatches between datasets (e.g. it is challenging to ensure that images sampled along the same route precisely overlap).

Milford and Wyeth [11] were the first to demonstrate that such issues could be overcome by matching sequences of images using a global search to overcome individual image mismatches. This has led to a family of models that improve VPR performance by exploiting the temporal relationships inherent in images sampled along routes [11], [12], [13], [14], [15], [16], [17]. While achieving state-of-the-art performance on challenging real-world datasets (e.g. Oxford RobotCar [18], Extended CMU Seasons [19], Pittsburgh [20], Tokyo24/7 [4], Nordland [21], Mapillary Street Level Sequences [22]), such models often include an explicit encoding of non-visual information to limit the image search space, require a cache of input images for global searches, and can be computationally expensive in

CHALLENGING TASKS

Visual Place Recognition: The Norland Test

Ozdemir, Scerri et al (2022), IEEE Robotics and Automation Letters

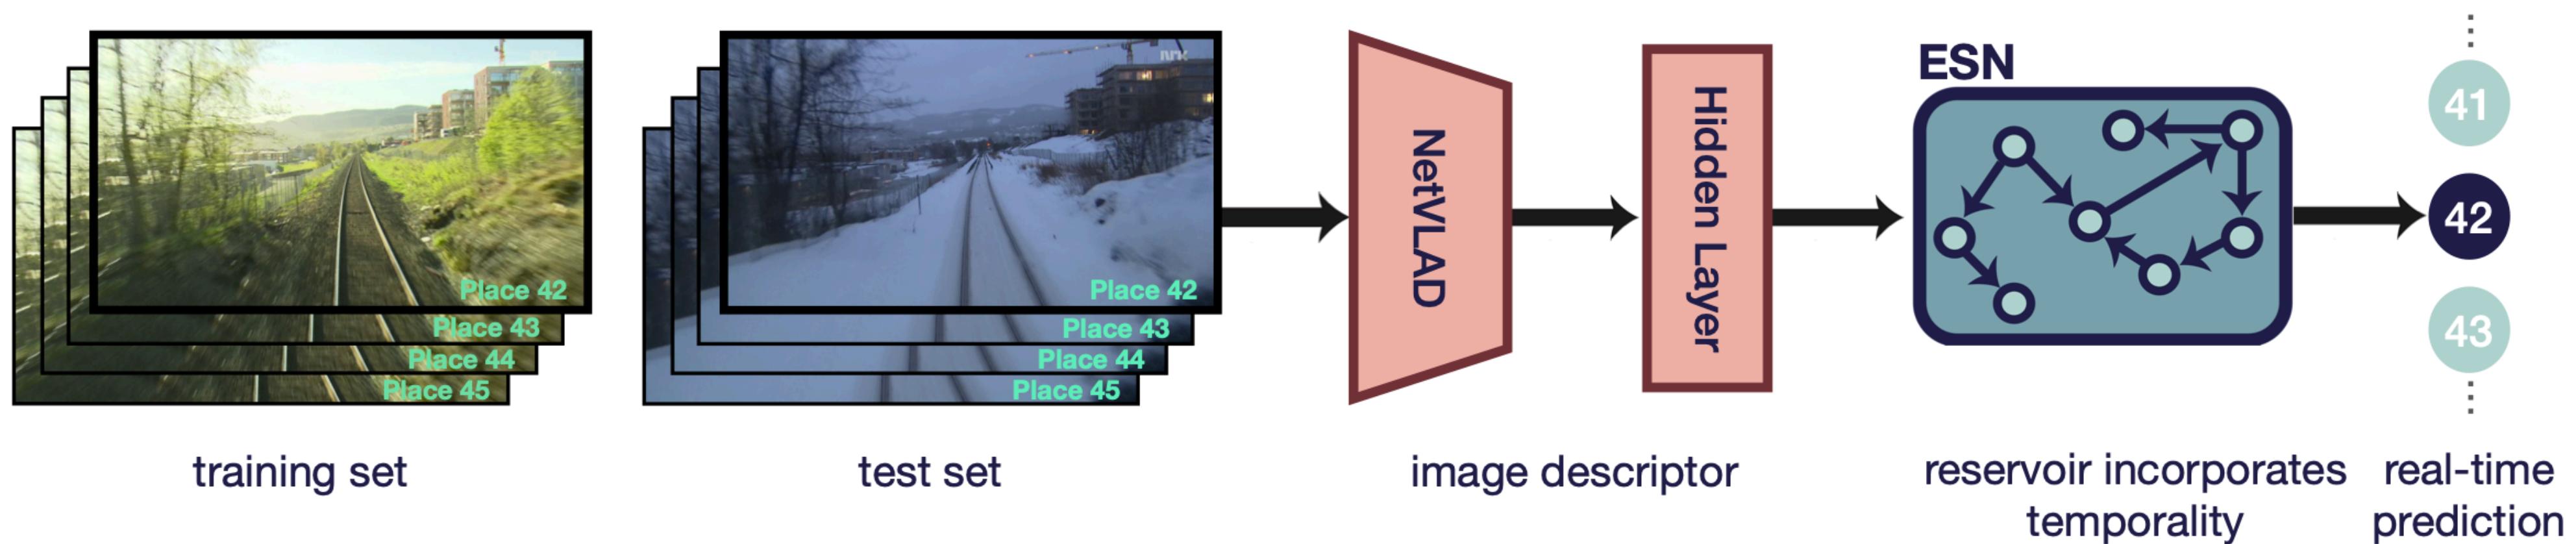
CHALLENGING TASKS

Visual Place Recognition: The Norland Test



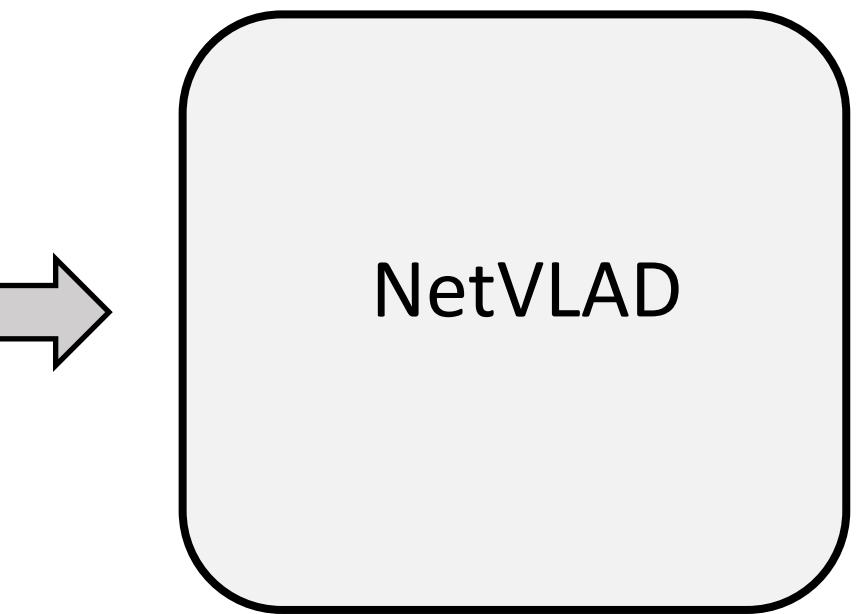
Ozdemir, Scerri et al (2022), IEEE Robotics and Automation Letters

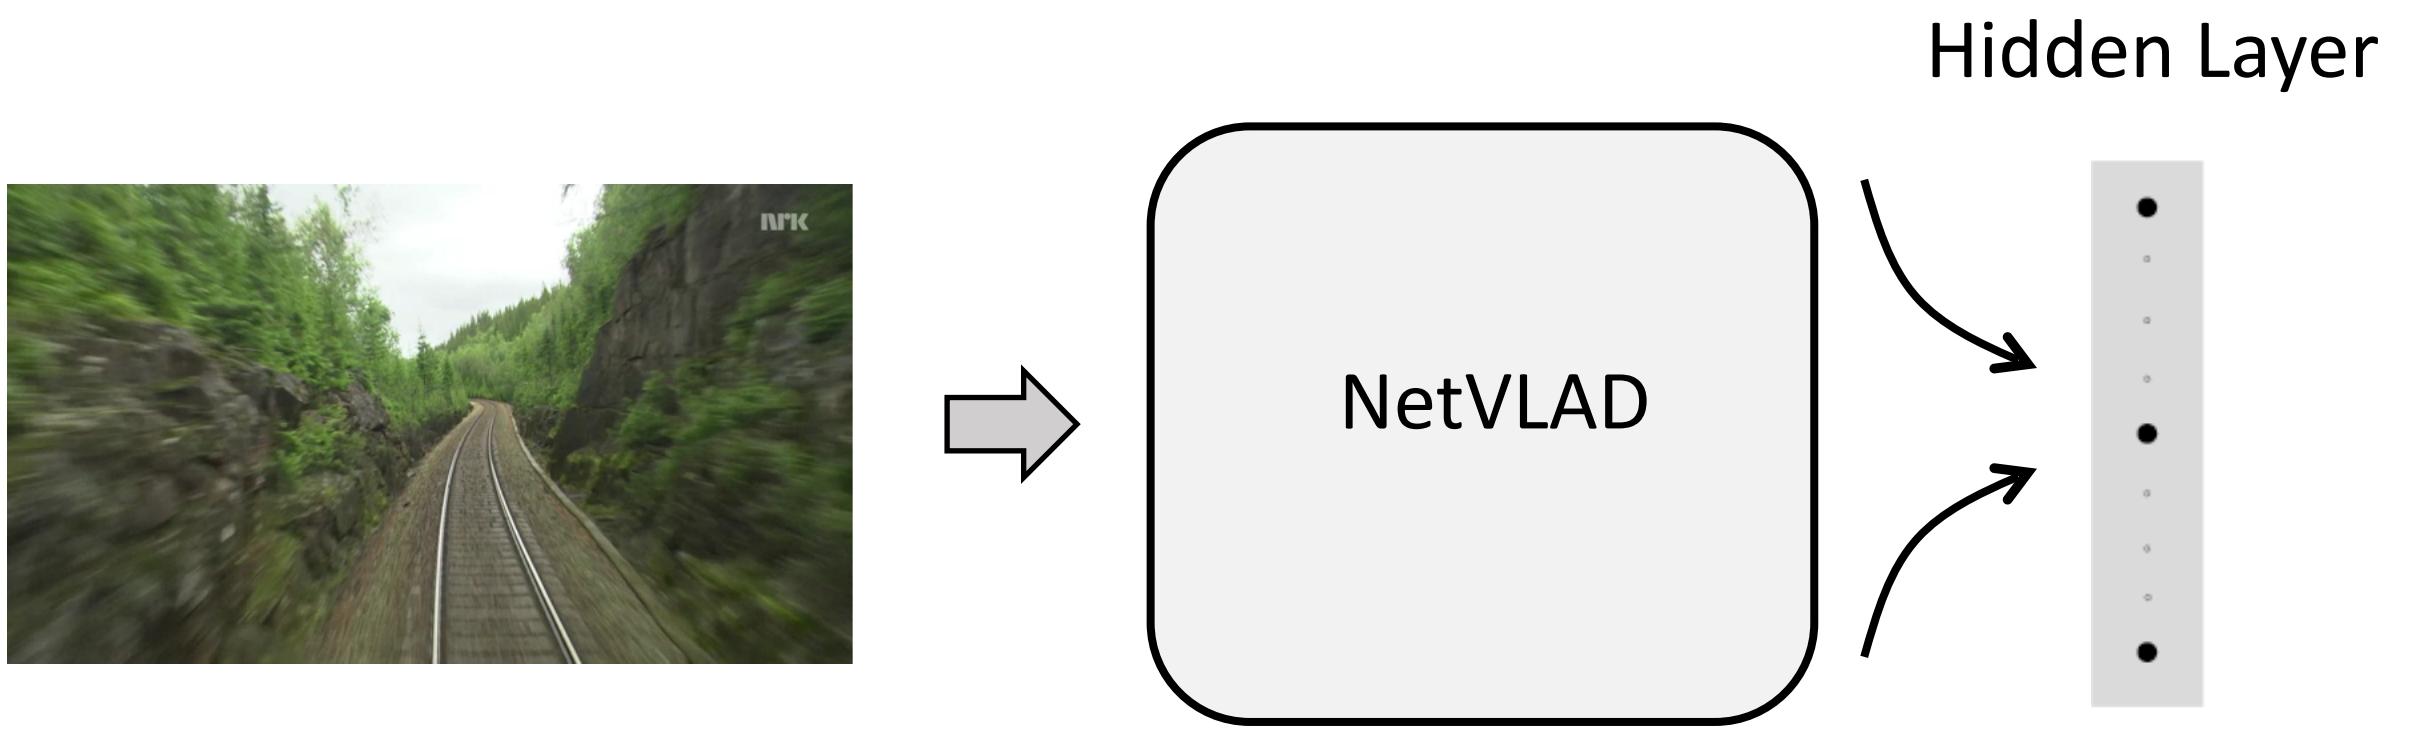
ECHO-VPR FRAMEWORK

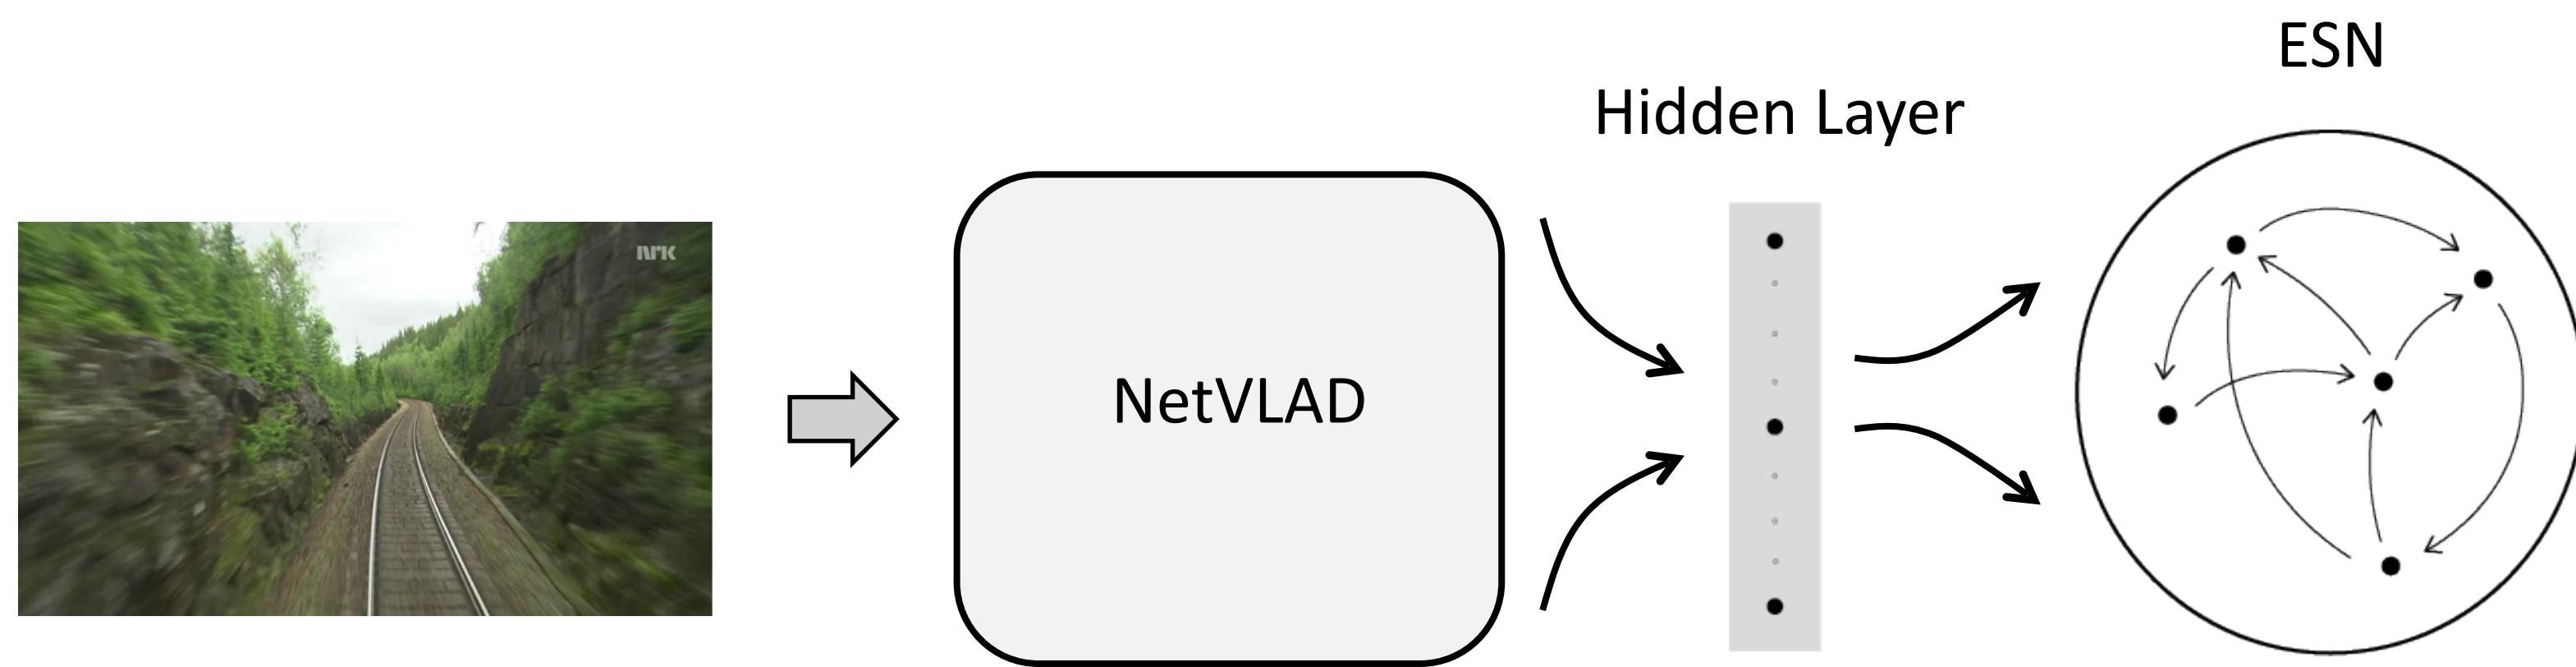


*We used a small subset removed from the test set
to tune hyper-parameters without overfitting.*

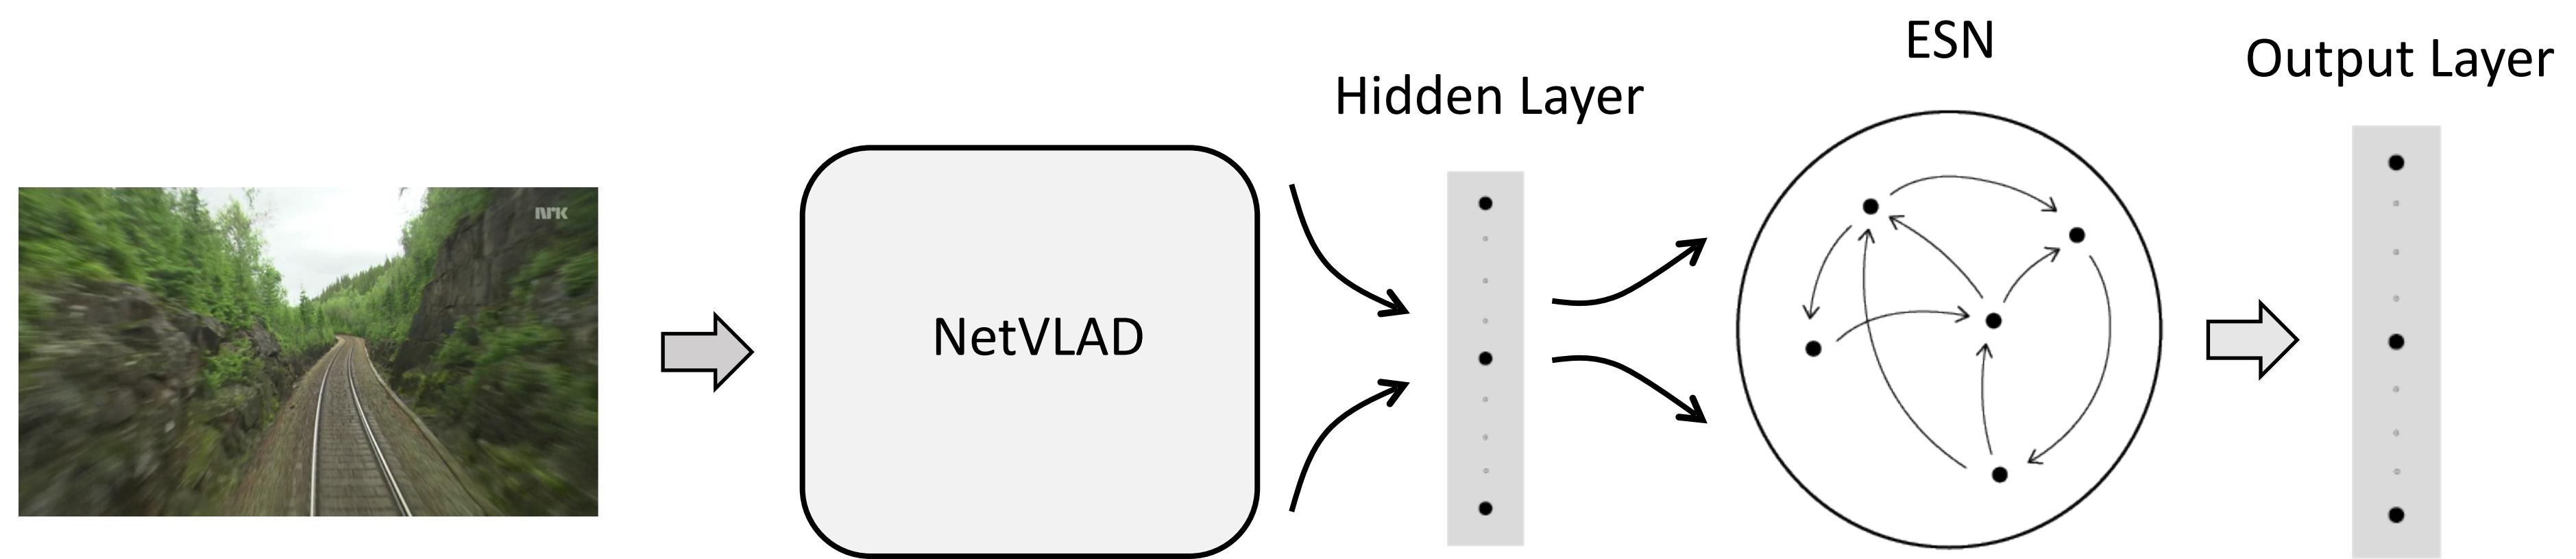
Ozdemir, Scerri et al (2022), IEEE Robotics and Automation Letters

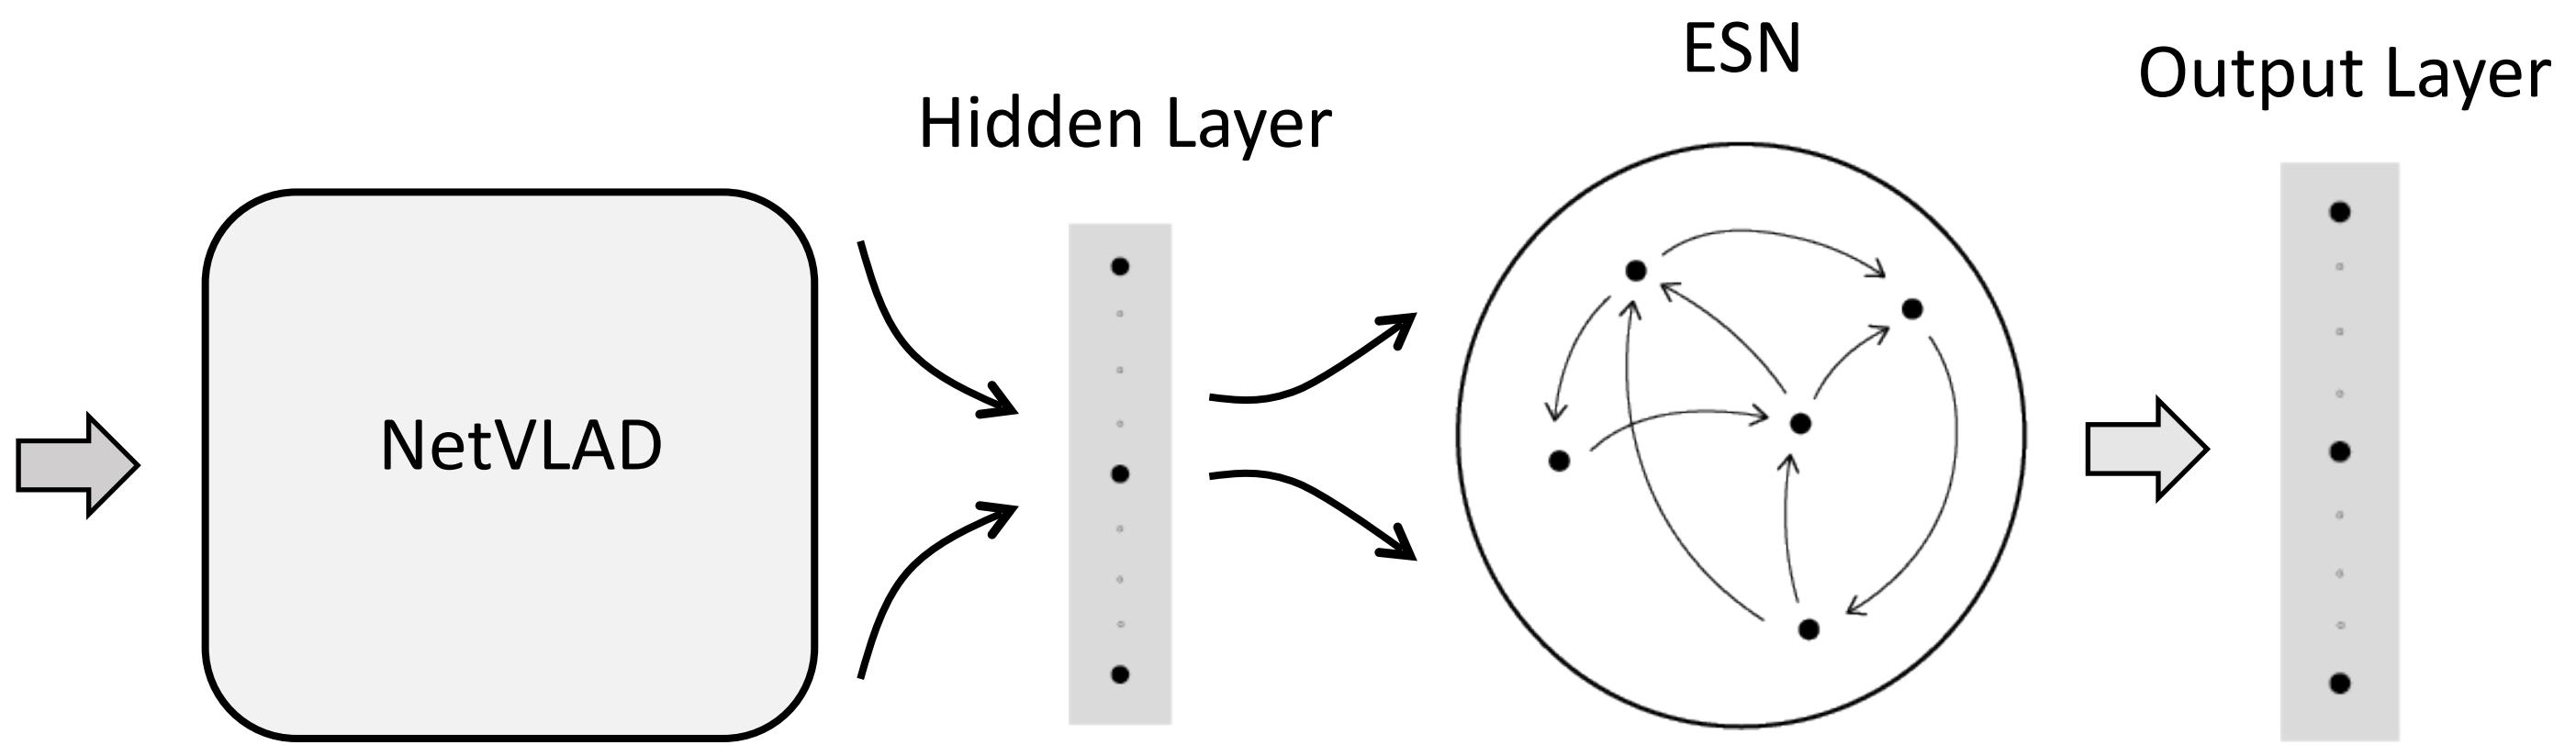






Ozdemir, Scerri et al (2022), IEEE Robotics and Automation Letters

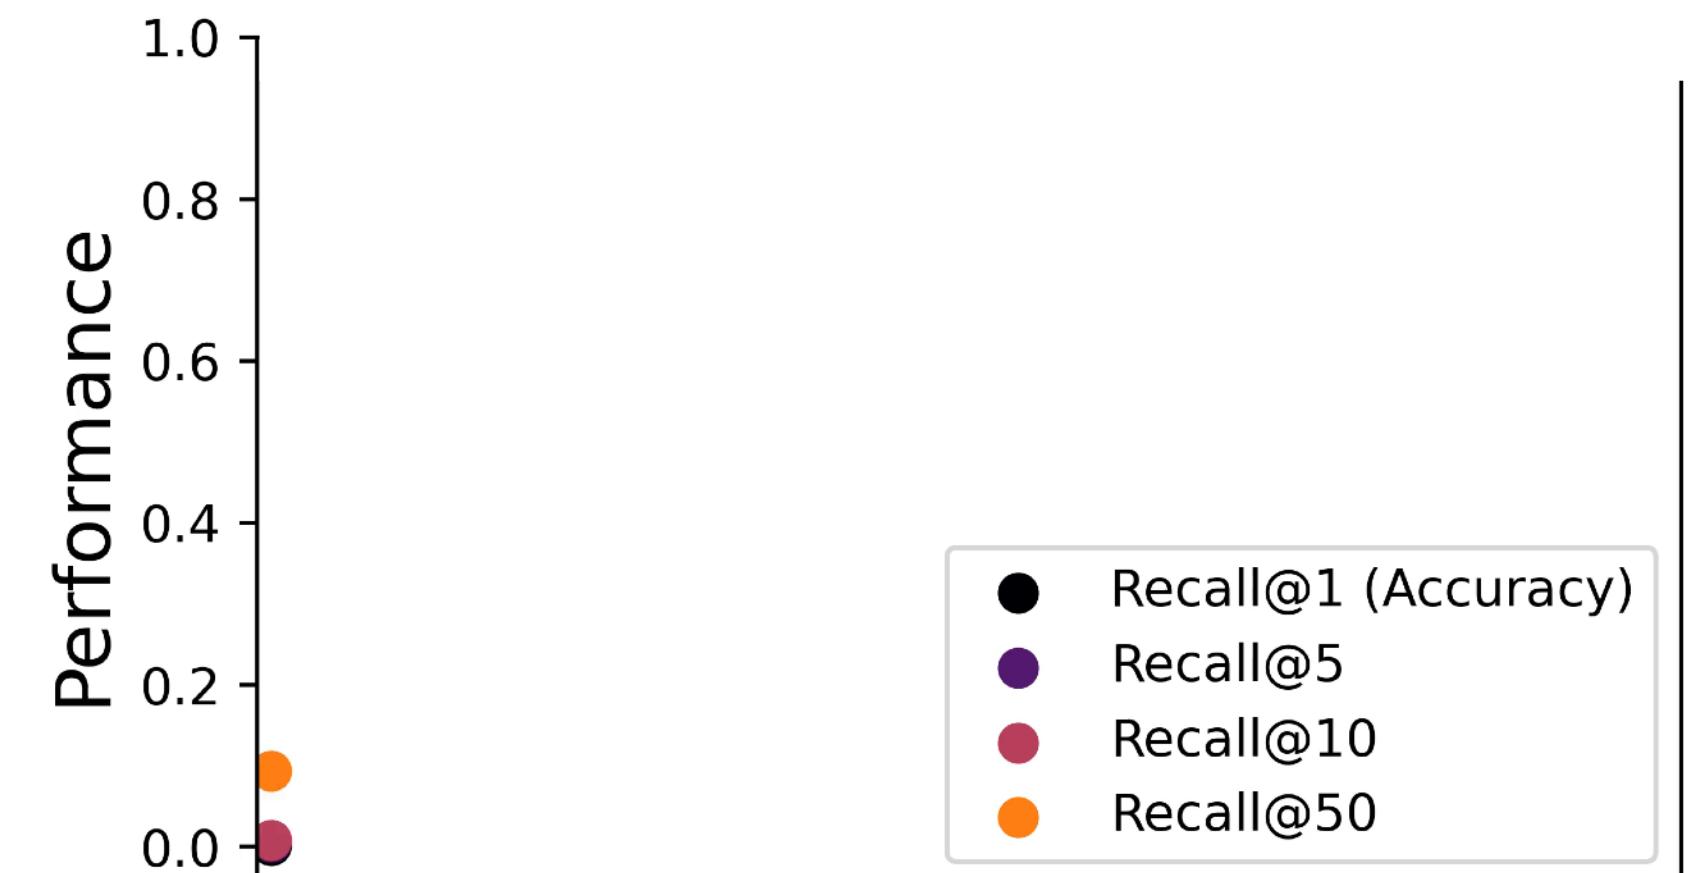




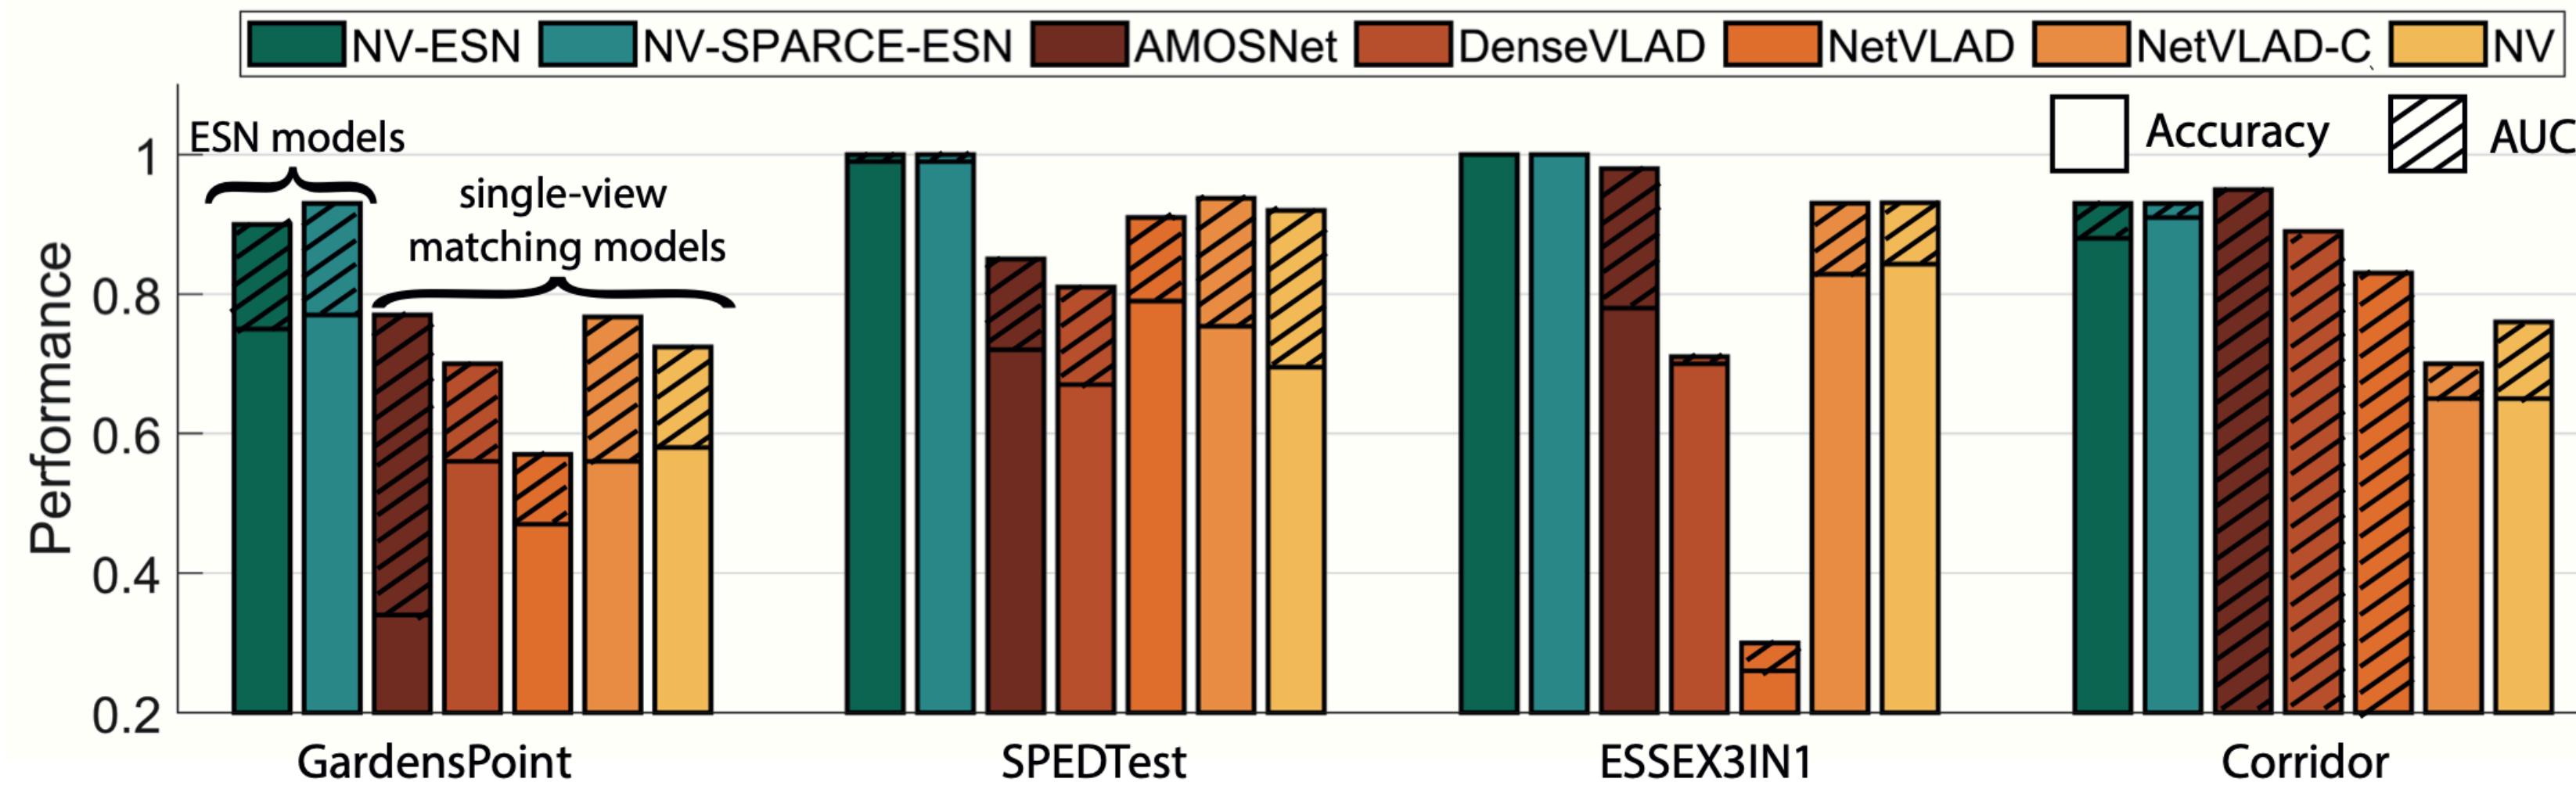
ESN Activities example

ESN Activities example

Output Layer example



RESULTS IN “SMALL” DATASETS (200–600 IMAGES)



RESULTS IN “LARGE” DATASETS (~30,000 IMAGES)

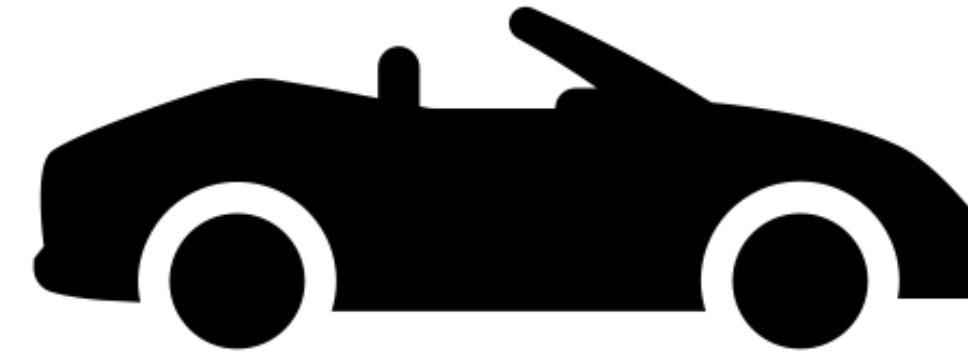
Method	Nordland Summer vs Winter			Oxford RobotCar Day vs Night		
	R@1	R@5	R@10	R@1	R@5	R@10
NetVLAD [5]	13.0	20.6	25.0	31.2	48.6	58.3
SuperGlue [7]	29.1	33.4	35.0	—	—	—
DELG global [8]	23.4	35.4	41.7	—	—	—
DELG local [8]	60.1	63.5	64.6	—	—	—
Patch-NetVLAD [9]	44.9	50.2	52.2	80.7	84.8	86.4
NV-ESN	47.5	61.0	66.4	44.0	58.1	65.5
NV-SPARCE-ESN	46.7	60.3	65.9	80.4	87.9	91.5
PatchL-ESN	66.4	76.0	78.7	75.9	85.2	87.7
PatchL-SPARCE-ESN	66.3	75.7	78.6	88.8	97.0	98.2

WHY DO WE CARE ABOUT EFFICIENT ML MODELS?

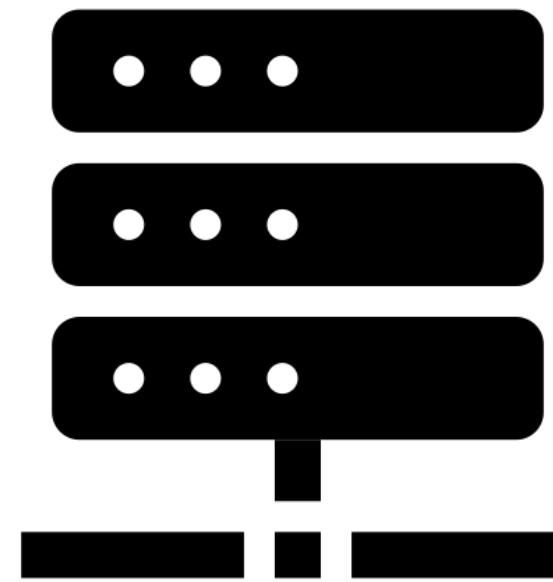
2019

Energy and Policy Considerations for Deep Learning in NLP

Emma Strubell Ananya Ganesh Andrew McCallum
College of Information and Computer Sciences
University of Massachusetts Amherst
`{strubell, aganesh, mccallum}@cs.umass.edu`



VS



WHY DO WE CARE ABOUT EFFICIENT ML MODELS?

2019

Energy and Policy Considerations for Deep Learning in NLP

Emma Strubell Ananya Ganesh Andrew McCallum
College of Information and Computer Sciences
University of Massachusetts Amherst
`{strubell, aganesh, mccallum}@cs.umass.edu`



Consumption	CO ₂ e (lbs)
Air travel, 1 passenger, NY↔SF	1984
Human life, avg, 1 year	11,023
American life, avg, 1 year	36,156
Car, avg incl. fuel, 1 lifetime	126,000

Training one model (GPU)	
NLP pipeline (parsing, SRL)	39
w/ tuning & experimentation	78,468
Transformer (big)	192
w/ neural architecture search	626,155

VS

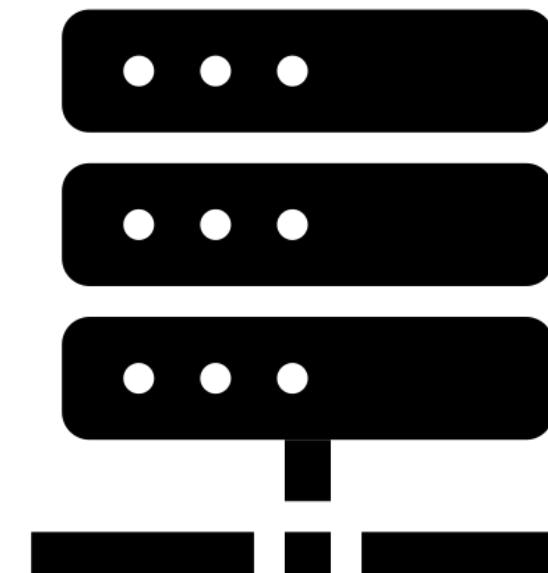


Table 1: Estimated CO₂ emissions from training common NLP models, compared to familiar consumption.¹

SUMMARY

SUMMARY

- SPARCe: SParse Reservoir Computing, inspired by the insect brain.

SUMMARY

- SPARCe: SParse Reservoir Computing, inspired by the insect brain.
- Key ingredients: threshold learning and “democratic” threshold initialisation.

SUMMARY

- SPARCe: SParse Reservoir Computing, inspired by the insect brain.
- Key ingredients: threshold learning and “democratic” threshold initialisation.
- Sparsity leads to neuronal specialisation and alleviates catastrophic forgetting.

SUMMARY

- SPARCe: SParse Reservoir Computing, inspired by the insect brain.
- Key ingredients: threshold learning and “democratic” threshold initialisation.
- Sparsity leads to neuronal specialisation and alleviates catastrophic forgetting.
- Threshold learning “kills” correlated neurons in the reservoir and helps increase neuronal specialisation.

SUMMARY

- SPARCe: SParse Reservoir Computing, inspired by the insect brain.
- Key ingredients: threshold learning and “democratic” threshold initialisation.
- Sparsity leads to neuronal specialisation and alleviates catastrophic forgetting.
- Threshold learning “kills” correlated neurons in the reservoir and helps increase neuronal specialisation.
- It helps increase performance in problems where memory is required, such as the VPR.

Thank you!